Covers
Hibernate 3.2

Beginning
Hibernate

From Novice to Professional

*An introduction to all the new features
of the Hibernate 3.2 persistence API*

Dave Minter **and** Jeff Linwood

APRESS®

Beginning Hibernate

From Novice to Professional

Dave Minter and
Jeff Linwood

Apress®

Beginning Hibernate: From Novice to Professional

Copyright © 2006 by Dave Minter, Jeff Linwood

ISBN-13 (pbk): 978-1-59059-693-7

ISBN-10 (pbk): 1-59059-693-5

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

The source code for this book is available to readers at `http://www.apress.com` in the Source Code section.

Contents at a Glance

Contents

About the Authors

■ DAVE MINTER has adored computers since he was small enough to play in the boxes they came in. He built his first PC from discarded, faulty, and obsolete components, and considers that to be the foundation of his career as an integration consultant. Dave is based in London, where he helps large and small companies build systems that work the way they're supposed to. He co-authored *Building Portals with the Java Portlet API* (Apress, 2004) and *Pro Hibernate 3* (Apress, 2005).

■ JEFF LINWOOD has been involved in software programming since he had a 286 in high school. He became caught up with the Internet when he got access to a UNIX shell account, and it has been downhill ever since. Jeff has published articles on several Jakarta Apache open source projects in *Dr. Dobb's Journal*, CNET's Builder.com, and *JavaWorld*. Jeff has a bachelor's degree in chemical engineering from Carnegie Mellon University. He currently works for the Gossamer Group in Austin, Texas, on content management and web application syndication systems. He gets to play with all the latest open source projects there. Jeff also co-authored *Professional Struts Applications* (Apress, 2003), *Building Portals with the Java Portlet API* (Apress, 2004), and *Pro Hibernate 3* (Apress, 2005). He was a technical reviewer for *Enterprise Java Development on a Budget* (Apress, 2003) and *Extreme Programming with Ant* (Sams Publishing, 2003).

About the Technical Reviewer

■**SUMIT PAL** has about 12 years of experience with software architecture, design, and development on a variety of platforms including Java J2EE. Sumit has worked with the SQL Server replication group while with Microsoft for two years, and with Oracle's OLAP Server group while with Oracle for seven years.

In addition to certifications like IEEE CSDP and J2EE Architect, Sumit also has a master's degree in computer science from the Asian Institute of Technology, Thailand.

Sumit has a keen interest in database internals, algorithms, and search engine technology. He has invented some basic generalized algorithms to find divisibility between numbers, and has also invented divisibility rules for prime numbers less than 100.

Currently, he loves to play as much as he can with his 22-month-old daughter.

Acknowledgments

Introduction

Hibernate is an amazing piece of software. With a little experience and the power of Java 5 annotations, you can build a complex database-backed system with disturbing ease. Once you have built a system using Hibernate, you will never want to go back to the traditional approaches.

While Hibernate is incredibly powerful, it presents a steep learning curve when you first encounter it—steep learning curves are actually a good thing, as they impart profound insight once you have scaled them. Yet gaining that insight takes some perseverance and assistance.

Our aim in this book is to help you up that learning curve by presenting you with the minimal requirements of a discrete Hibernate application, explaining the basis of those requirements, and walking you through an example application built according to them. We then provide additional material to be digested once the fundamentals are firmly understood. Throughout, we provide examples rather than relying upon pure discourse.

We hope that you will continue to find this book useful as a reference text long after you have become an expert on the subject.

Who This Book Is For

This book assumes a good understanding of Java fundamentals and some familiarity with database programming using the Java Database Connectivity (JDBC) API. We don't expect you to know anything about Hibernate—but if you buy this book, it will probably be because you have some exposure to the painful process of building a large database-based system.

All of our examples use open source software—primarily the Hibernate API itself—so you will not need to purchase any software to get started with Hibernate development.

This book is not an academic text. Our focus is instead on providing extensive examples and taking a pragmatic approach to the technology that it covers.

To true newcomers to the Hibernate API, we recommend that you read at least the first three chapters in order before diving into the juicy subjects of later chapters. Very experienced developers or those with experience with tools similar to Hibernate will want to skim through the latter half of the book for interesting chapters. Readers familiar with Hibernate will want to turn to the appendixes for discussion of more arcane topics.

How This Book Is Structured

This book is informally divided into three parts. Chapters 1 through 8 describe the fundamentals of Hibernate, including configuration, the creation of mapping files, and the basic APIs.

Chapters 9 through 11 then describe the use of queries, criteria, and filters to access the persistent information in more sophisticated ways.

Finally, the appendixes discuss features that you will use less often, or that are peripheral to the core Hibernate functionality. The following list describes more fully the contents of each chapter:

Chapter 1 outlines the purpose of persistence tools and presents excerpts from a simple example application to show how Hibernate can be applied. It also introduces core terminology and concepts.

Chapter 2 discusses the fundamentals of configuring a Hibernate application. It presents the basic architecture of Hibernate and discusses how a Hibernate application is integrated into an application.

Chapter 3 presents the example application from Chapter 1 in its entirety, walking you through the complete process of creating and running the application. It then looks at a slightly more complex example and introduces the notion of generating the database schema directly from the mapping files.

Chapter 4 covers the Hibernate life cycle in depth. It discusses the life cycle in the context of the methods available on the core interfaces. It also introduces key terminology and discusses the need for cascading and lazy loading.

Chapter 5 explains why mapping information must be retained by Hibernate, and demonstrates the various types of associations that can be represented by a relational database. It briefly discusses the other information that can be maintained within a Hibernate mapping.

Chapter 6 explains how Hibernate lets you use the Java 5 Annotations feature to represent mapping information. It provides detailed examples for the most important annotations, and discusses the distinctions between the standard EJB 3 annotations and the proprietary Hibernate 3 ones.

Chapter 7 explains how the XML-based mapping files can be used to represent mapping information in Hibernate. It provides examples for all of the most common mapping types and reference notes for the more obscure ones.

Chapter 8 revisits the Hibernate Session object in detail, explaining the various methods that it provides. The chapter also discusses the use of transactions, locking, and caching, and how to use Hibernate in a multithreaded environment.

Chapter 9 discusses how Hibernate can be used to make sophisticated queries against the underlying relational database using the built-in Hibernate Query Language (HQL).

Chapter 10 introduces the Criteria API, which is a programmatic analog of the query language discussed in Chapter 9.

Chapter 11 discusses how the filter API can be used to restrict the results of the queries introduced in Chapters 9 and 10.

Appendix A presents a large number of peripheral features that do not warrant more extensive coverage in a beginners' text. Basic discussion is given, with examples, of the use of the Hibernate `EntityManager` and EJB 3, the support for versioning and optimistic locking, the provision for persisting and retrieving Dom4J document models directly, the support for persisting and retrieving maps of information, and some of the obscure limitations of Hibernate and various ways that these can be worked around. It also discusses the use of events and interceptors.

Appendix B discusses how the Hibernate Tools toolset can be used to enhance development with the Eclipse development environment and the Ant build tool. It also explains how the Ant code-generation tasks can be customized.

Appendix C discusses how Hibernate can be integrated into the Spring API. The integration of Hibernate as the persistence layer of a Spring application is complex, so we present a working example, including the entire bean definition file, with discussions of the appropriate way to manage the session in the Spring MVC environment, and how Spring can enforce the proper transactional boundaries when using Hibernate.

Appendix D discusses some topics of interest to developers who are working with a pre-existing base of code that was built using version 2 of Hibernate. We present the various approaches to coexisting with Hibernate 3 code and to migrating a Hibernate 2 code base to the Hibernate 3 API.

Downloading the Code

The source code for this book is available to readers from `www.apress.com`, in the Source Code/Download section. Please feel free to visit the Apress web site and download all the code from there.

Contacting the Authors

We welcome feedback from our readers. If you have any queries or suggestions about this book, or technical questions about Hibernate, or if you just want to share a really good joke, you can e-mail Dave Minter at `dave@paperstack.com` and Jeff Linwood at `jlinwood@gmail.com`.

An Introduction to Hibernate 3

Most significant development projects involve a relational database. The mainstay of most commercial applications is the large-scale storage of ordered information, such as catalogs, customer lists, contract details, published text, and architectural designs.

With the advent of the World Wide Web, the demand for databases has increased. Though they may not know it, the customers of online bookshops and newspapers are using databases. Somewhere in the guts of the application a database is being queried and a response is offered.

While the demand for such applications has grown, their creation has not become noticeably simpler. Some standardization has occurred—the most successful being the Enterprise JavaBeans (EJB) standard of Java 2 Enterprise Edition (J2EE), which provides for container- and bean-managed persistence of entity bean classes. Unfortunately, this and other persistence models all suffer to one degree or another from the mismatch between the relational model and the object-oriented model. In short, database persistence is difficult.

There are solutions for which EJBs are appropriate, some for which some sort of object-relational mapping (ORM) like Hibernate is appropriate, and some for which the traditional approach of direct access via the Java Database Connectivity (JDBC) API is appropriate. We think that Hibernate represents a good first choice, as it does not preclude the simultaneous use of these alternative approaches.

To illustrate some of Hibernate's strengths, in this chapter we will show you a brief example using Hibernate and contrast this with the traditional JDBC approach.

Plain Old Java Objects (POJOs)

In our ideal world, it would be trivial to take any Java object and persist it to the database. No special coding would be required to achieve this, no performance penalty would ensue, and the result would be totally portable.

In this ideal world, we would perhaps perform such an operation in a manner like that shown in Listing 1-1.

Listing 1-1. *A Rose-Tinted View of Object Persistence*

```
POJO pojo = new POJO();
ORMSolution magic = ORMSolution.getInstance();
magic.save(pojo);
```

There would be no nasty surprises, no additional work to correlate the class with tables in the database, and no performance problems.

Hibernate comes remarkably close to this, at least when compared with the alternatives—but alas, there are configuration files to create and subtle performance issues to consider. Hibernate does, however, achieve its fundamental aim—it allows you to store POJOs in the database. Figure 1-1 shows how Hibernate fits into your application between the client code and the database.

Figure 1-1. *The role of Hibernate in a Java application*

The common term for the direct persistence of traditional Java objects is *object-relational mapping*—that is, mapping the objects in Java to the relational entities in a database.

Where entity beans have to follow a myriad of awkward naming conventions, POJOs can be any Java object at all. Hibernate allows you to persist POJOs with very few constraints. Listing 1-2 is an example of a simple POJO to represent a message.

Listing 1-2. *The POJO Used in this Chapter's Examples*

```
public class Message {
   private Message() {
   }

   public Message(String messageText) {
      this.messageText = messageText;
   }

   public String getMessageText() {
      return messageText;
   }
```

```
    public void setMessageText(String messageText) {
        this.messageText = messageText;
    }

    private String message;
}
```

The sole condescension to Hibernate here is the provision of a private default constructor. <mark>Hibernate demands that all POJOs to be stored should provide a default constructor;</mark> but even that can be worked around when third-party classes fail to satisfy this limited requirement (we will demonstrate this in Appendix A).

Origins of Hibernate and Object-Relational Mapping

If Hibernate is the solution, what was the problem? One answer is that doing things the right way when using JDBC requires a considerable body of code, and careful observation of various rules (such as those governing connection management) to ensure that your application does not leak resources. The gargantuan body of code in Listing 1-3 is required to populate the example Motd object from the database even when you know the appropriate message identifier.

Listing 1-3. *The JDBC Approach to Retrieving the POJO*

```
    public static List getMessages(int messageId) throws MessageException {
        Connection c = null;
        PreparedStatement p = null;
        List list = new ArrayList();

        try {

            Class.forName("org.postgresql.Driver");
            c = DriverManager.getConnection(
                "jdbc:hsqldb:testdb;shutdown=true",
                "hibernate",
                "hibernate");
            p = c.prepareStatement(
                "select message from motd");

            ResultSet rs = p.executeQuery();

            while(rs.next()) {
                String text = rs.getString(1);
                list.add(new Message(text));
            }
            return list;
```

```
        } catch (Exception e) {
          log.log(Level.SEVERE, "Could not acquire message", e);
          throw new MotdException(
              "Failed to retrieve message from the database.", e);
        } finally {
          if (p != null) {
            try {
              p.close();
            } catch (SQLException e) {
              log.log(Level.WARNING,
                   "Could not close ostensibly open statement.", e);
            }
          }

          if (c != null) {
            try {
              c.close();
            } catch (SQLException e) {
              log.log(Level.WARNING,
                   "Could not close ostensibly open connection.", e);
            }
          }
        }
      }
    }
```

Some of this can be trimmed down; there are various techniques that allow you to reduce the boilerplate code for opening connections and logging problems, but the basic logic that pulls the object instance from the ResultSet becomes more complex as the object itself does. Once the object includes references to other objects—or worse yet, other collections of objects—these "manual" techniques start to look more and more flawed.

EJBs As a Persistence Solution

So why not just use EJBs to retrieve data? Entity beans are, after all, designed to represent, store, and retrieve data in the database.

Strictly speaking, an entity bean is permitted two types of persistence in an EJB server: bean-managed persistence (BMP) and container-managed persistence (CMP). In BMP, the bean itself is responsible for carrying out all of the SQL associated with storing and retrieving its data—in other words, it requires the author to create the appropriate JDBC logic, complete with all the boilerplate from Listing 1-3. CMP, on the other hand, requires the container to carry out the work of storing and retrieving the bean data. So why doesn't that solve the problem? Here are just a few of the reasons:

- CMP entity beans require a one-to-one mapping to database tables.

- They do not directly support inheritance relationships.

- They are (by reputation, at least) slow.

- Someone has to determine which bean field maps to which table column.

- They require special method names. If these are not followed correctly, they will fail silently.

- Entity beans have to reside within a J2EE application server environment—they are a heavyweight solution.

- They cannot readily be extracted as "general purpose" components for other applications.

- They cannot be serializable.

- They rarely exist as portable components to be dropped into a foreign application— you generally have to roll your own EJB solution.

Hibernate As a Persistence Solution

Hibernate addresses a lot of these points, or alleviates some of the pain where it can't, so we'll address the points in turn.

Hibernate does not require you to map one POJO to one table. A POJO can be constructed out of a selection of table columns, or several POJOs can be persisted into a single table.

Hibernate directly supports inheritance relationships and the various other relationships between classes.

Though there is some performance overhead while Hibernate starts up and processes its configuration files, it is generally perceived as being a fast tool. This is very hard to quantify, and, to some extent, the poor reputation of entity beans may have been earned less from their own faults than from the mistakes of those designing and deploying such applications. As with all performance questions, you should carry out tests rather than relying on anecdotal evidence.

In Hibernate it is possible, but not necessary, to specify the mappings at deployment time. The EJB solution ensures portability of applications across environments, but the Hibernate approach tends to reduce the pain of deploying an application to a new environment.

Hibernate persistence has no requirement for a J2EE application server or any other special environment. It is, therefore, a much more suitable solution for stand-alone applications, client-side application storage, and other environments in which a J2EE server is not immediately available.

Hibernate uses POJOs that can very easily and naturally be generalized for use in other applications. There is no direct dependency upon the Hibernate libraries, so POJOs can be put to any use that does not require persistence; or they can be persisted using any other "POJO-friendly" mechanism.

Hibernate presents no problems when handling serializable POJOs.

There is a very large body of preexisting code. Any Java object capable of being persisted to a database is a candidate for Hibernate persistence. Therefore, Hibernate is a natural replacement for ad hoc solutions, or as the persistence engine for an application that has not yet had database persistence incorporated into it. Furthermore, by choosing Hibernate persistence, you are not tying yourself to any particular design decisions for the business objects in your application.

A Thin Solution?

One of the benefits often touted for Hibernate is that it is a "thin" solution. The problem with this description is that it is very much an informal term, so it doesn't really tell you anything about what attributes Hibernate has that could categorize it as thin.

Hibernate does not require an application server to operate (while EJBs do). It is therefore applicable in client-side applications in which EJBs are entirely inappropriate. So from this point of view, it is perhaps thin.

On the other hand, Hibernate makes use of an inordinate number of supporting libraries. So, if you are considering download times and disk space for an applet, Hibernate will look somewhat obese; though in these days of fast connections and copious disk space, it is unlikely to be a deciding factor.

A Hibernate Hello World Example

Listing 1-4 shows how much less boilerplate is required with Hibernate than with the JDBC approach from Listing 1-3.

Listing 1-4. *The Hibernate Approach to Retrieving the POJO*

```
public static List getMessages(int messageId)
   throws MessageException
{
   SessionFactory sessions =
      new Configuration().configure().buildSessionFactory();
   Session session = sessions.openSession();
   Transaction tx = null;
   try {
      tx = session.beginTransaction();

      List list = session.createQuery("from Message").list();

      tx.commit();
      tx = null;
      return list;

   } catch ( HibernateException e ) {
      if ( tx != null ) tx.rollback();
      log.log(Level.SEVERE, "Could not acquire message", e);
      throw new MotdException(
            "Failed to retrieve message from the database.",e);
   } finally {
      session.close();
   }
}
```

Even for this trivial example there would be a further reduction in the amount of code required in a real deployment—particularly in an application-server environment. For example, the `SessionFactory` would normally be created elsewhere and made available to the application as a Java Native Directory Interface (JNDI) resource.

Note that the manual coding to populate the message object has not been eradicated—rather, it has been moved into an external configuration file that isolates this implementation detail from the main logic.

Some of the additional code in the Hibernate 3 example given in Listing 1-4 actually provides functionality (particularly transactionality and caching) beyond that of the JDBC example.

Mappings

As we have intimated, Hibernate needs something to tell it which tables relate to which objects (this information is usually provided in an XML mapping file). While some tools inflict vast, poorly documented XML configuration files on their users, Hibernate offers a breath of fresh air—you create and associate a small, clear mapping file with each of the POJOs that you wish to map into the database. You're permitted to use a single monolithic configuration file if you prefer, but it's neither compulsory nor encouraged.

A document type definition (DTD) is provided for all of Hibernate's configuration files, so with a good XML editor you should be able to take advantage of autocompletion and autovalidation of the files as you create them. Java 5 annotations can be used to replace them entirely.

Listing 1-5 shows the file mapping the Message POJO into the database.

Listing 1-5. *The XML File That Maps the POJO to the Database*

```xml
<?xml version="1.0"?>
<!DOCTYPE hibernate-mapping PUBLIC
    "-//Hibernate/Hibernate Mapping DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">
<hibernate-mapping>
    <class name="Message" table="Message">
        <id type="int" column="id">
            <generator class="native"/>
        </id>
        <property name="message" column="message" type="string"/>
    </class>
</hibernate-mapping>
```

It would be reasonable to ask if the complexity has simply been moved from the application code into the XML mapping file. But, in fact, this isn't really the case for several reasons.

First, the XML file is much easier to edit than a complex population of a POJO from a result set—and far more easily changed after the fact should your mappings into the database change at a late stage of development.

Second, we have still done away with the complicated error handling that was required with the JDBC approach. This is probably the weakest reason, however, since there are various techniques to minimize this without resorting to Hibernate.

Finally, it is of particular note that Hibernate solves the problem in which the developer needs to extract from the database a single object that is related by references to a substantial number of objects in the database. Hibernate postpones such additional extractions until they are actually accessed, generally avoiding substantial memory and performance costs.

Database Generation

If you are creating a new Hibernate application around an existing database, creating the database schema is not an issue; it is presented to you on a platter. If you are starting a new application, you will need to create the schema, the POJOs, and the mapping directly.

Yet again, Hibernate makes life easy for you. A suite of tools has been made available that allows you to generate your database directly from the mapping file (generally using an Ant script). This is even better than it sounds! Hibernate comes with intimate knowledge of various different dialects of SQL, so the schema will be tailored to your particular database software—or it can be generated for each different database to which you want to deploy your application.

The Relationship of Hibernate 3 with EJB 3.0

The current version of the EJB standard does not provide an ORM mechanism. However, this will change. The forthcoming standard for the next (3.0) incarnation of EJB has been developed as Java Specification Request (JSR) 220, and includes an EJB ORM standard. Why learn Hibernate when a new standard blessed by Sun has just been released? Because the original developers of Hibernate are closely involved in the creation of the standard, and as one of the most successful ORM solutions, Hibernate has heavily influenced the design of EJB 3.0. Hibernate will support those parts of EJB 3.0 pertaining to ORM.

EJB 3.0 supports and encourages the use of transparent persistence of entity beans that conform to the same minimal set of requirements as the POJOs supported by Hibernate. As a result of this, Hibernate can be used as the implementation of the `EntityManager` component of an EJB container.

Hibernate's HQL (Hibernate Query Language) also has a strong correlation with the new EJB Query Language (EJB QL)—though this is probably due as much to their common ancestry in Structured Query Language (SQL) as it is to the Hibernate team's participation in the standards effort.

EJB 3.0 allows the relationship between entities/POJOs and the database to be described using annotations that Hibernate supports. The convenience of this approach suggests that it will become the standard way to maintain these mappings as Java 5 deployments become commonplace.

Hibernate 3 provides features above and beyond those mandated by the EJB 3.0 standard—and Hibernate 3 does not require the presence and burden of an application server, making it suitable for Swing applications, other client-side applications, and lightweight web apps running without a full J2EE stack (for example, those running on Tomcat).

In short, if you code to Hibernate 3 now, the effort involved in the transition to EJB 3.0 later will be minimal if at all necessary. Meanwhile, Hibernate 3 is available now and has a persuasive pedigree.

Summary

In this chapter, we have considered the problems and requirements that have driven the development of Hibernate. We have looked at some of the details of a trivial example application written with and without the aid of Hibernate. We have glossed over some of the implementation details, but we will discuss these in depth in Chapter 3.

In the next chapter, we will look at the architecture of Hibernate and how it is integrated into your applications.

■ ■ ■ ■

Integrating and Configuring Hibernate

Compared to other Java persistence solutions, integrating Hibernate into a Java application is easy. The designers of Hibernate avoided some of the more common pitfalls and problems with the existing Java persistence solutions, and created a clean but powerful architecture. In practice, this means that you do not have to run Hibernate inside any particular J2EE container or framework—Hibernate 3 only requires Java 2 Standard Edition (J2SE), version 1.3 or greater, although the new Annotations feature requires J2SE 5.0 (or later).

At first, adding Hibernate to your Java project looks intimidating—the distribution includes a large set of libraries. To get your first Hibernate application to work, you have to set up the database, the mapping files, the configuration, and your plain old Java objects (POJOs). After you have done all that, you need to write the logic in your application that uses the Hibernate session to actually do something! But once you learn how to integrate Hibernate with your application, the basics apply for any project that uses Hibernate.

If you already have an application that uses Hibernate 2, the migration path from Hibernate 2 to Hibernate 3 is easy. While Hibernate 3 is not completely backward-compatible, most of the changes are additional features that you can integrate into your existing application as you see fit. The Hibernate developers provided implementations of the core Hibernate 2 objects in Hibernate 3 with the Hibernate 2 methods for backward compatibility. We discuss the differences between Hibernate 2 and Hibernate 3 in more depth in Appendix D.

One of the key features of Hibernate's design is the principle of least intrusiveness—the Hibernate developers did not want Hibernate to intrude into your application more than was necessary. This led to several of the architectural decisions made for Hibernate. In Chapter 1 you saw how Hibernate can be applied to solve persistence problems using conventional Java objects. In this chapter, we explain some of the configuration details needed to support this behavior.

The Steps Needed to Integrate and Configure Hibernate

This chapter explains configuration and integration in detail, but for a quick overview, refer to the following bulleted list to determine what you need to do to get your first Hibernate application up and running. Chapter 3 leads you through the building of a pair of small example

applications that use Hibernate. The first of these is as simple as we could make it, so it is an excellent introduction to the following necessary steps:

- Identify the POJOs that have a database representation.

- Identify which properties of those POJOs need to be persisted.

- Create Hibernate XML mapping files for each of the POJOs that map properties to columns in a table (covered in more detail in Chapter 7).

- Create the database schema using the schema export tool, use an existing database, or create your own database schema.

- Add the Hibernate Java libraries to your application's classpath (covered in this chapter).

- Create a Hibernate XML configuration file that points to your database and your XML mapping files (covered in this chapter).

- In your Java application, create a Hibernate `Configuration` object that references your XML configuration file (covered in this chapter).

- Also in your Java application, build a Hibernate `SessionFactory` object from the `Configuration` object (covered in this chapter).

- Finally, retrieve Hibernate Session objects from the `SessionFactory`, and write your data access logic for your application (create, retrieve, update, and delete).

Don't worry if you don't understand every term or concept in the preceding list. After reading this chapter, and then going through the example in the next chapter, you will know what these terms mean and how they fit together.

Understanding Where Hibernate Fits in Your Java Application

You can call Hibernate from your Java application directly, or you can access Hibernate through another framework. You can call Hibernate from a Swing application, a servlet, a portlet, a JSP page, or any other Java application that has access to a database. Typically, you would use Hibernate to either create a data access layer for an application or replace an existing data access layer.

Hibernate supports Java Management Extensions (JMX), J2EE Connector Architecture (JCA), and Java Naming and Directory Interface (JNDI) Java language standards. Using JMX, you can configure Hibernate while it is running. Hibernate may be deployed as a JCA connector, and you can use JNDI to obtain a Hibernate session factory in your application. In addition, Hibernate uses standard Java Database Connectivity (JDBC) database drivers to access the relational database. Hibernate does not replace JDBC as a database connectivity layer—Hibernate sits on a level above JDBC.

In addition to the standard Java APIs, many Java web and application frameworks now integrate with Hibernate. Hibernate's simple, clean API makes it easy for these frameworks to support Hibernate in one way or another. The Spring framework provides excellent Hibernate integration, including generic support for persistence objects, a generic set of persistence

exceptions, and transaction management. Appendix C explains how Hibernate can be configured within a Spring application.

Regardless of the environment that you are integrating Hibernate into, certain requirements remain constant. You will need to define the configuration details that apply—these are then represented by a Configuration object. From the Configuration object, a single SessionFactory object is created; and from this, Session objects are instantiated, through which your application accesses Hibernate's representation of the database.

Deploying Hibernate

To integrate Hibernate with your Java application, you will need to use several Java libraries. The first library is the Java Archive (JAR) file for your JDBC driver, which you will need to find for your specific relational database. The Hibernate download does not include any JDBC drivers. You must obtain these yourself—typically, the database provider will offer them as a separate download; or they may be bundled with your database installation.

Because every relational database—and hence every JDBC driver—behaves slightly differently, the Hibernate team created *dialects* to abstract away the differences. These dialects define the SQL variant and the specific database features to use for each vendor's database. Every project that uses Hibernate must specify one dialect in the Hibernate configuration file. We discuss dialects in more detail further on in the chapter. The Hibernate web site also contains a platform-specific FAQ that offers some solutions to several vendor-specific questions.

If you encounter problems getting Hibernate to work with older JDBC versions, disable the following two JDBC 2–specific features: batch update and scrollable result sets. Use the following configuration values for Hibernate (we discuss the specifics of this configuration later in this chapter):

```
hibernate.jdbc.batch_size=0
hibernate.jdbc.use_scrollable_resultsets=false
```

Once you have configured the JDBC driver, your next step is to deploy hibernate3.jar with your application. This JAR file is provided with the Hibernate 3 binary distribution. The file contains the classes in the org.hibernate package, along with several DTD and XML Schema files. You will then need to deploy the other required libraries.

Required Libraries for Running Hibernate 3

Hibernate requires several libraries beyond hibernate3.jar. These libraries are included in the lib directory of your Hibernate 3 installation. The up-to-date list of libraries shipped with Hibernate 3 is in the lib/README.txt file in the Hibernate distribution.

There are several optional libraries included with the Hibernate 3 distribution. If you build Hibernate from source, a few of these are necessary for Hibernate to compile. Other libraries provide connection pools, additional caching functionality (the Session cache is mandatory), and the JCA API. The purpose of each library is detailed in the README file, which also states which libraries are optional and which are required.

Annotations and Enterprise JavaBeans 3

The newly released Enterprise JavaBeans 3 (EJB 3) specification includes a mandatory ORM component. Hibernate's design influenced many of the changes (including lightweight persistence and the query language). The configuration requirements for Hibernate in an EJB 3 environment are somewhat different from the default requirements—we discuss this in more depth in Appendix A. If you have used EJB 1.x or 2.x in the past, you will find EJB 3 to be a much-needed simplification.

Hibernate 3 permits you to take advantage of the new Annotations feature of Java 5. These annotations can be used in conjunction with, or in place of, some of the configuration files that previous versions of Hibernate demanded. These annotations are essentially an EJB 3 feature, although Hibernate supplies some additional proprietary extensions. In Chapter 6 we discuss how to use these persistence annotations.

JMX and Hibernate

JMX is a standard API for managing Java applications and components—mostly accessed through MBeans, which represent wrappers for services and resources. Hibernate provides two MBeans for JMX: `HibernateServiceMBean` and `StatisticsServiceMBean`. Both of these are interfaces that reside in the `org.hibernate.jmx` package. The `HibernateService` and `StatisticsService` classes implement the interfaces and reside within the same package. The `HibernateServiceMBean` provides getter and setter methods for many of the Hibernate configuration properties, including the data source, transaction strategy, caching, dialect, and other database options. It also provides a mechanism for adding mapping files to the configuration. When the `HibernateServiceMBean` starts, it creates a `Configuration` object from its properties and mapping files, and then builds a `SessionFactory` object. The `SessionFactory` object binds to the JNDI location specified on the JMX MBean, and your Java application can then use standard JNDI calls to obtain the session factory.

The other MBean supplies statistics. Hibernate can log statistics about the performance of query execution, caching, and object entity operations. Using a JMX console, an administrator can enable statistics and then access up-to-date performance indicators through the console.

The advantage of JMX over programmatic access to these features is that administrators or other non-developers may change properties at run time through a standardized JMX console that is independent of Hibernate and applies to a range of other frameworks and components.

Hibernate Configuration

Before you create a session factory, you must tell Hibernate where to find the mapping information that defines how your Java classes relate to the database tables. Hibernate also requires a set of configuration settings, which are usually supplied as a standard Java properties file called `hibernate.properties`, or as an XML file named `hibernate.cfg.xml`.

We recommend using the XML format. This allows you to specify the location of the mapping information from the configuration files—the alternative (when using properties files) being to programmatically supply this information to Hibernate through the `Configuration` class.

Listing 2-1 is a reprint of Listing 1-4 from the previous chapter, which shows a complete usage of Hibernate from within an application. The parts of this listing that deal with configuration and integration are highlighted.

Listing 2-1. *The Hibernate Approach to Retrieving the POJO*

```
public static List getMessages(int messageId)
   throws MessageException
{
   SessionFactory sessions =
      new Configuration().configure().buildSessionFactory();
   Session session = sessions.openSession();
   try {
      session.beginTransaction();
      List list = session.createQuery("from Message").list();
      session.getTransaction().commit();
      return list;

   } catch ( HibernateException e ) {
      if ( session.getTransaction() != null )
         session.getTransaction().rollback();

      log.log(Level.SEVERE, "Could not acquire message", e);
      throw new MotdException(
            "Failed to retrieve message from the database.",e);
   } finally {
      session.close();
   }
}
```

As you can see, we called the `configure()` method on the `org.hibernate.cfg.Configuration` class without any arguments. This tells Hibernate to look in the classpath for the configuration file. The default name for the configuration file is `hibernate.cfg.xml`—if you change it, you will need to pass that name explicitly as an argument to the `configure()` method. We discuss the `configure()` method and XML configuration in more detail later in this chapter.

The `configure()` method returns an instance of `Configuration`, which can be used to obtain a Hibernate `SessionFactory` instance by calling the `buildSessionFactory()` method, as follows:

```
public SessionFactory buildSessionFactory() throws HibernateException
```

The `SessionFactory` is a heavyweight object, and your application should use one Hibernate `SessionFactory` object for each discrete database instance that it interacts with. The `SessionFactory` relies on the Hibernate configuration properties, which we detail in the next section of this chapter.

After you have obtained the `SessionFactory`, you can retrieve Hibernate `org.hibernate.Session` objects. While the `SessionFactory` is a heavyweight object, the `Session` objects are lightweight. You perform your persistence operations using `Session` objects.

To sum up, there are three classes that you need to use: `Configuration`, `SessionFactory`, and `Session`.

- Use the `Configuration` class to read (and to set) configuration details.

- Use the `Configuration` object to create a `SessionFactory` object.

- Use the `SessionFactory` object to create `Session` objects as needed.

A typical application will have one `Configuration` object, which will only be used in initialization. There will be one `SessionFactory` object that will exist throughout the life cycle of the application. Your application will ask this `SessionFactory` object for a `Session` any time it needs to work with the database. The application could retrieve an object, make some property changes, and then persist it, all within one session, and then close down the `Session` object.

Hibernate Properties

Typically, you will specify your Hibernate configuration in a properties file called `hibernate.properties` in the root directory of your application's classpath, or as identified values in a `hibernate.cfg.xml` file. Hibernate has an extensive list of properties for configuration (see Table 2-1).

Unless you provide a JDBC connection programmatically in your application, you must either configure a JDBC connection here or specify the JNDI name of a container-provided JDBC connection. You must also configure the SQL dialect appropriate to the database that you are using. All the other properties take sensible default values, so they do not need to be explicitly stated.

Table 2-1. *Hibernate Configuration Property Names and Descriptions*

Property Name	Description
hibernate.c3p0.acquire_increment	After the connection pool is completely utilized, determines how many new connections are added to the pool.
hibernate.c3p0.idle_test_period	Determines how long to wait before a connection is validated.
hibernate.c3p0.max_size	The maximum size of the connection pool for C3PO.
hibernate.c3p0.max_statements	The upper limit for the SQL statement cache for C3PO.
hibernate.c3p0.min_size	The minimum size of the connection pool for C3PO.
hibernate.c3p0.timeout	The timeout for C3PO (in seconds).
hibernate.cache.provider_class	Specifies a class that implements the org.hibernate.cache.CacheProvider interface.
hibernate.cache.query_cache_factory	Specifies a class that implements the org.hibernate.cache.QueryCacheFactory interface for getting QueryCache objects.
hibernate.cache.region_prefix	The prefix to use for the name of the cache.
hibernate.cache.use_minimal_puts	Configures the cache to favor minimal puts over minimal gets.
hibernate.cache.use_query_cache	Specifies whether to use the query cache.

Property Name	Description
hibernate.cache.use_second_level_cache	Determines whether to use the Hibernate second-level cache.
hibernate.cglib.use_reflection_optimizer	Instead of using slower standard Java reflection, uses the CGLib code generation library to optimize access to business object properties. The application may be slower at startup if this is enabled, but with faster runtime performance.
hibernate.connection.autocommit	Allows autocommit mode to be used for the JDBC connection (not usually a good idea).
hibernate.connection.datasource	The DataSource name for a container-managed data source.
hibernate.connection.driver_class	The JDBC driver class.
hibernate.connection.isolation	The transaction isolation level for the JDBC connection.
hibernate.connection.*<JDBCpropertyname>*	Passes any JDBC property you like to the JDBC connection—for instance, hibernate.connection.debuglevel=info would pass a JDBC property called debuglevel.
hibernate.connection.password	The database password.
hibernate.connection.pool_size	Limits the number of connections waiting in the Hibernate database connection pool.
hibernate.connection.provider_class	The class that implements Hibernate's ConnectionProvider interface.
hibernate.connection.url	The JDBC URL to the database instance.
hibernate.connection.username	The database username.
hibernate.default_catalog	The default database catalog name that Hibernate uses to generate SQL for unqualified table names.
hibernate.default_schema	The default database owner name that Hibernate uses to generate SQL for unqualified table names.
hibernate.dialect	The SQL dialect to use for Hibernate; varies by database. See this chapter's "SQL Dialects" section.
hibernate.generate_statistics	Determines whether statistics are collected.
hibernate.hbm2ddl.auto	Automatically creates, updates, or drops the database schema on startup and shut down. There are three possible values: create, create-drop, and update. Be careful with create-drop!
hibernate.jdbc.batch_size	The maximum batch size for updates.
hibernate.jdbc.batch_versioned_data	Determines whether Hibernate batches versioned data, which depends on your JDBC driver properly implementing row counts for batch updates. Hibernate uses the row count to determine whether the update is successful.
hibernate.jdbc.factory_class	The class name of a custom implementation of the org.hibernate.jdbc.Batcher interface for controlling JDBC prepared statements.

Continued

Table 2-1. *Continued*

Property Name	Description
hibernate.jdbc.fetch_size	Determines how many rows the JDBC connection will try to buffer with every fetch. This is a balance between memory and minimizing database network traffic.
hibernate.jdbc.use_get_generated_keys	Determines Hibernate's behavior with respect to generated keys. If this property is set to true, and if the database driver supports the JDBC 3.0 generated keys API, Hibernate will retrieve generated keys from the statement after it executes an SQL query.
hibernate.jdbc.use_scrollable_resultset	Determines whether Hibernate will use JDBC scrollable result sets for a user-provided JDBC connection.
hibernate.jdbc.use_streams_for_binary	Determines whether binary data is read or written over JDBC as streams.
hibernate.jndi.class	The InitialContext class for JNDI.
hibernate.jndi.<JNDIpropertyname>	Passes any JNDI property you like to the JNDI InitialContext.
hibernate.jndi.url	Provides the URL for JNDI.
hibernate.max_fetch_depth	Determines how deep Hibernate will go to fetch the results of an outer join. Used by Hibernate's outer join loader.
hibernate.order_updates	Orders SQL update statements by each primary key.
hibernate.proxool	Prefix for the Proxool database connection pool.
hibernate.proxool.existing_pool	Configures Proxool with an existing pool.
hibernate.proxool.pool_alias	The alias to use for any of the configured Proxool pools previously mentioned.
hibernate.proxool.properties	Path to a Proxool properties file.
hibernate.proxool.xml	Path to a Proxool XML configuration file.
hibernate.query.factory_class	Specifies an HQL query factory class name.
hibernate.query.substitutions	Any possible SQL token substitutions that Hibernate should use.
hibernate.session_factory_name	If set, causes the Hibernate session factory to bind to this JNDI name.
hibernate.show_sql	Logs the generated SQL commands.
hibernate.sql_exception_converter	Specifies which SQLExceptionConverter to use to convert SQLExceptions into JDBCExceptions.
hibernate.transaction.auto_close_session	Automatically closes the session after a transaction.
hibernate.transaction.factory_class	Specifies a class that implements the org.hibernate.transaction.TransactionFactory interface.

Property Name	Description
`hibernate.transaction.flush_before_completion`	Automatically flushes before completion.
`hibernate.transaction.manager_lookup_class`	Specifies a class that implements the `org.hibernate.transaction.TransactionManagerLookup` interface.
`hibernate.use_identifier_rollback`	Determines whether Hibernate uses identifier rollback.
`hibernate.use_sql_comments`	Generates SQL with comments.
`hibernate.wrap_result_sets`	Turns on JDBC result set wrapping with column names.
`hibernate.xml.output_stylesheet`	Specifies an XSLT stylesheet for Hibernate's XML data binder. Requires `xalan.jar`.
`jta.UserTransaction`	The JNDI name for the `UserTransaction` object.

XML Configuration

As we have already mentioned, Hibernate offers XML configuration capabilities. To use them, you must create an XML configuration file, normally called `hibernate.cfg.xml`, and place it in the root of your application's classpath. The XML configuration file must conform to the Hibernate 3 Configuration DTD, which is available from `http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd`.

Listing 2-2 shows an example XML configuration for Hibernate.

Listing 2-2. *An XML Configuration for Hibernate*

```
<?xml version="1.0" encoding="utf-8"?>
<!DOCTYPE hibernate-configuration SYSTEM ➥
"http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <property name="dialect">org.hibernate.dialect.HSQLDialect</property>
        <mapping jar="hibernate-mappings.jar"/>
        <mapping resource="com/apress/hibernate/User.hbm.xml"/>
    </session-factory>
</hibernate-configuration>
```

When you use the XML configuration file, you do not need to use the `hibernate.` prefix for properties. As you can see in Listing 2-2, the `dialect` property is simply `dialect`, not `hibernate.dialect`. However, we usually elect to include the prefix for the sake of consistency. If you are already using `hibernate.properties`, `hibernate.cfg.xml` will override any settings in `hibernate.properties`.

In addition to specifying the properties listed in Table 2-1 to configure a session factory, with the `<property>` tag you can also configure mapping files, caching, listeners, and the JNDI name for the session factory in the XML configuration. Listing 2-2 includes two `<mapping>` elements, which identify a JAR file containing mapping files, and a specific mapping file available

on the classpath, respectively. We discuss mapping file configuration in the next section and caching in Chapter 9.

After placing your XML configuration file in the root directory of the classpath, you will need to call one of the `configure()` methods on your application's `Configuration` object. With the default file name (`hibernate.cfg.xml`), you can call `configure()` with no arguments. If you used a different file name (for instance, because you have production, staging, user acceptance test, and development environments, with different databases), use one of the following methods on a `Configuration` object:

- `public Configuration configure(String) throws HibernateException`: Loads the XML configuration from a resource accessible by the current thread's context class loader

- `public Configuration configure(URL) throws HibernateException`: Retrieves the XML configuration from a valid URL

- `public Configuration configure(File) throws HibernateException`: Uses an XML configuration file on the file system

Mapping Documents

Once you have created your mapping documents for Hibernate, it needs to know where to find them. Before you create the session factory, add them to your `Configuration` object, or specify them in the `hibernate.cfg.xml` XML configuration file. If you choose to add the mapping documents directly to an instance of `Configuration`, use one of the following methods:

- `addFile(String)`: Uses the path to an XML mapping document for Hibernate. An example of this would be `com/hibernatebook/config/Example.hbm.xml`

- `addFile(File)`: Uses a `File` object that represents an XML mapping document

- `addClass(Class)`: Translates a Java class name into a file name, which Hibernate then loads as an input stream from the Java class's class loader; for example, Hibernate would look for the file called `com/hibernatebook/config/Example.hbm.xml` for a class named `com.hibernatebook.config.Example`

- `addJar(File)`: Adds any mapping documents (`*.hbm.xml`) in the specified JAR file to the `Configuration` object

- `addDirectory(File)`: Adds any mapping documents (`*.hbm.xml`) in the specified directory to the `Configuration` object

The following methods also add mapping documents to the `Configuration` object, but you are unlikely to need them unless you have specialized application deployment issues:

- `addXML(String)`: Takes a `String` object that contains the Hibernate mapping XML.

- `addURL(URL)`: Requires a valid URL to the Hibernate mapping XML.

- `addCacheableFile(File)`: Saves time when Hibernate loads XML mapping files at startup by caching XML mapping documents on the file system as serialized DOM `Document` objects. This improves performance after the first load. Takes a `File` object that points to the XML mapping document, not the `.bin` cache file.

- `addCacheableFile(String)`: Same as `addCacheableFile(File)`, except this method takes a path to the file. Hibernate constructs a `File` object out of the path and then passes it to `addCacheableFile(File)`.

- `addDocument(Document)`: Takes a valid DOM `Document` object containing the XML.

The `addJar()` and `addDirectory()` methods are the most convenient, because they allow you to load all of your Hibernate mapping documents at one time. Both of these methods simplify code configuration, layout, and refactoring, because you do not have to separately maintain code that configures each document. We find that it is easy to create a mapping document and then forget to add it to your application's Hibernate initialization code; using either of these methods helps to avoid that problem.

As an alternative to specifying the locations of the mapping information in the code, you can instead use the `<mapping>` element in the `hibernate.cfg.xml` XML configuration file. The `<mapping>` element has four possible attributes—`jar`, `resource`, `file`, and `class`—which map to the `addJar()`, `addResource()`, `addFile()`, and `addClass()` methods on the `Configuration` object.

```
<?xml version="1.0" encoding="utf-8"?>
<!DOCTYPE hibernate-configuration SYSTEM ➥
"http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <mapping jar="hibernate-mappings.jar"/>
        <mapping resource="com/apress/hibernate/User.hbm.xml"/>
    </session-factory>
</hibernate-configuration>
```

Whether you use the XML configuration file or directly specify the mapping files in code is up to you—we suggest that you stick to the approach that you are most comfortable with.

Naming Strategy

If your project has an existing standard for naming database tables or columns, or you would like to specify exactly how Hibernate maps Java class names to database table names, you can use Hibernate's *naming strategy* functionality. Custom naming strategies specify how Hibernate maps Java class names to database table names, properties to column names, and the name of a table used to hold a collection of properties for a Java class. A naming strategy may also override the table names and column names specified in the Hibernate mapping documents—for instance, you might use this to enforce a consistent application-wide prefix to table names.

Although you can explicitly specify the names of all of the tables and columns in the mapping document, if you have clear and consistent rules for naming already, implementing a custom naming strategy can save a lot of time and frustration. Equally, if you decide to add

a prefix to all database table names after the fact, it is easy to do so with a naming strategy, while it would be a pain to correct these in every Hibernate mapping document.

■**Note** Using Hibernate with an existing well-specified database often means creating a custom naming strategy for Hibernate. If the database tables have a prefix, it may be cleaner to implement a naming strategy that adds that prefix than to specify the full table name with a prefix in every Hibernate mapping document.

A custom naming strategy must implement the `org.hibernate.cfg.NamingStrategy` inter-face or extend one of the two provided naming strategy classes, `org.hibernate.cfg.DefaultNamingStrategy` or `org.hibernate.cfg.ImprovedNamingStrategy`. The default naming strategy simply returns the unqualified Java class name as the database table name. For instance, the table name for the Java class `com.hibernatebook.AccessGroups` would be `AccessGroups`. The column name would be the same as the property name, and the collection table would have the same name as the property.

The improved naming strategy adds underscores in place of uppercase letters in mixed-case table and column names, and then lowercases the name. For instance, the same `com.hibernatebook.AccessGroups` Java class would correspond to a database table named `access_groups`.

Neither of these naming strategies takes into account the case in which you have two classes with the same name in different packages in your application. For instance, if you had two classes, `com.hibernatebook.webcast.Group` and `com.hibernatebook.security.Group`, both would default to a table named `Group`, which is not what you want. You would have to explic-itly set the table name in the mapping of at least one class.

Once you have created a naming strategy, pass an instance of it to the `Configuration` object's `setNamingStrategy()` method, as follows:

```
public Configuration setNamingStrategy(NamingStrategy namingStrategy)
```

You must call this method *before* building the session factory from the `Configuration`. For example, here's the code for using the `ImprovedNamingStrategy` naming strategy:

```
Configuration conf = new Configuration()
conf.setNamingStrategy(ImprovedNamingStrategy.INSTANCE);
```

Using a Container-Managed Data Source

When running in an environment with a JNDI server, Hibernate can obtain a data source through a JNDI lookup. You must use the `hibernate.connection.datasource` property to spec-ify the JNDI name, and then you may set the optional `hibernate.jndi.url` and `hibernate.jndi.class` properties to specify the location of the container's JNDI provider and the class name of the container's implementation of the JNDI `InitialContextFactory` interface. You may also use the `hibernate.connection.username` and `hibernate.connection.password` prop-erties to specify the database user your application uses. For example, your `hibernate.properties` file might have these lines for a WebLogic 7.0 managed data source:

```
hibernate.connection.datasource=java:/comp/env/jdbc/TestDB
hibernate.connection.username=dbuser
hibernate.connection.password=dbpassword
hibernate.jndi.url=t3://localhost:7001
hibernate.jndi.class=weblogic.jndi.WLInitialContextFactory
```

Typically only the mandatory `datasource` property is needed.

The Session Factory

You use the Hibernate session factory to create `Session` objects that manage connection data, caching, and mappings. Your application is responsible for managing the session factory. You should only have one session factory unless you are using Hibernate to connect to two or more database instances with different settings, in which case you should still have one session factory for each database instance.

In order to maintain backward compatibility with Hibernate 2, the Hibernate 3 session factory can also create `org.hibernate.classic.Session` session objects. These "classic" session objects implement all of the Hibernate 3 session functionality in addition to the deprecated Hibernate 2 session methods. We briefly discuss the changes in core functionality between Hibernate 2 and 3 in Appendix D.

You obtain a session from the `SessionFactory` object using one of the four `openSession()` methods. The no-argument `openSession()` method opens a session, with the database connection and interceptor specified in the `SessionFactory`'s original configuration. You can explicitly pass a JDBC connection to use, a Hibernate interceptor, or both as arguments to the remaining `openSession()` methods.

```
public org.hibernate.classic.Session openSession()
    throws HibernateException

public org.hibernate.classic.Session openSession(Interceptor interceptor)
    throws HibernateException

public org.hibernate.classic.Session openSession(
                Connection connection,
                Interceptor interceptor)

public org.hibernate.classic.Session openSession()
    throws HibernateException
```

We discuss Hibernate interceptors in Appendix A. You can also retrieve metadata and statistics from the `SessionFactory`.

The other important method on the session factory is `close()`. The `close()` method releases all the resource information used by the session factory and made available to the `Session` objects. It is therefore important that any related `Session` objects have been closed before invoking this to close the session factory.

When the session factory closes, it destroys the cache for the entity persisters and collection persisters, and also destroys the query cache and the timestamps cache. Then the session factory closes the JDBC connection provider and removes the current instance from its JNDI object factory binding.

```
public void close() throws HibernateException
```

The Hibernate developers designed their implementation of the SessionFactory interface to be scalable in a multithreaded application.

SQL Dialects

JDBC abstracts away many of the underlying connection details for each relational database, yet every relational database supports a different set of features and uses a slightly different version of SQL. Among the features that differ between relational databases are the syntax for marking identity columns, column data types, available SQL functions, foreign key constraint syntax, limits, GUID support, and support for cascade deletes.

Hibernate abstracts away all of these changes into *dialect* classes. Each supported database has its own dialect. When Hibernate constructs an SQL query, it obtains appropriate syntax information for the current database from the dialect. Hibernate 3 comes with over 20 different dialects. All of these standard dialects are supplied within the org.hibernate.dialect package. Table 2-2 shows the supported databases in Hibernate 3 and their corresponding dialect classes.

Table 2-2. *Supported Databases and Dialect Class Names for Hibernate 3*

Database Name	Dialect Class Name
DB2/390	DB2390Dialect
DB2/400	DB2400Dialect
DB2	DB2Dialect
Derby	DerbyDialect
Firebird	FirebirdDialect
FrontBase	FrontBaseDialect
HSQLDB	HSQLDialect
Informix	InformixDialect
Ingres	IngresDialect
InterBase	InterbaseDialect
JDataStore	JDataStoreDialect
Mimer SQL	MimerSQLDialect
Mckoi	MckoiDialect
MySQL 5	MySQL5Dialect
MySQL (< 5.x)	MySQLDialect
MySQL with InnoDB tables	MySQLInnoDBDialect
MySQL with MyISAM tables	MySQLMyISAMDialect

Database Name	Dialect Class Name
Oracle9*i*	`Oracle9Dialect`
Oracle9*i* (DataDirect drivers)	`DataDirectOracle9Dialect`
Oracle (< 9.x)	`OracleDialect`
PointBase	`PointbaseDialect`
PostgreSQL	`PostgreSQLDialect`
Progress	`ProgressDialect`
RDMS for Unisys OS2200	`RDMSOS2200Dialect`
SAP DB	`SAPDBDialect`
SQL Server	`SQLServerDialect`
Sybase	`SybaseDialect`
Sybase 11	`Sybase11Dialect`
Sybase Anywhere	`SybaseAnywhereDialect`
Times Ten 5.1	`TimesTenDialect`

Configure your chosen dialect by supplying the fully qualified dialect class name as the value for the `hibernate.dialect` configuration property.

Through Hibernate Query Language (HQL), Hibernate provides object-querying functionality that is database-independent. Hibernate translates the HQL queries into database-specific SQL using hints provided by the SQL dialect classes. We discuss HQL in more detail in Chapter 9.

Hibernate also provides a native SQL facility, which is especially useful for porting existing JDBC applications to Hibernate or for improving the performance of complicated queries.

Summary

In this chapter, we explained how to integrate Hibernate into your Java applications. We also detailed the configuration options for Hibernate, including the available Hibernate property settings. We discussed how naming strategies aid in the creation of consistent company- or application-wide database table-naming conventions, and how they help you to map your Hibernate classes to databases with existing naming conventions. Finally, we discussed how Hibernate uses dialects to manage the different behaviors of different database platforms.

In the next chapter, we build and configure a pair of simple Hibernate applications that illustrate the core Hibernate concepts discussed in the first two chapters.

CHAPTER 3

■ ■ ■

Building a Simple Application

In this chapter, you'll take another look at some of the steps necessary to get the example from Chapter 1 up and running. You'll also build a somewhat larger application from scratch. All of the code in this book is available for download from the Apress site (www.apress.com).

Installing the Tools

To run the examples in this chapter, you will need to install a number of tools. You will require a JDK, the Hibernate and Hibernate Tools distributions, the Ant build tool, and the HSQLDB database. Table 3-1 lists the specific tools you will need and where you can find them.

Table 3-1. *The Tools Used in This Book*

Tool	Version	Download Location
Hibernate	3.2.0	http://hibernate.org
Hibernate Tools	3.1	http://hibernate.org
Ant	1.6.5	http://ant.apache.org
HSQLDB	1.8.0.2	http://hsqldb.org

Hibernate and Hibernate Tools

The latest version of Hibernate is always available from http://hibernate.org, under the left-hand menu link named "Download." Various older versions and additional libraries are available from the resulting page, but you should select Hibernate Core 3.2.0 or a later version. At the time of writing, this is still a release-candidate version, but we expect the final release to be available by the time you read this book—if it is not, and you don't want to use a pre-release version, then most of the examples will work equally well with the previous 3.1.0 release of the Hibernate core. Download the archive and unpack it to a local directory. The unpacked archive contains all the source code for Hibernate itself, a JAR library built from this source, and all the library files that are necessary to run the sample.

You should then download Hibernate Tools from the same site. At the time of writing, it is currently at version 3.1 (again, this is currently in a late beta release, but we recommend using the beta version, rather that its inferior predecessors, if a final 3.1 version has not been released yet). Hibernate Tools provides various plug-ins for the Ant build tool and the free Eclipse IDE. In this chapter, we make use of the Ant plug-ins only, but we discuss the Eclipse

features in Appendix B. Again, the archive should be downloaded and unpacked to a local directory. This archive does not include the source code (which is available elsewhere on the `www.hibernate.org` site, if you decide to take a look at it).

HSQLDB 1.8.0

The database we will be using in our examples is the HSQL database. This is written in Java and is freely available open source software. While we used version 1.8.0.2 for our examples, we expect that any later version will be suitable. HSQL is derived from code originally released as "Hypersonic." You may encounter the term in some of the HSQL documentation and should treat it as synonymous with "HSQL." We may also refer to the product as HSQLDB when it might otherwise be mistaken for Hibernate Query Language (HQL), whose acronym is distressingly similar!

Our examples are tailored to HSQL because HSQL will run on any of the platforms that Hibernate will run on, and because HSQL is freely available with minimal installation requirements. However, if you want to run the examples with your own database, then the differences should boil down to the following:

- The Hibernate dialect class

- The JDBC driver

- The connection URL for the database

- The username for the database

- The password for the database

You will see where these can be specified later in this chapter. You will notice that where we specify the URL for connection to the database, we often append a `shutdown=true` attribute. This fixes a minor problem in which HSQLDB does not write its changes to disk until a `Connection` object is closed (something that may never happen when a connection is being managed by Hibernate's own connection pooling logic). This is not necessary on non-embedded databases.

Ant 1.6.5

You will want to install the Ant build tool. We will not attempt to explain the `build.xml` format in detail; if you are familiar with Ant, then the example build script provided in this chapter will be enough to get you started—if not, then Ant is a topic in its own right. We would recommend *Enterprise Java Development on a Budget*, by Christopher M. Judd and Brian Sam-Bodden (Apress, 2004), for good coverage of open source tools such as Ant.

While Ant in general lies outside the scope of this book, we will discuss the use of the Hibernate tasks used by our scripts.

Listing 3-1 provides the Ant script to build the example for this chapter.

Listing 3-1. *An Ant Script to Build the Chapter 3 Examples*

```
<project name="sample">

    <property file="build.properties"/>

    <property name="src" location="src"/>
    <property name="bin" location="bin"/>
    <property name="sql" location="sql"/>
    <property name="hibernate.tools"
        value="${hibernate.tools.home}${hibernate.tools.path}"/>

    <path id="classpath.base">
        <pathelement location="${src}"/>
        <pathelement location="${bin}"/>
        <pathelement location="${hibernate.home}/hibernate3.jar"/>
        <fileset dir="${hibernate.home}/lib" includes="**/*.jar"/>
        <pathelement location="${hsql.home}/lib/hsqldb.jar"/>
    </path>

    <path id="classpath.tools">
        <path refid="classpath.base"/>
        <pathelement
            location="${hibernate.tools}/hibernate-tools.jar"/>
    </path>

    <taskdef name="htools"
        classname="org.hibernate.tool.ant.HibernateToolTask"
        classpathref="classpath.tools"/>

    <target name="exportDDL" depends="compile">
        <htools destdir="${sql}">
            <classpath refid="classpath.tools"/>
            <configuration
                configurationfile="${src}/hibernate.cfg.xml"/>
            <hbm2ddl drop="true" outputfilename="sample.sql"/>
        </htools>
    </target>

    <target name="compile">
        <javac srcdir="${src}" destdir="${bin}" classpathref="classpath.base"/>
    </target>

    <target name="populateMessages" depends="compile">
        <java classname="sample.PopulateMessages" classpathref="classpath.base"/>
    </target>
```

```xml
<target name="listMessages" depends="compile">
    <java classname="sample.ListMessages" classpathref="classpath.base"/>
</target>

<target name="createUsers" depends="compile">
    <java classname="sample.CreateUser" classpathref="classpath.base">
        <arg value="dave"/>
        <arg value="dodgy"/>
    </java>
    <java classname="sample.CreateUser" classpathref="classpath.base">
        <arg value="jeff"/>
        <arg value="jammy"/>
    </java>
</target>

<target name="createCategories" depends="compile">
    <java classname="sample.CreateCategory" classpathref="classpath.base">
        <arg value="retro"/>
    </java>
    <java classname="sample.CreateCategory" classpathref="classpath.base">
        <arg value="kitsch"/>
    </java>
</target>

<target name="postAdverts" depends="compile">
    <java classname="sample.PostAdvert" classpathref="classpath.base">
        <arg value="dave"/>
        <arg value="retro"/>
        <arg value="Sinclair Spectrum for sale"/>
        <arg value="48k original box and packaging"/>
    </java>
    <java classname="sample.PostAdvert" classpathref="classpath.base">
        <arg value="dave"/>
        <arg value="kitsch"/>
        <arg value="Commemorative Plates"/>
        <arg value="Kitten and puppies design"/>
    </java>
    <java classname="sample.PostAdvert" classpathref="classpath.base">
        <arg value="jeff"/>
        <arg value="retro"/>
        <arg value="Atari 2600 wanted"/>
        <arg value="Must have original joysticks."/>
    </java>
</java>
```

```
        <java classname="sample.PostAdvert" classpathref="classpath.base">
            <arg value="jeff"/>
            <arg value="kitsch"/>
            <arg value="Inflatable Sofa"/>
            <arg value="Leopard skin pattern. Nice."/>
        </java>
    </target>

    <target name="listAdverts" depends="compile">
        <java classname="sample.ListAdverts" classpathref="classpath.base"/>
    </target>

</project>
```

The properties file imported in the first line provides the paths to your installed libraries, and you should adjust it as appropriate (as shown in Listing 3-2). If you unpack Hibernate 3.2.0, it will create a directory called `hibernate-3.2`, which we have renamed to the full version path; we have done something similar with the HSQL database directory.

The Hibernate Tools archive currently unpacks to two directories (`plugins` and `features`). We have created a parent directory to contain these. The path to the appropriate JAR file (`hibernate-tools.jar`) within the unpacked directory is dependent upon the specific Hibernate Tools version, so we have added the `hibernate.tools.path` property to point our build script at this.

Listing 3-2. *The* `build.properties` *File to Configure the Ant Script*

```
# Path to the hibernate install directory
hibernate.home=/hibernate/hibernate-3.2.0

# Path to the hibernate-tools install directory
hibernate.tools.home=/hibernate/hibernate-tools-3.1

# Path to hibernate-tools.jar relative to hibernate.tools.home
hibernate.tools.path=/plugins/org.hibernate.eclipse_3.1.0/lib/tools

# Path to the HSQL DB install directory
hsql.home=/hsqldb/hsqldb-1.8.0.2
```

Aside from the configuration settings, the only oddity in the `build.xml` file is the configuration and use of a Hibernate-specific Ant task. The `taskdef` (shown in Listing 3-3) makes this task available for use, using the appropriate classes from the `tools.jar` file.

Listing 3-3. *Defining the Hibernate Tools Ant Tasks*

```
<taskdef name="htools"
        classname="org.hibernate.tool.ant.HibernateToolTask"
        classpathref="classpath.tools"/>
```

This task provides several subtasks, but in this chapter we will only make use of the hbm2ddl subtask. This reads in the mapping and configuration files and generates Data Definition Language (DDL) scripts to create an appropriate schema in the database to represent our entities.

Table 3-2 shows the basic directories that our build script assumes, relative to the example project's root.

Table 3-2. *The Project Directories*

Directory	Contents
src	Source code and configuration files (excluding those directly related to the build)
bin	Compiled class files
sql	Generated DDL scripts

The root of the project contains the build script and build configuration file; it will also contain the database files generated by HSQL when the exportDDL task is run.

The Ant Tasks

Table 3-3 shows the tasks contained in the Ant build script.

Table 3-3. *The Tasks Available in the Example Ant Script*

Task	Action
exportDDL	Creates the appropriate database objects. It also generates a script that can be run against an HSQL database to re-create these objects if necessary.
compile	Builds the class files. This task is a dependency of all the tasks except exportDDL (which does not require the class files), so it is not necessary to invoke it directly.
populateMessages	Populates the database with a sample message.
listMessages	Lists all messages stored in the database by populateMessages.
createUsers	Creates a pair of users in the database for the Advert example.
createCategories	Creates a pair of categories in the database for the Advert example.
postAdverts	Creates several adverts in the database for the Advert example.
listAdverts	Lists the adverts in the database for the Advert example.

Enabling Logging

Before going on to run any of the examples in this chapter, you will want to create a log4j.properties file in the classpath. A suitable example is provided with the Hibernate tools in the etc directory of the unpacked archive.

Our example includes this file in the src directory of our project and places that directory itself on the classpath. In some circumstances—such as when building a JAR file for inclusion in other projects—it may be better to copy the appropriate properties file(s) into the target directory with the class files.

Creating a Hibernate Configuration File

There are several ways that Hibernate can be given all of the information that it needs to connect to the database and determine its mappings. For our Message example, we used the configuration file `hibernate.cfg.xml` placed in our project's `src` directory and given in Listing 3-4.

Listing 3-4. *The Message Application's Mapping File*

```xml
<?xml version='1.0' encoding='utf-8'?>
  <session-factory>
    <property name="hibernate.connection.url">
        jdbc:hsqldb:file:testdb;shutdown=true
    </property>
    <property name="hibernate.connection.driver_class">
        org.hsqldb.jdbcDriver
    </property>
    <property name="hibernate.connection.username">sa</property>
    <property name="hibernate.connection.password"></property>
    <property name="hibernate.connection.pool_size">0</property>
    <property name="hibernate.dialect">
        org.hibernate.dialect.HSQLDialect
    </property>
    <property name="hibernate.show_sql">false</property>

    <!-- "Import" the mapping resources here -->
    <mapping resource="sample/entity/Message.hbm.xml"/>

  </session-factory>
</hibernate-configuration>
```

The various database-related fields (`hibernate.connection.*`) should look pretty familiar from setting up JDBC connections, with the exception of the `hibernate.connection.pool` property, which is used to disable a feature (connection pooling) that causes problems when using the HSQL database. The `show_sql` value, set to `false` in our example, is extremely useful when debugging problems with your programs—when set to `true`, all of the SQL prepared by Hibernate is logged to the standard output stream (i.e., the console).

The SQL dialects, discussed in Chapter 2, allow you to select the database type that Hibernate will be talking to. You must select a dialect, even if it is `GenericDialect`—most database platforms accept a common subset of SQL, but there are inconsistencies and extensions specific to each. Hibernate uses the dialect class to determine the appropriate SQL to use when creating and querying the database. If you elect to use `GenericDialect`, then Hibernate will only be able to use a common subset of SQL to perform its operations, and will be unable to take advantage of various database-specific features to improve performance.

■**Caution** Hibernate looks in the classpath for the configuration file. If you place it anywhere else, Hibernate will complain that you haven't provided necessary configuration details.

Hibernate does not require you to use an XML configuration file. You have two other options. First, you can provide a normal Java properties file. The equivalent properties file to Listing 3-4 would be as follows:

```
hibernate.connection.driver_class=org.hsqldb.jdbcDriver
hibernate.connection.url=jdbc:hsqldb:file:testdb;shutdown=true
hibernate.connection.username=sa
hibernate.connection.password=
hibernate.connection.pool_size=0
hibernate.show_sql=false
hibernate.dialect=org.hibernate.dialect.HSQLDialect
```

As you'll notice, this does not contain the resource mapping from the XML file—and in fact, you cannot include this information in a properties file; if you want to configure Hibernate this way, you'll need to directly map your classes into the Hibernate `Configuration` at run time. Here's how this can be done:

```
Configuration config = new Configuration();
config.addClass( sample.entity.Message.class );
config.setProperties( System.getProperties() );
SessionFactory sessions = config.buildSessionFactory();
```

Note that the `Configuration` object will look in the classpath for a mapping file *in the same package* as the class it has been passed. So, in this example, where the fully qualified name of the class is `sample.entity.Message`, you should see the following pair of files from the root of the classpath:

```
/sample/entity/Message.class
/sample/entity/Message.hbm.xml
```

Here, `Message.class` is the compiled output from the `Message.java` code given in Listing 3-5 (and briefly discussed in Chapter 1), and `Message.hbm.xml` is the XML mapping file provided in Chapter 1 as Listing 1-5. If for some reason you want to keep your mapping files in a different directory, you can alternatively provide them as resources like this (note that this resource path must still be relative to the classpath):

```
Configuration config = new Configuration();
config.addResource( "Message.hbm.xml" );
config.setProperties( System.getProperties() );
SessionFactory sessions = config.buildSessionFactory();
```

You may have as many or as few mapping files as you wish, given any names you like—however, it is conventional to have one mapping file for each class that you are mapping, placed in the same directory as the class itself, and named similarly (for example, `Message.hbm.xml` in the default package to map the `Message` class also in the default package). This allows you to find any given class mapping quickly, and keeps the mapping files easily readable.

If you don't want to provide the configuration properties in a file, you can apply them directly using the `-D` flag. Here's an example:

```
java -classpath ...
   -Dhibernate.connection.driver_class=org.hsqldb.jdbcDriver
   -Dhibernate.connection.url= jdbc:hsqldb:file:testdb;shutdown=true
   -Dhibernate.connection.username=sa
   -Dhibernate.connection.password=
   -Dhibernate.connection.pool_size=0
   -Dhibernate.show_sql=false
   -Dhibernate.dialect=org.hibernate.dialect.HSQLDialect
    ...
```

Given its verbosity, this is probably the least convenient of the three methods, but it is occasionally useful when running tools and utilities on an ad hoc basis. For most other purposes, we think that the XML configuration file is the best choice.

Running the Message Example

With Hibernate and a database installed, and our configuration file created, all we need to do now is create the classes in full, and then build and run everything. Chapter 1 omitted the trivial parts of the required classes, so we provide them in full in Listings 3-5 through 3-7, after which we'll look at some of the details of what's being invoked.

Listing 3-5. *The* Message *POJO Class*

```
package sample.entity;

public class Message {
   private String message;

   public Message(String message) {
      this.message = message;
   }

   Message() {
   }

   public String getMessage() {
      return this.message;
   }

   public void setMessage(String message) {
      this.message = message;
   }
}
```

Listing 3-6 shows a simple application to populate the messages table with examples.

Listing 3-6. *The Code to Create a Sample Message*

```java
package sample;

import java.util.Date;

import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.cfg.Configuration;

import sample.entity.Message;

public class PopulateMessages {
    public static void main(String[] args) {
        SessionFactory factory =
            new Configuration().configure().buildSessionFactory();
        Session session = factory.openSession();
        session.beginTransaction();

        Message m1 = new Message(
            "Hibernated a message on " + new Date());

        session.save(m1);
        session.getTransaction().commit();
        session.close();
    }
}
```

Finally, Listing 3-7 shows the full text of the application to list all the messages in the database.

Listing 3-7. *The Message Application*

```java
package sample;

import java.util.Iterator;
import java.util.List;

import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.cfg.Configuration;

import sample.entity.Message;
```

```
public class ListMessages {
   public static void main(String[] args)
   {
      SessionFactory factory =
         new Configuration().configure().buildSessionFactory();
      Session session = factory.openSession();

      List messages = session.createQuery("from Message").list();
      System.out.println("Found " + messages.size() + " message(s):");

      Iterator i = messages.iterator();
      while(i.hasNext()) {
         Message msg = (Message)i.next();
         System.out.println(msg.getMessage());
      }

      session.close();
   }
}
```

The Ant target `exportDDL` will create an appropriate schema in the HSQLDB database files. Running the `populateMessages` task will create a message entry (this can be invoked multiple times). Running the `listMessages` task will list the messages that have been entered into the database so far.

■**Caution** Because we have selected the `drop="true"` option for the `hbm2ddl` subtask of our `exportDDL` target, running this script will effectively delete any data in the named tables. It is rarely a good idea to run such a script from a machine that has database access to the production environment because of the risk of accidentally deleting your production data!

The appropriate classpath entries have been set up in the Ant build script. To run a Hibernate application, you need the `hibernate.jar` file from the root of the Hibernate distribution, and a subset of the libraries provided in the `lib` subdirectory. The origin, purpose, and optionality of each of these libraries is explained in a `README` text file provided in the `lib` directory.

Most of the work required to get this example running is the sort of basic configuration trivia that any application requires (writing Ant scripts, setting classpaths, and so on). The real work consists of these steps:

1. Creating the Hibernate configuration file

2. Creating the mapping file

3. Writing the POJOs (introduced in Chapter 1)

Persisting Multiple Objects

Our example in Chapter 1 was as simple a persistence scenario as you can imagine. In the next few sections of this chapter, we will look at a slightly more complicated scenario.

Our example application will provide the persistence technology for an online billboard application, as shown in Figure 3-1.
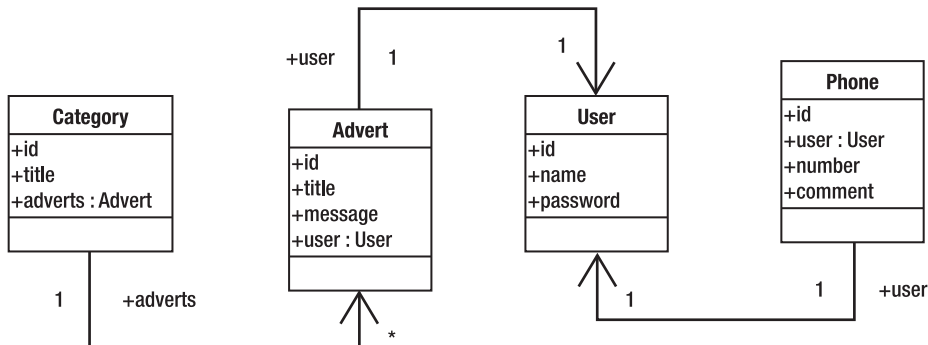

Figure 3-1. *The online billboard classes*

This is a gross simplification of the sort of classes that would be required in a production application. For example, we make no distinction between the roles of users of the application, but it should suffice to show some of the simpler relationships between classes.

Particularly interesting is the many-to-many relationship between categories and advertisements. We would like to be able to have multiple categories and adverts, and place any given advert in more than one category. For example, an electric piano should be listed in the "Instruments" category as well as the "Electronics" category.

Creating Persistence Classes

We will begin by creating the POJOs for the application. This is not strictly necessary in a new application, as they can be generated directly from the mapping files, but since this will be familiar territory, it should help to provide some context for our subsequent creation of the mapping files.

From the class diagram, we know that three classes will be persisted into the database (see Listings 3-8, 3-9, and 3-10). Each class that will be persisted by Hibernate is required to have a default constructor with at least package scope. They should have get and set methods for all of the attributes that are to be persisted. We will provide each with an id field, allowing this to be the primary key in our database (we prefer the use of surrogate keys, as changes to business rules can make the use of direct keys risky).

■**Note** A surrogate key is an arbitrary value (usually numeric), with the data type depending on the number of objects expected (e.g., 32-bit, 64-bit, etc.). The surrogate key has no meaning outside the database—it is not a customer number, a phone number, or anything else. As such, if a business decision causes previously unique business data to be duplicated, this will not cause problems since the business data does not form the primary key.

As well as the default constructor for each class, we provide a constructor that allows the fields other than the primary key to be assigned directly. This allows us to create and populate an object in one step instead of several, but we let Hibernate take care of the allocation of our primary keys.

The classes shown in Figure 3-1 are our POJOs. Their implementation is shown in Listings 3-8, 3-9, and 3-10.

Listing 3-8. *The Class Representing Users*

```java
package sample.entity;

public class User {
    private long id;
    private String name;
    private String password;

    public User(String name, String password) {
        this.name = name;
        this.password = password;
    }

    User() {
    }

    public String getName() {
        return name;
    }

    public void setName(String name) {
        this.name = name;
    }

    public String getPassword() {
        return password;
    }
```

```
   public void setPassword(String password) {
      this.password = password;
   }

   protected long getId() {
      return id;
   }

   protected void setId(long id) {
      this.id = id;
   }
}
```

Listing 3-9. *The Class Representing Categories (Each Having an Associated Set of* Advert *Objects)*

```
package sample.entity;

import java.util.HashSet;
import java.util.Set;

public class Category {
   private long id;
   private String title;
   private Set adverts = new HashSet();

   public Category(String title) {
      this.title = title;
      this.adverts = new HashSet();
   }

   Category() {
   }

   public Set getAdverts() {
      return adverts;
   }

   void setAdverts(Set adverts) {
      this.adverts = adverts;
   }

   public void addAdvert(Advert advert) {
      getAdverts().add(advert);
   }
```

```java
    public String getTitle() {
        return title;
    }

    public void setTitle(String title) {
        this.title = title;
    }

    protected long getId() {
        return id;
    }

    protected void setId(long id) {
        this.id = id;
    }
}
```

Listing 3-10. *The Class Representing Adverts (Each Instance Has an Associated User Who Placed the Advert)*

```java
package sample.entity;

public class Advert {
    private long id;
    private String title;
    private String message;
    private User user;

    public Advert(String title, String message, User user) {
        this.title = title;
        this.message = message;
        this.user = user;
    }

    Advert() {
    }

    public String getMessage() {
        return message;
    }

    public void setMessage(String message) {
        this.message = message;
    }
```

```
    public String getTitle() {
        return title;
    }

    public void setTitle(String title) {
        this.title = title;
    }

    public User getUser() {
        return user;
    }

    public void setUser(User user) {
        this.user = user;
    }

    protected long getId() {
        return id;
    }

    protected void setId(long id) {
        this.id = id;
    }
}
```

We have not had to add any unusual features to these classes in order to support the Hibernate tool, though we have chosen to provide package-scoped default constructors to support use of the (optional) lazy-loading feature of Hibernate. Most existing applications will contain POJOs "out of the box" that are compatible with Hibernate.

Creating the Object Mappings

Now that we have our POJOs, we need to map them to the database, representing the fields of each directly or indirectly as values in the columns of the associated tables. We take each in turn.

The fully qualified name of the type that we are mapping is specified, and the table in which we would like to store it is specified (we used aduser because user is a keyword in many databases).

The class has three fields, as follows:

The id *field*: Corresponds to the surrogate key to be used in, and generated by, the database. This special field is handled by the <id> element. The name of the field is specified by the name attribute (so that name="id" corresponds as it must with the method name of "getId"). It is identified as being of long type, and we would like to store its values in the database in the long column. We specify that it should be generated by the database, rather than by Hibernate.

The name *field*: Represents the name of the user. It should be stored in a column called name. It has type String. We do not permit duplicate names to be stored in the table.

The password *field*: Represents a given user's password. It should be stored in a column called password. It has type String.

Bearing these features in mind, the mapping file in Listing 3-11 should be extremely easy to follow.

Listing 3-11. *The Mapping of the* User *Class into the Database*

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping>
    <class name="com.hibernatebook.chapter3.User" table="aduser">

        <id name="id" type="long" column="id">
            <generator class="native"/>
        </id>

        <property name="name" column="name" type="string" unique="true"/>

        <property name="password" column="password" type="string"/>

    </class>
</hibernate-mapping>
```

The Category mapping presents another type of relationship: many-to-many. Each Category object is associated with a set of adverts, while any given advert can be associated with multiple categories.

The <set> element indicates that the field in question has a java.util.Set type with the name adverts. This sort of relationship requires the creation of an additional link table, so we specify the name of the table containing that information.

We state that the primary key (used to retrieve items) for the objects contained in the link table is represented by the id column, and provide the fully qualified name of the class type contained in the table. We specify the column in the link table representing the adverts associated with each category.

Again, this is complicated when described, but if you look at the example table from Listing 3-14, the need for each field in the mapping becomes clear (see Listing 3-12).

Listing 3-12. *The Mapping of the* Category *Class into the Database*

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">
```

```
<hibernate-mapping>
    <class name="sample.entity.Category" table="category">

        <id name="id" type="long" column="id">
            <generator class="native"/>
        </id>

        <property
                name="title"
                column="title"
                type="string"
                unique="true"/>

        <set name="adverts" table="link_category_advert" >
            <key column="category" foreign-key="fk_advert_category"/>
            <many-to-many class="sample.entity.Advert"
                column="advert"
                foreign-key="fk_category_advert"/>
        </set>

    </class>
</hibernate-mapping>
```

Finally, we represent the Advert class (see Listing 3-13). This class introduces the many-to-one association, in this case with the User class. Any given advertisement must belong to a single user, but any given user can place many different advertisements.

Listing 3-13. *The Mapping of the* Advert *Class into the Database*

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping>
    <class name="sample.entity.Advert" table="advert">

        <id name="id" type="long" column="id">
            <generator class="native"/>
        </id>

        <property name="message" column="message" type="string"/>
        <property name="title" column="title" type="string"/>
```

```
    <many-to-one
        name="user"
        column="aduser"
        class="sample.entity.User"
        not-null="true"
        foreign-key="fk_advert_user"/>

    </class>
</hibernate-mapping>
```

Once you have created the individual mapping files, you need to tell Hibernate where to find them. If you're using a Hibernate configuration file, as in the Chapter 1 example, the simplest thing to do is include links to the mapping files directly within this.

For our example, take the configuration file described for Chapter 1 (Listing 1-5) and add the following three mapping resource entries:

```
<mapping resource="sample/entity/Advert.hbm.xml"/>
<mapping resource="sample/entity/Category.hbm.xml"/>
<mapping resource="sample/entity/User.hbm.xml"/>
```

after the following line:

```
<mapping resource="sample/entity/Message.hbm.xml"/>
```

This section may seem confusing, as it is something of a flying visit to the subject of mappings and some of their whys and wherefores. We provide a more in-depth discussion of mapping in later chapters—specifically, general mapping concepts in Chapter 5, and XML-based mapping files in Chapter 7. We also discuss how you can use the new Java 5 Annotations features to represent mappings directly in your source code in Chapter 6.

Creating the Tables

With the object mapping in place and our Hibernate configuration file set up correctly, we have everything we need to generate a script to create the database for our application by invoking the exportDDL task. This builds the entities shown in Figure 3-2.

Even though we can generate the database directly, we also recommend taking some time to work out what schema you would expect your mappings to generate. This allows you to "sanity check" the script to make sure it corresponds with your expectations. If you and the tool both agree on what things should look like, then all is well and good; if not, your mappings may be wrong or there may be a subtle error in the way that you have related your data types.

Figure 3-2. *The database entity relationships*

The script in Listing 3-14 is generated by the `exportDDL` task. It could easily have been written by hand, and it is easy to compare it against your prior expectations of the database schema (we have changed the formatting slightly, but otherwise this is identical to the output of the task).

Listing 3-14. *The Script Generated by the* `exportDDL` *Task*

```
alter table advert
   drop constraint fk_advert_user;

alter table link_category_advert
   drop constraint fk_advert_category;

alter table link_category_advert
   drop constraint fk_category_advert;

drop table aduser if exists;
drop table advert if exists;
drop table category if exists;
drop table link_category_advert if exists;
drop table message if exists;

create table aduser (
   id bigint generated by default as identity (start with 1),
   name varchar(255),
   password varchar(255),
   primary key (id),
   unique (name));
```

```
create table advert (
    id bigint generated by default as identity (start with 1),
    message varchar(255),
    title varchar(255),
    aduser bigint not null,
    primary key (id));

create table category (
    id bigint generated by default as identity (start with 1),
    title varchar(255),
    primary key (id),
    unique (title));

create table link_category_advert (
    category bigint not null,
    advert bigint not null,
    primary key (category, advert));

create table message (
    id bigint generated by default as identity (start with 1),
    message varchar(255),
    primary key (id));

alter table advert
    add constraint fk_advert_user
    foreign key (aduser) references aduser;

alter table link_category_advert
    add constraint fk_advert_category
    foreign key (category) references category;

alter table link_category_advert
    add constraint fk_category_advert
    foreign key (advert) references advert;
```

Note the foreign key constraints and the link table representing the many-to-many relationship.

Sessions

Chapter 4 will discuss the full life cycle of persistence objects in detail—but you need to understand the basics of the relationship between the session and the persistence objects if you are to build even a trivial application in Hibernate.

The Session and Related Objects

The session is always created from a `SessionFactory`. The `SessionFactory` is a heavyweight object, and there would normally be a single instance per application. In some ways, it is a little like a connection pool in a connected application. In a J2EE application, it would typically be retrieved as a JNDI resource. It is created from a `Configuration` object, which in turn acquires the Hibernate configuration information and uses this to generate an appropriate `SessionFactory` instance.

The session itself has a certain amount in common with a JDBC `Connection` object. To read an object from the database, you must use a session directly or indirectly. An example of a direct use of the session to do this would be, as in Chapter 1, calling the `session.get()` method, or creating a `Query` object from the session (a `Query` is very much like a `PreparedStatement`).

An indirect use of the session would be using an object itself associated with the session. For example, if we have retrieved a `Phone` object from the database using a session directly, we can retrieve a `User` object by calling `Phone`'s `getUser()` method, even if the associated `User` object has not yet been loaded (as a result of lazy loading).

An object that has not been loaded via the session can be explicitly associated with the session in several ways, the simplest of which is to call the `session.update()` method passing in the object in question.

The session does a lot more than this, however, as it provides some caching functionality, manages the lazy loading of objects, and watches for changes to associated objects (so that the changes can be persisted to the database).

A Hibernate transaction is typically used in much the same way as a JDBC transaction. It is used to batch together mutually dependent Hibernate operations, allowing them to be completed or rolled back atomically, and to isolate operations from external changes to the database. Hibernate can also take advantage of a transaction's scope to limit unnecessary JDBC "chatter," queuing SQL to be transmitted in a batch at the end of the transaction when possible.

We will discuss all of this in much greater detail in Chapter 4, but for now it suffices that we need to maintain a single `SessionFactory` for the entire application. However, a session should only be accessed within a single thread of execution. Because a session also represents information cached from the database, it is desirable to retain it for use within the thread until anything (specifically any Hibernate exception) causes it to become invalid.

We present in Listing 3-15 a pattern from which Data Access Objects (DAOs) can be derived, providing an efficient way for a thread to retrieve and (if necessary) create its sessions with a minimal impact on the clarity of the code.

Listing 3-15. *The Base Class Used to Manage the Session in the Example*

```
package sample.dao;

import java.util.logging.Level;
import java.util.logging.Logger;
```

```java
import org.hibernate.HibernateException;
import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.cfg.Configuration;

public class DAO {

    protected DAO() {
    }

    public static Session getSession() {
        Session session = (Session) DAO.session.get();
        if (session == null) {
            session = sessionFactory.openSession();
            DAO.session.set(session);
        }
        return session;
    }

    protected void begin() {
        getSession().beginTransaction();
    }

    protected void commit() {
        getSession().getTransaction().commit();
    }

    protected void rollback() {
        try {
        getSession().getTransaction().rollback();
        } catch( HibernateException e ) {
            log.log(Level.WARNING,"Cannot rollback",e);
        }

        try {
            getSession().close();
        } catch( HibernateException e ) {
            log.log(Level.WARNING,"Cannot close",e);
        }
        DAO.session.set(null);
    }

    public static void close() {
        getSession().close();
        DAO.session.set(null);
    }
```

```
   private static final Logger log = Logger.getAnonymousLogger();

   private static final ThreadLocal session = new ThreadLocal();

   private static final SessionFactory sessionFactory =
      new Configuration().configure().buildSessionFactory();
}
```

Using the Session

The most common use cases for our POJOs will be to create them and delete them. In both cases, we want the change to be reflected in the database.

For example, we want to be able to create a user, specifying the username and password, and have this information stored in the database when we are done.

The logic to create a user (and reflect this in the database) is incredibly simple, as shown in Listing 3-16.

Listing 3-16. *Creating a* User *Object and Reflecting This in the Database*

```
try {
   begin();
   User user = new User(username,password);
   getSession().save(user);
   commit();
   return user;
} catch( HibernateException e ) {
   rollback();
   throw new AdException("Could not create user " + username,e);
}
```

We begin a transaction, create the new User object, ask the session to save the object, and then commit the transaction. If a problem is encountered (if, for example, a User entity with that username has already been created in the database), then a Hibernate exception will be thrown, and the entire transaction will be rolled back.

To retrieve the User object from the database, we will make our first excursion into HQL. HQL is somewhat similar to SQL, but you should bear in mind that it refers to the names used in the mapping files, rather than the table names and columns of the underlying database.

The appropriate HQL query to retrieve the users having a given name field is as follows:

```
from User where name= :username
```

where User is the class name and :username is the HQL named parameter that our code will populate when we carry out the query. This is remarkably similar to the SQL for a prepared statement to achieve the same end:

```
select * from user where name = ?
```

The complete code to retrieve a user for a specific username is shown in Listing 3-17.

Listing 3-17. *Retrieving a* User *Object from the Database*

```
try {
   begin();
   Query q = getSession().createQuery("from User where name = :username");
   q.setString("username",username);
   User user = (User)q.uniqueResult();
   commit();
   return user;
} catch( HibernateException e ) {
   rollback();
   throw new AdException("Could not get user " + username,e);
}
```

We begin a transaction, create a Query object (similar in purpose to PreparedStatement in connected applications), populate the parameter of the query with the appropriate username, and then list the results of the query. We extract the user (if one has been retrieved successfully) and commit the transaction. If there is a problem reading the data, the transaction will be rolled back.

The key line used to obtain the User entity is:

```
User user = (User)q.uniqueResult();
```

We use the uniqueResult()method because it is guaranteed to throw an exception if somehow our query identifies more than one User object for the given username. In principle, this could happen if the underlying database's constraints don't match our mapping constraint for a unique username field, and an exception is an appropriate way to handle the failure.

The logic to delete a user from the database (Listing 3-18) is even more trivial than that required to create one.

Listing 3-18. *Deleting a* User *Object and Reflecting This in the Database*

```
try {
   begin();
   getSession().delete(user);
   commit();
} catch( HibernateException e ) {
   rollback();
   throw new AdException("Could not delete user " + user.getName(),e);
}
```

We simply instruct the session to delete the User object from the database, and commit the transaction. The transaction will roll back if there is a problem—for example, if the user has already been deleted.

You have now seen all the basic operations that we want to perform on our data, so we will now take a look at the architecture we are going to use to do this.

Building DAOs

The DAO pattern is well known to most developers. The idea is to separate out the POJOs from the logic used to persist them into, and retrieve them from, the database. The specifics of the implementation vary—at one extreme, they can be provided as interfaces instantiated from a factory class, allowing a completely pluggable database layer. For our example, we have selected a compromise of concrete DAO classes. Each DAO class represents the operations that can be performed on a POJO type.

We have already described the base class DAO in Listing 3-15, and the preceding examples made use of this.

To help encapsulate the specifics of the database operations that are being carried out, we catch any HibernateException that is thrown and wrap it in a business AdException instance, as shown in Listing 3-19.

Listing 3-19. *The* AdException *Class for the Example*

```
package sample;

public class AdException extends Exception {
    public AdException(String message) {
        super(message);
    }

    public AdException(String message, Throwable cause) {
        super(message,cause);
    }
}
```

The UserDAO provides all the methods required to retrieve an existing User object, delete an existing User object, or create a new User object (see Listing 3-20). Changes to the object in question will be persisted to the database at the end of the transaction.

Listing 3-20. *The* UserDAO *Class for the Example*

```
package sample.dao;

import org.hibernate.HibernateException;
import org.hibernate.Query;

import sample.AdException;
import sample.entity.User;

public class UserDAO extends DAO {
    public UserDAO() {
    }
```

```java
    public User get(String username)
       throws AdException
    {
       try {
          begin();
          Query q = getSession().createQuery("from User where name = :username");
          q.setString("username",username);
          User user = (User)q.uniqueResult();
          commit();
          return user;
       } catch( HibernateException e ) {
          rollback();
          throw new AdException("Could not get user " + username,e);
       }
    }

    public User create(String username,String password)
       throws AdException
    {
       try {
          begin();
          User user = new User(username,password);
          getSession().save(user);
          commit();
          return user;
       } catch( HibernateException e ) {
          rollback();
          throw new AdException("Could not create user " + username,e);
       }
    }

    public void delete(User user)
       throws AdException
    {
       try {
          begin();
          getSession().delete(user);
          commit();
       } catch( HibernateException e ) {
          rollback();
          throw new AdException("Could not delete user " + user.getName(),e);
       }
    }
}
```

CategoryDAO provides all the methods required to retrieve all of the Category objects, delete an existing Category object, or create a new Category object (see Listing 3-21). Changes to the object in question will be persisted to the database at the end of the transaction.

Listing 3-21. *The* CategoryDAO *Class for the Example*

```java
package sample.dao;

import java.util.List;

import org.hibernate.HibernateException;
import org.hibernate.Query;

import sample.AdException;
import sample.entity.Category;

public class CategoryDAO extends DAO {
   public Category get(String title) throws AdException {
      try {
         begin();
         Query q = getSession().createQuery(
               "from Category where title = :title");
         q.setString("title", title);
         Category category = (Category) q.uniqueResult();
         commit();
         return category;
      } catch (HibernateException e) {
         rollback();
         throw new AdException("Could not obtain the named category " + title, e);
      }
   }

   public List list() throws AdException {
      try {
         begin();
         Query q = getSession().createQuery("from Category");
         List list = q.list();
         commit();
         return list;
      } catch (HibernateException e) {
         rollback();
         throw new AdException("Could not list the categories", e);
      }
   }
}
```

```java
    public Category create(String title) throws AdException {
        try {
            begin();
            Category cat = new Category(title);
            getSession().save(cat);
            commit();
            return null;
        } catch (HibernateException e) {
            rollback();
            throw new AdException("Could not create the category", e);
        }
    }

    public void save(Category category) throws AdException {
        try {
            begin();
            getSession().update(category);
            commit();
        } catch (HibernateException e) {
            rollback();
            throw new AdException("Could not save the category", e);
        }
    }

    public void delete(Category category) throws AdException {
        try {
            begin();
            getSession().delete(category);
            commit();
        } catch (HibernateException e) {
            rollback();
            throw new AdException("Could not delete the category", e);
        }
    }
}
```

AdvertDAO provides all the methods required to delete an existing Advert object or create a new Advert object (adverts are always retrieved by selecting them from a category, and are thus indirectly loaded by the CategoryDAO class). Changes to the object in question will be persisted to the database at the end of the transaction (see Listing 3-22).

Listing 3-22. *The* AdvertDAO *Class for the Example*

```java
package sample.dao;

import org.hibernate.HibernateException;
```

```java
import sample.AdException;
import sample.entity.Advert;
import sample.entity.User;

public class AdvertDAO extends DAO {
    public Advert create(String title, String message, User user)
            throws AdException {
        try {
            begin();
            Advert advert = new Advert(title, message, user);
            getSession().save(advert);
            commit();
            return advert;
        } catch (HibernateException e) {
            rollback();
            throw new AdException("Could not create advert", e);
        }
    }

    public void delete(Advert advert)
        throws AdException
    {
        try {
            begin();
            getSession().delete(advert);
            commit();
        } catch (HibernateException e) {
            rollback();
            throw new AdException("Could not delete advert", e);
        }
    }
}
```

If you compare the amount of code required to create our DAO classes here with the amount of code that would be required to implement them using the usual JDBC approach, you will see that Hibernate's logic is admirably compact.

The Example Client

Listing 3-23 shows the example code tying this together. Of course, this isn't a full application, but you now have all the DAOs necessary to manage the advertisement database. This example gives a flavor of how they can be used.

The code should be run with the tasks in the Ant script delivered in Listing 3-1. After running the exportDDL task to create the empty database, you should run the createUsers and createCategories tasks to provide initial users and categories, and then the postAdverts task to place advertisements in the database. Finally, run the listAdverts task to display the saved data.

The code invoking the DAOs to perform the tasks in question is shown in Listing 3-23.

Listing 3-23. *The Class to Create the Example Users*

```
package sample;

import sample.dao.DAO;
import sample.dao.UserDAO;

public class CreateUser {
   public static void main(String[] args) {

      if (args.length != 2) {
         System.out.println("params required: username, password");
         return;
      }
      String username = args[0];
      String password = args[1];

      try {
         UserDAO userDao = new UserDAO();

         System.out.println("Creating user " + username);
         userDao.create(username, password);
         System.out.println("Created user");
         DAO.close();
      } catch (AdException e) {
         System.out.println(e.getMessage());
      }

   }
}
```

The CreateUser class uses the UserDAO class to create and persist an appropriate User object. The specifics of the (two) users created are drawn from the command-line parameters provided in the createUsers Ant task.

In Listing 3-24, we create Category objects via the CategoryDAO class—and again we draw the specific details from the command line provided by the Ant script.

Listing 3-24. *The Class to Create the Example Categories*

```
package sample;

import sample.dao.CategoryDAO;
import sample.dao.DAO;

public class CreateCategory {
   public static void main(String[] args) {
```

```
        if (args.length != 1) {
            System.out.println("param required: categoryTitle");
            return;
        }

        CategoryDAO categories = new CategoryDAO();
        String title = args[0];
        try {
            System.out.println("Creating category " + title);
            categories.create(title);
            System.out.println("Created category");
            DAO.close();
        } catch (AdException e) {
            System.out.println(e.getMessage());
        }

    }
}
```

The code in Listing 3-25 allows us to create an advert for a preexisting user in a pre-existing category. Note our use of UserDAO and CategoryDAO to obtain User and Category objects from the database. As with the user and category, the advert details are supplied by the Ant task.

Listing 3-25. *The Class to Create the Example Adverts*

```
package sample;

import sample.dao.AdvertDAO;
import sample.dao.CategoryDAO;
import sample.dao.DAO;
import sample.dao.UserDAO;
import sample.entity.Advert;
import sample.entity.Category;
import sample.entity.User;

public class PostAdvert {
    public static void main(String[] args) {

        if (args.length != 4) {
System.out.println("params required: username, categoryTitle, title, message");
            return;
        }
```

```
        String username = args[0];
        String categoryTitle = args[1];
        String title = args[2];
        String message = args[3];

        try {
            UserDAO users = new UserDAO();
            CategoryDAO categories = new CategoryDAO();
            AdvertDAO adverts = new AdvertDAO();

            User user = users.get(username);
            Category category = categories.get(categoryTitle);
            Advert advert = adverts.create(title, message, user);

            category.addAdvert(advert);
            categories.save(category);

            DAO.close();
        } catch (AdException e) {
            e.printStackTrace();
        }

    }
}
```

Finally, in Listing 3-26, we make use of `CategoryDAO` to iterate over the categories, and within these, the adverts drawn from the database. It is easy to see how this logic could now be incorporated into a JSP file or servlet. It could even be used from within an EJB session bean.

Listing 3-26. *The Class to Display the Contents of the Database*

```
package sample;

import java.util.Iterator;
import java.util.List;

import sample.dao.CategoryDAO;
import sample.dao.DAO;
import sample.entity.Advert;
import sample.entity.Category;

public class ListAdverts {
    public static void main(String[] args) {
        try {
            List categories = new CategoryDAO().list();
```

```
        Iterator ci = categories.iterator();
        while(ci.hasNext()) {
            Category category = (Category)ci.next();
            System.out.println("Category: " + category.getTitle());
            System.out.println();
            Iterator ai = category.getAdverts().iterator();
            while(ai.hasNext()) {
                Advert advert = (Advert)ai.next();
                System.out.println();
                System.out.println("Title: " + advert.getTitle());
                System.out.println(advert.getMessage());
                System.out.println(" posted by " + advert.getUser().getName());
            }
        }

        DAO.close();
    } catch( AdException e ) {
        System.out.println(e.getMessage());
    }


    }
}
```

A large part of the logic here is either output information, or concerned with accessing the collections themselves. Java 5 devotees will see an obvious opportunity to make use of generics and enhanced `for` loops in this example. A quick taste of the simplified version of this code might look like Listing 3-27.

Listing 3-27. *Enhancing Your DAOs with Java 5 Features*

```
List<Category> categories = new CategoryDAO().list();
for( Category category : categories ) {
   // ...
   for( Advert advert : category.getAdverts() ) {
      // ...
   }
}
DAO.close();
```

When you run the example applications, you will see a considerable amount of "chatter" from the logging API, and from the Ant tool when you run these tasks, much of which can be controlled or eliminated in a production application.

You will also notice that because you are starting each of these applications as new tasks (several times in the case of the tasks to create data), the tasks proceed relatively slowly. This is an artifact of the repeated creation of `SessionFactory`—a heavyweight object—from each invocation of the JVM from the Ant `java` task, and is not a problem in "real" applications.

Summary

In this chapter, we've shown how to acquire the Hibernate tools, how to create and run the example from Chapter 1, and how to create a slightly larger application from scratch, driving the database table generation from the `hbm2ddl` Ant task. All of the files described in this chapter and the others can be downloaded from the Apress web site (`www.apress.com`).

In the next chapter, we will look at the architecture of Hibernate and the lifecycle of a Hibernate-based application.

■ ■ ■

The Persistence Life Cycle

In this chapter, we discuss the life cycle of persistent objects in Hibernate. These persistent objects are POJOs without any special marker interfaces or inheritance related to Hibernate. Part of Hibernate's popularity comes from its ability to work with a normal object model.

We also discuss the methods of the `Session` interface that are used for creating, retrieving, updating, and deleting persistent objects from Hibernate.

Introduction to the Life Cycle

After adding Hibernate to your application, you do not need to change your existing Java object model to add persistence marker interfaces or any other type of hint for Hibernate. Instead, Hibernate works with normal Java objects that your application creates with the `new` operator, or that other objects create. For Hibernate's purposes, these can be drawn up into two categories: objects for which Hibernate has entity mappings, and objects that are not directly recognized by Hibernate. A correctly mapped entity object will consist of fields and properties that are mapped, and that are themselves either references to correctly mapped entities, references to collections of such entities, or "value" types (primitives, primitive wrappers, strings, or arrays of these).

Given an instance of an object that is mapped to Hibernate, it can be in any one of three different states: transient, persistent, or detached.

Transient objects exist in memory, as illustrated in Figure 4-1. Hibernate does not manage transient objects or persist changes to transient objects.

Transient Object

Figure 4-1. *Transient objects are independent of Hibernate.*

To persist the changes to a transient object, you would have to ask the session to save the transient object to the database, at which point Hibernate assigns the object an identifier.

Persistent objects exist in the database, and Hibernate manages the persistence for persistent objects. We show this relationship between the objects and the database in Figure 4-2. If fields or properties change on a persistent object, Hibernate will keep the database representation up-to-date.

Figure 4-2. *Persistent objects are maintained by Hibernate.*

Detached objects have a representation in the database, but changes to the object will not be reflected in the database, and vice versa. This temporary separation of the object and the database is shown in Figure 4-3. A detached object can be created by closing the session that it was associated with, or by evicting it from the session with a call to the session's `evict()` method.

Figure 4-3. *Detached objects exist in the database but are not maintained by Hibernate.*

In order to persist changes made to a detached object, the application must reattach it to a valid Hibernate session. A detached instance can be associated with a new Hibernate session when your application calls one of the `load()`, `refresh()`, `merge()`, `update()`, or `save()` methods on the new session with a reference to the detached object. After the call, the detached object would be a persistent object managed by the new Hibernate session.

Versions prior to Hibernate 3 had support for the `Lifecycle` and `Validatable` interfaces. These allowed your objects to listen for `save`, `update`, `delete`, `load`, and `validate` events using methods on the object. In Hibernate 3, this functionality moved into events and interceptors, and the old interfaces were removed.

Entities, Classes, and Names

Entities represent Java objects with mappings that permit them to be stored in the database. The mappings indicate how the fields and properties of the object should be stored in the database tables. However, it is possible that you will want objects of a particular type to be represented in two different ways in the database. In this case, how does Hibernate choose which to use?

An object representing an entity will have a normal Java class type. It will also have an entity name. By default, the name of the entity will be the same as the name of the class type. You have the option, however, to change this via the mappings, and thus distinguish between objects of the same type that are mapped to different tables. There are therefore methods in the Session API that require an entity name to be provided to determine the appropriate mapping. If this is omitted, it will either be because no such distinction is needed, or because, for convenience, the method assumes the most common case—in which the entity name is the same as the class name—and duplicates the functionality of another more specific method that permits the entity name to specified explicitly.

Identifiers

Hibernate requires all entities to have an identifier, which represents the primary key column(s) of the table to which it will be persisted. When an entity is persisted, a suitable identifier can be assigned to it automatically by Hibernate, or a suitable identifier may be explicitly assigned by the user (see Listing 4-1).

Listing 4-1. *A Typical Identifier Field*

```
public int id;
```

Usually, the entity will provide a suitable identifier field or property, and Hibernate will use this value to correlate entities in memory with those persisted to the tables of the database. However, if no such field or property is available (as will likely be the case with legacy code), then Hibernate itself can manage the identifier value internally. The type of the identifier must be defined in the mapping information.

Entities and Associations

Entities can contain references to other entities—either directly as a property or field, or indirectly via a collection of some sort (arrays, sets, lists, etc.). These associations are represented using foreign key relationships in the underlying tables.

When only one of the pair of entities contains a reference to the other, the association is unidirectional. If the association is mutual, then it is referred to as bidirectional.

■**Tip** A common mistake when designing entity models using Hibernate is to try to make all associations bidirectional. Associations that are not a natural part of the object model should not be forced into it. HQL often presents a more natural way to access the same information.

If both ends of the association managed the foreign keys, then we would encounter a problem when client code called the appropriate set method on both ends of the association. Should two foreign key columns be maintained—one in each direction (risking circular dependencies)—or only one? (And if only one, should altering either side affect it, or only one?)

Ideally, we would like to dictate that only changes to one end of the relationship will result in any updates to the foreign key; and indeed, Hibernate allows us to do this by marking one end of the association as being managed by the other (in the XML mapping files, this is known as the "`inverse`" of the parent, whereas in the EJB 3 terminology used by the annotation mappings, it is marked as being "`mappedBy`" the parent).

■**Caution** `inverse` and `mappedBy` are purely about how the foreign key relationships between entities are saved. They have nothing to do with saving the entities themselves. Despite this, they are often confused with the entirely orthogonal cascade functionality (described in the "Cascading Operations" section of this chapter).

While Hibernate lets us specify that changes to one association will result in changes to the database, it does *not* allow us to cause changes to one end of the association to be automatically reflected in the other end in the Java POJOs. For example, in a one-to-one bidirectional association between an `Email` class and a `Message` class, the code in Listing 4-2 is incomplete even if the `Message` entity is the inverse of the `Email` entity:

Listing 4-2. *A Common Misconception About Bidirectional Associations*

```
Email email = new Email("Test Email");
Message message = new Message("Test Message");
email.setMessage(message);
// Incorrectly managed
session.save(email);
session.save(message);
System.out.println(message.getEmail());
```

The final call to `message.getEmail()` will return `null` (assuming simple getters and setters are used). To get the desired effect, both entities must be updated—If the `Email` entity owns the association, this merely ensures the proper assignment of a foreign key column value. There is *no* implicit call of `message.setEmail(email)`. This must be explicitly given as in Listing 4-3.

Listing 4-3. *The Correct Maintenance of a Bidirectional Association*

```
Email email = new Email("Test Email");
Message message = new Message("Test Message");
email.setMessage(message);
message.setEmail(email); // Correctly managed
session.save(email);
session.save(message);
System.out.println(message.getEmail());
```

It is common for users new to Hibernate to get confused about this point—the confusion arises from two origins, which are described in the following paragraphs.

EJB container-managed persistence (CMP) *does* work in this way—when a reference is assigned in one entity, the corresponding reference in the other entity will be updated to reflect this. Hibernate does not take this approach because it was designed to work in other environments and with POJOs from other bodies of code, where such behavior would be unexpected. If you pass a pair of objects to some third-party API, mysterious side effects should not occur. Since Hibernate is precisely that—a third-party API—from the perspective of most client code, Hibernate cannot safely cause their references to become connected in this way!

Though Hibernate does not produce this behavior automatically, there *is* a side effect of persistence to the database that can make it appear that it does (see Listing 4-4).

Listing 4-4. *Misleading Behavior*

```
openSession();
beginTransaction();

Email email = new Email("Test Email");
Message message = new Message("Test Message");
email.setMessage(message);

save(email,message);

System.out.println("Stored...");
System.out.println(email);
System.out.println(email.getMessage());
System.out.println(message);
System.out.println(message.getEmail());

Serializable emailPrimaryKey = session.getIdentifier(email);
Serializable messagePrimaryKey = session.getIdentifier(message);
endTransaction();

closeSession();

System.out.println();

openSession();

beginTransaction();

email = (Email)session.get(Email.class,emailPrimaryKey);
message = (Message)session.get(Message.class,messagePrimaryKey);
```

```
System.out.println("Retrieved...");
System.out.println(email);
System.out.println(email.getMessage());
System.out.println(message);
System.out.println(message.getEmail());
endTransaction();
closeSession();
```

If you run the code from Listing 4-4, you will see the following output:

```
Stored...
Test Email
Test Message
Test Message
null

Retrieved...
Test Email
Test Message
Test Message
Test Emails
```

When the entities are initially stored, the Message object's reference to its associated Email is null, even after Hibernate has stored the data. The entity in memory is not updated to reflect the change to the Email entity. However, after we have closed the session, opened a new one, and loaded the entities from the database, the entity has been updated.

Because the session has been closed, the session is forced to reload the entities from the database when we request them by primary key. Because the Email entity is the owner of the association, the association exists in the database purely in the form of a foreign key relationship from the Email table onto the Message table's primary key. When we altered the Email entity and saved it, this foreign key relationship was therefore updated. So, when we reload the entities, the Message entity's association details are (correctly) obtained from the same foreign key.

If we alter this code to make the association in the Message entity instead of the Email entity, but leave the Email entity the owner of the association, we will see the reverse effect, as follows:

```
Email email = new Email("Test Email");
Message message = new Message("Test Message");
//email.setMessage(message);
message.setEmail(email);
```

Because we have not made the association in the Email entity (the owner), the foreign key of the Email table is not pointed at the Message table. When we reload the entities, we do not see the "automatic" association behavior—quite the opposite:

```
Stored...
Test Email
null
Test Message
Test Email
```

```
Retrieved...
Test Email
null
Test Message
null
```

As you can see, although the two entities have been saved, the attempt to associate the Email with the Message entity by calling a method on the Message entity has not been honored in the database, because the Message entity does not own the association. The following list recaps the points made so far:

- You must explicitly manage both ends of an association.

- Only changes to the owner of an association will be honored in the database.

- When you load a detached entity from the database, it will reflect the foreign key relationships persisted into the database.

Table 4-1 shows how you can select the side of the relationship that should be made the owner of a bidirectional association. Remember that to make an association the owner, you must mark the *other* end as inverse="true" (the choice of terminology is poor, but entrenched).

Table 4-1. *Marking the Owner of an Association*

Type of Association	Options
One-to-one	Either end can be made the owner, but one (and only one) of them should be—if you don't specify this, you will end up with a circular dependency.
One-to-many	The *many* end must be made the owner of the association.
Many-to-one	This is the same as the one-to-many relationship, viewed from the opposite perspective, so the same rule applies—the *many* end must be made the owner of the association.
Many-to-many	Either end of the association can be made the owner.

If this all seems rather confusing, just remember that association ownership is concerned exclusively with the management of the foreign keys in the database, and things should become clearer as you use Hibernate further. Associations and mappings are discussed in detail in the next three chapters.

Saving Entities

Creating an instance of a class you mapped with a Hibernate mapping does not automatically persist the object to the database. Until you explicitly save the object with a valid Hibernate session, the object is transient, like any other Java object. In Hibernate, we use one of the save() methods on the Session interface to store a transient object in the database, as follows:

```
public Serializable save(Object object) throws HibernateException
public void save(Object object, Serializable id) throws HibernateException
public Serializable save(String entityName,Object object) throws HibernateException
```

All of the `save()` methods take a transient object reference (which must not be null) as an argument. Hibernate expects to find a mapping for the transient object's class—Hibernate cannot persist arbitrary unmapped objects.

When persisting an object, it is possible for the user to override the identifier value generated by Hibernate by setting the `id` field property or calling an alternative `save()` method that also takes an object `id` as an argument. Obviously, this is most useful when the entity does not have its own `id` field or property, but it can be used even if it does. The `id` must not be null, and must be a `Serializable` value. If your entity uses a primitive to represent the `id` field, you can use the appropriate wrapper object. For example, `int` identifier values can be wrapped in `java.lang.Integer` objects.

The `save()` methods all create a new `org.hibernate.event.SaveOrUpdateEvent` event. We discuss events in more detail in Appendix A, although you do not have to worry about these implementation details to use Hibernate effectively.

At its simplest, we create a new object in Java, set a few of its properties, and then save it through the session, as follows:

```
Supplier superCorp = new Supplier();
superCorp.setName("SuperCorp");
session.save(superCorp);
```

It is not appropriate to save an object that has already been persisted. Equally, it is not appropriate to update a transient object. If it is impossible or inconvenient to determine the state of the object from your application code, you may use the `saveOrUpdate()` method. Hibernate uses the identifier of the object to determine whether to insert a new row into the database or update an existing row. The method signature is as follows:

```
public void saveOrUpdate(Object object) throws HibernateException
```

Once an object is in a persistent state, Hibernate manages updates to the database itself as you change the fields and properties of the object.

In order to support the behavior required by the new EJB 3 `EntityManager`, a `persist()` method has been added to the API. This behaves the same as the `save()` methods, except it is not guaranteed that the entities will immediately be assigned a primary key value. We discuss EJB 3 and the `EntityManager` in Appendix A.

Object Equality and Identity

When we discuss persistent objects in Hibernate, we also need to consider the role that object equality and identity plays with Hibernate. When we have a persistent object in Hibernate, that object represents both an instance of a class in a particular Java virtual machine (JVM) and a row (or rows) in a database table (or tables).

Requesting a persistent object again from the same Hibernate session returns the same Java instance of a class, which means that you can compare the objects using the standard Java `==` equality syntax. If, however, you request a persistent object from more than one Hibernate session, Hibernate will provide distinct instances from each session, and the `==` operator will return `false` if you compare these object instances.

Taking this into account, if you are comparing objects in two different sessions, you will need to implement the `equals()` method on your Java persistence objects.

Loading Entities

Hibernate's Session interface provides several load() methods for loading entities from your database. Each load() method requires the object's primary key as an identifier.

In addition to the id, Hibernate also needs to know which class or entity name to use to find the object with that id. Last, you will need to cast the object returned by load() to the class you desire. The basic load() methods are as follows:

```
public Object load(Class theClass, Serializable id) throws HibernateException
public Object load(String entityName, Serializable id) throws HibernateException
public void load(Object object, Serializable id) throws HibernateException
```

The last load() method takes an object as an argument. The object should be of the same class as the object you would like loaded, and it should be empty. Hibernate will populate that object with the object you requested. We find this syntax to be somewhat confusing when put into applications, so we do not tend to use it ourselves.

The other load() methods take a lock mode as an argument. The lock mode specifies whether Hibernate should look into the cache for the object, and which database lock level Hibernate should use for the row (or rows) of data that represent this object. The Hibernate developers claim that Hibernate will usually pick the correct lock mode for you, although we have seen situations in which it is important to manually choose the correct lock. In addition, your database may choose its own locking strategy—for instance, locking down an entire table rather than multiple rows within a table. In order of least restrictive to most restrictive, the various lock modes you can use are as follows:

- NONE: Uses no row-level locking, and uses a cached object if available; this is the Hibernate default.

- READ: Prevents other SELECT queries from reading data that is in the middle of a transaction (and thus possibly invalid) until it is committed.

- UPGRADE: Uses the SELECT FOR UPDATE SQL syntax to lock the data until the transaction is finished.

- UPGRADE_NOWAIT: Uses the NOWAIT keyword (for Oracle), which returns an error immediately if there is another thread using that row. Otherwise this is similar to UPGRADE.

All of these lock modes are static fields on the org.hibernate.LockMode class. We discuss locking and deadlocks with respect to transactions in more detail in Chapter 8. The load() methods that use lock modes are as follows:

```
public Object load(Class theClass, Serializable id, LockMode lockMode)
          throws HibernateException

public Object load(String entityName, Serializable id, LockMode lockMode)
          throws HibernateException
```

You should not use a load() method unless you are sure that the object exists. If you are not certain, then use one of the get() methods. The load() methods will throw an exception if the unique id is not found in the database, whereas the get() methods will merely return a null reference.

Much like `load()`, the `get()` methods take an identifier and either an entity name or a class. There are also two `get()` methods that take a lock mode as an argument. The `get()` methods are as follows:

```
public Object get(Class clazz, Serializable id)
        throws HibernateException

public Object get(String entityName, Serializable id)
        throws HibernateException

public Object get(Class clazz, Serializable id, LockMode lockMode)
        throws HibernateException

public Object get(String entityName, Serializable id, LockMode lockMode)
        throws HibernateException
```

If you need to determine the entity name for a given object (by default, this is the same as the class name), you can call the `getEntityName()` method on the `Session` interface, as follows:

```
public String getEntityName(Object object) throws HibernateException
```

Using the `get()`and `load()` methods is straightforward. For the following code sample, we would be getting the `Supplier` id from another Java class. For instance, through a web application, someone may select a `Supplier` details page for the supplier with the id 1. If we are not sure that the supplier exists, we use the `get()` method, with which we could check for `null`, as follows:

```
// get an id from some other Java class, for instance, through a web application
Supplier supplier = (Supplier) session.get(Supplier.class,id);
if (supplier == null) {
    System.out.println("Supplier not found for id " + id);
    return;
}
```

We can also retrieve the entity name from Hibernate and use it with either the `get()` or `load()` method. The `load()` method will throw an exception if an object with that `id` cannot be found.

```
String entityName = session.getEntityName(supplier);
Supplier secondarySupplier = (Supplier) session.load(entityName,id);
```

Refreshing Entities

Hibernate provides a mechanism to refresh persistent objects from their database representation. Use one of the `refresh()` methods on the `Session` interface to refresh an instance of a persistent object, as follows:

```
public void refresh(Object object)
            throws HibernateException

public void refresh(Object object, LockMode lockMode)
            throws HibernateException
```

Updating Entities

Hibernate automatically persists into the database changes made to persistent objects. If a property changes on a persistent object, the associated Hibernate session will queue the change for persistence to the database using SQL. From a developer's perspective, you do not have to do any work to store these changes, unless you would like to force Hibernate to commit all of its changes in the queue. You can also determine whether the session is dirty and changes need to be committed. When you commit a Hibernate transaction, Hibernate will take care of these details for you.

The flush() method forces Hibernate to flush the session, as follows:

```
public void flush() throws HibernateException
```

You can determine if the session is dirty with the isDirty() method, as follows:

```
public boolean isDirty() throws HibernateException
```

You can also instruct Hibernate to use a flushing mode for the session with the setFlushMode() method. The getFlushMode() method returns the flush mode for the current session, as follows:

```
public void setFlushMode(FlushMode flushMode)
public FlushMode getFlushMode()
```

The possible flush modes are the following:

- ALWAYS: Every query flushes the session before the query is executed.

- AUTO: Hibernate manages the query flushing to guarantee that the data returned by a query is up-to-date.

- COMMIT: Hibernate flushes the session on transaction commits.

- NEVER: Your application needs to manage the session flushing with the flush() method. Hibernate never flushes the session itself.

By default, Hibernate uses the AUTO flush mode. Generally, you should use transaction boundaries to ensure that appropriate flushing is taking place, rather than trying to "manually" flush at the appropriate times.

Deleting Entities

In order to allow convenient removal of entities from the database, the Session interface provides a delete() method, as follows:

```
public void delete(Object object) throws HibernateException
```

This method takes a persistent object as an argument. The argument can also be a transient object with the identifier set to the id of the object that needs to be erased.

In the simplest form, in which you are simply deleting an object with no associations to other objects, this is straightforward; but many objects do have associations with other objects. To allow for this, Hibernate can be configured to allow deletes to cascade from one object to its associated objects.

For instance, consider the situation in which you have a parent with a collection of child objects, and you would like to delete them all. The easiest way to handle this is to use the cascade attribute on the collection's element in the Hibernate mapping. If you set the cascade attribute to delete or all, the delete will be cascaded to all of the associated objects. Hibernate will take care of deleting these for you—deleting the parent erases the associated objects.

Hibernate 3 also supports bulk deletes (see Listing 4-5), where your application executes a DELETE HQL statement against the database. These are very useful for deleting more than one object at a time because each object does not need to be loaded into memory just to be deleted.

Listing 4-5. *A Bulk Delete Using a Hibernate Query*

```
session.createQuery("delete from User").executeUpdate();
```

Network traffic is greatly reduced, as are the memory requirements compared to those for individually issuing a delete() call against each entity identifier.

■**Caution** Bulk deletes do not cause cascade operations to be carried out. If cascade behavior is needed, you will need to carry out the appropriate deletions yourself, or use the session's delete() method.

Cascading Operations

When you perform one of the operations described in this chapter on an entity, the operations will not be performed on the associated entities unless you explicitly tell Hibernate to perform them.

For example, the following code will fail when we try to commit the transaction because the message entity that is associated with the Email entity has not been persisted into the database—and so the Email entity cannot be accurately represented (with its foreign key onto the appropriate message row) in its table.

```
Session session = factory.openSession();
session.beginTransaction();

Email email = new Email("Email title");
Message message = new Message("Message content");
email.setMessage(message);
message.setEmail(email);

session.save(email);

session.getTransaction().commit();
session.close();

session = factory.openSession();
session.beginTransaction();
```

Ideally, we would like the save operation to be propagated from the Email entity to its associated Message object. We do this by setting the cascade operations for the properties and fields of the entity (or assigning an appropriate default value for the entity as a whole). So, the preceding code will perform correctly if at least the save cascade operation is set for the Email entity's message property. All of the basic life cycle operations discussed in this chapter have associated cascade values, as follows:

- create

- merge

- delete

- save-update

- evict

- replicate

- lock

- refresh

These values can be concatenated in a comma-separated list to allow cascading for any combination of these operations. When all operations should be cascaded, Hibernate provides a shortcut value named all that tells Hibernate to cascade all of these operations from the parent to each child object (for that relationship).

As part of the Hibernate mapping process, you can tell Hibernate to use one of these cascading types for a relationship between two objects (the parent and the child). On the collection or property element in the mapping file, set the cascade attribute to the type (or types) you would like to use.

By default, Hibernate does not cascade any operations—the default behavior can be overridden at the entity level via the XML mapping files using the default-cascade attribute on the <hibernate-mapping> XML element or in the annotated source files.

The last possible cascading type is `delete-orphan`. Use `delete-orphan` to remove a child object from the database when you remove the child from the parent's collection. This cascading type only works on one-to-many associations. The `all` cascading type does not include `delete-orphan`—you will have to use `"all,delete-orphan"`, as in the following excerpt from a Hibernate mapping file:

```
<bag name="products" inverse="true" cascade="all,delete-orphan">
    <key column="supplierId"/>
    <one-to-many class="Product"/>
</bag>
```

Simply remove a child object from a parent object's collection after you have added the `delete-orphan` cascading type. Hibernate will remove the child object from the database itself, without any additional calls. The following example removes a child object from the collection:

```
supplier.getProducts().remove(product);
```

Lazy Loading, Proxies, and Collection Wrappers

Consider the stereotypical Internet web application: the online store. The store maintains a catalog of products. At the crudest level, this can be modeled as a catalog entity managing a series of product entities. In a large store, there may be tens of thousands of products grouped into various overlapping categories.

When a customer visits the store, the catalog must be loaded from the database. We probably don't want the implementation to load every single one of the entities representing the tens of thousands of products to be loaded into memory. For a sufficiently large retailer, this might not even be possible given the amount of physical memory available on the machine. Even if this were possible, it would probably cripple the performance of the site.

Instead, we want only the catalog to load, possibly with the categories as well. Only when the user drills down into the categories should a subset of the products in that category be loaded from the database.

To manage this problem, Hibernate provides a facility called lazy loading. When enabled (this is the default using XML mappings, but not when using annotations), an entity's associated entities will only be loaded when they are directly requested. For example, the following code loads only a single entity from the database:

```
Email email = (Email)session.get(Email.class,new Integer(42));
```

whereas if an association of the class is accessed, and lazy loading is in effect, the associations are pulled from the database as needed. For instance, in the following snippet, the associated `Message` object will be loaded since it is explicitly referenced.

```
Email email = (Email)session.get(Email.class,new Integer(42));
String text = email.getMessage().getContent();
```

The simplest way that Hibernate can force this behavior upon our entities is by providing a proxy implementation of them. Hibernate intercepts calls to the entity by substituting for it a proxy derived from the entity's class. Where the requested information is missing, it will be

loaded from the database before control is ceded to the parent entity's implementation. Where the association is represented as a collection class, a wrapper (essentially a proxy for the collection, rather than for the entities that it contains) is created and substituted for the original collection.

Hibernate can only access the database via a session. If an entity is detached from the session when we try to access an association (via a proxy or collection wrapper) that has not yet been loaded, Hibernate throws an exception: the infamous `LazyInitializationException`. The cure is to either ensure that the entity is made persistent again, or ensure that all of the fields that will be required are accessed *before* the entity is detached from the session.

If you need to determine whether a proxy, a persistence collection, or an attribute has been lazy loaded or not, you can call the `isInitialized()` and `isPropertyInitialized()` methods on the `org.hibernate.Hibernate` class. You can also force a proxy or collection to become fully populated by calling the `initialize()` method on this class.

Querying Objects

Hibernate provides several different ways to query for objects stored in the database. The Criteria Query API is a Java API for constructing a query as an object. HQL is an object-oriented query language, similar to SQL, that you may use to retrieve objects that match the query. Hibernate provides a way to execute SQL directly against the database to retrieve objects—if you have legacy applications that use SQL or if you need to use SQL features that are not supported through HQL and the Criteria Query API (discussed in Chapter 11).

Summary

Hibernate provides a simple API for creating, retrieving, updating, and deleting objects from a relational database through the `Session` interface. Understanding the difference between transient, persistent, and detached objects in Hibernate will allow you to understand how changes to the objects update database tables.

We have touched upon the need to create mappings to correlate the database tables with the fields and properties of the Java objects that you want to persist. The next chapter covers these in detail, and discusses why they are required and what they can contain.

An Overview of Mapping

The purpose of Hibernate is to allow you to treat your database as if it stores Java objects. However, databases in practice do not store objects—they store data in tables and columns. Unfortunately, there is no simple way to correlate the data stored in a database with the data represented by Java objects.

The difference between an object-oriented association and a relational one is fundamental. Consider a simple class to represent a user, and another to represent an e-mail address, as shown in Figure 5-1.

Figure 5-1. *An object-oriented association*

User objects contain fields referring to Email objects. The association has a direction—given a User object, you can determine its associated Email object. For example, consider Listing 5-1.

Listing 5-1. *Acquiring the* Email *Object from the* User *Object*

```
User user = ...
Email email = user.email;
```

The reverse, however, is not true. The natural way to represent this relationship in the database, as illustrated in Figure 5-2, is superficially similar.

Figure 5-2. *A relational association*

Despite that similarity, the direction of the association is effectively reversed. Given an `Email` row, you can immediately determine which `User` row it belongs to in the database; this relationship is mandated by a foreign key constraint. It is possible to reverse the relationship in the database world through suitable use of SQL—another difference.

Given the differences between the two worlds, it is necessary to manually intervene to determine how your Java classes should be represented in database tables.

Why Mapping Cannot Be Automated

It is not immediately obvious why you cannot create simple rules for storing your Java objects in the database so that they can be easily retrieved. For example, the most immediately obvious rule would be that a Java class must correlate to a single table. For example, instances of the `User` class defined in Listing 5-2 could surely be represented by a simple table like the one for a user shown in Figure 5-1.

Listing 5-2. *A Simple* `User` *Class with a Password Field*

```java
public class User {
    String name;
    String password;
}
```

And indeed it could, but some questions present themselves:

- How many rows should you end up with if you save a user twice?

- Are you allowed to save a user without a name?

- Are you allowed to save a user without a password?

When you start to think about classes that refer to other classes, there are additional questions to consider. Have a look at the `Customer` and `Email` classes defined in Listing 5-3.

Listing 5-3. `Customer` *and* `Email` *Classes*

```java
public class Customer {
    int customerId;
    int customerReference;
    String name;
    Email email;
}

public class Email {
    String address;
}
```

Based on this, the following questions arise:

- Is a unique customer identified by their customer ID, or their customer reference?

- Can an e-mail address be used by more than one customer?

- Should the relationship be represented in the `Customer` table?

- Should the relationship be represented in the `Email` table?

- Should the relationship be represented in some third (link) table?

Depending upon the answers to these questions, your database tables could vary considerably. You could take a stab at a reasonable design, such as that given in Figure 5-3, based upon your intuition about likely scenarios in the real world.

Figure 5-3. *Tables in which the customer is identified by* `customerId`. *Here, e-mail address entities can only be used by a single customer, and the relationship is maintained by the* `Email` *table.*

As soon as you take away the context provided by the variable and class names, it becomes much harder to form any useful decision about these classes (see Listing 5-4). It would be an impossible task to design an automated tool that could make this sort of decision.

Listing 5-4. *A Class Identical in Structure to Listing 5-3, but with All Contextual Information Removed*

```
public class Foo {
    int x;
    int y;
    String s;
    Bar bar;
}

public class Bar {
    String a;
}
```

Primary Keys

Most "relational" databases that provide SQL access are prepared to accept tables that have no predefined primary key. Hibernate is not so tolerant—even if your table has been created without a primary key, Hibernate will require you to specify one. This often seems perverse to users who are familiar with SQL and databases, but who are not familiar with ORM tools. As such, we will examine in more depth the problems that arise without a primary key.

Without a primary key, it is impossible to uniquely identify a row in a table. For example, consider Table 5-1.

Table 5-1. *A Table in Which the Rows Cannot Be Uniquely Identified*

User	Age
dminter	35
dminter	40
dminter	55
dminter	40
jlinwood	57

This table clearly contains information about users and their respective ages. However, there are four users with the same name (Dave Minter, Denise Minter, Daniel Minter, and Dashiel Minter). There is probably a way of distinguishing them somewhere else in the system—perhaps by an e-mail address or a user number. But if, for example, you want to know the ages of Dashiel Minter with user ID 32, there is no way to obtain it from Table 5-1.

While Hibernate will not let you omit the primary key, it will permit you to form the primary key from a collection of columns. For example, Table 5-2 could be keyed by Usernumber and User.

Table 5-2. *A Table in Which the Rows Can Be Uniquely Identified*

User	Usernumber	Age
dminter	1	35
dminter	2	40
dminter	3	55
dminter	32	42
jlinwood	1	57

Neither User nor Usernumber contains unique entries, but in combination they uniquely identify the age of a particular user, and so they are acceptable to Hibernate as a primary key.

Why does Hibernate need to uniquely identify entries when SQL doesn't? Hibernate is representing Java objects, which are *always* uniquely identifiable. This is why the classic mistake made by new Java developers is to compare strings using the == operator instead of the equals() method. You can distinguish between references to two String objects that represent the same text and two references to the same String object. SQL has no such obligation, and there are arguably cases in which it is desirable to give up the ability to make the distinction.

If Hibernate could not uniquely identify an object with a primary key, then the following code could have several possible outcomes in the underlying table.

```
String customer = getCustomerFromHibernate("dcminter");
customer.setAge(10);
saveCustomerToHibernate(customer);
```

For example, let's say the table originally contained the data shown in Table 5-3.

Table 5-3. *Updating an Ambiguous Table*

User	Age
dcminter	30
dcminter	42

Which of the following should be contained in the resulting table?

- A single row for the user dcminter, with the age set to 10

- Two rows for the user, with both ages set to 10

- Two rows for the user, with one age set to 10 and the other to 42

- Two rows for the user, with one age set to 10 and the other to 30

- Three rows for the user, with one age set to 10 and the others to 30 and 42

In short, the Hibernate developers made a decision to enforce the use of primary keys when creating mappings so that this problem does not arise. Hibernate does provide facilities that will allow you to work around this if it is absolutely necessary (you can create views or stored procedures to "fake" the appropriate key, or you can use conventional JDBC to access the table data), but when using Hibernate, it is always more desirable to work with tables that have correctly specified primary keys if at all possible.

Lazy Loading

When you load classes into memory from the database, you don't necessarily want *all* the information to *actually* be loaded. To take an extreme example, loading a list of e-mails should not cause the full body text and attachments of every e-mail to be loaded into memory. First, they might demand more memory than is actually available. Second, even if they fit, it could take a long time for all of this information to be obtained.

If you were to tackle this problem in SQL, you would probably select a subset of the appropriate fields for the query to obtain the list; for example:

```
SELECT from, to, date, subject FROM email WHERE username = 'dcminter';
```

Hibernate will allow you to fashion queries that are rather similar to this, but it also offers a more flexible approach, known as *lazy loading*. Certain relationships can be marked as being "lazy," and they will not be loaded from disk until they are actually required.

The default in Hibernate 3 is that classes (including collections like Set and Map) should be lazily loaded. For example, when an instance of the User class given in the next listing is loaded from the database, the only fields initialized will be userId and username.

```
public class User {
    int userId;
    String username;
    EmailAddress emailAddress;
    Set roles;
}
```

However, as long as the object is still associated with Hibernate in the appropriate way (see Chapter 9), the appropriate objects for emailAddress and roles will be loaded from the database if they are accessed.

This is the default behavior only; the mapping file can be used to specify which classes and fields should behave in this way.

Associations

When we looked at why the mapping process could not be automated, we discussed the following example classes:

```
public class Customer {
    int customerId;
    int customerReference;
    String name;
    Email email;
}

public class Email {
    String address;
}
```

We also gave the following five questions that it raised:

- Is a unique customer identified by their customer ID, or their customer reference?

- Can a given e-mail address be used by more than one customer?

- Should the relationship be represented in the Customer table?

- Should the relationship be represented in the Email table?

- Should the relationship be represented in some third (link) table?

The first question can be answered simply—it depends on what column you specify as the primary key. The remaining four questions are related, and their answers depend upon the object relationships. Furthermore, if your Customer class represents the relationship with the EmailAddress using a Collection class or an array, it would be possible for a user to have multiple e-mail addresses.

```
public class Customer {
    int customerId;
    int customerReference;
    String name;
    Set email;
}
```

So, you should add another question: can a customer have more than one e-mail address? The set could contain a single entry, so you can't automatically infer that this is the case.

The key questions from the previous options are as follows:

- *Q1*: Can an e-mail address belong to more than one user?

- *Q2*: Can a customer have more than one e-mail address?

The answers to these questions can be formed into a truth table, as in Table 5-4.

Table 5-4. *Deciding the Cardinality of an Entity Relationship*

Q1 Answer	Q2 Answer	Relationship Between Customer and Email
No	No	One-to-one
Yes	No	One-to-many
No	Yes	Many-to-one
Yes	Yes	Many-to-many

These are the four ways in which the cardinality of the relationship between the objects can be expressed. Each relationship can then be represented within the mapping table(s) in various ways.

The One-to-One Association

A one-to-one association between classes can be represented in a variety of ways. At its simplest, the properties of both classes are maintained in the same table. For example, a one-to-one association between a User and an Email class might be represented as a single table, as in Table 5-5.

Table 5-5. *A Combined* User/Email *Table*

ID	Username	Email
1	dcminter	dcminter@example.com
2	jlinwood	jlinwood@example.com
3	tjkitchen	tjkitchen@example.com

The single database entity representing this combination of a User and an Email class is shown in Figure 5-4.

	User
PK	<u>id</u>
	username email

Figure 5-4. *A single entity representing a one-to-one relationship*

Alternatively, the entities can be maintained in distinct tables with identical primary keys, or with a key maintained from one of the entities into the other, as in Tables 5-6 and 5-7.

Table 5-6. *The* User *Table*

ID	Username
1	dcminter
2	jlinwood
3	tjkitchen

Table 5-7. *The* Email *Table*

ID	Username
1	dcminter@example.com
2	jlinwood@example.com
3	tjkitchen@example.com

It is possible to create a mandatory foreign key relationship from one of the entities to the other, but this should not be applied in both directions because a circular dependency would be created. It is also possible to omit the foreign key relationships entirely (as shown in Figure 5-5) and rely upon Hibernate to manage the key selection and assignment.

	User
PK	<u>id</u>
	username

	Email
PK	<u>id</u>
	email

Figure 5-5. *Entities related by primary keys*

If it is not appropriate for the tables to share primary keys, then a foreign key relationship between the two tables can be maintained, with a "unique" constraint applied to the foreign key column. For example, reusing the User table from Table 5-6, the Email table can be suitably populated, as shown in Table 5-8.

Table 5-8. *An* Email *Table with a Foreign Key to the* User *Table*

ID	Email	UserID (Unique)
34	dcminter@example.com	1
35	jlinwood@example.com	2
36	tjkitchen@example.com	3

This has the advantage that the association can easily be changed from one-to-one to many-to-one by removing the unique constraint. Figure 5-6 shows this type of relationship.

Figure 5-6. *Entities related by a foreign key relationship*

The One-to-Many and Many-to-One Association

A one-to-many association (or from the perspective of the other class, a many-to-one association) can most simply be represented by the use of a foreign key, with no additional constraints.

The relationship can also be maintained by the use of a link table. This will maintain a foreign key into each of the associated tables, which will itself form the primary key of the link table. An example of this is shown in Tables 5-9, 5-10, and 5-11.

Table 5-9. *A Simple* User *Table*

ID	Username
1	dcminter
2	jlinwood

Table 5-10. *A Simple* Email *Table*

ID	Email
1	dcminter@example.com
2	dave@example.com
3	jlinwood@example.com
4	jeff@example.com

Table 5-11. *A Link Table Joining* User *and* Email *in a One-to-Many Relationship*

UserID	EmailID
1	1
1	2
2	3
2	4

Additional columns can be added to the link table to maintain information on the ordering of the entities in the association.

A unique constraint must be applied to the "one" side of the relationship (the userId column of the UserEmailLink table in Figure 5-7); otherwise, the link table can represent the set of all possible relationships between User and Email entities—a many-to-many set association.

Figure 5-7. *A relationship represented by a link table (duplicates are not permitted because of the use of a compound primary key)*

The Many-to-Many Association

As noted at the end of the previous section, if a unique constraint is not applied to the "one" end of the relationship when using a link table, it becomes a limited sort of many-to-many relationship. All of the possible combinations of User and Email can be represented, but it is not possible for the same User to have the same e-mail address entity associated twice, because that would require the compound primary key to be duplicated.

If instead of using the foreign keys together as a compound primary key, we give the link table its own primary key (usually a surrogate key), the association between the two entities can be transformed into a full many-to-many relationship, as shown in Table 5-12.

Table 5-12. *A Many-to-Many* User/Email *Link Table*

ID	UserID	EmailID
1	1	1
2	1	2
3	1	3
4	1	4
5	2	1
6	2	2

Table 5-12 might describe a situation in which the user dcminter receives all e-mail sent to any of the four addresses, whereas jlinwood receives only e-mail sent to his own accounts.

When the link table has its own independent primary key, as with the association shown in Figure 5-8, thought should be given to the possibility that a new class should be created to represent the contents of the link table as an entity in its own right.

Figure 5-8. *A many-to-many relationship represented by a link table (duplicates are permitted because of the use of a surrogate key)*

Applying Mappings to Associations

The mappings are applied to express the various different ways of forming associations in the underlying tables—there is no automatically correct way to represent them.

In addition to the basic choice of the approach to take, the mappings are used to specify the minutiae of the tables' representations. While Hibernate tends to use sensible default values when possible, it is often desirable to override these. For example, the foreign key names generated automatically by Hibernate will be effectively random—whereas an informed developer can apply a name (e.g., FK_USER_EMAIL_LINK) to aid in the debugging of constraint violations at run time.

Types of Mapping

At present, Hibernate supports two standard ways to express the mappings.

The technique that has been available the longest is the use of XML mapping files. As the most mature approach, this is currently the best way to control Hibernate, and gives the most sophisticated control over the Hibernate feature set. You have seen examples of simple mapping files in Chapters 1 and 3.

These files can be created directly with a text editor or with the help of various tools created by the Hibernate team and others. We discuss the details of XML mapping files in Chapter 8.

Hibernate now also supports the Annotations feature introduced in Java 5. This permits the use of a special syntax to include metadata directly in the source code for the application. While this allows the core features of Hibernate to be controlled, many of the additional features cannot be specified in annotations. There is therefore something of a trade-off between the advantages of maintaining the mapping information directly within the associated source code, and the more flexible features available from the XML-based mappings. We discuss the details of annotation-based mapping in Chapter 6.

Other Information Represented in Mappings

While Hibernate can determine a lot of sensible default values for the mappings, most of these can be overridden by one or both of the file- and XML-based approaches. Some apply directly to mapping; others, such as the foreign key names, are really only pertinent when the mapping is used to create the database schema. Lastly, some mappings can also provide a place to configure some features that are perhaps not "mappings" in the purest sense. The final sections of this chapter discuss the features that Hibernate supports in addition to those already mentioned.

Specification of (Database) Column Types and Sizes

Java provides the primitive types and allows user declaration of interfaces and classes to extend these. Relational databases generally provide a small subset of "standard" types, and then provide additional proprietary types.

Restricting yourself to the proprietary types will still cause problems, as there are only approximate correspondences between these and the Java primitive types.

A typical example of a problematic type is `java.lang.String` (treated by Hibernate as if it were a primitive type since it is used so frequently), which by default will be mapped to a fixed-size character data database type. Typically, the database would perform poorly if a character field of unlimited size were chosen—but lengthy `String` fields will be truncated as they are persisted into the database.

By overriding the default type mappings, the developer can make appropriate trade-offs between storage space, performance, and fidelity to the original Java representation.

The Mapping of Inheritance Relationships to the Database

There is no SQL standard for representing inheritance relationships for the data in tables; and while some database implementations provide a proprietary syntax for this, not all do. Hibernate provides several configurable ways in which to represent inheritance relationships, and the mapping file permits users to select a suitable approach for their model.

Primary Key

Hibernate demands that a primary key be used to identify entities. The choice of a surrogate key, a key chosen from the business data, and/or a compound primary key can be made via the mapping file.

When a surrogate key is used, Hibernate also permits the key-generation technique to be selected—from a range of techniques that vary in portability and efficiency.

The Use of SQL Formula–Based Properties

It is sometimes desirable that a property of an entity should be maintained not as data directly stored in the database, but rather as a function performed on that data—for example, a subtotal field should not be managed directly by the Java logic, but instead maintained as an aggregate function of some other property.

Mandatory and Unique Constraints

As well as the implicit constraints of a primary or foreign key relationship, you can specify that a field must not be duplicated—for example, a `username` field should often be unique.

Fields can also be made mandatory—for example, requiring a message entity to have both a subject and message text.

The generated database schema will contain corresponding `NOT NULL` and `UNIQUE` constraints so that it is literally impossible to corrupt the table with invalid data (rather, the application logic will throw an exception if any attempt to do so is made).

Note that primary keys are implicitly both mandatory and unique.

Cascading of Operations

As alterations are made to the object model, operations on some objects should cascade through to related objects. For example, deleting a stocked item should perhaps cause any associated catalog entries to be deleted. The reverse—deleting a single catalog entry—should not necessarily cause the stocked item to be deleted from the database entirely!

It would be awkward to manage the appropriate cascading rules from code alone, so cascading rules can be specified at a fine level of detail within the mappings.

Summary

This chapter has given you an overview of the reason why mapping files are needed, and what features they support beyond these absolute requirements. It has discussed the various types of associations, and the circumstances under which you would choose to use them.

The next two chapters look at how mappings are specified using annotations and XML files respectively.

■ ■ ■

Mapping with Annotations

In Chapter 5, we discussed the need to create mappings between the database model and the object model. Mappings can be created as separate XML files, or as Java 5 annotations inline with the source code for your POJOs. In this chapter, we discuss the use of annotations, and in the next chapter, we will discuss the use of XML files.

Java 5 Features

Java 5 was introduced in late 2004 as a major new release of the language. Annotations are not supported by versions of Java prior to this, so while core Hibernate 3 is compatible with earlier versions, you will not be able to take advantage of the features described in this chapter unless your development, compilation, and runtime tools support at least version 5 of the language (version 6 of Java, codenamed Mustang, is expected some time in late 2006).

Since we must perforce assume that you have a Java 5 environment available to you, the examples in this chapter will also take advantage of some of the other enhanced language features introduced in Java 5, as follows:

- Generics

- Enhanced `for` loops

- Static imports

- Enumerations

- Autoboxing

- Variable parameter lists

Using these features will make the source code for this chapter noticeably more compact. Similarly, annotation-based mappings are significantly terser than their XML-based counterparts.

Creating Hibernate Mappings with Annotations

Prior to annotations, the only way to create mappings was through XML files—although tools from Hibernate and third-party projects allowed part or all of these to be generated from Java source code. Although using annotations is the newest way to define mappings, it is not automatically the best way to do so. We will briefly discuss the drawbacks and benefits of annotations before discussing when and how to apply them.

Cons of Annotations

Using annotations immediately restricts your code to a Java 5 environment. This immediately rules out the use of annotations for some developers, as some application servers do not yet support this version of the JVM. Even when there are no technical reasons why a current JVM could not be used, many shops are quite conservative in the deployment of new technologies.

If you are migrating from a Hibernate 2 environment or an existing Hibernate 3 environment, you will already have XML-based mapping files to support your code base. All else being equal, you will not want to re-express these mappings using annotations just for the sake of it.

If you are migrating from a legacy environment, you may not want to alter the preexisting POJO source code, in order to avoid contaminating known-good code with possible bugs.

If you do not have the source code to your POJOs (because it has been lost, or because it was generated by an automated tool), you may prefer the use of external XML-based mappings to the decompilation of class files to obtain Java source code for alteration.

Maintaining the mapping information as external XML files allows the mapping information to be changed to reflect business changes or schema alterations without forcing you to rebuild the application as a whole.

Hibernate 3 support for annotation-based mappings is not yet as mature as its support for XML-based mapping files. For example, while you can still make appropriate foreign key relationships for use in schema generation, you cannot generally name the foreign keys used.

Pros of Annotations

Having considered the drawbacks, there are some powerful benefits to contrast against them.

First, and perhaps most persuasively, we find annotations-based mappings to be far more intuitive than their XML-based alternatives, as they are immediately in the source code along with the properties that they are associated with.

Partly as a result of this, annotations are less verbose than their XML equivalents, as evidenced by the contrast between Listings 6-1 and 6-2.

Listing 6-1. *A Minimal Class Mapped Using Annotations*

```
import javax.persistence.Entity;
import javax.persistence.Id;

@Entity
public class Sample {
    @Id
    public Integer id;
    public String name;
}
```

Listing 6-2. *A Minimal Class Mapped Using XML*

```xml
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE
    hibernate-mapping
    PUBLIC
    "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping default-access="field">
    <class name="Sample">
        <id type="int" column="id">
            <generator class="native"/>
        </id>
        <property name="name" type="string"/>
    </class>
</hibernate-mapping>
```

Some of this verbosity is in the nature of XML itself (the tag names and the boilerplate document type declaration), and some of it is due to the closer integration of annotations with the source code. Here, for example, the XML file must explicitly declare that field access is used in place of property access (i.e., the fields are accessed directly rather than through their get/set methods), but the annotation infers this from the fact that it has been applied to the id field rather than the getId() method.

Hibernate uses and supports the EJB 3 persistence annotations. If you elect not to use Hibernate-specific features in your code and annotations, you will have the freedom to deploy your entities to environments using other ORM tools that support EJB 3.

Finally—and perhaps a minor point—because the annotations are compiled directly into the appropriate class files, there is less risk of a missing or stale mapping file causing problems at deployment (this point will perhaps prove most persuasive to those who already have some experience with this hazard of the XML technique).

Choosing Which to Use

When you are creating a Hibernate application that has complete or primary ownership of its database, and that is a new project, we would generally recommend the use of annotations.

If you intend to make your application portable to other EJB 3–compliant ORM applications, you *must* use annotations to represent the mapping information. Hibernate 3 XML file–based mapping is a proprietary format. However, you may lose this benefit if you rely upon any of the Hibernate 3–specific annotations (that is to say, annotations taken from the org.hibernate package tree rather than the javax.persistence package tree).

If you are migrating an existing application to Hibernate, or creating a new project reliant upon a database primarily owned by other applications, you should use the greater flexibility of XML-based mappings to ensure that your project will not be unduly inconvenienced by changes to the database schema.

Using Annotations in Your Application

You will need to install the Hibernate 3 annotations toolset, available from the Hibernate annotations page (`http://annotations.hibernate.org`). If you do not already use JDK 5.0, you will need to upgrade to take advantage of the native support for annotations.

■**Tip** If you want to declare your mappings inline with your source code, but cannot use a Java 5 environment, the XDoclet tool allows you to use javadoc-style comments to achieve a similar effect. XDoclet can be obtained from `http://xdoclet.sourceforge.net`.

Your application needs the `hibernate-annotations.jar` and `ejb3-persistence.jar` files provided in the annotations toolset.

If you are using a `hibernate.cfg.xml` file to establish the mapping configuration, you will need to provide the fully qualified name of the annotated class with the `<mapping>` element:

```
<mapping class="com.hibernatebook.annotations.Book"/>
```

When you are configuring the `SessionFactory`, you will need to make use of an `AnnotationConfiguration` object instead of the `Configuration` object used with XML mappings, as follows:

```
SessionFactory factory =
   new AnnotationConfiguration().configure().buildSessionFactory();
```

If you prefer to configure the mappings manually rather than through the `hibernate.cfg.xml` file, you can do this through the `AnnotationConfiguration` object, as follows:

```
AnnotationConfiguration config = new AnnotationConfiguration();
config.addAnnotatedClass(Book.class);
SessionFactory factory = config.configure().buildSessionFactory();
```

If you need to use your annotated entities from within an EJB 3 container, you must use the standard `EntityManager` instead of the Hibernate-specific `Session`. Hibernate provides an implementation of `EntityManager` as a separate download. At the time of writing, this is still a beta version, but as it closely follows the state of the EJB 3 specification, you should have little trouble migrating code from the current implementation over to any final release (or to third-party `EntityManagers`). See Appendix A for details of how to use the Hibernate `EntityManager`.

EJB 3 Persistence Annotations

In Chapter 3, we walked you through the creation of a very simple application using the basic XML mapping files—annotations might have been simpler, but as already noted, only Java 5 users would be able to use an annotation-based example.

When you develop using annotations, you start with a Java class, and then annotate the source code listing with metadata notations. In J2SE 5.0, the Java Runtime Environment (JRE) parses these annotations. Hibernate uses Java reflection to read the annotations and apply the

mapping information. If you want to use the Hibernate tools to generate your database schema, you must compile your entity classes containing their annotations first.

The full set of persistence annotations available in the EJB 3 API is listed in Table 6-1. In this section, we are going to introduce the significant core of these annotations alongside a simple set of classes to illustrate how they are applied.

Table 6-1. *The EJB 3 Annotations*

Attribute Name	Target	Purpose
AttributeOverride/ AttributeOverrides	T, M, and F	Overrides the default column details of embedded (component) entities.
Basic	M and F	Overrides the default fetch strategy and nullability of basic fields and properties.
Column	M and F	Associates a field or property of the class with a column in the mapped table.
ColumnResult	Pm	Used as a parameter of the @SqlResultSetMapping annotation; permits the fields of an entity to be returned as columns in a conventional JDBC ResultSet.
DiscriminatorColumn	T	Overrides the default behavior of the discriminator column in single or joined table inheritance strategies.
DiscriminatorValue	T	Determines the value associated with the entity in the discriminator column of the root of the entity's inheritance hierarchy.
Embeddable	T	Marks an entity as being an embeddable (component) entity.
Embedded	M and F	Marks a field or property as consisting of an embedded (component) entity.
EmbeddedId	M and F	Marks a primary key field as consisting of an embedded (component) entity. This is mutually exclusive with the @Id annotation.
Entity	T	Identifies an entity and allows attributes, such as its name, to be overridden from the defaults.
EntityListeners	T	Allows appropriate javax.persistence. EntityListener classes to be invoked during the life cycle of the marked entity.
EntityResult	Pm	Used as a parameter of the @SqlResultSetMapping annotation; permits the fields of an entity to be returned as columns in a conventional JDBC ResultSet.
Enumerated	M and F	Defines a field or property as being an enumerated type.
ExcludeDefaultListeners	T	Prevents the default EntityListeners from being invoked during the life cycle of the marked entity.
ExcludeSuperclassListeners	T	Prevents the EntityListeners of the superclass from being invoked during the life cycle of the marked entity.

Continued

Table 6-1. *Continued*

Attribute Name	Target	Purpose
FieldResult	Pm	Used as a parameter of the @SqlResultSetMapping annotation, permits the fields of an entity to be returned as columns in a conventional JDBC ResultSet.
GeneratedValue	M and F	Allows generation strategies to be specified for the marked entity's primary key value(s).
Id	M and F	Identifies the primary key of the entity. Placement of the @Id attribute also determines whether the default access mode for the entity class is field or property access.
IdClass	T	Applied to indicate that an entity's primary key is represented with columns corresponding to the fields of another entity. The appropriate fields forming the primary key will be marked with the @Id attribute.
Inheritance	T	Marks an entity as being the root of an entity inheritance hierarchy (i.e., the highest persistent class in the class inheritance hierarchy).
JoinColumn/JoinColumns	T, M, and F	Defines the column(s) being used as a foreign key into another table.
JoinTable	M and F	Allows the details of the link table to be specified in a one-to-many or many-to-many relationship.
Lob	M and F	Marks a field or property as being stored as a large object data type—typically a binary large object (BLOB). This can be used to remove the length limitations on strings and binary data, but usually implies reduced scope for querying the data so marked.
ManyToMany	M and F	Allows a many-to-many association to be defined between entities.
ManyToOne	M and F	Allows a many-to-one association to be defined between entities.
MapKey	M and F	Allows a key to be specified when making an association with a Map object.
MappedSuperclass	T	Allows a non-persistence class to be used as the basis of the mapping information for its derived classes.
NamedNativeQuery/ NamedNativeQueries	Pk and T	Allows a named SQL query to be stored in the annotations.
NamedQuery/NamedQueries	Pk and T	Allows a named EJB QL query to be stored in the annotations.
OneToMany	M and F	Allows a one-to-many association to be defined between entities.
OneToOne	M and F	Allows a one-to-one association to be defined between entities.
OrderBy	MF	Allows the ordering of a collection to be defined as it is retrieved.

Attribute Name	Target	Purpose
PersistenceContext/ PersistenceContexts	T, M, and F	For use with @EntityManager; marks a field or property as representing the EntityManager to be injected by the container.
PersistenceUnit/ PersistenceUnits	T, M, and F	For use with @EntityManager; marks a field or property as representing the EntityManagerFactory to be injected by the container.
PostLoad	M	Marks a method for invocation after performing a load operation on the entity.
PostPersist	M	Marks a method for invocation after performing a persist operation on the entity.
PostRemove	M	Marks a method for invocation after performing a remove operation on the entity.
PostUpdate	M	Marks a method for invocation after performing an update operation on the entity.
PrePersist	M	Marks a method for invocation prior to performing a persist operation on the entity.
PreRemove	M	Marks a method for invocation prior to performing a remove operation on the entity.
PreUpdate	M	Marks a method for invocation prior to performing an update operation on the entity.
PrimaryKeyJoinColumn/ PrimaryKeyJoinColumns	T, M, and F	Allows the columns joining a secondary table to a primary table to be specified.
QueryHint	Pm	Allows implementation-specific "hints" to be provided as a parameter of named queries and named native queries.
SecondaryTable/ SecondaryTables	T	Allows an entity's basic fields and properties to be persisted to more than one table.
SequenceGenerator	Pk, T, M, and F	Allows a named primary key generator to be defined for use by one or more entities.
SqlResultSetMapping	Pk and T	Allows an entity to be mapped as if it were a named native query (i.e., so that it can be retrieved as a conventional JDBC ResultSet).
Table T	Pk, T, M, and F	Allows the default details of an entity's primary table to be overridden.
TableGenerator	Pk, T, M, and F	Overrides the default properties of the table used to generate primary keys when the table generation strategy of the generated value annotation is used on the primary key field or property (the field or property marked with the @Id annotation).
Temporal	M and F	Specifies the behavior of Date and Calendar fields or properties (if omitted, such fields will be treated as TIMESTAMP values).
Transient	M and F	Allows a field or property to be marked so that it will not be persisted.
UniqueConstraint	Pm	Enforces a unique constraint at schema generation time as a parameter of @Table.
Version	M and F	Marks the field or property serving as the optimistic lock value of the entity.

Key to the Target column: Pk = package, T = type, M = method, F = field, Pm = parameter

The set of example classes represents a publisher's catalog of books. You'll start with a single class, Book, which has no annotations or mapping information. For this example's purposes, you do not have an existing database schema to work with, so you need to define your relational database schema as you go.

At the beginning of the example, the Book class is very simple. It has two fields, title and pages; and an identifier, id, which is an integer. The title is a String object, and pages is an integer. As we go through this example, we will add annotations, fields, and methods to the Book class. The complete source code listing for the Book and Author classes is given at the end of this chapter—the source files for the rest are available in the source code download for this chapter on the Apress web site (www.apress.com).

Listing 6-3 gives the source code of the Book class, in its unannotated form, as a starting point for the example.

Listing 6-3. *The* Book *Class, Unannotated*

```
package com.hibernatebook.annotations;

public class Book {

    private String title;
    private int pages;
    private int id;

    // Getters...

    public int getId() {
        return id;
    }

    public String getTitle() {
        return title;
    }

    public int getPages() {
        return pages;
    }

    // Setters...

    public void setId(int id) {
        this.id = id;
    }

    public void setTitle(String title) {
        this.title = title;
    }
```

```
    public void setPages(int pages) {
        this.pages = pages;
    }
}
```

As you can see, this is a POJO. We are going to annotate this class as we go along, explaining the concepts behind annotation.

Entity Beans with @Entity

The first step is to annotate the Book class as an EJB 3 entity bean. With traditional EJB, we would have added an EJB marker interface to mark the class as an entity bean. Instead, we add the @Entity annotation to the Book class, as follows:

```
package com.hibernatebook.annotations;
import javax.persistence.*;

@Entity
public class Book
```

The EJB 3 standard annotations are contained in the javax.persistence package, so we import the appropriate annotations (here we will use wildcard imports to keep the listings short, but in the downloadable source code accompanying this chapter, we use explicit imports such as import javax.persistence.Entity;—annotations are imported in exactly the same way as the ordinary interfaces that they resemble).

The @Entity annotation marks this class as an entity bean, so it must have a no-argument constructor that is visible with at least protected scope. Hibernate supports package scope as the minimum, but you lose portability to other EJB 3 containers if you take advantage of this. Other EJB 3 rules for an entity bean class are that the class must not be final, and that the entity bean class must be concrete. Many of the rules for EJB 3 entity bean classes and Hibernate 3 persistent objects are the same—partly because the Hibernate team had much input into the EJB 3 design process, and partly because there are only so many ways to design a relatively unobtrusive object-relational persistence solution.

As you can see, although we did have to add the import statement and the annotations, we have not had to change the rest of the code. The POJO is essentially unchanged.

Primary Keys with @Id and @GeneratedValue

Each entity bean has to have a primary key, which you annotate on the class with the @Id annotation. Typically, the primary key will be a single field, though it can also be a composite of multiple fields.

The placement of the @Id annotation determines the default access strategy that Hibernate will use for the mapping. If the annotation is applied to a field as shown in Listing 6-4, then field access will be used.

Listing 6-4. *A Class with Field Access*

```java
import javax.persistence.*;

@Entity
public class Sample {
    @Id
    int id;

    public int getId() {
        return this.id;
    }

    public void setId(int id) {
        this.id = id;
    }
}
```

If instead the annotation is applied to the getter for the field, as shown in Listing 6-5, then property access will be used.

Listing 6-5. *The Same Class with Property Access*

```java
import javax.persistence.*;

@Entity
public class Sample {
    int id;

    @Id
    public int getId() {
        return this.id;
    }

    public void setId(int id) {
        this.id = id;
    }
}
```

Here you can see one of the strengths of the annotations approach—because the annotations are placed inline with the source code, information can be extracted from the context of the mapping in the code, allowing many mapping decisions to be inferred rather than stated explicitly—which helps to further reduce the verbosity of the annotations.

By default, the @Id annotation will automatically determine the most appropriate primary key generation strategy to use—you can override this by also applying the @GeneratedValue annotation. This takes a pair of attributes: strategy and generator. The strategy attribute must be a value from the javax.persistence.GeneratorType enumeration. If you do not specify a generator type, the default is AUTO. There are four different types of primary key generators on GeneratorType, as follows:

- AUTO: Hibernate decides which generator type to use, based on the database's support for primary key generation.

- IDENTITY: The database is responsible for determining and assigning the next primary key.

- SEQUENCE: Some databases support a SEQUENCE column type. See the "Generating Primary Key Values with @SequenceGenerator" section later in the chapter.

- TABLE: This type keeps a separate table with the primary key values. See the "Generating Primary Key Values with @TableGenerator" section later in the chapter.

You will notice that the available values for the strategy attribute do not exactly match the values for Hibernate's primary key generators for XML mapping. If you need to use Hibernate-specific primary key generation strategies, you can use some of the Hibernate extensions described at the end of this chapter—but as always, you risk forfeiting portability of your application to other EJB 3 environments when taking advantage of Hibernate-specific features.

For the Book class, we are going to use the default key generation strategy. Letting Hibernate determine which generator type to use makes your code portable between different databases. Because we want Hibernate to use property access to our POJO, we must annotate the getter method for the identifier, not the field that it accesses:

```
@Id
@GeneratedValue
 public int getId() {
     return id;
 }
```

Generating Primary Key Values with @SequenceGenerator

As noted in the section on the @Id tag, we can declare the primary key property as being generated by a database sequence. A sequence is a database object that can be used as a source of primary key values. It is similar to the use of an identity column type, except that a sequence is independent of any particular table and can therefore be used by multiple tables.

To declare the specific sequence object to use and its properties, you must include an @SequenceGenerator annotation on the annotated field. Here's an example:

```
@Id
@SequenceGenerator(name="seq1",sequenceName="HIB_SEQ")
@GeneratedValue(strategy=SEQUENCE,generator="seq1")
public int getId() {
    return id;
}
```

Here, a sequence generation annotation named seq1 has been declared. This refers to the database sequence object called HIB_SEQ. The name seq1 is then referenced as the generator attribute of the @GeneratedValue annotation.

Only the sequence generator name is mandatory—the other attributes will take sensible default values—but you should provide an explicit value for the `sequenceName` attribute as a matter of good practice anyway. If not specified, the `sequenceName` value to be used is selected by the <mark>persistence provider</mark> (i.e., Hibernate or EJB 3). The other (optional) attributes are `initialValue` and `allocationSize`, which default to values of 1 and 50, respectively.

Generating Primary Key Values with @TableGenerator

The `@TableGenerator` annotation is used in a very similar way to the `@SequenceGenerator` annotation—but because `@TableGenerator` manipulates a standard database table to obtain its primary key values, instead of using a vendor-specific sequence object, it is guaranteed to be portable between database platforms.

■**Note** For optimal portability *and* optimal performance, you should not specify the use of a table generator, but instead use the <mark>`@GeneratorValue(strategy=GeneratorType.AUTO)`</mark> configuration, which allows the persistence provider to select the most appropriate strategy for the database in use.

As with the sequence generator, the name attributes of `@TableGenerator` are mandatory and the other attributes are optional, with the table details being selected by the persistence provider.

```
@Id
@TableGenerator(name="tablegen",
                table="ID_TABLE",
                pkColumnName="ID",
                valueColumnName="NEXT_ID")
@GeneratedValue(strategy=TABLE,generator="tablegen")
public int getId() {
    return id;
}
```

The optional attributes are as follows:

- `allocationSize`: Allows the increment on the primary key value to be specified.

- `catalog`: Allows the catalog that the table resides within to be specified.

- `initialValue`: Allows the starting primary key value to be specified.

- `pkColumnName`: Allows the primary key column of the table to be identified. The table can contain the details necessary for generating primary key values for multiple entities.

- `pkColumnValue`: Allows the primary key for the row containing the primary key generation information to be identified.

- `schema`: Allows the schema that the table resides within to be specified.

- `table`: The name of the table containing the primary key values.

- `uniqueConstraints`: Allows additional constraints to be applied to the table for schema generation.

- `valueColumnName`: Allows the column containing the primary key generation information for the current entity to be identified.

Because the table can be used to contain the primary key values for a variety of entries, it is likely to contain a single row for each of the entities using it. It therefore needs its own primary key (`pkColumnName`), as well as a column containing the next primary key value to be used (`pkColumnValue`) for any of the entities obtaining their primary keys from it.

Compound Primary Keys with @Id, @IdClass, or @EmbeddedId

While the use of single column surrogate keys is advantageous for various reasons, you may sometimes be forced to work with business keys. When these are contained in a single column, you can use `@Id` without specifying a generation strategy (forcing the user to assign a primary key value before the entity can be persisted). However, when the primary key consists of multiple columns, you need to take a different strategy to group these together in a way that allows the persistence engine to manipulate the key values as a single object.

You must create a class to represent this primary key. It will not require a primary key of its own, of course, but it must be a public class, must have a default constructor, must be serializable, and must implement `hashCode()` and `equals()` methods to allow the Hibernate code to test for primary key collisions (i.e., they must be implemented with the appropriate database semantics for the primary key values).

Your three strategies for using this primary key class once it has been created are as follows:

- Mark it as `@Embeddable` and add to your entity class a normal property for it, marked with `@Id`.

- Add to your entity class a normal property for it, marked with `@EmbeddableId`.

- Add properties to your entity class for all of its fields, mark them with `@Id`, and mark your entity class with `@IdClass`, supplying the class of your primary key class.

All these techniques require the use of an `id` class because Hibernate must be supplied with a primary key object when various parts of its persistence API are invoked. For example, you can retrieve an instance of an entity by invoking the `Session` object's `get()` method, which takes as its parameter a single serializable object representing the entity's primary key.

The use of `@Id` with a class marked as `@Embeddable`, as shown in Listing 6-6, is the most natural approach. The `@Embeddable` tag can be used for non–primary key embeddable values anyway (`@Embeddable` is discussed in more detail later in the chapter). It allows you to treat the compound primary key as a single property, and it permits the reuse of the `@Embeddable` class in other tables.

Listing 6-6. *Using the* @Id *and* @Embeddable *Annotations to Map a Compound Primary Key*

```java
package com.hibernatebook.annotations;

import javax.persistence.*;

@Entity
public class Account {
   private String description;
   private AccountPk id;

   public Account (String description) {
      this.description = description;
   }

   protected Account() {
   }

   @Id
   public AccountPk getId() {
      return this.id;
   }

   public String getDescription() {
      return this.description;
   }

   public void setId(AccountPk id) {
      this.id = id;
   }

   public void setDescription(String description) {
      this.description = description;
   }

   @Embeddable
   public static class AccountPk {
      private String code;
      private Integer number;

      public AccountPk() {
      }

      public String getCode() {
         return this.code;
      }
```

```
    public Integer getNumber() {
       return this.number;
    }

    public void setNumber(Integer number) {
       this.number = number;
    }

    public void setCode(String code) {
       this.code = code;
    }


    public int hashCode() {
       int hashCode = 0;
       if( code != null ) hashCode ^= code.hashCode();
       if( number != null ) hashCode ^= number.hashCode();
       return hashCode;
    }

    public boolean equals(Object obj) {
       if( !(obj instanceof AccountPk) ) return false;
       AccountPk target = (AccountPk)obj;
       return ((this.code == null) ?
              (target.code == null) :
                 this.code.equals(target.code))
           && ((this.number == null) ?
              (target.number == null) :
                 this.number.equals(target.number));
    }
  }
}
```

The next most natural approach is the use of the @EmbeddedId tag. Here, the primary key class cannot be used in other tables since it is not an @Embeddable entity, but it does allow us to treat the key as a single attribute of the Account class (in Listings 6-7 and 6-8, the implementation of AccountPk is identical to that in Listing 6-6, and is thus omitted for brevity). Note that in Listings 6-7 and 6-8, the AccountPk class is *not* marked as @Embeddable.

Listing 6-7. *Using the* @EmbeddedId *Annotation to Map a Compound Primary Key*

```
package com.hibernatebook.annotations;

import javax.persistence.*;

@Entity
public class Account {
   private String description;
   private AccountPk id;
```

```java
    public Account(String description) {
        this.description = description;
    }

    protected Account() {
    }

    @EmbeddedId
    public AccountPk getId() {
        return this.id;
    }

    public String getDescription() {
        return this.description;
    }

    public void setId(AccountPk id) {
        this.id = id;
    }

    public void setDescription(String description) {
        this.description = description;
    }

    public static class AccountPk {
        // ...
    }
}
```

Finally, the use of the @IdClass and @Id annotations allows us to map the compound primary key class using properties of the entity itself corresponding to the names of the properties in the primary key class. The names must correspond (there is no mechanism for overriding this), and the primary key class must honor the same obligations as with the other two techniques. The only advantage to this approach is its ability to "hide" the use of the primary key class from the interface of the enclosing entity. The @IdClass annotation takes a value parameter of Class type, which must be the class to be used as the compound primary key. The fields that correspond to the properties of the primary key class to be used must all be annotated with @Id—note in Listing 6-8 that the getCode() and getNumber() methods of the Account class are so annotated, and the AccountPk class is not mapped as @Embeddable, but it is supplied as the value of the @IdClass annotation.

Listing 6-8. *Using the* @IdClass *and* @Id *Annotations to Map a Compound Primary Key*

```java
package com.hibernatebook.annotations;

import javax.persistence.*;

@Entity
@IdClass(Account.AccountPk.class)
public class Account {
   private String description;
   private String code;
   private Integer number;

   public Account(String description) {
      this.description = description;
   }

   protected Account() {
   }

   @Id
   public String getCode() {
      return this.code;
   }

   @Id
   public Integer getNumber() {
      return this.number;
   }

   public String getDescription() {
      return this.description;
   }

   public void setDescription(String description) {
      this.description = description;
   }

   public void setNumber(Integer number) {
      this.number = number;
   }

   public void setCode(String code) {
      this.code = code;
   }
```

```
    public static class AccountPk {
        // ...
    }
}
```

Regardless of which of these approaches we take to declare our compound primary key, the table that will be used to represent it will require the same set of columns. Listing 6-9 shows the DDL that will be generated from any of Listings 6-6, 6-7, or 6-8.

Listing 6-9. *The DDL Generated from the Annotated* Account *Class (Regardless of the Approach Used)*

```
create table Account (
    code varchar(255) not null,
    number integer not null,
    description varchar(255),
    primary key (code, number)
);
```

Database Table Mapping with @Table and @SecondaryTable

The @Table annotation allows you to specify the details of the table that will be used to persist the entity in the database. If you omit the annotation, Hibernate will default to using the class name for the table name, so you only need to provide this annotation if you want to override that behavior. The @Table annotation provides four attributes, allowing you to override the name of the table, its catalog, and its schema, and enforce unique constraints on columns in the table. Typically, you would only provide a substitute table name thus: @Table(name=" ORDER_HISTORY"). The unique constraints will be applied if the database schema is generated from the annotated classes, and will supplement any column-specific constraints (see discussions of @Column and @JoinColumn later in this chapter). They are not otherwise enforced.

The @SecondaryTable annotation provides a way to model an entity bean that is persisted across several different database tables. Here, in addition to providing an @Table annotation for the primary database table, your entity bean can have an @SecondaryTable annotation, or an @SecondaryTables annotation in turn containing zero or more @SecondaryTable annotations. The @SecondaryTable annotation takes the same basic attributes as the @Table annotation, with the addition of the join attribute. The join attribute defines the join column for the primary database table. It accepts an array of javax.persistence.PrimaryKeyJoinColumn objects. If you omit the join attribute, then it will be assumed that the tables are joined on identically named primary key columns.

When an attribute in the entity is drawn from the secondary table, it must be marked with the @Column annotation, with a table attribute identifying the appropriate table. Listing 6-10 shows how a property of the Customer entity could be drawn from a second table mapped in this way.

Listing 6-10. *An Example of a Field Access Entity Mapped Across Two Tables*

```
package com.hibernatebook.annotations;

import javax.persistence.*;

@Entity
@Table(name="CUSTOMER")
@SecondaryTable(name="CUSTOMER_DETAILS")
public class Customer {

    @Id
    public int id;

    public String name;

    @Column(table="CUSTOMER_DETAILS")
    public String address;
}
```

Columns in the primary or secondary tables can be marked as having unique values within their tables by adding one or more appropriate @UniqueConstraint annotations to @Table or @SecondaryTable's uniqueConstraints attribute. For example, to mark the name field in the preceding declaration as being unique, use the following:

```
@Entity
@Table(
    name="CUSTOMER",
    uniqueConstraints={@UniqueConstraint(columnNames="name")}
)
@SecondaryTable(name="CUSTOMER_DETAILS")
public class Customer {
    ...
}
```

Persisting Basic Types with @Basic

By default, properties and instance variables in your POJO are persistent—Hibernate will store their values for you. The simplest mappings are therefore for the "basic" types. These include primitives, primitive wrappers, arrays of primitives or wrappers, enumerations, and any types that implement Serializable but are not themselves mapped entities. These are all mapped implicitly—no annotation is needed. By default, such fields are mapped to a single column, and eager fetching is used to retrieve them (i.e., when the entity is retrieved from the database, all the basic fields and properties are retrieved). Also, when the field or property is not a primitive, it can be stored and retrieved as a null value.

This default behavior can be overridden by applying the @Basic annotation to the appropriate class member. This annotation takes two optional attributes, and is itself entirely

optional. The first attribute is named `optional` and takes a `boolean`. Defaulting to `true`, this can be set to `false` to provide a hint to schema generation that the associated column should be created `NOT NULL`. The second is named `fetch` and takes a member of the enumeration `FetchType`. This is `EAGER` by default, but can be set to `LAZY` to permit loading on access of the value.

The use of lazy loading is unlikely to be valuable, except when large serializable objects have been mapped as basic types (rather than given entity mappings of their own) and retrieval time may become significant. While the (default) `EAGER` value must be honored, the `LAZY` flag is considered to be a hint, and can be ignored by the persistence engine.

The `@Basic` attribute is usually omitted, with the `@Column` attribute being used where the `@Basic` annotation's optional attribute might otherwise be used to provide the `NOT NULL` behavior.

Omitting Persistence with @Transient

Some fields may be used at run time only, and should be discarded from objects as they are persisted into the database. The EJB 3 specification provides the `@Transient` annotation for these transient fields. The `@Transient` annotation does not have any attributes—you just add it to the instance variable or the getter method as appropriate for the entity bean's property access strategy.

For our example, we contrive to add a `Date` field named `publicationDate`, which will not be stored in the database to our `Book` class. We mark this field transient thus:

```
@Transient
public Date getPublicationDate() {
    return publicationDate;
}
```

Because we are using a property access strategy for our `Book` class, we must put the `@Transient` annotation on the getter method.

Mapping Properties and Fields with @Column

The `@Column` annotation is used to specify the details of the column to which a field or property will be mapped. Some of the details are schema related, and therefore apply only if the schema is generated from the annotated files. Others apply and are enforced at run time by Hibernate (or the EJB 3 persistence engine). It is optional, with an appropriate set of default behaviors, but is often useful when overriding default behavior, or when you need to fit your object model into a preexisting schema. It is more commonly used than the similar `@Basic` annotation, with the following attributes commonly being overridden:

name permits the name of the column to be explicitly specified—by default, this would be the name of the property. However, it is often necessary to override the default behavior when it would otherwise result in an SQL keyword being used as the column name (e.g., user).

length permits the size of the column used to map a value (particularly a `String` value) to be explicitly defined. The column size defaults to 255, which might otherwise result in truncated `String` data, for example.

`nullable` permits the column to be marked `NOT NULL` when the schema is generated. The default is that fields should be permitted to be null; however, it is common to override this when a field is, or ought to be, mandatory.

`unique` permits the column to be marked as containing only unique values. This defaults to `false`, but commonly would be set for a value that might not be a primary key but would still cause problems if duplicated (such as `username`).

We have marked up the title field of our `Book` entity using the `@Column` entity to show how three of these attributes would be applied:

```
@Column(name="working_title",length=200,nullable=false)
public String getTitle() {
    return title;
}
```

The remaining attributes, less commonly used, are as follows:

`table` is used when the owning entity has been mapped across one or more secondary tables. By default, the value is assumed to be drawn from the primary table, but the name of one of the secondary tables can be substituted here (see the `@SecondaryTable` annotation example earlier in this chapter).

`insertable` defaults to `true`, but if set to `false`, the annotated field will be omitted from insert statements generated by Hibernate (i.e., it won't be persisted).

`updatable` defaults to `true`, but if set to `false`, the annotated field will be omitted from update statements generated by Hibernate (i.e., it won't be altered once it has been persisted).

`columnDefinition` can be set to an appropriate DDL fragment to be used when generating the column in the database. This can only be used during schema generation from the annotated entity, and should be avoided if possible, since it is likely to reduce the portability of your application between database dialects.

`precision` permits the precision of decimal numeric columns to be specified for schema generation, and will be ignored when a non-decimal value is persisted. The value given represents the number of digits in the number (usually requiring a minimum length of n+1, where n is the scale).

`scale` permits the scale of decimal numeric columns to be specified for schema generation and will be ignored where a non-decimal value is persisted. The value given represents the number of places after the decimal point.

Modeling Entity Relationships

Naturally, annotations also allow you to model associations between entities. EJB 3 supports one-to-one, one-to-many, many-to-one, and many-to-many associations. Each of these has its corresponding annotation.

We discussed the various ways in which these mappings can be established in the tables in Chapter 5. In this section, we will show how the various mappings are requested using the annotations.

Mapping an Embedded (Component) One-to-One Association

When all the fields of one entity are maintained within the same table as another, the enclosed entity is referred to in Hibernate as a *component*. The EJB 3 standard refers to such an entity as being *embedded*.

The @Embedded and @Embeddable attributes are used to manage this relationship. In this book's database example, we associate an AuthorAddress class with an Author class in this way.

The AuthorAddress class is marked with the @Embeddable annotation. An embeddable entity must be composed entirely of basic fields and attributes. An embeddable entity can only use the @Basic, @Column, @Lob, @Temporal, and @Enumerated annotations. It cannot maintain its own primary key with the Id tag because its primary key is the primary key of the enclosing entity.

The @Embeddable annotation itself is purely a marker annotation, and takes no additional attributes, as demonstrated in Listing 6-11. Typically, the fields and properties of the embeddable entity need no further markup.

Listing 6-11. *Marking an Entity for Embedding Within Other Entities*

```
@Embeddable
public class AuthorAddress {
...
}
```

The enclosing entity then marks appropriate fields or getters in entities, making use of the embeddable class with the @Embedded annotation, as shown in Listing 6-12.

Listing 6-12. *Marking an Embedded Property*

```
@Embedded
public AuthorAddress getAddress() {
    return this.address;
}
```

The @Embedded annotation draws its column information from the embedded type, but permits the overriding of a specific column or columns with the @AttributeOverride and @AttributeOverrides tags (the latter to enclose an array of the former if multiple columns are being overridden). For example, Listing 6-13 shows how to override the default column names of the address and country attributes of AuthorAddress with columns named ADDR and NATION.

Listing 6-13. *Overriding Default Attributes of an Embedded Property*

```
@Embedded
@AttributeOverrides({
    @AttributeOverride(name="address",column=@Column(name="ADDR")),
    @AttributeOverride(name="country",column=@Column(name="NATION"))
})
public AuthorAddress getAddress() {
    return this.address;
}
```

Neither Hibernate nor the EJB 3 standard supports mapping an embedded object across more than one table. In practice, if you want this sort of persistence for your embedded entity, you will usually be better off making it a first-class entity (i.e., *not* embedded) with its own `@Entity` marker and `@Id` annotations, and then mapping it via a conventional one-to-one association, as explained in the next section.

Mapping a Conventional One-to-One Association

There is nothing intrinsically wrong with mapping a one-to-one association between two entities where one is not a component of (i.e., embedded into) the other. The relationship is often somewhat suspect, however. You should give some thought to using the embedded technique described previously before using the `@OneToOne` annotation.

Assuming that you are resolute on declaring the association in this way (perhaps because you anticipate converting it to a one-to-many or many-to-one relationship in the foreseeable future), applying the annotation is quite simple—all of the attributes are optional. Listing 6-14 shows how simply a relationship like this might be declared.

Listing 6-14. *Declaring a Simple One-to-One Relationship*

```
@OneToOne
public Address getAddress() {
    return this.address;
}
```

The `@OneToOne` annotation permits the following optional attributes to be specified:

`targetEntity` can be set to the class of an entity storing the association. If left unset, the appropriate type will be inferred from the field type, or the return type of the property's getter.

`cascade` can be set to any of the members of the `javax.persistence.CascadeType` enumeration. It defaults to none being set. See the "Cascading Operations" sidebar for a discussion of these values.

`fetch` can be set to the `EAGER` or `LAZY` members of `FetchType`.

`optional` indicates whether the value being mapped can be null.

`mappedBy` indicates that a bidirectional one-to-one relationship is owned by the named entity.[1] The owning entity contains the primary key of the subordinate entity.

Mapping a Many-to-One or One-to-Many Association

A many-to-one association and a one-to-many association are the same association seen from the perspective of the owning and subordinate entities, respectively.

---

1. An association is bidirectional if each entity maintains a property or field representing its end of the same relationship. For example, if our `Address` class maintained a reference to the `Publisher` located there, and the `Publisher` class maintained a reference to its `Address`, then the association would be bidirectional.

CASCADING OPERATIONS

When an association between two entities is established (for example, a one-to-one association between Human and Pet), it is common to want certain persistence operations on one entity to also be applied to the entity that it is linked to. Take, for example, the following code:

```
Human dave = new Human("dave");
Pet cat = new PetCat("Tibbles");
dave.setPet(cat);
session.save(dave);
```

In the last line, highlighted in bold, we are likely to want to save the Pet object associated with the Human object. In a one-to-one relationship, we usually expect all operations on the owning entity to be propagated through—that is, to be *cascaded*—to the dependent entity. In other associations this is not true, and even in a one-to-one relationship we may have special reasons for wanting to spare the dependent entity from delete operations (perhaps for auditing reasons).

We are therefore able to specify the types of operations that should be cascaded through an association to another entity using the cascade annotation, which takes an array of members of the CascadeType enumeration. The members correspond with the names of the key methods of the EntityManager class used for EJB 3 persistence, and have the following rough correspondence with operations on entities:

- ALL requires all operations to be cascaded to dependent entities. This is the same as including MERGE, PERSIST, REFRESH, and REMOVE.

- MERGE cascades updates to the entity's state in the database (i.e., UPDATE . . .).

- PERSIST cascades the initial storing of the entity's state in the database (i.e., INSERT. . .).

- REFRESH cascades the updating of the entity's state from the database (i.e., SELECT . . .).

- REMOVE cascades deletion of the entity from the database (i.e., DELETE . . .).

- If no cascade type is specified, no operations will be cascaded through the association.

In the light of these options, the appropriate annotation for the relationship between a publisher and its address would be as follows:

```
@OneToOne(cascade=CascadeType.ALL)
public Address getAddress() {
    return this.address;
}
```

The simplest way to maintain a many-to-one relationship between two entities is by managing the foreign key of the entity at the "one" end of the one-to-many relationship as a column in the "many" entity's table.

The @OneToMany annotation can be applied to a field or property value for a collection or an array representing the mapped "many" end of the association.

COLLECTION ORDERING

EJB 3 does not provide for maintaining the ordering of an ordered collection such as an array or a list—see the "Ordering Collections with `@IndexColumn`" section of the chapter for discussion on how the `@IndexColumn` annotation can be used in a non-portable way to remedy this deficiency. The collection, however, can be specified in terms of the fields of the associated entity at retrieval time by means of the `@OrderBy` annotation. For example, if we were to retrieve a list ordered by the books' names in ascending order, we could annotate a suitable method.

The following code snippet specifies a retrieval order for an ordered collection.

```
@OneToMany(cascade = ALL, mappedBy = "publisher"
@OrderBy("name ASC")
public List<Book> getBooks() {
    return books
}
```

The value of the `@OrderBy` annotation is an ordered list of the field names to sort by, each one optionally appended with `ASC` (for ascending order, as in the preceding code) or `DESC` (for descending order). If neither `ASC` nor `DESC` is appended to one of the field names, the order will default to ascending. `@OrderBy` can be applied to any collection-valued association.

The `mappedBy` attribute is mandatory on a bidirectional association and optional (being implicit) on a unidirectional association.

`cascade` is optional, taking a member of the `javax.persistence.CascadeType` enumeration and dictating the cascade behavior of the mapped entity.

`targetEntity` is optional, as it can usually be deduced from the type of the field or property, as in Listing 6-15, where the property represents a `Set` of `Book` entities, making the target entity implicitly `Book`. However, if necessary (if generics are not being used, for example), the class of the target entity can be provided here.

`fetch` is optional, allowing lazy or eager fetching to be specified as a member of the `javax.persistence.FetchType` enumeration.

Listing 6-15. *Mapping a One-to-Many Relationship from the* Book *Entity to the* Publisher *Entity*

```
@OneToMany(cascade = ALL,mappedBy = "publisher")
public Set<Book> getBooks() {
    return books;
}
```

The many-to-one end of this relationship is expressed in similar terms to the one-to-many end, as shown in Listing 6-16.

Listing 6-16. *Mapping a Many-to-One Relationship from the* Publisher *Entity to the* Book *Entity*

```
@ManyToOne
@JoinColumn(name = "publisher_id")
public Publisher getPublisher() {
    return publisher;
}
```

The @ManyToOne annotation takes a similar set of attributes to @OneToMany. The following list describes the attributes, all of which are optional.

cascade indicates the appropriate cascade policy for operations on the association; it defaults to none.

fetch indicates the fetch strategy to use; it defaults to LAZY.

optional indicates whether the value can be null; it defaults to true.

targetEntity indicates the entity that stores the primary key—this is normally inferred from the type of the field or property (Publisher in the preceding example).

We have also supplied the optional @JoinColumn attribute to name the foreign key column required by the association something other than the default (publisher)—this is not necessary, but it illustrates the use of the annotation.

When a unidirectional one-to-many association is to be formed, it is possible to express the relationship using a link table. This is achieved by adding the @JoinTable annotation as shown in Listing 6-17.[2]

Listing 6-17. *A Simple Unidirectional One-to-Many Association with a Join Table*

```
@OneToMany(cascade = ALL)
@JoinTable
public Set<Book> getBooks() {
    return books;
}
```

The @JoinTable annotation provides attributes that allow various aspects of the link table to be controlled. These attributes are as follows:

name is the name of the join table to be used to represent the association.

catalog is the name of the catalog containing the join table.

schema is the name of the schema containing the join table.

joinColumns is an array of @JoinColumn attributes representing the primary key of the entity at the "one" end of the association.

inverseJoinColumns is an array of @JoinColumn attributes representing the primary key of the entity at the "many" end of the association.

---

2. When a join table is being used, the foreign key relationship is maintained within the join table itself—it is therefore not appropriate to combine the mappedBy attribute of the @OneToMany annotation with the use of an @JoinTable annotation.

Listing 6-18 shows a fairly typical application of the @JoinTable annotation to specify the name of the join table and its foreign keys into the associated entities.

Listing 6-18. *A Unidirectional One-to-Many Association with a More Fully Specified Join Table*

```
@OneToMany(cascade = ALL)
@JoinTable(
        name="PublishedBooks",
        joinColumns = { @JoinColumn( name = "publisher_id") },
        inverseJoinColumns = @JoinColumn( name = "book_id")
)
public Set<Book> getBooks() {
    return books;
}
```

Mapping a Many-to-Many Association

When a many-to-many association does not involve a first-class entity joining the two sides of the relationship, a link table must be used to maintain the relationship. This can be generated automatically, or the details can be established in much the same way as with the link table described in the "Mapping a Many-to-One or One-to-Many Association" section of the chapter.

The appropriate annotation is naturally @ManyToMany, and takes the following attributes:

mappedBy is the field that owns the relationship—this is only required if the association is bidirectional. If an entity provides this attribute, then the other end of the association is the owner of the association, and the attribute must name a field or property of that entity.

targetEntity is the entity class that is the target of the association. Again, this may be inferred from the generic or array declaration, and only needs to be specified if this is not possible.

cascade indicates the cascade behavior of the association, which defaults to none.

fetch indicates the fetch behavior of the association, which defaults to LAZY.

The example maintains a many-to-many association between the Book class and the Author class. The Book entity owns the association, so its getAuthors() method must be marked with an appropriate @ManyToMany attribute, as shown in Listing 6-19.

Listing 6-19. *The Book Side of the Many-to-Many Association*

```
@ManyToMany(cascade = ALL)
public Set<Author> getAuthors() {
    return authors;
}
```

The Author entity is managed by the Book entity. The link table is not explicitly managed, so, as shown in Listing 6-20, we mark it with a @ManyToMany annotation and indicate that the foreign key is managed by the authors attribute of the associated Book entity.

Listing 6-20. *The Author Side of the Many-to-Many Association*

```
@ManyToMany(mappedBy = "authors")
public Set<Book> getBooks() {
    return books;
}
```

Alternatively, we could specify the link table in full, as in Listing 6-21.

Listing 6-21. *Specifying the Link Table in Full Using the* Book *Entity Annotations*

```
@ManyToMany(cascade = ALL)
@JoinTable(
        name="Books_to_Author",
        joinColumns={@JoinColumn(name="book_ident")},
        inverseJoinColumns={@JoinColumn(name="author_ident")}
)
public Set<Author> getAuthors() {
    return authors;
}
```

Inheritance

The EJB 3 standard and Hibernate both support three approaches to mapping inheritance hierarchies into the database. These are as follows:

- Single table (`SINGLE_TABLE`)

- Joined (`JOINED`)

- Table-per-class (`TABLE_PER_CLASS`)

Persistent entities that are related by inheritance must be marked up with the `@Inheritance` annotation. This takes a single strategy attribute, which is set to one of three `javax.persistence.InheritanceType` enumeration values corresponding to these approaches (shown in brackets in the preceding bulleted list).

The single table approach manages one class for the superclass and all its subtypes. There are columns for each mapped field or property of the superclass, and for each distinct field or property of the derived types. When following this strategy, you will need to ensure that columns are appropriately renamed when any field or property names collide in the hierarchy.

To determine the appropriate type to instantiate when retrieving entities from the database, an `@DiscriminatorColumn` annotation should be provided in the root (and only in the root) of the persistent hierarchy.[3] This defines a column containing a value that distinguishes between each of the types used. The attributes permitted by the `@DiscriminatorColumn` annotation are as follows:

---

3. That is to say, the highest class in the hierarchy that is mapped to the database as an entity should be annotated in this way.

name is the name of the discriminator column.

discriminatorType is the type of value to be stored in the column as selected from the javax.persistence.DiscriminatorType enumeration of STRING, CHAR, or INTEGER.

columnDefinition is a fragment of DDL defining the column type. Using this is liable to reduce the portability of your code across databases.

length is the column length of STRING discriminator types. It is ignored for CHAR and INTEGER types.

All of these (and the annotation itself) are optional, but we recommend supplying at least the name attribute. If no @DiscriminatorColumn is specified in the hierarchy, a default column name of DTYPE and type of STRING will be used.

Hibernate will supply an appropriate discriminator value for each of your entities. For example, if the STRING discriminator type is used, the value this column contains will be the name of the entity (which defaults to the class name). You can also override this behavior with specific values using the @DiscriminatorValue annotation. If the discriminator type is INTEGER, any value provided via the @DiscriminatorValue annotation must be convertible directly into an integer.

In Listing 6-22, we specify that an INTEGER discriminator type should be stored in the column named DISCRIMINATOR. Rows representing Book entities will have a value of 1 in this column, whereas the following mapping in Listing 6-23 requires that rows representing ComputerBook entities should have a value of 2 in the same column.

Listing 6-22. *The Root of the Inheritance Hierarchy Mapped with the* SINGLE_TABLE *Strategy*

```
@Entity
@Inheritance(strategy = SINGLE_TABLE)
@DiscriminatorColumn(
    name="DISCRIMINATOR",
    discriminatorType=INTEGER
)
@DiscriminatorValue("1")
public class Book {
...
}
```

Listing 6-23. *A Derived Entity in the Inheritance Hierarchy*

```
@Entity
@DiscriminatorValue("2")
public class ComputerBook extends Book {
...
}
```

An alternative to the monolithic single table approach is the otherwise similar joined table approach. Here a discriminator column is used, but the fields of the various derived types are stored in distinct tables. Other than the differing strategy, this inheritance type is specified in the same way (as shown in Listing 6-24).

Listing 6-24. *The Root of the Inheritance Hierarchy Mapped with the* JOINED *Strategy*

```
@Entity
@Inheritance(strategy = JOINED)
@DiscriminatorColumn(
    name="DISCRIMINATOR"
)
public class Book {
...
}
```

Finally, there is the table-per-class approach, in which all of the fields of each type in the inheritance hierarchy are stored in distinct tables. Because of the close correspondence between the entity and its table, the @DiscriminatorColumn annotation is not applicable to this inheritance strategy. Listing 6-25 shows how our Book class could be mapped in this way.

Listing 6-25. *The Root of the Inheritance Hierarchy Mapped with the* TABLE_PER_CLASS *Strategy*

```
@Entity
@Inheritance(strategy = TABLE_PER_CLASS)
public class Book {
...
}
```

Other EJB 3 Persistence Annotations

Although we have now covered most of the core EJB 3 persistence annotations, there are a few others that you will encounter fairly frequently. We cover some of these in passing in the following sections.

Temporal Data

Fields or properties of an entity that have java.util.Date or java.util.Calendar types represent temporal data. By default, these will be stored in a column with the TIMESTAMP data type, but this default behavior can be overridden with the @Temporal annotation.

The annotation accepts a single value attribute from the javax.persistence. TemporalType enumeration. This offers three possible values: DATE, TIME, and TIMESTAMP. These correspond respectively to java.sql.Date, java.sql.Time, and java.sql.Timestamp. The table column is given the appropriate data type at schema generation time. Listing 6-26 shows an example mapping a java.util.Date property as a TIME type—the java.sql.Date and java.sql.Time classes are both derived from the java.util.Date class, so confusingly, both are capable of representing dates *and* times!

Listing 6-26. *A* Date *Property Mapped as a* Time *Temporal Field*

```
@Temporal(TIME)
public java.util.Date getStartingTime() {
    return this.startingTime;
}
```

Large Objects

A persistent property or field can be marked for persistence as a database-supported large object type by applying the `@Lob` annotation.

The annotation takes no attributes, but the underlying large object type to be used will be inferred from the type of the field or parameter. String- and character-based types will be stored in an appropriate character-based type. All other objects will be stored in a BLOB. Listing 6-27 maps a `String` into a large object column type.

Listing 6-27. *An Example of a Large Object Property*

```
@Lob
public String getTitle() {
    return this.title;
}
```

The `@Lob` annotation can be used in combination with the `@Basic` annotation.

Mapped Superclasses

A special case of inheritance occurs when the root of the hierarchy is not itself a persistent entity, but various classes derived from it are. Such a class can be abstract or concrete. The `@MappedSuperclass` annotation allows you to take advantage of this circumstance.

The class marked with `@MappedSuperclass` is not an entity, and is not queryable (it cannot be passed to methods that expect an entity in the `Session` or `EntityManager` objects). It cannot be the target of an association.

The mapping information for the columns of the superclass will be stored in the same table as the details of the derived class (in this way, the annotation resembles the use of the an `@Inheritance` tag with the `SINGLE_TABLE` strategy).

In other respects, the superclass can be mapped as a normal entity, but the mappings will apply to the derived classes only (since the superclass itself does not have an associated table in the database). When a derived class needs to deviate from the superclass's behavior, the `@AttributeOverride` annotation can be used (much as with the use of an embeddable entity).

For example, if in our example model `Book` was a superclass of `ComputerBook`, but `Book` objects themselves were never persisted directly, then `Book` could be marked as `@MappedSuperclass`, as in Listing 6-28.

Listing 6-28. *Marking the* Book *Class As a Mapped Superclass*

```
@Entity
@MappedSuperclass
public class Book {
...
}
```

The fields of the `ComputerBook` entity derived from `Book` would then be stored in the `ComputerBook` entity class's table. Classes derived directly from `Book` but not mapped as entities in their own right, such as a hypothetical `MarketingBook` class, would not be persistable. In this respect alone, the mapped superclass approach behaves differently from the conventional `@Inheritance` approach with a `SINGLE_TABLE` strategy.

Named Queries (HQL or EJB QL)

`@NamedQuery` and `@NamedQueries` allow one or more EJB QL queries to be associated with an entity. The required attributes are as follows:

`name` is the name by which the query is retrieved.

`query` is the EJB QL (or HQL) query associated with the name.

Listing 6-29 shows an example associating a named query with the `Author` entity. The query would retrieve `Author` entities by name, so it is natural to associate it with that entity—however, there is no actual requirement that a named query be associated in this way with the entity that it concerns.

Listing 6-29. *An EJB QL Named Query Annotation*

```
@Entity
@NamedQuery(
        name="findAuthorsByName",
        query="from Author where name = :author"
)
public class Author {
...
}
```

There is also a `hints` attribute, taking a `QueryHint` annotation name/value pair, which allows caching mode, timeout value, and a variety of other platform-specific tweaks to be applied (this can also be used to comment the SQL generated by the query).

You do not need to directly associate the query with the entity against which it is declared, but it is normal to do so. If a query has no natural association with any of the entity declarations, it is possible to make the `@NamedQuery` annotation at the package level.

There is no natural place to put a package-level annotation, so Java annotations allow for a specific file, called `package-info.java`, to contain them. Listing 6-30 gives an example of this.

Listing 6-30. *A* `package-info.java` *File*

```
@javax.annotations.NamedQuery(
    name="findAuthorsByName",
    query="from Author where name = :author"
)
package com.hibernatebook.annotations;
```

Hibernate's session allows named queries to be accessed directly, as shown in Listing 6-31.

Listing 6-31. *Invoking a Named Query via the Session*

```
Query query = session.getNamedQuery("findAuthorsByName");
query.setParameter("author", "Dave");
List booksByDave = query.list();
System.out.println("There is/are " + booksByDave.size()
        + " author(s) called Dave in the catalog");
```

If you have multiple @NamedQuery annotations to apply to an entity, they can be provided as an array of values of the @NamedQueries annotation.

Named Native Queries (SQL)

EJB 3 also allows the database's native query language (usually a dialect of SQL) to be used in place of EJB QL. You risk losing portability here if you use a database-specific feature, but as long as you use reasonably generic SQL, you should be OK. The @NamedNativeQuery annotation is declared in almost exactly the same manner as the @NamedQuery annotation. The following block of code shows a simple example of the use of a named native query.

```
@NamedNativeQuery(
  name="nativeFindAuthorNames",
  query="select name from author"
)
```

Multiple @NamedNativeQuery annotations can be grouped with the @NamedNativeQueries annotation.

■**Note** Hibernate does not currently fully support named native queries.

Configuring the Annotated Classes

Once you have an annotated class, you will need to provide the class to your application's Hibernate configuration, just as if it were an XML mapping. With annotations, you can use either the declarative configuration in the hibernate.cfg.xml XML configuration document, or you can programmatically add annotated classes to Hibernate's org.hibernate.cfg.AnnotationConfiguration object. Your application may use both annotated entities and XML mapped entities in the same configuration.

To provide declarative mapping, we use a normal hibernate.cfg.xml XML configuration file and add the annotated classes to the mapping using the mapping element (see Listing 6-32). Notice that we have specified the name of the annotated class as a mapping.

Listing 6-32. *A Hibernate XML Configuration File with an Annotated Class*

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <property name="connection.driver_class">
            org.hsqldb.jdbcDriver
        </property>
```

```
            <property name="connection.url">
                jdbc:hsqldb:file:annotationsdb;shutdown=true
            </property>
            <property name="connection.username">sa</property>
            <property name="connection.password"></property>
            <property name="hibernate.connection.pool_size">0</property>
            <property name="show_sql">false</property>
            <property name="dialect">
                org.hibernate.dialect.HSQLDialect
            </property>

            <!-- Mapping files -->
            <mapping class="com.hibernatebook.annotations.Author"/>
            <mapping class="com.hibernatebook.annotations.AuthorAddress"/>
            <mapping class="com.hibernatebook.annotations.Book"/>
            <mapping class="com.hibernatebook.annotations.Address"/>
            <mapping class="com.hibernatebook.annotations.Publisher"/>
            <mapping class="com.hibernatebook.annotations.ComputerBook"/>

    </session-factory>
</hibernate-configuration>
```

You can also add an annotated class to your Hibernate configuration programmatically. The annotations toolset comes with an `org.hibernate.cfg.AnnotationConfiguration` object that extends the base Hibernate `Configuration` object for adding mappings. The methods on `AnnotationConfiguration` for adding annotated classes to the configuration are as follows:

```
addAnnotatedClass(Class persistentClass) throws MappingException
addAnnotatedClasses(List<Class> classes)
addPackage(String packageName) throws MappingException
```

Using these methods, you can add one annotated class, a list of annotated classes, or an entire package (by name) of annotated classes. As with the Hibernate XML configuration file, the annotated entities are interoperable with XML mapped entities.

Hibernate 3–Specific Persistence Annotations

Table 6-2 lists the Hibernate-specific annotations. We will now look at some of these Hibernate-specific annotations in more detail—however, bear in mind that using any of this functionality potentially reduces the portability of your application to other EJB 3 solutions.

Annotations that are not recognized by an EJB 3 environment will be ignored, rather than causing a runtime exception directly—however, this may result in different runtime application behavior that may not be desirable. In some cases, Hibernate 3 annotations can be used to prepare resources that are referenced by standard EJB 3 annotations, in which case the application will fail when the EJB 3 environment attempts to use the missing resource.

Tip It is possible to overstate the importance of portability—most bespoke applications are never deployed to an environment other than the one for which they were originally developed. As a mature product, Hibernate 3 has numerous features to offer above and beyond the base EJB 3 specification. You should not waste too much time trying to achieve a portable solution in preference to these proprietary features unless you have a definite requirement for portability.

Table 6-2. *The Hibernate Annotations*

Attribute Name	Target	Purpose
AccessType	T, M, and F	Allows the default access type (normally determined by the placement of the @javax.persistence.Id annotation) for the annotated object to be overridden.
BatchSize	T, M, and F	Allows the batch size for a query to be specified.
Cache	T, M, and F	Allows a cache concurrency strategy (NONE, NONSTRICT_READ_WRITE, READ_ONLY, READ_WRITE, or TRANSACTIONAL) to be selected.
Cascade	M and F	Applies a Hibernate-specific cascade strategy to an association.
Check	T, M, and F	Allows an arbitrary SQL constraint to be specified during the schema generation.
CollectionOfElements	M and F	Allows a collection field or an attribute to be marked as a collection of elements or embedded objects, rather than a full-fledged association with an entity.
Columns	M and F	Allows an array of @javax.persistence.Column annotations to be applied to an annotated Hibernate composite user type.
DiscriminatorFormula	T	Allows the discriminator type to be determined with an HQL formula instead of the default EJB 3 mechanisms.
Entity	T	Allows Hibernate-specific attributes to be applied in addition to the javax.persistence.Entity annotation information.
Filter/Filters	T, M, and F	Adds named filters to an entity or a collection.
FilterDef/FilterDefs	Pk and T	Allows named filters to be declared.
Formula	MF	Allows an SQL formula to be used in place of values drawn from a column.
GenericGenerator	Pk, T, M, and F	Allows a Hibernate-specific generator to be used when creating a primary key value for the entity.
Index	M and F	Allows a database index to be defined for a column or columns.
IndexColumn	M and F	Allows a collection to maintain an order on the basis of an index column maintaining the element ordering (i.e., collection ordering rather than database ordering).

Continued

Table 6-2. *Continued*

Attribute Name	Target	Purpose
NamedNativeQuery/ NamedNativeQueries	Pk and T	Extends the corresponding EJB 3 named native query functionality with various Hibernate-specific query hints.
NamedQuery/NamedQueries	Pk and T	Extends the corresponding EJB 3 named query functionality with various Hibernate-specific query hints.
NotFound	M and F	Allows the behavior to be defined for circumstances in which an expected entity is missing. The options drawn from the NotFoundAction enumeration are the self-explanatory EXCEPTION and IGNORE values. The default is EXCEPTION.
OnDelete	T, M, and F	Allows Hibernate-specific behavior on deletion of collections, arrays, and joined subclasses.
OrderBy	M and F	Allows a collection to be ordered by SQL rather than HQL (as with the EJB 3 annotation) ordering.
ParamDef	Pm	Used to define parameters for Filter annotations.
Parameter	Pm	Used to declare parameters for GenericGenerator annotations.
Proxy	T	Allows the proxy behavior for an entity to be configured or disabled.
Sort	M and F	Allows a collection to be sorted using a comparator.
Table/Tables	T	Allows indexes to be applied to a table (see the "Applying Indexes with @Table and @Index" section later in the chapter).
Type	M and F	Marks a field or an attribute as being a composite user type.
TypeDef/TypeDefs	Pk and T	Allows a composite user type to be defined.
Where	T, M, and F	Applies a Where clause to an entity or association.

Key to the Target column: Pk = package, T = type, M = method, F = field, Pm = parameter

All the annotations and enumerations described here fall into the org.hibernate.annotations package. When we refer to an EJB 3 annotation or enumeration, we will use the fully qualified javax.persistence.* class name.

@Entity

The Hibernate-specific @Entity annotation extends the basic details of the @javax.persistence.Entity annotation, but is otherwise used in the same contexts. It allows the following additional attributes to be specified:

dynamicInsert is used to flag that insert statements should be generated at run time (not at startup), allowing only the altered columns to be inserted. By default this is disabled.

dynamicUpdate is used to flag that update statements should be generated at run time, allowing only the altered columns to be updated. By default this is disabled.

`mutable` is `true` by default, but if set to `false`, it allows the persistence engine to cache the values read from the database, and the persistence engine will make no attempt to update them in response to changes (changes that should not be made if this flag is set to `false`).

`optimisticLock` allows an optimistic lock strategy to be selected from the `OptimisticLockType` enumeration values of `ALL`, `DIRTY`, `NONE`, and `VERSION`. This defaults to `VERSION`.

`persister` allows a persister class other than the default Hibernate one to be selected for the entity (for example, allowing serialization to be used instead of relational persistence).

`polymorphism` allows the polymorphism strategy to be selected from the `PolymorphismType` enumeration values of `EXPLICIT` and `IMPLICIT`. This defaults to `IMPLICIT`.

`selectBeforeUpdate` allows the user to request that a `SELECT` be performed to retrieve the entity before any potential update.

Sorting Collections with @Sort

The Hibernate-specific `@Sort` annotation allows a collection managed by Hibernate to be sorted by a standard Java comparator. The following code gives an example.

```
@javac.persistence.OneToMany
@org.hibernate.annotations.Sort(
    type=org.hibernate.annotations.SortType.COMPARATOR,
    comparator=EmployeeComparator.class
)
public Set<Employee> getEmployees() {
    return this.employees;
}
```

Ordering Collections with @IndexColumn

While `@Sort` allows data to be sorted once it has been retrieved from the database, Hibernate also provides a non-standard persistence feature that allows the ordering of appropriate collection types such as `List` to be maintained within the database by maintaining an index column to represent that order. Here's an example:

```
@javax.persistence.OneToMany
@org.hibernate.annotations.IndexColumn(
    name="employeeNumber"
)
public List<Employee> getEmployees() {
    return this.employees;
}
```

Here, we are declaring that an `employeeNumber` column will maintain a value, starting at `0` and incrementing as each entry is added to the list. The default starting value can be overridden by the `base` attribute. By default, the column can contain null (unordered) values. This can be overridden by setting the `nullable` attribute to `false`. By default, when the schema is generated

from the annotations, the column is assumed to be an integer type, but this can be overridden by supplying a columnDefinition attribute specifying a different column definition string.

Applying Indexes with @Table and @Index

The Hibernate-specific @Table annotation supplements the standard table annotation and allows additional index hints to be provided to Hibernate. These will be used at schema generation time to apply indexes to the columns specified. The following code gives an example.

```
// Standard persistence annotations:
@javax.persistence.Entity
@javax.persistence.Table(name="FOO")

// Hibernate-specific table annotation:
@Table(
  appliesTo="FOO", indexes = {
      @Index(name="FOO_FROM_TO_IDX",columnNames={"FIRST","LAST"}),
      @Index(name="FOO_EMPLOYEE_IDX",columnNames={"EMPLOYEE_NUM"}))
public class Foo {
...
}
```

Restricting Collections with @Where

The contents of a collection that will be retrieved from the database can be restricted with a Hibernate-specific @Where annotation. This simply adds a Where clause to the query that will be used to obtain the entities contained within the collection. Here's an example:

```
@javax.persistence.OneToMany
@org.hibernate.annotations.Where(clause="grade > 2")
public Set<Employee> getEmployees() {
    return this.employees;
}
```

Alternative Key Generation Strategies with @GenericGenerator

As mentioned in the "Primary Keys with @Id and @GeneratedValue" section, the full gamut of Hibernate primary key value generators is not supported by the standard set of annotations. Hibernate therefore supplies the @GenericGenerator annotation to fill the void.

The attributes that can be supplied to the annotation are as follows:

name is mandatory, and is used to identify the generic generator in the @GeneratedValue annotation.

strategy is mandatory, and determines the generator type to be used. This can be a standard Hibernate generator type or the name of a class implementing the org.hibernate.id.IdentifierGenerator interface.

parameters is a list of @Parameter annotations defining any parameter values required by the generator strategy.

The available standard Hibernate strategies are `increment`, `identity`, `sequence`, `hilo`, `seqhilo`, `uuid`, `guid`, `native`, `assigned`, `select`, and `foreign`. For example, the non-standard `uuid` strategy for a primary key is configured as follows:

```
@Id
@GenericGenerator(name="unique_id",strategy="uuid")
@GeneratedValue(generator="unique_id")
public String getId() {
    return this.id;
}
```

Alternatively, to configure the `sequence` strategy (equivalent to specifying a strategy of `SEQUENCE` in the `@GeneratedValue` annotation), you can supply the following parameters:

```
@Id
@GenericGenerator(name="seq_id",strategy="sequence",
    parameters= {
        @Parameter(name="sequence",value="HIB_SEQ")
    }
)
@GeneratedValue(generator="seq_id")
public Integer getId() {
    return this.id;
}
```

Using Ant with Annotation-Based Mappings

When using the Hibernate Ant tasks in conjunction with the annotation-based mappings, you operate under one important constraint: the Ant task cannot read the mapping information from the raw source files. The annotated files must be compiled before you can perform any operations on them (including schema generation). You should therefore ensure that any Hibernate Ant tasks are granted a dependency upon the compile task for the entities.

The Ant task will also need access to the classes via the configuration object—you will therefore need to explicitly include any annotated classes in the `hibernate.cfg.xml` file as described in the first part of the previous "Configuring the Annotated Classes" section. You cannot use programmatic configuration of the classes in conjunction with tasks such as `hbm2ddl`, so this is an important step.

The various Hibernate JAR files, including `hibernate-annotations.jar`, will need to be in the classpath of the task definition.

Finally, you will need to specify an `<annotationconfiguration .../>` element, rather than the usual `<configuration .../>` element. An example Ant target to build a DDL script from annotated classes is shown in Listing 6-33.

Listing 6-33. *A Sample Excerpt from this Chapter's Task to Perform Schema Generation*

```
<target name="exportDDL" depends="compile">
    <mkdir dir="${sql}"/>
    <htools destdir="${sql}">
        <classpath refid="classpath.tools"/>
        <annotationconfiguration
           configurationfile="${src}/hibernate.cfg.xml"/>
        <hbm2ddl
           create="true"
           drop="true"
           format="true"
           export="true"
           outputfilename="${ant.project.name}.dll"/>
    </htools>
</target>
```

A full Ant script is provided with the online source code for this chapter (at www.apress.com).

Code Listings

Listings 6-34 and 6-35 contain the completed annotated source code for the Author and Book classes described in this chapter. The database schema also follows in Listing 6-36.

Listing 6-34 illustrates use of the @Entity, @Inheritance, @Id, @GeneratedValue, @Column, @Transient, @ManyToOne, @JoinColumn, and @ManyToMany annotations.

Listing 6-34. *The Fully Annotated* Book *Class*

```
package com.hibernatebook.annotations;

import static javax.persistence.CascadeType.ALL;
import static javax.persistence.InheritanceType.JOINED;

import java.util.*;
import javax.persistence.*;

@Entity
@Inheritance(strategy = JOINED)
public class Book {

    private String title;
    private Publisher publisher;
    private Set<Author> authors = new HashSet<Author>();
    private int pages;
    private int id;
    protected Date publicationDate;
```

```java
// Constructors...

protected Book() {
}

public Book(String title, int pages) {
    this.title = title;
    this.pages = pages;
}

// Getters...

@Id
@GeneratedValue
public int getId() {
    return id;
}

@Column(name = "working_title", length = 200, nullable = false)
public String getTitle() {
    return title;
}

public int getPages() {
    return pages;
}

@Transient
public Date getPublicationDate() {
    return publicationDate;
}

@ManyToOne
@JoinColumn(name = "publisher_id")
public Publisher getPublisher() {
    return publisher;
}

@ManyToMany(cascade = ALL)
public Set<Author> getAuthors() {
    return authors;
}

// Setters...
```

```
    public void setId(int id) {
        this.id = id;
    }

    public void setTitle(String title) {
        this.title = title;
    }

    public void setPages(int pages) {
        this.pages = pages;
    }

    public void setPublicationDate(Date publicationDate) {
        this.publicationDate = publicationDate;
    }

    public void setPublisher(Publisher publisher) {
        this.publisher = publisher;
    }

    public void setAuthors(Set<Author> authors) {
        this.authors = authors;
    }

    // Helpers...
    public void addAuthor(Author author) {
        authors.add(author);
    }
}
```

Listing 6-35 demonstrates the use of the `@NamedQuery`, `@Embedded`, `@AttributeOverrides`, and `@AttributeOverride` annotations.

Listing 6-35. *The Fully Annotated* Author *Class*

```
package com.hibernatebook.annotations;

import java.util.HashSet;
import java.util.Set;

import javax.persistence.*;

@Entity
@NamedQuery(
  name="findAuthorsByName",
  query="from Author where name = :author"
)
public class Author {
```

```java
private int id;
private String name;
private Set<Book> books = new HashSet<Book>();
private AuthorAddress address;

// Constructors...

protected Author() {
}

public Author(String name, AuthorAddress address) {
    this.name = name;
    this.address = address;
}

// Getters...

@Id
@GeneratedValue
public int getId() {
    return id;
}

public String getName() {
    return name;
}

@ManyToMany(mappedBy = "authors")
public Set<Book> getBooks() {
    return books;
}

@Embedded
@AttributeOverrides({
    @AttributeOverride(name="address",column=@Column(name="ADDR")),
    @AttributeOverride(name="country",column=@Column(name="NATION"))
})
public AuthorAddress getAddress() {
    return this.address;
}

// Setters...

public void setId(int id) {
    this.id = id;
}
```

```
    public void setName(String name) {
        this.name = name;
    }

    public void setBooks(Set<Book> books) {
        this.books = books;
    }

    public void setAddress(AuthorAddress address) {
        this.address = address;
    }
}
```

Finally, Listing 6-36 shows the database schema for the classes supporting this chapter,
as generated by the Ant `hbm2ddl` export task for the HSQL database. You will note that we are
unable to control the names selected for the foreign key relationships. This is one area in
which the Hibernate XML mapping files are superior to the EJB 3 annotations.

Listing 6-36. *The Database Schema for the Example*

```
    create table Address (
        id integer not null,
        city varchar(255),
        country varchar(255),
        primary key (id)
    );

    create table Author (
        id integer generated by default as identity (start with 1),
        ADDR varchar(255),
        NATION varchar(255),
        name varchar(255),
        primary key (id)
    );

    create table Book (
        id integer generated by default as identity (start with 1),
        pages integer not null,
        working_title varchar(200) not null,
        publisher_id integer,
        primary key (id)
    );

    create table Book_Author (
        books_id integer not null,
        authors_id integer not null,
        primary key (books_id, authors_id)
    );
```

```
create table ComputerBook (
    BOOK_ID integer not null,
    softwareName varchar(255),
    primary key (BOOK_ID)
);

create table Publisher (
    id integer generated by default as identity (start with 1),
    name varchar(255),
    address_id integer,
    primary key (id)
);

alter table Book
    add constraint FK1FAF0990BF1C70
    foreign key (publisher_id)
    references Publisher;

alter table Book_Author
    add constraint FK1A9A0FA1B629DD87
    foreign key (authors_id)
    references Author;

alter table Book_Author
    add constraint FK1A9A0FA1D3BA8BC3
    foreign key (books_id)
    references Book;

alter table ComputerBook
    add constraint FK98D97CC4600B1724
    foreign key (BOOK_ID)
    references Book;

alter table Publisher
    add constraint FKCDB7C1DC158ECEF0
    foreign key (address_id)
    references Address;
```

Summary

In this chapter, we used EJB 3 annotations to add metadata to our POJOs for Hibernate, and we looked at the Hibernate-specific annotations that can enhance these at the cost of reduced portability.

In the next chapter, we discuss the alternative approach of using XML mapping documents to express the mapping requirements.

CHAPTER 7

■ ■ ■

Creating Mappings with Hibernate XML Files

In the simple example programs in Chapters 1 and 3, we demonstrated how a mapping file could be used to establish the relationship between the object model and the database schema. A mapping file can map a single class or multiple classes to the database. The mapping can also describe standard queries (in HQL and SQL) and filters.

Hibernate Types

Although we have referred to the Hibernate types in passing, we have not discussed the terminology in any depth. In order to express the behavior of the mapping file elements, we need to make these fine distinctions explicit.

Hibernate types fall into three broad categories: entities, components, and values.

Entities

Generally, an entity is a POJO class that has been mapped into the database using the `<class>` or `<subclass>` elements.

An entity can also be a dynamic map (actually a `Map` of `Maps`). These are mapped against the database in the same way as a POJO, but with the default entity mode of the `SessionFactory` set to `dynamic-map`.

The advantage of POJOs over the `dynamic-map` approach is that compile-time type safety is retained. Conversely, dynamic maps are a quick way to get up and running when building prototypes.

It is also possible to represent your entities as Dom4J `Document` objects. This is a useful feature when importing and exporting data from a preexisting Hibernate database, but it is not really central to the everyday use of Hibernate.

We recommend that you use the standard entity mode unless you need to sacrifice accuracy for timeliness, so the alternate approaches are not discussed in this chapter—however, we give some simple examples of the Dom4J- and `Map`-based mappings in Appendix A.

Components

Lying somewhere between entities and values are component types. When the class representation is simple and its instances have a strong one-to-one relationship with instances of another class, then it is a good candidate to become a component of that other class.

The component will normally be mapped as columns in the same table that represents most of the other attributes of the owning class, so the strength of the relationship must justify this inclusion. In the following code, the MacAddress class might a good candidate for a component relationship.

```
public class NetworkInterface {
    public int id;
    public String name;
    public String manufacturer;
    public MacAddress physicalAddress;
}
```

The advantage of this approach is that it allows you to dispense with the primary key of the component and the join to its containing table. If a poor choice of component is made (for example, when a many-to-one relationship actually holds), then data will be duplicated unnecessarily in the component columns.

Values

Everything that is not an entity or a component is a value. Generally, these correspond to the data types supported by your database, the collection types, and, optionally, some user-defined types.

The details of these mappings will be vendor-specific, so Hibernate provides its own value type names; the Java types are defined in terms of these (see Table 7-1).

Table 7-1. *The Standard Hibernate 3 Value Names*

Hibernate 3 Type	Corresponding Java Type
Primitives and Wrappers	
integer	int, java.lang.Integer
long	long, java.lang.Long
short	short, java.lang.Short
float	float, java.lang.Float
double	double, java.lang.Double
character	char, java.lang.Character
byte	byte, java.lang.Byte
boolean, yes_no, true_false	boolean, java.lang.Boolean
Other Classes	
string	java.lang.String
date, time, timestamp	java.util.Date
calendar, calendar_date	java.util.Calendar

Hibernate 3 Type	Corresponding Java Type
big_decimal	java.math.BigDecimal
big_integer	java.math.BigInteger
locale	java.util.Locale
timezone	java.util.TimeZone
currency	java.util.Currency
class	java.lang.Class
binary	byte[]
text	java.lang.String
serializable	java.io.Serializable
clob	java.sql.Clob
blob	java.sql.Blob

In addition to these standard types, you can create your own. Your user type class should implement either the `org.hibernate.usertype.UserType` interface or the `org.hibernate.usertype.CompositeUserType` interface. Once implemented, a custom type can behave identically to the standard types; though depending on your requirements, it may be necessary to specify multiple column names to contain its values, or to provide initialization parameters for your implementation.

For one-off cases, we recommend that you use components—these have similar behavior, but they can be "created" in the mapping file without needing to write Hibernate-specific code. Unless you propose to make substantial use of a custom type throughout your application, it will not be worth the effort. We do not discuss this feature further in this book.

The Anatomy of a Mapping File

A mapping file is a normal XML file. It is validated against a DTD, which can be downloaded from `http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd`. You can also look through the annotated version at `http://hibernatebook.com`.

The terminology used in the naming of elements and attributes is somewhat confusing at first because it is the point of contact between the jargon of the object-oriented and relational worlds.

The <hibernate-mapping> Element

The root element of any mapping file is `<hibernate-mapping>`. As the top-level element, its attributes mostly define default behaviors and settings to apply to the child elements (see Table 7-2).

Table 7-2. *The* `<hibernate-mapping>` *Attributes*

Attribute	Values	Default	Description
auto-import	true, false	true	By default, allows you to use the unqualified class names in Hibernate queries. You would normally only set this to false if the class name would otherwise be ambiguous.
catalog			The database catalog against which queries should apply.
default-access		property	The default access type. If set to property, then get and set methods are used to access the data. If set to field, then the data is accessed directly. Alternatively, you can provide the class name of a PropertyAccessor implementation defining any other access mechanism.
default-cascade			Defines how (and whether) direct changes to data should affect dependent data by default.
default-lazy	true, false	true	Defines whether lazy instantiation is used by default. Generally, the performance benefits are such that you will want to use lazy instantiation whenever possible.
package			The package from which all implicit imports are considered to occur.
schema			The database schema against which queries should apply.

The default cascade modes available for the `default-cascade` attribute (and for the cascade attributes in all other elements) are as follows:

```
create, merge, delete, save-update, evict, replicate, lock, refresh
```

These correspond to the various possible changes in the lifestyle of the parent object. When set (you can include combinations of them as comma-separated values), the relevant changes to the parent will be cascaded to the relation. For example, you may want to apply the `save-update` cascade option to a class that includes `Set` attributes, so that when new persistent classes are added to these, they will not have to be saved explicitly in the session.

There are also three special options:

```
all, delete-orphan, none
```

`all` specifies that all changes to the parent should be propagated to the relation, and `none` specifies that none should. `delete-orphan` applies only to one-to-many associations, and specifies that the relation should be deleted when it is no longer referenced by the parent.

The required order and cardinality of the child elements of `<hibernate-mapping>` are as follows:

```
(meta*,
 typedef*,
 import*,
 (class | subclass | joined-subclass | union-subclass)*,
 (query | sql-query)*,
 filter-def*)
```

THE ORDER AND CARDINALITY INFORMATION FROM THE DTD

The mapping files used by Hibernate have a great many elements and are somewhat self-referential. For example, the `<component>` element permits you to include within it further `<component>` elements, and within those further `<component>` elements—and so on, ad infinitum.

While we do not quote exhaustively from the mapping file's DTD, we sometimes quote the part of it that specifies the permitted ordering and cardinality (number of occurrences) of the child elements of a given element.

The cardinality is expressed by a symbol after the end of the name of the element: * means "zero or more occurrences," ? means "zero or one occurrences," and no trailing symbol means "exactly one occurrence."

The elements can be grouped using brackets, and where the elements are interchangeable, | (the pipe symbol) means "or."

In practical terms, this allows us to tell from the order and cardinality information quoted for the `hibernate-mapping` file that all of the elements immediately below it are, in fact, optional. We can also see that there is no limit to the number of `<class>` elements that can be included.

You can look up this ordering and cardinality information in the DTD for the mapping file for all the elements, including the ones that we have omitted from this chapter. You will also find within the DTD the specification of which attributes are permitted to each element, the values they may take (when they are constrained), and their default values when provided. We recommend that you look at the DTD for enlightenment whenever you are trying to work out whether a specific mapping file should be syntactically valid.

Throughout this book, we have assumed that the mappings are defined in one mapping file for each significant class that is to be mapped to the database. We suggest that you follow this practice in general, but there are some exceptions to this rule. You may, for instance, find it useful to place query and sql-query entries into an independent mapping file, particularly when they do not fall clearly into the context of a single class.

The <class> Element

The child element that you will use most often—indeed, in nearly all of your mapping files—is `<class>`. As you have seen in earlier chapters, we generally describe the relationships between Java objects and database entities in the body of the `<class>` element. The `<class>` element permits the following attributes to be defined (see Table 7-3).

Table 7-3. *The* `<class>` *Attributes*

Attribute	Values	Default	Description
abstract	true, false	false	The flag that should be set if the class being mapped is abstract.
batch-size		1	Specifies the number of items that can be batched together when retrieving instances of the class by identifier.
catalog			The database catalog against which the queries should apply.

Continued

Table 7-3. *Continued*

Attribute	Values	Default	Description
check			Defines an additional row-level check constraint, effectively adding this as a SQL CHECK(...) clause during table generation (for example, check="salary < 1000000").
discriminator-value			A value used to distinguish between otherwise identical subclasses of a common type persisted to the same table. is null and is not null are permissible values. To distinguish between a Cat and a Dog derivative of the Mammal abstract class, for example, you might use discriminator values of C and D, respectively.
dynamic-insert	true, false	false	Indicates whether all columns should appear in INSERT statements. If the attribute is set to true, null columns will not appear in generated INSERT commands. On very wide tables, this may improve performance; but because insert statements are cached, dynamic-insert can easily produce a performance hit.
dynamic-update	true, false	false	Indicates whether all columns should appear in UPDATE statements. If the attribute is set to true, unchanged columns will not appear in generated UPDATE commands. As with dynamic-insert, this can be tweaked for performance reasons. You must enable dynamic-update if you want to use version-based optimistic locking (discussed in Appendix A).
entity-name			The name of the entity to use in place of the class name (therefore required if dynamic mapping is used).
lazy	true, false		Used to disable or enable lazy fetching against the enclosing mapping's default.
mutable	true, false	true	Used to flag that a class is immutable (allowing Hibernate to make some performance optimizations when dealing with these classes).
name			The fully qualified Java name, or optionally unqualified if the <hibernate-mapping> element declares a package attribute, of the class (or interface) that is to be made persistent.
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features.

Attribute	Values	Default	Description
optimistic-lock	none, version	version	Specifies the optimistic locking dirty, all strategy to use. The strategy applies at a class level, but in Hibernate 3 can also be specified (or overridden) at an attribute level. Optimistic locking is discussed in Appendix A.
persister			Allows a custom ClassPersister object to be used when persisting the entity.
polymorphism	implicit, explicit	implicit	Determines how polymorphism is to be used. The default implicit behavior will return instances of the class if superclasses or implemented interfaces are named in the query, and will return subclasses if the class itself is named in the query.
proxy			Specifies a class or an interface to use as the proxy for lazy initialization. Hibernate uses runtime-generated proxies by default, but you can specify your own implementation of org.hibernate.HibernateProxy in their place.
rowid			Flags that row IDs should be used (a database-implementation detail allowing Hibernate to optimize updates).
schema			Optionally overrides the schema specified by the <hibernate-mapping> element.
select-before-update	true, false	false	Flags that Hibernate should carry out extra work to avoid issuing unnecessary UPDATE statements. If set to true, Hibernate issues a SELECT statement before attempting to issue an UPDATE statement in order to ensure that the UPDATE statement is actually required (i.e., that columns have been modified). While this is likely to be less efficient, it can prevent database triggers from being invoked unnecessarily.
subselect			A subselection of the contents of the underlying table. A class can only use a subselect if it is immutable and read-only (because the SQL defined here cannot be reversed). Generally, the use of a database view is preferable.
table			The table name associated with the class (if unspecified, the unqualified class name will be used).
where			An arbitrary SQL where condition to be used when retrieving objects of this class from the table.

Many of these attributes in the <class> element are designed to support preexisting database schemas. In practice, the name attribute is very often the only one set.

The required order and cardinality of the child elements of `<class>` are as follows:

```
(meta*,
 subselect?,
 cache?,
 synchronize*,
 comment?,
 tuplizer*,
 (id | composite-id),
 discriminator?,
 (version | timestamp)?,
 (property | many-to-one | one-to-one | component | dynamic-component |
  properties | any | map | set | list | bag | idbag |
  array | primitive-array)*,
 ((join*, subclass*) | joined-subclass* | union-subclass*),
 loader?,
 sql-insert?,
 sql-update?,
 sql-delete?,
 filter*
 resultset,
 (query | sql-query)
)
```

The <id> Element

All entities need to define their primary key in some way. Any class directly defined by the `<class>` element (not a derived or component class) must therefore have an `<id>` or a `<composite-id>` element to define this (see Table 7-4). Note that while it is not a requirement that your class implementation itself should implement the primary key attribute, it is certainly advisable. If you cannot alter your class design to accommodate this, you can instead use the getIdentifier() method on the Session object to determine the identifier of a persistent class independently.

Table 7-4. *The* `<id>` *Attributes*

Attribute	Values	Default	Description
access			Defines how the properties should be accessed: through field (directly), through property (calling the get/set methods), or through the name of a PropertyAccessor class to be used. The value from the `<hibernate-mapping>` element will be inherited if this is not specified.
column			The name of the column in the table containing the primary key. The value given in the name attribute will be used if this is not specified.
length			The column length to be used.

Attribute	Values	Default	Description
name			The name of the attribute in the class representing this primary key. If this is omitted, it is assumed that the class does not have an attribute directly representing this primary key. Naturally, the column attribute must be provided if the name attribute is omitted.
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features.
type			The Hibernate type of the column.
unsaved-value			The value that the attribute should take when an instance of the class has been created but not yet persisted to the database. This attribute is mandatory.

The `<id>` element requires a `<generator>` element to be specified, which defines how to generate a new primary key for a new instance of the class. The generator takes a class attribute, which defines the mechanism to be used. The class should be an implementation of `org.hibernate.id.IdentifierGenerator`. Optional `<param>` elements can be provided if the identifier needs additional configuration information, each having the following form:

```
<param name="parameter name">parameter value</param>
```

Hibernate provides several default `IdentifierGenerator` implementations, which can be referenced by convenient short names, as shown in Table 7-5. These are fairly comprehensive, so you are unlikely to need to implement your own `IdentifierGenerator`.

Table 7-5. *The Default* `IdentiferGenerator` *Implementations*

Short Name	Description
guid	Uses a database-generated "globally" unique identifier. This is not portable to databases that do not have a guid type. The specific implementation, and hence the quality of the uniqueness of this key, may vary from vendor to vendor.
hilo	Uses a database table and column to efficiently and portably maintain and generate identifiers that are unique to that database. The Hibernate int, short, and long types are supported.
identity	Supports the identity column type available in some, but not all, databases. This is therefore not a fully portable option. The Hibernate int, short, and long types are supported.
increment	Generates a suitable key by adding 1 to the current highest key value. Can apply to int, short, or long hibernate types. This only works if other processes are not permitted to update the table at the same time. If multiple processes are running, then depending on the constraints enforced by the database, the result may be an error in the application(s) or data corruption.
native	Selects one of sequence, identity, or hilo, as appropriate. This is a good compromise option since it uses the innate features of the database and is portable to most platforms. This is particularly appropriate if your code is likely to be deployed to a number of database implementations with differing capabilities.

Continued

Table 7-5. *Continued*

Short Name	Description
seqhilo	Uses a sequence to efficiently generate identifiers that are unique to that database. The Hibernate `int`, `short`, and `long` types are supported. This is not a portable technique (see `sequence`, following).
sequence	Supports the `sequence` column type (essentially a database-enforced increment) available in some, but not all, databases. This is, therefore, not a fully portable option. The Hibernate `int`, `short`, and `long` types are supported.
uuid	Attempts to portably generate a (cross-database) unique primary key. The key is composed of the local IP address, the startup time of the JVM (accurate to ¼ of a second), the system time, and a counter value (unique within the JVM). This cannot guarantee absolutely that a given key is unique, but it will be good enough for most clustering purposes.

The child elements of the `<id>` element are as follows:

```
(meta*, column*, type?, generator?)
```

While this is all rather complex, Listing 7-1 shows a typical `<id>` element from Chapter 3, which illustrates the simplicity of the usual case.

Listing 7-1. *A Typical* `<id>` *Element*

```
<id name="id" type="long" column="id">
   <generator class="native"/>
</id>
```

■**Note** When the `<id>` element cannot be defined, a compound key can instead be defined using the `<composite-id>` element. This is provided purely to support existing database schemas. A new Hibernate project with a clean database design does not require this.

In addition to using the standard and custom generator types, you have the option of using the special `assigned` generator type. This allows you to explicitly set the identifier for the entities that you will be persisting—Hibernate will not then attempt to assign any identifier value to such an entity. If you use this technique, you will not be able to use the `saveOrUpdate()` method on a transient entity—instead, you will have to call the appropriate `save()` or `update()` method explicitly.

The <property> Element

While it is not absolutely essential, almost all classes will also maintain a set of properties in the database in addition to the primary key. These must be defined by a `<property>` element (see Table 7-6).

Table 7-6. *The* `<property>` *Attributes*

Attribute	Values	Default	Description
access			Defines how the properties should be accessed: through `field` (directly), through `property` (calling the get/set methods), or through the name of a `PropertyAccessor` class to be used. The value from the `<class>` element or `<hibernate-mapping>` element will be inherited if this is not specified.
column			The column in which the property will be maintained. If omitted, this will default to the name of the attribute; or it can be specified with nested `<column>` elements (see Listing 7-2).
formula			An arbitrary SQL query representing a computed property (i.e., one that is calculated dynamically, rather than represented in a column).
index			The name of an index to be maintained for the column.
insert	true, false	true	Specifies whether creation of an instance of the class should result in the column associated with this attribute being included in `insert` statements.
lazy	true, false	false	Defines whether lazy instantiation is used by default for this column.
length			The column length to be used.
name			The (mandatory) name of the attribute. This should start with a lowercase letter.
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features.
not-null	true, false	false	Specifies whether the column is permitted to contain `null` values.
optimistic-lock	true, false	true	Determines whether optimistic locking should be used when the attribute has been updated.
precision			Allows the precision (the number of digits) to be specified for numeric data.
scale			Allows the scale (the number of digits to the right of the decimal point) to be specified for numeric data.
type			The Hibernate type of the column.
unique	true, false	false	Indicates whether duplicate values are permitted for this column/attribute.
unique-key			Groups the columns together by this attribute value. Represents columns across which a unique key constraint should be generated (not yet supported in the schema generation).
update	true, false	true	Specifies whether changes to this attribute in instances of the class should result in the column associated with this attribute being included in `update` statements.

The child elements of the `<property>` element are as follows:

```
(meta*, (column | formula)*, type?)
```

Any element accepting a `column` attribute, as is the case for the `<property>` element, will also accept `<column>` elements in its place. For an example, see Listing 7-2.

Listing 7-2. *Using the* `<column>` *Element*

```
<property name="message"/>
   <column name="message" type="string"/>
</property>
```

This particular example does not really give us anything beyond the use of the `column` attribute directly; but the `<column>` element comes into its own with custom types and some of the more complex mappings that we will be looking into later in the chapter.

The <component> Element

The `<component>` element is used to map classes that will be represented as extra columns within a table describing some other class. We have already discussed how components fit in as a compromise between full entity types and mere value types.

The `<component>` element can take the attributes listed in Table 7-7.

Table 7-7. *The* `<component>` *Attributes*

Attribute	Values	Default	Description
access			Defines how the properties should be accessed: through `field` (directly), through `property` (calling the get/set methods), or through the name of a `PropertyAccessor` class to be used
class			The class that the parent class incorporates by composition
insert	true, false	true	Specifies whether creation of an instance of the class should result in the column associated with this attribute being included in `insert` statements
lazy	true, false	false	Defines whether lazy instantiation is used by default for this mapped entity
name			The name of the attribute (component) to be persisted
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features
optimistic-lock	true, false	true	Specifies the optimistic locking strategy to use
unique	true, false	false	Indicates that the values that represent the component must be unique within the table
update	true, false	true	Specifies whether changes to this attribute in instances of the class should result in the column associated with this attribute being included in `update` statements

The child elements of the `<component>` element are as follows:

```
(meta*,
 tuplizer*,
 parent?,
 (property | many-to-one | one-to-one |
  component | dynamic-component | any |
  map | set | list | bag |
  array | primitive-array)* )
```

We provide a full example of the use of the `<component>` element in the "Mapping Composition" section later in this chapter.

The `<one-to-one>` Element

The `<one-to-one>` element expresses the relationship between two classes, where each instance of the first class is related to a single instance of the second, and vice versa. Such a one-to-one relationship can be expressed either by giving each of the respective tables the same primary key values, or by using a foreign key constraint from one table onto a unique identifier column of the other. Table 7-8 shows the attributes that apply to the `<one-to-one>` element.

Table 7-8. *The* `<one-to-one>` *Attributes*

Attribute	Values	Default	Description
access			Specifies how the class member should be accessed: `field` for direct field access or `attribute` for access via the get and set methods.
cascade			Determines how changes to the parent entity will affect the linked relation.
check			The SQL to create a multirow check constraint for schema generation.
class			The property type of the `attribute` or `field` (if omitted, this will be determined by reflection).
constrained	true, false		Indicates that a foreign key constraint on the primary key of this class references the table of the associated class.
embed-xml	true, false		When using XML relational persistence, indicates whether the XML tree for the associated entity itself, or only its identifier, will appear in the generated XML tree.
entity-name			The entity name of the associated class.
fetch	join, select		The mode in which the element will be retrieved (outer join, a series of `select`s, or a series of subselects). Only one member of the enclosing class can be retrieved by outer join.
foreign-key			The name to assign to the foreign key enforcing the relationship.
formula			Allows the value to which the associated class maps its foreign key to be overridden using an SQL formula.

Continued

Table 7-8. *Continued*

Attribute	Values	Default	Description
lazy	true, false		Overrides the entity-loading mode.
name			Assigns a name to the entity (required in dynamic mappings).
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features.
outer join	true, false, auto		Specifies whether an outer join should be used.
property-ref			Specifies the column in the target entity's table that the foreign key references. If the referenced table's foreign key does not reference the primary key of the "many" end of the relationship, then property-ref can be used to specify the column that it references. This should only be the case for legacy designs—when creating a new schema, your foreign keys should always reference the primary key of the related table.

You would select a primary key association when you do not want an additional table column to relate the two entities. The master of the two entities takes a normal primary key generator, and its one-to-one mapping entry will typically have the attribute name and associated class specified only. The slave entity will be mapped similarly, but must have the constrained attribute setting applied to ensure that the relationship is recognized.

Because the slave class's primary key must be identical to that allocated to the master, it is given the special id generator type of foreign. On the slave end, the <id> and <one-to-one> elements will therefore look like this:

```
<id name="id" column="product">
   <generator class="foreign">
      <param name="property">campaign</param>
   </generator>
</id>

<one-to-one name="campaign"
   class="com.hibernatebook.xmlmapping.Campaign"
   constrained="true"/>
```

There are some limitations to this approach: it cannot be used on the receiving end of a many-to-one relationship (even when the "many" end of the association is limited by a unique constraint), and the slave entity cannot be the slave of more than one entity.

In these circumstances, you will need to declare the master end of the association as a uniquely constrained one-to-many association. The slave entity's table will then need to take a foreign key column associating it with the master's primary key. The property-ref attribute setting is used to declare this relationship, like so:

```
<one-to-one
    name="campaign"
    class="com.hibernatebook.xmlmapping.Campaign"
    property-ref="product"/>
```

The format used in this example is the most common. The body of the element consists of an infrequently used optional element:

```
(meta* | formula*)
```

We discuss the `<many-to-many>` element and the alternative approach of composition in some detail in the "Mapping Collections" section later in this chapter.

The <many-to-one> Element

The many-to-one association describes the relationship in which multiple instances of one class can reference a single instance of another class. This enforces a relational rule for which the "many" class has a foreign key into the (usually primary) unique key of the "one" class. Table 7-9 shows the attributes permissible for the `<many-to-one>` element.

Table 7-9. *The* `<many-to-one>` *Attributes*

Attribute	Values	Default	Description
access			Specifies how the class member should be accessed: `field` for direct field access, or `attribute` for access via the get and set methods.
cascade			Determines how changes to the parent entity will affect the linked relation.
class			The property type of the attribute or field (if omitted, this will be determined by reflection).
column			The column containing the identifier of the target entity (i.e., the foreign key from this entity into the mapped one).
embed-xml	true, false		When using XML relational persistence, indicates whether the XML tree for the associated entity itself, or only its identifier, will appear in the generated XML tree.
entity-name			The name of the associated entity.
fetch	join, select		The mode in which the element will be retrieved (`outer join`, a series of `selects`, or a series of `subselects`). Only one member of the enclosing class can be retrieved by `outer join`.
foreign-key			The name of the foreign key constraint to generate for this association.
formula			An arbitrary SQL expression to use in place of the normal primary key relationship between the entities.
index			The name of the index to be applied to the foreign key column in the parent table representing the "many" side of the association.

Continued

Table 7-9. *Continued*

Attribute	Values	Default	Description
insert	true, false	true	Indicates whether the field can be persisted. When set to false, this prevents inserts if the field has already been mapped as part of a composite identifier or some other attribute.
lazy	false, proxy, noproxy		Overrides the entity-loading mode.
name			The (mandatory) name of the attribute. This should start with a lowercase letter.
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features.
not-found	exception, ignore	exception	The behavior to exhibit if the related entity does not exist (either throw an exception or ignore the problem).
not-null	true, false	false	Specifies whether a not-null constraint should be applied to this column.
optimistic-lock	true, false	true	Specifies whether optimistic locking should be used.
outer-join	true, false, auto		Specifies whether an outer join should be used.
property-ref			Specifies the column in the target entity's table that the foreign key references. If the referenced table's foreign key does not reference the primary key of the "many" end of the relationship, then property-ref can be used to specify the column that it references. This should only be the case for legacy designs—when creating a new schema, your foreign keys should always reference the primary key of the related table.
unique	true, false	false	Specifies whether a unique constraint should be applied to the column.
unique-key			Groups the columns together by this attribute value. Represents columns across which a unique key constraint should be generated (not yet supported in the schema generation).
update	true, false	true	When set to false, prevents updates if the field has already been mapped elsewhere.

If a unique constraint is specified on a many-to-one relationship, it is effectively converted into a one-to-one relationship. This approach is preferred over creating a one-to-one association, both because it results in a simpler mapping and because it requires less intrusive changes to the database should it become desirable to relax the one-to-one association into a many-to-one.

This element has a small number of optional daughter elements—the `<column>` element will be required when a composite key has to be specified:

```
(meta*, (column | formula)*)
```

The following mapping illustrates the creation of a simple many-to-one association between a `User` class and an `Email` class: each user can have only one e-mail address—but an e-mail address can belong to more than one user.

```
<many-to-one
    name="email"
    class="com.hibernatebook.xmlmapping.Email"
    column="email"
    cascade="all" unique="true"/>
```

The simplest approach to creating a many-to-one relationship, as shown in the previous example, requires two tables and a foreign key dependency. An alternative is to use a link table to combine the two entities. The link table contains the appropriate foreign keys referencing the two tables associated with both of the entities in the association. The following code shows the mapping of a many-to-one relationship via a link table.

```
<join table="link_email_user" inverse="true" optional="false">
    <key column="user_id"/>
    <many-to-one name="email" column="email_id" not-null="true"/>
</join>
```

The disadvantage of the link table approach is its slightly poorer performance (it requires a join of three tables to retrieve the associations, rather than one). Its benefit is that it requires less extreme changes to the schema if the relationship is modified—typically, changes would be made to the link table, rather than to one of the entity tables.

The Collection Elements

These are the elements that are required for you to include an attribute in your class that represents any of the collection classes. For example, if you have an attribute of type `Set`, then you will need to use a `<bag>` or `<set>` element to represent its relationship with the database.

Because of the simplicity of the object-oriented relationship involved, where one object has an attribute capable of containing many objects, it is a common fallacy to assume that the relationship must be expressed as a one-to-many. In practice, however, this will almost always be easiest to express as a many-to-many relationship, where an additional link table closely corresponds with the role of the collection itself. See the "Mapping Collections" section later in this chapter for a more detailed illustration of this.

All the collection mapping elements share the attributes shown in Table 7-10.

Table 7-10. *The Attributes Common to the Collection Elements*

Attribute	Values	Default	Description
access			Specifies how the class member should be accessed: `field` for direct field access or `attribute` for access via the get and set methods.
batch-size			Specifies the number of items that can be batched together when retrieving instances of the class by identifier.

Continued

Table 7-10. *Continued*

Attribute	Values	Default	Description
cascade			Determines how changes to the parent entity will affect the linked relation.
catalog			The database catalog against which the queries should apply.
collection-type			The name of a UserCollectionType class describing the collection type to be used in place of the defaults.
check			The SQL to create a multirow check constraint for schema generation.
embed-xml	true, false		When using XML relational persistence, indicates whether the XML tree for the associated entity itself, or only its identifier, will appear in the generated XML tree.
fetch	join, select		The mode in which the element will be retrieved (outer-join, a series of selects, or a series of subselects). Only one member of the enclosing class can be retrieved by outer-join.
lazy	true, false		Can be used to disable or enable lazy fetching against the enclosing mapping's default.
mutable	true, false	true	Can be used to flag that a class is mutable (allowing Hibernate to make some performance optimizations when dealing with these classes).
name			The (mandatory) name of the attribute. This should start with a lowercase letter.
node			Specifies the name of the XML element or attribute that should be used by the XML relational persistence features.
optimistic-lock	true, false	true	Specifies the optimistic locking strategy to use.
outer-join	true, false, auto		Specifies whether an outer join should be used.
persister			Allows a custom ClassPersister object to be used when persisting this class.
schema			The database schema against which queries should apply.
subselect			A query to enforce a subselection of the contents of the underlying table. A class can only use a subselect if it is immutable and read-only (because the SQL defined here cannot be reversed). Generally, the use of a database view is preferable.
table			The name of the table in which the associated entity is stored.
where			An arbitrary SQL where clause limiting the linked entities.

The set Collection

A set collection allows collection attributes derived from the Set interface to be persisted.

In addition to the common collection mappings, the <set> element offers the inverse, order-by, and sort attributes, as shown in Table 7-11.

Table 7-11. *The Additional* <set> *Attributes*

Attribute	Values	Default	Description
inverse	true, false	false	Specifies that an entity is the opposite navigable end of a relationship expressed in another entity's mapping.
order-by			Specifies an arbitrary SQL order by clause to constrain the results returned by the SQL query that populates the set collection.
sort			Specifies the collection class sorting to be used. The value can be unsorted, natural, or any Comparator class.

The child elements of the <set> element are as follows:

```
(meta*,
 subselect?,
 cache?,
 synchronize*,
 comment?,
 key,
 (element | one-to-many | many-to-many |
  composite-element | many-to-any),
 loader?,
 sql-insert?,
 sql-update?,
 sql-delete?,
 sql-delete-all?,
 filter*)
```

The following code shows an implementation of mapping a set of strings into a property called titles:

```
<set name="titles" table="nameset">
   <key column="titleid"/>
   <element type="string" column="name" not-null="true"/>
</set>
```

A typical implementation, however, maps other entities into the collection. Here we map Phone entities from the "many" side of a one-to-many association into a Set property, called phoneNumbers, that belongs to a User entity:

```
<set name="phoneNumbers">
   <key column="aduser"/>
   <one-to-many class="sample.Phone"/>
</set>
```

If the Phone class contains a reference to a User object, it is not automatically clear whether this constitutes a pair of unrelated associations or two halves of the same association—a bidirectional association. When a bidirectional association is to be established, one side must be selected as the owner (in a one-to-many or many-to-one association, it must always be the "many" side), and the other will be marked as being the inverse half of the relationship. See the discussion of unidirectional and bidirectional associations at the end of Chapter 4. The following code shows a mapping of a one-to-many relationship as a reverse association.

```
<set name="phoneNumbers" inverse="true">
   <key column="aduser"/>
   <one-to-many class="sample.Phone"/>
</set>
```

The list Collection

A list collection allows collection attributes derived from the List interface to be persisted.

In addition to the common collection mappings, the <list> element offers the inverse attribute, as shown in Table 7-12.

Table 7-12. *The Additional* <list> *Attribute*

Attribute	Values	Default	Description
inverse	true, false	false	Specifies that an entity is the opposite navigable end of a relationship expressed in another entity's mapping

The child elements of the <list> element are as follows:

```
(meta*,
 subselect?,
 cache?,
 synchronize*,
 comment?,
 key,
 (index | list-index),
 (element | one-to-many | many-to-many |
  composite-element | many-to-any),
 loader?,
 sql-insert?,
 sql-update?,
 sql-delete?,
 sql-delete-all?,
 filter*)
```

A typical implementation of a `list` mapping is as follows:

```
<list name="list" table="namelist">
   <key column="fooid"/>
   <index column="position"/>
   <element type="string" column="name" not-null="true"/>
</list>
```

The idbag Collection

An `idbag` collection allows for appropriate use of collection attributes derived from the `List` interface. A `bag` data structure permits unordered storage of unordered items, and permits duplicates. Because the collection classes do not provide a native `bag` implementation, classes derived from the `List` interface tend to be used as a substitute. The imposition of ordering imposed by a `list` is not itself a problem, but the implementation code can become dependent upon the ordering information.

`idbag` usually maps to a `List`. However, by managing its database representation with a surrogate key, you can make the performance of updates and deletions of items in a collection defined with `idbag` dramatically better than with an unkeyed `bag` (described at the end of this section). Hibernate does not provide a mechanism for obtaining the identifier of a row in the `bag`.

In addition to the common collection mappings, the `<idbag>` element offers the `order-by` element, as shown in Table 7-13.

Table 7-13. *The Additional* `<idbag>` *Attribute*

Attribute	Values	Default	Description
order-by			Specifies an arbitrary SQL `order by` clause to constrain the results returned by the SQL query that populates the collection

The child elements of the `<idbag>` element are as follows:

```
(meta*,
 subselect?,
 cache?,
 synchronize*,
 comment?,
 collection-id,
 key,
 (element | many-to-many |
  composite-element | many-to-any),
 loader?,
 sql-insert?,
 sql-update?,
 sql-delete?,
 sql-delete-all?,
 filter*)
```

A typical implementation of an idbag mapping is as follows:

```
<idbag name="idbag" table="nameidbag">
   <collection-id column="id" type="int">
      <generator class="native"/>
   </collection-id>

   <key column="fooid"/>
   <element type="string" column="name" not-null="true"/>
</idbag>
```

The map Collection

A map collection allows collection attributes derived from the Map interface to be persisted.

In addition to the common collection mappings, the <map> element offers the inverse, order-by, and sort attributes, as shown in Table 7-14.

Table 7-14. *The Additional* <map> *Attributes*

Attribute	Values	Default	Description
inverse	true, false	false	Specifies that this entity is the opposite navigable end of a relationship expressed in another entity's mapping
order-by			Specifies an arbitrary SQL order by clause to constrain the results returned by the SQL query that populates the map
sort		unsorted	Specifies the collection class sorting to be used. The value can be unsorted, natural, or any Comparator class

The child elements of the <map> element are as follows:

```
(meta*,
 subselect?,
 cache?,
 synchronize*,
 comment?,
 key,
 (map-key | composite-map-key | map-key-many-to-many |
  index | composite-index | index-many-to-many |
  index-many-to-any),
 (element | one-to-many | many-to-many | composite-element |
  many-to-any),
 loader?,
 sql-insert?,
 sql-update?,
 sql-delete?,
 sql-delete-all?,
 filter*)
```

A typical implementation of the mapping is as follows:

```
<map name="map" table="namemap">
   <key column="fooid"/>
   <index column="name" type="string"/>
   <element column="value" type="string" not-null="true"/>
</map>
```

The bag Collection

If your class represents data using a class derived from the List interface, but you do not want to maintain an index column to keep track of the order of items, you can optionally use the bag collection mapping to achieve this. The order in which the items are stored and retrieved from a bag is completely ignored.

Although the bag's table does not contain enough information to determine the order of its contents prior to persistence into the table, it *is* possible to apply an order by clause to the SQL used to obtain the contents of the bag so that it has a natural sorted order as it is acquired. This will not be honored at other times during the lifetime of the object.

If the <bag> elements lack a proper key, there will be a performance impact that will manifest itself when update or delete operations are performed on the contents of the bag.

In addition to the common collection mappings, the <bag> element therefore offers the order-by as well as the inverse attribute, as shown in Table 7-15.

Table 7-15. *The Additional* <bag> *Attributes*

Attribute	Values	Default	Description
inverse	true, false	false	Specifies that an entity is the opposite navigable end of a relationship expressed in another entity's mapping
order-by			Specifies an arbitrary SQL order by clause to constrain the results returned by the SQL query that populates the collection

The child elements of the <bag> element are as follows:

```
(meta*,
 subselect?,
 cache?,
 synchronize*,
 comment?,
 key,
 (element | one-to-many | many-to-many |
  composite-element | many-to-any),
 loader?,
 sql-insert?,
 sql-update?,
 sql-delete?,
 sql-delete-all?,
 filter*)
```

A typical implementation of a bag mapping is as follows:

```
<bag name="bag" table="namebag">
   <key column="fooid"/>
   <element column="value" type="string" not-null="true"/>
</bag>
```

Mapping Simple Classes

Figure 7-1 shows the class diagram and entity relationship diagram for a simple class. They are as straightforward as you would expect.

Figure 7-1. *Representing a simple class*

The elements discussed so far are sufficient to map a basic class into a single table, as shown in Listing 7-3.

Listing 7-3. *A Simple Class to Represent a User*

```
package com.hibernatebook.xmlmapping;

public class User {

   public User(String username) {
      this.username = username;
   }

   User() {
   }

   public int getId() {
      return id;
   }

   public String getUsername() {
      return username;
   }
```

```
   public void setId(int id) {
      this.id = id;
   }

   public void setUsername(String username) {
      this.username = username;
   }

   // We will map the id to the table's primary key
   private int id = -1;

   // We will map the username into a column in the table
   private String username;
}
```

It's pretty easy to see that we might want to represent the class in Listing 7-3 in a table with the format shown in Table 7-16.

Table 7-16. *Mapping a Simple Class to a Simple Table*

Column	Type
Id	Integer
Username	Varchar(32)

The mapping between the two is, thus, similarly straightforward:

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
   PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
   "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping>
   <class name="book.hibernatebook.chapter06.User">

      <id name="id" type="int">
         <generator class="native"/>
      </id>

      <property name="username" type="string" length="32"/>

   </class>
</hibernate-mapping>
```

Aside from the very limited number of properties maintained by the class, this is a pretty common mapping type, so it is reassuring to see that it can be managed with a minimal number of elements (<hibernate-mapping>, <class>, <id>, <generator>, and <property>).

Mapping Composition

Figure 7-2 shows the class diagram and the entity relationship diagram for a composition relationship between two classes. Here, the Advert class is composed of a Picture class in addition to its normal value types.

Figure 7-2. *Representing composition*

Composition is the strongest form of aggregation—in which the life cycle of each object is dependent upon the life cycle of the whole. Although Java does not make the distinction between other types of aggregation and composition, it becomes relevant when we choose to store the components in the database, because the most efficient and natural way to do this is to store them in the same table.

In our example, we will look at an Advert class that has this relationship with a Picture class. The idea is that our advert is always going to be associated with an illustration (see Listings 7-4 and 7-5). In these circumstances, there is a clear one-to-one relationship that could be represented between two distinct tables, but which is more efficiently represented with one.

Listing 7-4. *The Class Representing the Illustration*

```
package com.hibernatebook.xmlmapping;

public class Picture {
   public Picture(String caption, String filename) {
      this.caption = caption;
      this.filename = filename;
   }

   Picture() {
   }

   public String getCaption() {
      return this.caption;
   }
```

```java
    public String getFilename() {
        return this.filename;
    }

    public void setCaption(String title) {
        this.caption = title;
    }

    public void setFilename(String filename) {
        this.filename = filename;
    }

    private String caption;
    private String filename;
}
```

Listing 7-5. *The Class Representing the Advert*

```java
package com.hibernatebook.xmlmapping;

public class Advert {
    public Advert(String title, String content, Picture picture) {
        this.title = title;
        this.content = content;
        this.picture = picture;
    }

    Advert() {
    }

    public int getId() {
        return id;
    }

    public String getTitle() {
        return this.title;
    }

    public String getContent() {
        return this.content;
    }

    public Picture getPicture() {
        return this.picture;
    }
```

```
    public void setId(int id) {
       this.id = id;
    }

    public void setTitle(String title) {
       this.title = title;
    }

    public void setContent(String content) {
       this.content = content;
    }

    public void setPicture(Picture picture) {
       this.picture = picture;
    }

    private int id = -1;
    private String title;
    private String content;
    private Picture picture;
}
```

Again, Hibernate manages to express this simple relationship with a correspondingly sim-
ple mapping file. We introduce the component entity for this association. Here it is in use:

```xml
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<class name="com.hibernatebook.xmlmapping.Advert">
    <id name="id" type="int">
        <generator class="native"/>
    </id>
    <property name="title" type="string" length="255"/>
    <property name="content" type="text"/>
    <component name="picture" class="com.hibernatebook.xmlmapping.Picture">
        <property name="caption" type="string" length="255"/>
        <property name="filename" type="string" length="32"/>
    </component>
</class>
```

In this example, we use the <property> element to describe the relationship between
Picture and its attributes. In fact, this is true of all of the rest of the elements of <class>—
a <component> element can even contain more <component> elements. Of course, this makes
perfect sense, since a component usually corresponds with a Java class.

Mapping Other Associations

In Figure 7-3, the Advert class includes an instance of a Picture class. The relationship in the tables is represented with the Picture table having a foreign key onto the Advert table.
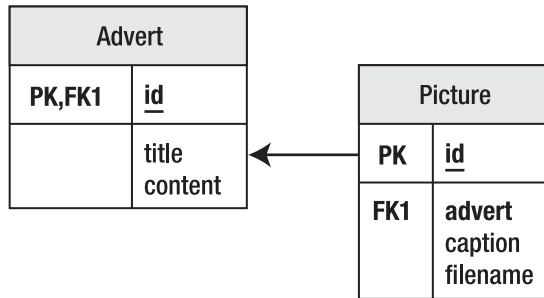
Figure 7-3. *Mapping an aggregation or composition relationship*

A one-to-one correspondence does not absolutely require you to incorporate both parties into the same table. There are often good reasons not to. For instance, in the Picture example, it is entirely possible that while the initial implementation will permit only one Picture per Advert, a future implementation will relax this relationship. Consider this scenario from the perspective of the database for a moment (see Table 7-17).

Table 7-17. *The* Advert *Table*

ID	Title	Contents	PictureCaption	PictureFilename
1	Bike	Bicycle for sale	My bike (you can ride it if you like)	advert001.jpg
2	Sofa	Sofa, comfy but used	Chesterfield sofa	advert002.jpg
3	Car	Shabby MGF for sale	MGF VVC (BRG)	advert003.jpg

If we want to allow the advert for the sofa to include another picture, we would have to duplicate some of the data, or include null columns. It would probably be preferable to set up a pair of tables: one to represent the adverts, and one to represent the distinct tables (as shown in Tables 7-18 and 7-19).

Table 7-18. *The Refined* Advert *Table*

ID	Title	Contents
1	Bike	Bicycle for sale
2	Sofa	Sofa, comfy but used
3	Car	Shabby MGF for sale

Table 7-19. *The* `Picture` *Table*

ID	Advert	Caption	Filename
1	1	My bike (you can ride it if you like)	advert001.jpg
2	2	Chesterfield sofa	advert002.jpg
3	3	MGF VVC (BRG)	advert003.jpg

If we decide (considering the database only) to allow additional pictures, we can then include extra rows in the `Picture` table without duplicating any data unnecessarily (see Table 7-20).

Table 7-20. *The* `Picture` *Table with Multiple Pictures per Advert*

ID	Advert	Caption	Filename
1	1	My bike (you can ride it if you like)	advert001.jpg
2	2	Chesterfield sofa	advert002.jpg
3	2	Back of sofa	advert003.jpg
4	3	MGF VVC (BRG)	advert004.jpg

With the single `Advert` table, the query to extract the data necessary to materialize an instance of the `Advert` consists of something like this:

```
select id,title,contents,picturecaption,picturefilename from advert where id = 1
```

It is obvious here that a single row will be returned, since we are carrying out the selection on the primary key.

Once we split things into two tables, we have a slightly more ambiguous pair of queries:

```
select id,title,contents from advert where id = 1
select id,caption,filename from picture where advert = 1
```

While Hibernate is not under any particular obligation to use this pair of SQL instructions to retrieve the data (it could reduce it to a join on the table pair), it is the easiest way of thinking about the data we are going to retrieve. While the first query of the two is required to return a single row, this is not true for the second query—if we have added multiple pictures, we will get multiple rows back.

In these circumstances, there is very little difference between a one-to-one relationship and a one-to-many relationship, except from a business perspective. That is to say, we choose not to associate an advert with multiple pictures, even though we have that option.

This, perhaps, explains why the expression of a one-to-one relationship in Hibernate is usually carried out via a many-to-one mapping. If you do not find that persuasive, remember that a foreign key relationship, which is the relationship that the `advert` column in the `Picture` table has with the `id` column in the `Advert` table, is a many-to-one relationship between the entities.

In our example, the `Picture` table will be maintaining the `advert` column as a foreign key into the `Advert` table, so this must be expressed as a many-to-one relationship with the `Advert` object (see Listing 7-6).

Listing 7-6. *The New Picture Mapping*

```xml
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<class name="com.hibernatebook.xmlmapping.Picture">
    <id name="id" type="int">
        <generator class="native"/>
    </id>
    <many-to-one
      name="advert"
      class="com.hibernatebook.xmlmapping.Advert"
      column="advert"/>
    <property name="caption" type="string" length="255"/>
    <property name="filename" type="string" length="32"/>
</class>
```

If you still object to the many-to-one relationship, you will probably find it cathartic to note that we have explicitly constrained this relationship with the `unique` attribute. You will also find it reassuring that in order to make navigation possible directly from the `Advert` to its associated `Picture`, we can in fact use a one-to-one mapping entry. We need to be able to navigate in this direction because we expect to retrieve adverts from the database, and then display their associated pictures (see Listing 7-7).

Listing 7-7. *The Revised* Advert *Mapping*

```xml
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<class name="com.hibernatebook.xmlmapping.Advert">
    <id name="id" type="int">
        <generator class="native"/>
    </id>
    <property name="title" type="string" length="255"/>
    <property name="content" type="text"/>
    <one-to-one name="picture"
                class="com.hibernatebook.xmlmapping.Picture"
                property-ref="picture">
</class>
```

Now that we have seen how one-to-one and many-to-one relationships are expressed, we will see how a many-to-many relationship can be expressed.

Mapping Collections

In Figure 7-4, we show the User objects as having an unknown number of Advert instances. In the database, this is then represented with three tables, one of which is a link table between the two entity tables.
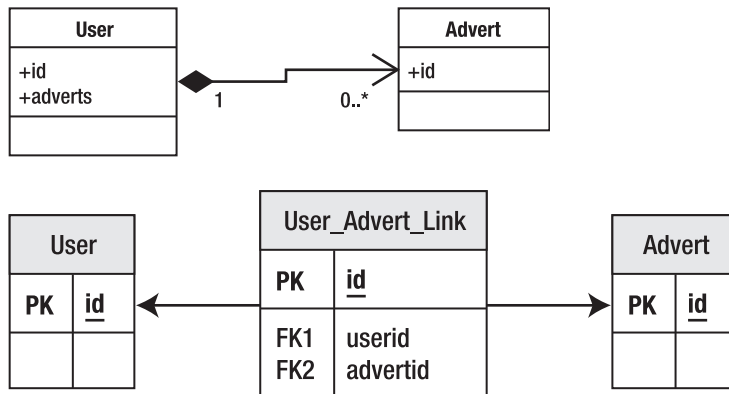
Figure 7-4. *Mapping collections*

The Java collection classes provide the most elegant mechanism for expressing the "many" end of a many-to-many relationship in our own classes:

```
public Set getAdverts();
```

If we use generics, we can give an even more precise specification:

```
public Set<Advert> getAdverts();
```

■**Note** A lot of legacy code will not use generics. However, if you have the opportunity you should do so, as it allows you to make this sort of distinction clear at the API level, instead of at the documentation level. Hibernate 3 is compatible with Java 5 generics.

Of course, we can place values (of Object type) into collections as well as entities, and Java 5 introduced autoboxing so that we have the illusion of being able to place primitives into them as well.

```
List<Integer> ages = getAges();
int first = ages.get(0);
```

The only catch with collection mapping is that an additional table may be required to correctly express the relationship between the owning table and the collection. Table 7-21 shows how it should be done; the entity table contains only its own attributes.

Table 7-21. *The* Entity *Table*

ID	Name
1	Entity 1

A separate collection table, on the other hand, contains the actual values (see Table 7-22). In this case, we are linking a List to the owning entity, so we need to include a column to represent the position of the values in the list, as well as the foreign key into the owning entity and the column for the actual values that are contained within the collection.

Table 7-22. ListTable

entityid	positionInList	listValue
1	1	Good
1	2	Bad
1	3	Indifferent

In a legacy schema, you may quite often encounter a situation in which all the values have been retained within a single table (see Table 7-23).

Table 7-23. EntityTable

ID	Name	positionInList	listValue
1	Entity 1	1	Good
1	Entity 1	2	Bad
1	Entity 1	3	Indifferent

It should be obvious that this is not just poor design from Hibernate's perspective—it's also bad relational design. The values in the entity's name attribute have been duplicated needlessly, so this is not a properly normalized table. We also break the foreign key of the table, and need to form a compound key of id and positionInList. Overall, this is a poor design, and we encourage you to use a second table if at all possible. If you must work with such an existing design, see Chapter 13 for some techniques for approaching this type of problem.

If your collection is going to contain entity types instead of value types, the approach is essentially the same, but your second table will contain keys into the second entity table instead of value types. This changes the combination of tables into the situation shown in the entity relationship diagram (see Figure 7-4), in which we have a link table joining two major tables into a many-to-many relationship. This is a very familiar pattern in properly normalized relational schemas.

The following code shows a mapping of a Set attribute representing the adverts with which the User class is associated:

```
<set name="adverts"
    table="user_advert_link"
    cascade="save-update">
  <key column="userid"/>
  <many-to-many
     class="com.hibernatebook.xmlmapping.Advert"
     column="advertid"/>
</set>
```

Hibernate's use of collections tends to expose the lazy loading issues more than most other mappings. If you enable lazy loading, the collection that you retrieve from the session will be a proxy implementing the relevant collection interface (in our example, Set), rather than one of the usual Java concrete collection implementations.

This allows Hibernate to retrieve the contents of the collection only as they are required by the user. If you load an entity, consult a single item from the collection, and then discard it, often only a handful of SQL operations will be required. If the collection in question represents hundreds of entity instances, the performance advantages of lazy loading (compared with the massive task of reading in *all* of the entities concerned) are massive.

However, you will need to ensure that you do not try to access the contents of a lazily loaded collection at a time when it is no longer associated with the session, unless you can be certain that the contents of the collection that you are accessing have already been loaded.

Mapping Inheritance Relationships

Figure 7-5 shows a simple class hierarchy. The superclass is Advert, and there are two classes derived from this: a Personal class to represent personal advertisements and a Property class to represent property advertisements.
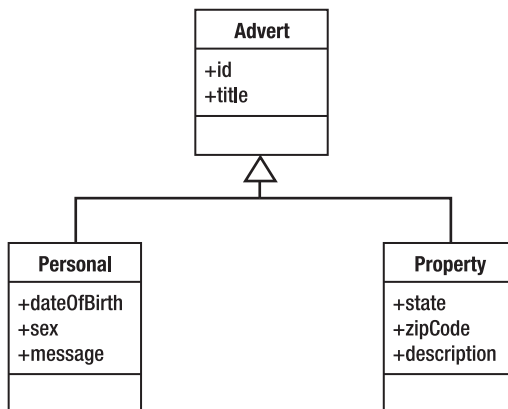
Figure 7-5. *A simple inheritance hierarchy*

Hibernate can represent inheritance relationships in a relational schema in three ways, each mapped in a slightly different way. These are as follows:

- One table for each concrete class implementation

- One table for each subclass (including interfaces and abstract classes)

- One table for each class hierarchy

Each of these techniques has different costs and benefits, so we will show you an example mapping from each and discuss some of these issues.

One Table per Concrete Class

This approach is the easiest to implement. You map each of the concrete classes as normal, writing mapping elements for each of its persistent properties (including those that are inherited). No mapping files are required for interfaces and abstract classes.

Figure 7-6 shows the schema required to represent the hierarchy from Figure 7-5 using this technique.

Advert	
PK	**id**
	title

Personal	
PK	**id**
	title dateOfBirth sex message

Property	
PK	**id**
	title state zipCode description
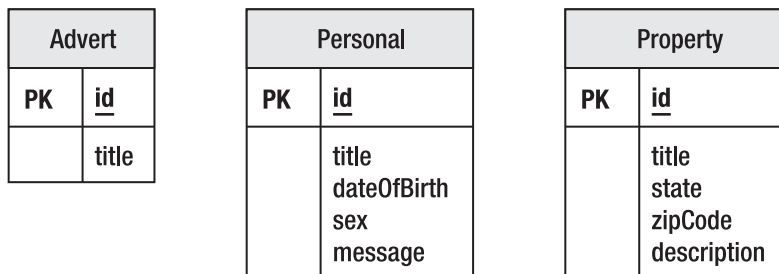
Figure 7-6. *Mapping one table per concrete class*

While this is easy to create, there are several disadvantages; the data belonging to a parent class is scattered across a number of different tables, so a query couched in terms of the parent class is likely to cause a large number of select operations. It also means that changes to a parent class can touch an awful lot of tables. We suggest that you file this approach under "quick-and-dirty solutions."

Listing 7-8 demonstrates how a derived class (Property) can be mapped to a single table independently of its superclass (Advert).

Listing 7-8. *Mapping a* Property *Advert with the One-Table-per-Concrete-Class Approach*

```
<hibernate-mapping>
    <class name="com.hibernatebook.xmlmapping.Property">
        <id name="id" type="int">
            <generator class="native"/>
        </id>
        <property name="title" type="string" length="255"/>
```

```
        <property name="state" type="string"/>
        <property name="zipCode" type="string"/>
        <property name="description" type="string"/>
    </class>
</hibernate-mapping>
```

One Table per Subclass

A slightly more complex mapping is to provide one table for each class in the hierarchy, including the abstract and interface classes. The pure "is a" relationship of our class hierarchy is then converted into a "has a" relationship for each entity in the schema.

Figure 7-7 shows the schema required to represent the hierarchy from Figure 7-5 using this technique.
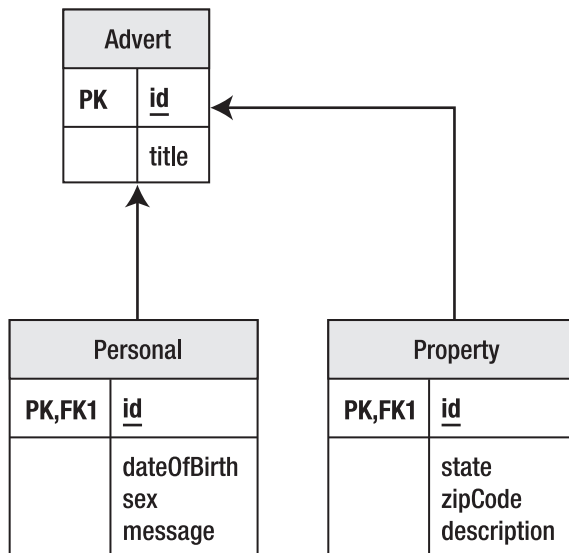
Figure 7-7. *Mapping one table per subclass*

We like this approach, as it is conceptually easy to manage, does not require complex changes to the schema when a single parent class is modified, and is similar to how most JVMs manage the same data behind the scenes.

The disadvantage of this approach is that while it works well from an object-oriented point of view, and is correct from a relational point of view, it can result in poor performance. As the hierarchy grows, the number of joins required to construct a leaf class also grows.

The technique works well for shallow inheritance hierarchies. Deep inheritance hierarchies are often a symptom of poorly designed code, so you may want to reconsider your application architecture before abandoning this technique. In our opinion, it should be preferred until performance issues are substantially proven to be an issue.

Listing 7-9 shows how you can map a derived class (Property) as a table joined to another representing the superclass (Advert).

Listing 7-9. *Mapping a* `Property` *Advert with the One-Table-per-Subclass Approach*

```
<hibernate-mapping>
   <joined-subclass
       name="com.hibernatebook.xmlmapping.Property"
       extends="com.hibernatebook.xmlmapping.Advert">

       <key column="advertid"/>

       <property name="state" type="string"/>
       <property name="zipCode" type="string"/>
       <property name="description" type="string"/>
   </joined-subclass>
</hibernate-mapping>
```

Note in the mapping that we replace `class` with `joined-subclass` to associate our mapping explicitly with the parent. You specify the entity that is being extended and replace the `id` and `title` classes from the subclass with a single key element that maps the foreign key column to the parent class table's primary key. Otherwise, the `<joined-subclass>` element is virtually identical to the `<class>` element. Note, however, that a `<joined-subclass>` cannot contain `<subclass>` elements and vice versa—the two strategies are not compatible.

One Table per Class Hierarchy

The last of the inheritance mapping strategies is to place each inheritance hierarchy in its own table. The fields from each of the child classes are added to this table, and a discriminator column contains a key to identify the base type represented by each row in the table.

Figure 7-8 shows the schema required to represent the hierarchy from Figure 7-5 using this technique.

Advert	
PK	<u>id</u>
	title advertType dateOfBirth sex message state zipCode description

Figure 7-8. *Mapping one table per hierarchy*

This technique offers the best performance—for simple queries on simple classes even in the deepest of inheritance hierarchies, a single select may suffice to gather all the fields to populate the entity.

Conversely, this is not a satisfying representation of the attribute. Changes to members of the hierarchy will usually require a column to be altered, added, or deleted from the table. This will often be a very slow operation. As the hierarchy grows (horizontally as well as vertically), so too will the number of columns required by this table.

Each mapped subclass must specify the class that it extends and a value that can be used to discriminate this subclass from the other classes held in the same table. Thus, this is known as the discriminator value, and is mapped with a `discriminator-value` attribute in the `<subclass>` element (see Listing 7-10).

Listing 7-10. *Mapping a* `Property` *Advert with the One-Table-per-Class-Hierarchy Approach*

```
<hibernate-mapping>
   <subclass
      name="com.hibernatebook.xmlmapping.Property"
      extends="com.hibernatebook.xmlmapping.Advert"
      discriminator-value="property">

      <property name="state" type="string"/>
      <property name="zipCode" type="string"/>
      <property name="description" type="string"/>
   </subclass>
</hibernate-mapping>
```

Note that this also requires the specification of a discriminator column for the root of the class hierarchy, from which the discriminator values identifying the types of the child classes can be obtained (see Listing 7-11).

Listing 7-11. *The Addition to* `Advert.hbm.xml` *Required to Support a One-Table-per-Class-Hierarchy Approach*

```
<discriminator column="advertType" type="string"/>
```

A subclass mapping cannot contain `<joined-subclass>` elements and vice versa—the two strategies are not compatible.

More Exotic Mappings

The Hibernate mapping DTD is large. We have discussed the core set of mappings that you will use on a day-to-day basis; but before we move on, we will take a very quick tour around four of the more interesting remaining mapping types.

The any Tag

The any tag represents a polymorphic association between the attribute and several entity classes. The mapping is expressed in the schema with a column to specify the type of the related entity, and then columns for the identifier of the related entity.

Because a proper foreign key cannot be specified (being dependent upon multiple tables), this is not the generally recommended technique for making polymorphic associations. When possible, use the techniques described in the previous "Mapping Inheritance Relationships" section.

The array Tag

The array tag represents the innate array feature of the Java language. The syntax of this is virtually identical to that used for the List collection class, and we recommend the use of List except when primitive values are to be stored, or when you are constrained by an existing application architecture to work with arrays.

The <dynamic-component> Element

While the full-blown dynamic class approach (discussed briefly in the "Entities" section at the beginning of the chapter) is really only suitable for prototyping exercises, the dynamic component technique allows some of that flexibility in a package that reflects some legitimate techniques.

The <dynamic-component> element permits you to place any of the items that can be mapped with the normal <component> element into a map with a given key. For example, we could obtain and combine several items of information relating to an entity's ownership into a single Map with named elements, as follows:

```
<dynamic-component name="ownership">
    <property name="user" type="string" column="user"/>
    <many-to-one
       name="person"
       class="com.hibernatebook.xmlmapping.Person"
       column="person_id"/>
</dynamic-component>
```

The code to access this information in the entity is then very familiar:

```
Map map = entity.getOwnership();
System.out.println(map.get("user"));
System.out.println(map.get("person"));
```

The output would then be as follows:

```
dcminter
person: { "Dave Minter", 33, "5'10" }
```

Summary

This chapter has covered the data types supported by Hibernate 3: entities, values, and components. You have seen how all three can be expressed in a mapping file, and how each relates to the underlying database schema. We have listed the attributes available to the major mapping elements, and we have discussed some detailed examples of the elements that you will use most frequently when working with Hibernate.

In the next chapter, we will look at how a client application communicates with the database representation of the entities by using the `Session` object.

Using the Session

You may have noticed that the `Session` object is the central point of access to Hibernate functionality. We will now look at what it embodies and what that implies about how you should use it.

Sessions

From the examples in the earlier chapters, you will have noticed that a small number of classes dominate our interactions with Hibernate. Of these, `Session` is the linchpin.

The `Session` object is used to create new database entities, read in objects from the database, update objects in the database, and delete objects from the database. It allows you to manage the transaction boundaries of database access, and (in a pinch) it allows you to obtain a traditional JDBC connection object so that you can do things to the database that the Hibernate developers have not already considered in their existing design (precious little).

If you are familiar with the JDBC approach, it helps to think of a `Session` object as being somewhat like a JDBC connection, and the `SessionFactory`, which provides `Session` objects, as being somewhat like a `ConnectionPool`, which provides `Connection` objects. These similarities in roles are illustrated in Figure 8-1.
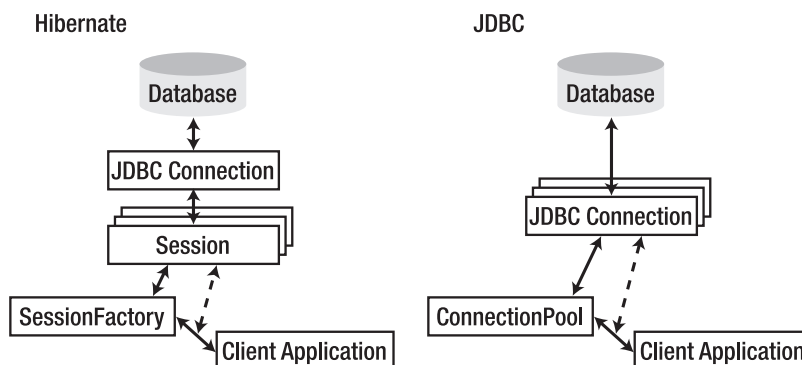
Figure 8-1. *Similarities between Hibernate and JDBC objects*

SessionFactory objects are expensive objects—needlessly duplicating them will cause problems quickly, and creating them is a relatively time-consuming process. Ideally, you should have a single SessionFactory for each database your application will access. SessionFactory objects are threadsafe, so it is not necessary to obtain one for each thread. However, you will create numerous Session objects—at least one for each thread using Hibernate. Sessions in Hibernate are *not* threadsafe, so sharing Session objects between threads could cause data loss or deadlock. In fact, you will often want to create multiple Session instances even during the lifetime of a specific thread (see the "Threads" section for concurrency issues).

■**Caution** The analogy between a Hibernate session and a JDBC connection only goes so far. One important difference is that if a Hibernate Session object throws an exception of any sort, you must discard it and obtain a new one. This prevents data in the session's cache from becoming inconsistent with the database.

We've already covered the core methods in Chapter 4, so we won't discuss all the methods available to you through the Session interface. For an exhaustive look at what's available, you should read the API documentation on the Hibernate web site or in the Hibernate 3 download. Table 8-1 gives an overview of the various categories of methods available to you.

Table 8-1. *Hibernate Method Summary*

Method	Description
Create, Read, Update, and Delete	
save()	Saves an object to the database. This should not be called for an object that has already been saved to the database.
saveOrUpdate()	Saves an object to the database, or updates the database if the object already exists. This method is slightly less efficient than the save() method since it may need to perform a SELECT statement to check whether the object already exists, but it will not fail if the object has already been saved.
merge()	Merges the fields of a nonpersistent object into the appropriate persistent object (determined by ID). If no such object exists in the database, then one is created and saved.
persist()	Reassociates an object with the session so that changes made to the object will be persisted.
get()	Retrieves a specific object from the database by the object's identifier.
getEntityName()	Retrieves the entity name (this will usually be the same as the class name of the POJO).
getIdentifier()	Determines the identifier—the object(s) representing the primary key—for a specific object associated with the session.
load()	Loads an object from the database by the object's identifier (you should use the get() methods if you are not certain that the object is in the database).
refresh()	Refreshes the state of an associated object from the database.
update()	Updates the database with changes to an object.
delete()	Deletes an object from the database.
createFilter()	Creates a filter (query) to narrow operations on the database.

Method	Description
enableFilter()	Enables a named filter in queries produced by createFilter().
disableFilter()	Disables a named filter.
getEnabledFilter()	Retrieves a currently enabled filter object.
createQuery()	Creates a Hibernate query to be applied to the database.
getNamedQuery()	Retrieves a query from the mapping file.
cancelQuery()	Cancels execution of any query currently in progress from another thread.
createCriteria()	Creates a criteria object for narrowing search results.
Transactions and Locking	
beginTransaction()	Begins a transaction.
getTransaction()	Retrieves the current transaction object. This does not return null when no transaction is in progress. Instead, the active property of the returned object is false.
lock()	Gets a database lock for an object (or can be used like persist() if LockMode.NONE is given).
Managing Resources	
contains()	Determines whether a specific object is associated with the database.
clear()	Clears the session of all loaded instances and cancels any saves, updates, or deletions that have not been completed. Retains any iterators that are in use.
evict()	Disassociates an object from the session so that subsequent changes to it will not be persisted.
flush()	Flushes all pending changes into the database—all saves, updates, and deletions will be carried out; essentially, this synchronizes the session with the database.
isOpen()	Determines whether the session has been closed.
isDirty()	Determines whether the session is synchronized with the database.
getCacheMode()	Determines the caching mode currently employed.
setCacheMode()	Changes the caching mode currently employed.
getCurrentLockMode()	Determines the locking mode currently employed.
setFlushMode()	Determines the approach to flushing currently used. The options are to flush after every operation, flush when needed, never flush, or flush only on commit.
setReadOnly()	Marks a persistent object as read-only (or as writable). There are minor performance benefits from marking an object as read-only, but changes to its state will be ignored until it is marked as writable.
close()	Closes the session, and hence, the underlying database connection; releases other resources (such as the cache). You must not perform operations on the Session object after calling close().
getSessionFactory()	Retrieves a reference to the SessionFactory object that created the current Session instance.
The JDBC Connection	
connection()	Retrieves a reference to the underlying database connection.
disconnect()	Disconnects the underlying database connection.
reconnect()	Reconnects the underlying database connection.
isConnected()	Determines whether the underlying database connection is connected.

Transactions and Locking

Transactions and locking are intimately related—the locking techniques chosen to enforce a transaction can determine both the performance and likelihood of success of the transaction. The type of transaction selected dictates, to some extent, the type of locking that it must use.

You are not obliged to use transactions if they do not suit your needs, but there is rarely a good reason to avoid them. If you decide to avoid them, you will need to invoke the flush() method on the session at appropriate points to ensure that your changes are persisted to the database.

Transactions

A transaction is a unit of work guaranteed to behave as if you have exclusive use of the database. If you wrap your work in a transaction, the behavior of other system users will not affect your data. A transaction can be started, committed to write data to the database, or rolled back to remove all changes from the beginning onward (usually as the result of an error). To achieve this, you obtain a Transaction object from the database (beginning the transaction) and manipulate the session as shown in the following code:

```
Session session = factory.openSession();
try {
  session.beginTransaction();

  // Normal session usage here...

  session.getTransaction().commit();
} catch (HibernateException e) {
  Transaction tx = session.getTransaction();
  if (tx.isActive()) tx.rollback();
} finally {
  session.close();
}
```

In the real world, it's not actually desirable for all transactions to be fully ACID (see the sidebar entitled "The ACID Tests") because of the performance problems that this can cause.

Different database suppliers support and permit you to break the ACID rules to a lesser or greater extent, but the degree of control over the isolation rule is actually mandated by the SQL-92 standard. There are important reasons that you might want to break this rule, so both JDBC and Hibernate also make explicit allowances for it.

<div>

THE ACID TESTS

- *Atomicity*: A transaction should be all or nothing. If it fails to complete, the database will be left as if none of the operations had ever been performed—this is known as a *rollback*.

- *Consistency*: A transaction should be incapable of breaking any rules defined for the database. For example, foreign keys must be obeyed. If for some reason this is impossible, the transaction will be rolled back.

- *Isolation*: The effects of the transaction will be completely invisible to all other transactions until it has completed successfully. This guarantees that the transaction will always see the data in a sensible state. For example, an update to a user's address should only contain a correct address (i.e., it will never have the house name for one location but the ZIP code for another); without this rule, a transaction could easily see when another transaction had updated the first part but had not yet completed.

- *Durability*: The data should be retained intact. If the system fails for any reason, it should always be possible to retrieve the database up to the moment of the failure.

</div>

The isolation levels permitted by JDBC and Hibernate are listed in Table 8-2.

Table 8-2. *JDBC Isolation Levels*

Level	Name	Transactional Behavior
0	None	Anything is permitted; the database or driver does not support transactions.
1	Read Uncommitted	Dirty, nonrepeatable, and phantom reads are permitted.
2	Read Committed	Nonrepeatable reads and phantom reads are permitted.
4	Repeatable Read	Phantom reads are permitted.
8	Serializable	The rule must be obeyed absolutely.

A *dirty read* may see the in-progress changes of an uncommitted transaction. As with the isolation example discussed in the preceding sidebar, it could see the wrong ZIP code for an address.

A *nonrepeatable read* sees different data for the same query. For example, it might determine a specific user's ZIP code at the beginning of the transaction and again at the end, and get a different answer both times without making any updates.

A *phantom read* sees different numbers of rows for the same query. For example, it might see 100 users in the database at the beginning of the query and 105 at the end without making any updates.

The HSQLDB database that we are using in this book only supports the first level of isolation here: Read Uncommitted. While this means that deadlocks cannot occur (see the "Deadlocks" section later in the chapter), the three undesirable behaviors of dirty, nonrepeatable, and phantom reads are permitted.

Hibernate treats the isolation as a global setting—you apply the configuration option `hibernate.connection.isolation` in the usual manner, setting it to one of the values permitted in Table 8-2. This is not always ideal. You will sometimes want to treat one particular transaction at a high level of isolation (usually Serializable), while permitting lower degrees of isolation for others. To do so, you will need to obtain the JDBC connection directly, alter the isolation level, begin the transaction, roll back or clean up the transaction as appropriate, and reset the isolation level back to its original value before releasing the connection for general usage. Hibernate does not provide a more direct way to alter the isolation level of the connection in a localized way. The implementation of the `createUser()` method, shown in Listing 8-1, demonstrates the additional complexity that the connection-specific transaction isolation involves.

Listing 8-1. *Using a Specific Isolation Level*

```
public static void createUser(String username)
    throws HibernateException
{
    Session session = factory.openSession();
    int isolation = -1;
    try {
        isolation = session.connection().getTransactionIsolation();
        session.connection().setTransactionIsolation(
            Connection.TRANSACTION_SERIALIZABLE);
        session.beginTransaction();

        // Normal usage of the Session here...
        Publisher p = new Publisher(username);
        Subscriber s = new Subscriber(username);
        session.saveOrUpdate(p);
        session.saveOrUpdate(s);

        // Commit the transaction
        session.getTransaction().commit();
    } catch (SQLException e1) {
        rollback(session);
        throw new HibernateException(e1);
    } catch (HibernateException e1) {
        rollback(session);
        throw e1;
    } finally {
        // reset isolation
        reset(session,isolation);

        // Close the session
        close(session);
    }
}
```

Fortunately, the normal case for a transaction using the global isolation level is much simpler. We provide a more standard implementation of the `createUser()` method for comparison in Listing 8-2.

Listing 8-2. *Using the Global (Default) Isolation Level*

```
public static void createUser(String username) throws HibernateException {
    Session session = factory.openSession();
    try {
        session.beginTransaction();

        // Normal usage of the Session here...
        Publisher p = new Publisher(username);
        Subscriber s = new Subscriber(username);
        session.saveOrUpdate(p);
        session.saveOrUpdate(s);

        // Commit the transaction
        session.getTransaction().commit();
    } catch (HibernateException e1) {
        rollback(session);
        throw e1;
    } finally {
        // Close the session
        close(session);
    }
}
```

Locking

A database can conform to these various levels of isolation in a number of ways, and you will need a working knowledge of locking to elicit the desired behavior and performance from your application in all circumstances.

To prevent simultaneous access to data, the database itself will acquire a lock on that data. This can be acquired for the momentary operation on the data only, or it can be retained until the end of the transaction. The former is called *optimistic locking* and the latter is called *pessimistic locking*.

The Read Uncommitted isolation level always acquires optimistic locks, whereas the Serializable isolation level will only acquire pessimistic locks. Some databases offer a feature that allows you to append the `FOR UPDATE` query to a select operation, which requires the database to acquire a pessimistic lock even in the lower isolation levels.

Hibernate provides some support for this feature when it is available, and takes it somewhat further by adding facilities that describe additional degrees of isolation obtainable from Hibernate's own cache.

The `LockMode` object controls this fine-grained isolation (see Table 8-3). It is only applicable to the `get()` methods, so it is limited—however, when possible, it is preferable to the direct control of isolation mentioned previously.

Table 8-3. *Lock Modes*

Mode	Description
NONE	Reads from the database only if the object is not available from the caches.
READ	Reads from the database regardless of the contents of the caches.
UPGRADE	Obtains a dialect-specific upgrade lock for the data to be accessed (if this is available from your database).
UPGRADE_NOWAIT	Behaves like UPGRADE, but when support is available from the database and dialect, the method will fail with a locking exception immediately. Without this option, or on databases for which it is not supported, the query must wait for a lock to be granted (or for a timeout to occur).

An additional lock mode, WRITE, is acquired by Hibernate automatically when it has written to a row within the current transaction. This mode cannot be set explicitly, but calls to getLockMode may return it.

Having discussed locking in general, we need to touch on some of the problems that locks can cause.

Deadlocks

Even if you have not encountered a deadlock (sometimes given the rather louche name of "deadly embrace") in databases, you have probably encountered the problem in multithreaded Java code. The problem arises from similar origins.

Two threads of execution can get into a situation in which each is waiting for the other to release a resource that it needs. The most common way to create this situation in a database is shown in Figure 8-2.

Each thread obtains a lock on its table when the update begins. Each thread proceeds until the table held by the other user is required. Neither thread can release the lock on its own table until the transaction completes—so something has to give.

A deadlock can also occur when a single thread of execution is carrying out an equivalent sequence of operations using two Session objects connected to the same database. In practice, the multiple-thread scenario is more common.

Fortunately, a database management system (DBMS) can detect this situation automatically, at which point the transaction of one or more of the offending processes will be aborted by the database. The resulting deadlock error will be received and handled by Hibernate as a normal HibernateException. Now you must roll back your transaction, close the session, and then (optionally) try again.

Listing 8-3 demonstrates how four updates from a pair of sessions can cause a deadlock. If you look at the output from the threads, you will see that one of them completes while the other fails with a deadlock error.

■**Caution** The HSQL database does not support sufficient levels of isolation to run this example—you will never get a deadlock, and you will get inconsistent data. If you want to see the deadlock in action, you will need to use a more sophisticated database product (e.g., PostgreSQL).

You should also be aware that if you want to run the test using the MySQL database, you must use MySQLInnoDBDialect to ensure that the appropriate level of transactions is supported.
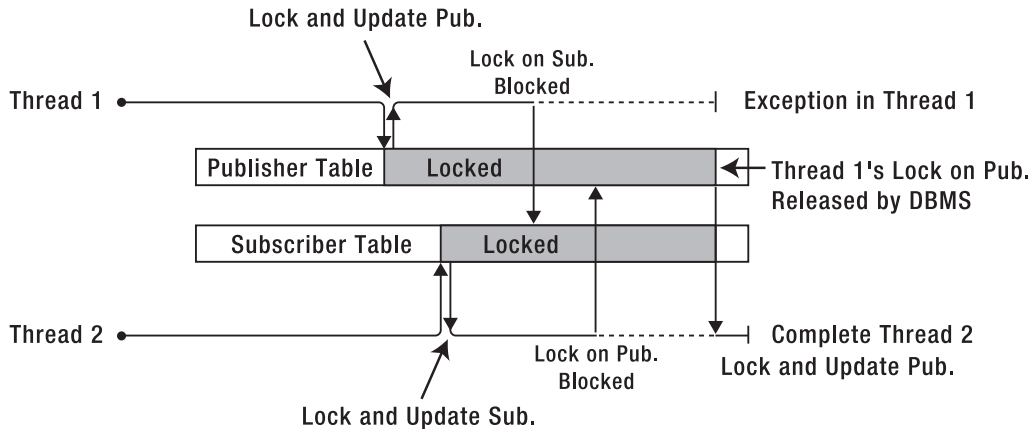
Figure 8-2. *The anatomy of a deadlock*

Looking at the database after completion, you will see that the test user has been replaced with either jeff or dave in both tables (you will never see dave from one thread and jeff from the other). Though it is not necessary here, because we close the session regardless, in a more extensive application it is important to ensure that the session associated with a deadlock or any other Hibernate exception is closed and never used again because the cache may be left in a corrupted state.

It is worth building and running Listing 8-3 to ensure that you are familiar with the symptoms of a deadlock when they occur.

Listing 8-3. *Code to Generate a Deadlock*

```
package com.hibernatebook.session.deadlock;

import java.sql.Connection;
import java.sql.SQLException;

import org.hibernate.HibernateException;
import org.hibernate.Query;
import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.Transaction;
import org.hibernate.cfg.Configuration;

public class GenerateDeadlock {

    private static SessionFactory factory = new Configuration().configure()
            .buildSessionFactory();

    public static void createUser(String username) throws HibernateException {
        Session session = factory.openSession();
        try {
            session.beginTransaction();
```

```java
        // Normal usage of the Session here...
        Publisher p = new Publisher(username);
        Subscriber s = new Subscriber(username);
        session.saveOrUpdate(p);
        session.saveOrUpdate(s);

        // Commit the transaction
        session.getTransaction().commit();
    } catch (HibernateException e1) {
        rollback(session);
        throw e1;
    } finally {
        // Close the session
        close(session);
    }
}

public static void reset(Session session, int isolation) {
    if (isolation >= 0) {
        try {
            session.connection().setTransactionIsolation(isolation);
        } catch (SQLException e) {
            System.err.println("Could not reset the isolation level: " + e);
        } catch (HibernateException e) {
            System.err.println("Could not reset the isolation level: " + e);
        }
    }
}

public static void close(Session session) {
    try {
        session.close();
    } catch (HibernateException e) {
        System.err.println("Could not close the session: " + e);
    }
}

public static void rollback(Session session) {
    try {
        Transaction tx = session.getTransaction();
        if (tx.isActive())
            tx.rollback();
    } catch (HibernateException e) {
        System.err.println("Could not rollback the session: " + e);
    }
}
```

```java
public static void main(String[] argv) {

    System.out.println("Creating test user...");
    createUser("test");

    System.out.println("Proceeding to main test...");
    Session s1 = factory.openSession();
    Session s2 = factory.openSession();

    try {
        s1.beginTransaction();
        s2.beginTransaction();

        System.out.println("Update 1");
        Query q1 = s1.createQuery("from Publisher");
        Publisher pub1 = (Publisher) q1.uniqueResult();
        pub1.setUsername("jeff");
        s1.flush();

        System.out.println("Update 2");
        Query q2 = s2.createQuery("from Subscriber");
        Subscriber sub1 = (Subscriber) q2.uniqueResult();
        sub1.setUsername("dave");
        s2.flush();

        System.out.println("Update 3");
        Query q3 = s1.createQuery("from Subscriber");
        Subscriber sub2 = (Subscriber) q3.uniqueResult();
        sub2.setUsername("jeff");
        s1.flush();

        System.out.println("Update 4");
        Query q4 = s2.createQuery("from Publisher");
        Publisher pub2 = (Publisher) q4.uniqueResult();
        pub2.setUsername("dave");
        s2.flush();

        s1.getTransaction().commit();
        s2.getTransaction().commit();

    } catch (RuntimeException e1) {
        e1.printStackTrace();
        // Run the boilerplate to roll back the sessions
        rollback(s1);
        rollback(s2);
        throw e1;
    } finally {
```

```
        // Run the boilerplate to close the sessions
        close(s1);
        close(s2);
      }
    }
}
```

Caching

Accessing a database is an expensive operation, even for a simple query. The request has to be sent (usually over the network) to the server. The database server may have to compile the SQL into a query plan. The query plan has to be run and is limited largely by disk performance. The resulting data has to be shuttled back (again, usually across the network) to the client, and only then can the application program begin to process the results.

Most good databases will cache the results of a query if it is run multiple times, eliminating the disk I/O and query compilation time; but this will be of limited value if there are large numbers of clients making substantially different requests. Even if the cache generally holds the results, the time taken to transmit the information across the network is often the larger part of the delay.

Some applications will be able to take advantage of in-process databases, but this is the exception rather than the rule—and such databases have their own limitations.

The natural and obvious answer is to have a cache at the client end of the database connection. This is not a feature provided or supported by JDBC directly, but Hibernate provides one cache (the first-level, or L1, cache) through which all requests must pass. A second-level cache is optional and configurable.

The L1 cache ensures that within a session, requests for a given object from a database will always return the same object instance, thus preventing data from conflicting and preventing Hibernate from trying to load an object multiple times.

Items in the L1 cache can be individually discarded by invoking the `evict()` method on the session for the object that you wish to discard. To discard all items in the L1 cache, invoke the `clear()` method.

In this way, Hibernate has a major advantage over the traditional JDBC approach: with no additional effort from the developer, a Hibernate application gains the benefits of a client-side database cache.

Figure 8-3 shows the two caches available to the session: the compulsory L1 cache, through which all requests must pass, and the optional level-two (L2) cache. The L1 cache will always be consulted before any attempt is made to locate an object in the L2 cache. You will notice that the L2 cache is external to Hibernate; and although it is accessed via the session in a way that is transparent to Hibernate users, it is a pluggable interface to any one of a variety of caches that are maintained on the same JVM as your Hibernate application or on an external JVM. This allows a cache to be shared between applications on the same machine, or even between multiple applications on multiple machines.

In principle, any third-party cache can be used with Hibernate. An `org.hibernate.cache.CacheProvider` interface is provided, which must be implemented to provide Hibernate with a handle to the cache implementation. The cache provider is then specified by giving the implementation class name as the value of the `hibernate.cache.provider_class` property.
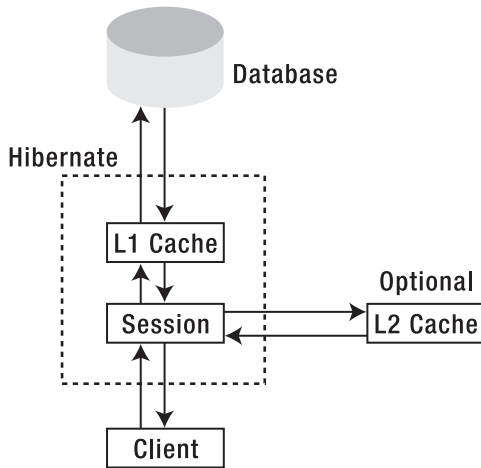
Figure 8-3. *The session's relationship to the caches*

In practice, the four production-ready caches, which are already supported, will be adequate for most users (see Table 8-4).

Table 8-4. *L2 Cache Implementations Supported by Hibernate Out of the Box*

Cache Name	Description
EHCache	An in-process cache
OSCache	An alternative in-process cache
SwarmCache	A multicast distributed cache
TreeCache	A multicast distributed transactional cache

The type of access to the L2 cache can be configured on a per-session basis by selecting a CacheMode option (see Table 8-5) and applying it with the setCacheMode() method.

Table 8-5. CacheMode *Options*

Mode	Description
NORMAL	Data is read from and written to the cache as necessary.
GET	Data is never added to the cache (although cache entries are invalidated when updated by the session).
PUT	Data is never read from the cache, but cache entries will be updated as they are read from the database by the session.
REFRESH	This is the same as PUT, but the use_minimal_puts Hibernate configuration option will be ignored if it has been set.
IGNORE	Data is never read from or written to the cache (except that cache entries will still be invalidated when they are updated by the session).

The CacheMode setting does not affect the way in which the L1 cache is accessed.

The decision to use an L2 cache is not clear-cut. Although it has the potential to greatly reduce access to the database, the benefits depend on the type of cache and the way in which it will be accessed.

A distributed cache will cause additional network traffic. Some types of database access may result in the contents of the cache being flushed before they are used—in which case, it will be adding unnecessary overhead to the transactions.

The L2 cache cannot account for the changes in the underlying data, which are the result of actions by an external program that is not cache-aware. This could potentially lead to problems with stale data, which is not an issue with the L1 cache.

In practice, as with most optimization problems, it is best to carry out performance testing under realistic load conditions. This will let you determine if a cache is necessary and help you select which one will offer the greatest improvement.

Threads

Having considered the caches available to a Hibernate application, you may now be concerned about the risk of a conventional Java deadlock if two threads of execution were to contend for the same object in the Hibernate session cache.

In principle, this is possible—and unlike database deadlocks, Java thread deadlocks do not time out with an error message. Fortunately, there is a very simple solution:

Patient: Doctor, it hurts when I do this.

Doctor: Don't do that then.

Do not share the Session object between threads. This will eliminate any risk of deadlocking on objects contained within the session cache.

The easiest way to ensure that you do not use the same Session object outside the current thread is to use an instance local to the current method. If you absolutely must maintain an instance for a longer duration, maintain the instance within a ThreadLocal object. For most purposes, however, the lightweight nature of the Session object makes it practical to construct, use, and destroy an instance, rather than to store a session.

Summary

In this chapter, we have discussed the nature of Session objects and how they can be used to obtain and manage transactions. We have looked at the two levels of caching that are available to applications, and how concurrent threads should manage sessions.

In the next chapter, we discuss the various ways in which you can retrieve objects from the database. We also show you how to perform more complicated queries against the database using HQL.

■ ■ ■

Searches and Queries

In the last chapter, we discussed how the Hibernate session is used to interact with the database. Some of the session's methods take query strings in their parameter lists or return `Query` objects. These methods are used to request arbitrary information from the database. In order to fully show how they're used, we must introduce you to the HQL used to phrase these requests. As well as extracting information (with `SELECT`), HQL can be used to alter the information in the database (with `INSERT`, `UPDATE`, and `DELETE`). We cover all of this basic functionality in this chapter. Hibernate's query facilities do not allow you to alter the database structure.

HQL is an object-oriented query language, similar to SQL, but instead of operating on tables and columns, HQL works with persistent objects and their properties.

HQL is a language with its own syntax and grammar. HQL is written as strings, like `from Product p`, as opposed to Hibernate's criteria queries (discussed in the next chapter), which take the form of a conventional Java API. Ultimately, your HQL queries are translated by Hibernate into conventional SQL queries, and Hibernate also provides an API that allows you to directly issue SQL queries.

HQL

While most ORM tools and object databases offer an object query language, Hibernate's HQL stands out as being complete and easy to use. Although you can use SQL statements directly with Hibernate (which is covered in detail in the "Using Native SQL" section of this chapter), we recommend that you use HQL (or criteria) whenever possible to avoid database portability hassles, and to take advantage of Hibernate's SQL-generation and caching strategies. In addition to its technical advantages over traditional SQL, HQL is a more compact query language than SQL because it can make use of the relationship information defined in the Hibernate mappings.

We realize that not every developer trusts Hibernate's generated SQL to be perfectly optimized. If you do encounter a performance bottleneck in your queries, we recommend that you use SQL tracing on your database during performance testing of your critical components. If you see an area that needs optimization, we suggest trying first to optimize using HQL, and only later dropping into native SQL. Hibernate 3 provides statistics information through a JMX MBean, which you can use for analyzing Hibernate's performance. Hibernate's statistics also give you insight into how caching is performing.

■**Note** If you would like to execute HQL statements through a GUI-based tool, the Hibernate team provides a Hibernate console for Eclipse in the Hibernate Tools subproject. This console is a plug-in for recent versions of Eclipse. This tool is described in detail in Appendix B.

Syntax Basics

HQL is inspired by SQL and is the inspiration for the new EJB Query Language (EJB QL). The EJB QL specification is included in the standard for EJB 3 available from the Java Community Process web site (`www.jcp.org/en/jsr/detail?id=220`). HQL's syntax is defined as an ANTLR grammar; the grammar files are included in the `grammar` directory of the Hibernate core download (ANTLR is a tool for building language parsers).

As the ANTLR grammar files are somewhat cryptic, and as not every statement that is permissible according to the ANTLR grammar's rules can be used in Hibernate, we outline the syntax for the four fundamental HQL operations in this section. Note that the following descriptions of syntax are not comprehensive—there are some deprecated or more obscure usages (particularly for `SELECT` statements) that are not covered here.

UPDATE

`UPDATE` alters the details of existing objects in the database. In-memory entities will not be updated to reflect changes resulting from issuing `UPDATE` statements. Here's the syntax of the `UPDATE` statement:

```
UPDATE [VERSIONED]
   [FROM] path [[AS] alias] [, ...]
   SET property = value [, ...]
   [WHERE logicalExpression]
```

`path` is the fully qualified name of the entity or entities. The `alias` names may be used to abbreviate references to specific entities or their properties, and must be used when property names used in the query would otherwise be ambiguous.

The `property` names are the names of properties of entities listed in the `FROM` path.

The syntax of logical expressions is discussed later, in the "Using Restrictions with HQL" section.

DELETE

`DELETE` removes the details of existing objects from the database. In-memory entities will not be updated to reflect changes resulting from `DELETE` statements. This also means that cascade rules will not be followed for deletions carried out using HQL. This approach to deletion is commonly referred to as "bulk deletion" since it is the most efficient way to remove large numbers of entities from the database. Here's the syntax of the `DELETE` statement:

```
DELETE
    [FROM] path [[AS] alias]
    [WHERE logicalExpression]
```

   path is the fully qualified name of an entity. The alias names may be used to abbreviate references to specific entities or their properties, and must be used when property names used in the query would otherwise be ambiguous.

INSERT

An HQL INSERT cannot be used to directly insert arbitrary entities—it can only be used to insert entities constructed from information obtained from SELECT queries (unlike ordinary SQL, in which an INSERT command can be used to insert arbitrary data into a table, as well as insert values selected from other tables). Here's the syntax of the INSERT statement:

```
INSERT
    INTO path ( property [, ...])
    select
```

   path is the fully qualified name of an entity. The property names are the names of properties of entities listed in the FROM path of the incorporated SELECT query.
   The select query is an HQL SELECT query (as described in the next section).

SELECT

An HQL SELECT is used to query the database for classes and their properties. As noted previously, this is very much a summary of the full expressive power of HQL SELECT queries—however, for more complex joins and the like, you may find that the use of the Criteria API described in the next chapter is more appropriate. Here's the syntax of the SELECT statement:

```
[SELECT [DISTINCT] property [, ...]]
    FROM path [[AS] alias] [, ...] [FETCH ALL PROPERTIES]
    WHERE logicalExpression
    GROUP BY property [, ...]
    HAVING logicalExpression
    ORDER BY property [ASC | DESC] [, ...]
```

   path is the fully qualified name of an entity. The alias names may be used to abbreviate references to specific entities or their properties, and must be used when property names used in the query would otherwise be ambiguous.
   The property names are the names of properties of entities listed in the FROM path.
   If FETCH ALL PROPERTIES is used, then lazy loading semantics will be ignored, and all the immediate properties of the retrieved object(s) will be actively loaded (this does not apply recursively).
   When the properties listed consist only of the names of aliases in the FROM clause, the SELECT clause can be omitted.

The First Example with HQL

The simplest HQL query returns all objects for a given class in the database. In a syntax similar to that of SQL, we use the HQL clause `from`. As noted, when retrieving objects with HQL, you do not have to use the leading `select` clause for this query—instead, you can use the following simple shortcut query to select all objects from the `Product` table:

```
from Product
```

■**Note** Like all SQL syntax, you can write `from` in lowercase or uppercase (or mixed case). However, any Java classes or properties that you reference in an HQL query have to be specified in the proper case. For example, when you query for instances of a Java class named `Product`, the HQL query `from Product` is the equivalent of `FROM Product`. However, the HQL query `from product` is not the same as the HQL query `from Product`. Because Java class names are case-sensitive, Hibernate is case-sensitive about class names as well.

Embedding the following HQL statement into our application is straightforward. The `org.hibernate.Session` object contains a method named `createQuery()`:

```
public Query createQuery(String queryString) throws HibernateException
```

The `createQuery()` method takes a valid HQL statement, and returns an `org.hibernate.Query` object. The `Query` class provides methods for returning the query results as a Java `List`, as an `Iterator`, or as a unique result. Other functionality includes named parameters, results scrolling, JDBC fetch sizes, and JDBC timeouts. You can also add a comment to the SQL that Hibernate creates, which is useful for tracing which HQL statements correspond to which SQL statements.

In order to fully illustrate our examples, we must first introduce the sample application that we are using in this chapter and the next (which discusses criteria). The sample application has three classes: `Supplier`, `Product`, and `Software`. The `Supplier` class, shown in Listing 9-1, has a `name` property and a `List` collection of `Product` objects.

Listing 9-1. *The* `Supplier` *Class*

```java
package com.hibernatebook.queries;

import java.util.ArrayList;
import java.util.List;

public class Supplier
{
    private int id;
    private String name;
    private List products = new ArrayList();
```

```
    public int getId() {
        return id;
    }

    public void setId(int id) {
        this.id = id;
    }

    public String getName() {
        return name;
    }

    public void setName(String name) {
        this.name = name;
    }

    public List getProducts() {
        return products;
    }

    public void setProducts(List products) {
        this.products = products;
    }
}
```

The Product class, shown in Listing 9-2, has name, price, and description properties, along with a reference to its parent supplier.

Listing 9-2. *The* Product *Class*

```
package com.hibernatebook.queries;

public class Product
{
    private int id;
    private Supplier supplier;

    private String name;
    private String description;
    private double price;

    public Product() {
    }
```

```java
    public Product(String name, String description, double price) {
        this.name = name;
        this.description = description;
        this.price = price;
    }

    public String getDescription() {
        return description;
    }

    public void setDescription(String description) {
        this.description = description;
    }

    public int getId() {
        return id;
    }

    public void setId(int id) {
        this.id = id;
    }

    public String getName() {
        return name;
    }

    public void setName(String name) {
        this.name = name;
    }

    public Supplier getSupplier() {
        return supplier;
    }

    public void setSupplier(Supplier supplier) {
        this.supplier = supplier;
    }

    public double getPrice() {
        return price;
    }

    public void setPrice(double price) {
        this.price = price;
    }
}
```

The Software class, shown in Listing 9-3, extends the Product class and adds a version property—we added this subclass so that we could demonstrate polymorphism with Hibernate's queries.

Listing 9-3. *The* Software *Class*

```
package com.hibernatebook.queries;

public class Software extends Product
{
    private String version;

    public Software() {
    }

    public Software(String name, String description,
                    double price, String version)
    {
        super(name, description, price);
        this.setVersion(version);
    }

    public String getVersion() {
        return version;
    }

    public void setVersion(String version) {
        this.version = version;
    }
}
```

The Hibernate mapping files for these three classes are in the source directory for the book, along with a test harness for populating the database and running the examples in this chapter and the next.

The first example executes our HQL statement, from Product, and then retrieves a List of Product objects.

```
Query query = session.createQuery("from Product");
List results = query.list();
```

Many of the other examples in this chapter use the same supporting Java code as this example. We are going to provide just the HQL for these examples—you can execute them the same way we did here, substituting that HQL for the from Product HQL statement. This should make each example clearer as to what you should be looking at. You could also execute these HQL statements in the Hibernate Tools scratch pad.

Logging the Underlying SQL

Hibernate can output the underlying SQL behind your HQL queries into your application's log file. This is especially useful if the HQL query does not give the results you expect, or the query takes longer than you wanted. You can run the SQL that Hibernate generates directly against your database in the database's query analyzer at a later date to determine the causes of the problem. This is not a feature you will have to use frequently, but it is useful should you need to turn to your database administrators for help in tuning your Hibernate application.

The easiest way to see the SQL for a Hibernate HQL query is to enable SQL output in the logs with the `hibernate.show_sql` property. Set this property to `true` in your `hibernate.properties` or `hibernate.cfg.xml` configuration files, and Hibernate will output the SQL into the logs. You do not need to enable any other logging settings—although setting logging for Hibernate to `debug` also outputs the generated SQL statements, along with a lot of other verbiage.

After enabling SQL output in Hibernate, you should rerun the previous example. Here is the generated SQL statement for the HQL statement `from Product`:

```
select product0_.id as id, product0_.name as name0_, product0_.description ➥
as descript3_0_, product0_.price as price0_, product0_.supplierId ➥
as supplierId0_, product0_1_.version as version1_, ➥
case when product0_1_.productId is not null then 1 ➥
when product0_.id is not null then 0 end ➥
as clazz_ from Product product0_ left outer join Software product0_1_ ➥
on product0_.id=product0_1_.productId
```

As an aside, remember that the `Software` class inherits from `Product`, which complicates Hibernate's generated SQL for this simple query. When we select all objects from our simple `Supplier` class, the generated SQL for the HQL query `from Supplier` is much simpler:

```
select supplier0_.id as id, supplier0_.name as name2_ from Supplier supplier0_
```

When you look in your application's output for the Hibernate SQL statements, they will be prefixed with `Hibernate:`. The previous SQL statement would look like this:

```
Hibernate: select supplier0_.id as id, supplier0_.name as name2_ ➥
from Supplier supplier0_
```

If you turn your log4j logging up to `debug` for the Hibernate classes, you will see SQL statements in your log files, along with lots of information about how Hibernate parsed your HQL query and translated it into SQL.

Commenting the Generated SQL

Tracing your HQL statements through to the generated SQL can be difficult, so Hibernate provides a commenting facility on the `Query` object that lets you apply a comment to a specific query. The `Query` interface has a `setComment()` method that takes a `String` object as an argument, as follows:

```
public Query setComment(String comment)
```

Use this to identify the SQL output in your application's logs if SQL logging is enabled. For instance, if we add a comment to this example, the Java code would look like this:

```
String hql = "from Supplier";
Query query = session.createQuery(hql);
query.setComment("My HQL: " + hql);
List results = query.list();
```

The output in your application's log will have the comment in a Java-style comment before the SQL:

```
Hibernate: /*My HQL: from Supplier*/ select supplier0_.id as id, supplier0_.name ➡
as name2_ from Supplier supplier0_
```

We have found this useful for identifying SQL in our logs, especially because the generated SQL is a little difficult to follow when you are scanning large quantities of it in logs.

The from Clause and Aliases

We have already discussed the basics of the `from` clause in HQL in the earlier section, "The First Example with HQL." The most important feature to note is the *alias*. Hibernate allows you to assign aliases to the classes in your query with the `as` clause. Use the aliases to refer back to the class inside the query. For instance, our previous simple example would be the following:

```
from Product as p
```

or the following:

```
from Product as product
```

You'll see either alias-naming convention in applications. The `as` keyword is optional—you can also specify the alias directly after the class name, as follows:

```
from Product product
```

If you need to fully qualify a class name in HQL, just specify the package and class name. Hibernate will take care of most of this behind the scenes, so you only really need this if you have classes with duplicate names in your application. If you need to do this in Hibernate, use syntax such as the following:

```
from com.hibernatebook.criteria.Product
```

The `from` clause is very basic and useful for working directly with objects. However, if you want to work with the object's properties without loading the full objects into memory, you must use the `select` clause.

The select Clause and Projection

The `select` clause provides more control over the result set than the `from` clause. If you want to obtain the properties of objects in the result set, use the `select` clause. For instance, we could

run a projection query on the products in the database that only returned the names, instead of loading the full object into memory, as follows:

```
select product.name from Product product
```

The result set for this query will contain a `List` of Java `String` objects. Additionally, we can retrieve the prices and the names for each product in the database, like so:

```
select product.name, product.price from Product product
```

This result set contains a `List` of `Object` arrays—each array represents one set of properties (in this case, a name and price pair).

If you're only interested in a few properties, this approach can allow you to reduce network traffic to the database server and save memory on the application's machine.

Using Restrictions with HQL

As with SQL, you use the `where` clause to select results that match your query's expressions. HQL provides many different expressions that you can use to construct a query. In the HQL language grammar, there are the following possible expressions:

- *Logic operators*: `OR`, `AND`, `NOT`

- *Equality operators*: `=`, `<>`, `!=`, `^=`

- *Comparison operators*: `<`, `>`, `<=`, `>=`, `like`, `not like`, `between`, `not between`

- *Math operators*: `+`, `-`, `*`, `/`

- *Concatenation operator*: `||`

- *Cases*: Case when `<logical expression>` then `<unary expression>` else `_<unary expression>` end

- *Collection expressions*: `some`, `exists`, `all`, `any`

In addition, you may also use the following expressions in the `where` clause:

- *HQL named parameters*: `:date`, `:quantity`

- *JDBC query parameter*: `?`

- *Date and time SQL-92 functional operators*: `current_time()`, `current_date()`, `current_timestamp()`

- *SQL functions (supported by the database)*: `length()`, `upper()`, `lower()`, `ltrim()`, `rtrim()`, etc.

Using these restrictions, you can build a `where` clause in HQL that is as powerful as an SQL query. For many queries, HQL syntax is more compact and elegant than the Criteria Query API syntax (discussed in Chapter 10). For instance, here is an example of a criteria query that uses logical expressions:

```
Criteria crit = session.createCriteria(Product.class);
Criterion price = Restrictions.gt("price",new Double(25.0));
Criterion name = Restrictions.like("name","Mou%");
LogicalExpression orExp = Restrictions.or(price,name);
crit.add(orExp);
crit.add(Restrictions.ilike("description","blocks%"));
List results = crit.list();
```

The equivalent HQL would be the following:

```
from Product where price > 25.0 and name like 'Mou%'
```

We would have to wrap that HQL in a couple of lines of Java code, but even so, we find this particular example to be clearer in HQL. In the previous HQL example, you can see that we used the `where` clause with a `>` (greater than) comparison operator, an `and` logical operator, and a `like` comparison operator. You do have to enclose literal strings in quotes in HQL. To find names that have the literal `Mou` at the beginning of the string, we used `%` in the query.

Using Named Parameters

Hibernate supports named parameters in its HQL queries. This makes writing queries that accept input from the user easy—and you do not have to defend against SQL injection attacks.

■**Note** SQL injection is an attack against applications that create SQL directly from user input with string concatenation. For instance, if we accept a name from the user through a web application form, then it would be very bad form to construct an SQL (or HQL) query like this:

```
String sql = "select p from products where name = '" + name + "'";
```

A malicious user could pass a name to the application that contained a terminating quote and semicolon, followed by another SQL command (such as `delete from products`) that would let them do whatever they wanted. They would just need to end with another command that matched the SQL statement's ending quote. This is a very common attack, especially if the malicious user can guess details of your database structure.

You could escape the user's input yourself for every query, but it is much less of a security risk if you let Hibernate manage all of your input with named parameters. Hibernate's named parameters are similar to the JDBC query parameters (?) you may already be familiar with, but Hibernate's parameters are less confusing. It is also more straightforward to use Hibernate's named parameters if you have a query that uses the same parameter in multiple places.

When you use JDBC query parameters, any time you add, change, or delete parts of the SQL statement, you need to update your Java code that sets its parameters, because the parameters are indexed based on the order they appear in the statement. Hibernate lets you provide names for the parameters in the HQL query, so you do not have to worry about accidentally moving parameters further up or back in the query.

The simplest example of named parameters uses regular SQL types for the parameters:

```
String hql = "from Product where price > :price";
Query query = session.createQuery(hql);
query.setDouble("price",25.0);
List results = query.list();
```

Normally, you do not know the values that are to be substituted for the named parameters—and if you did, you would probably encode them directly into the query string. When the value to be provided will be known only at run time, you can use some of HQL's object-oriented features to provide objects as values for named parameters. The Query interface has a setEntity() method that takes the name of a parameter and an object. Using this functionality, we could retrieve all the products that have a supplier whose object we already have:

```
String supplierHQL = "from Supplier where name='MegaInc'";
Query supplierQuery = session.createQuery(supplierHQL);
Supplier supplier = (Supplier) supplierQuery.list().get(0);

String hql = "from Product as product where product.supplier=:supplier";
Query query = session.createQuery(hql);
query.setEntity("supplier",supplier);
List results = query.list();
```

You can also use regular JDBC query parameters in your HQL queries. We do not particularly see any reason why you would want to, but they do work.

Paging Through the Result Set

Pagination through the result set of a database query is a very common application pattern. Typically, you would use pagination for a web application that returned a large set of data for a query. The web application would page through the database query result set to build the appropriate page for the user. The application would be very slow if the web application loaded all of the data into memory for each user. Instead, you can page through the result set and retrieve the results you are going to display one chunk at a time.

There are two methods on the Query interface for paging: setFirstResult() and setMaxResults(), just as with the Criteria interface. The setFirstResult() method takes an integer that represents the first row in your result set, starting with row 0. You can tell Hibernate to only retrieve a fixed number of objects with the setMaxResults() method. Your HQL is unchanged—you only need to modify the Java code that executes the query. Excuse our tiny dataset for this trivial example of pagination:

```
Query query = session.createQuery("from Product");
query.setFirstResult(1);
query.setMaxResults(2);
List results = query.list();
displayProductsList(results);
```

You can change the numbers around and play with the pagination. If you turn on SQL logging, you can see which SQL commands Hibernate uses for pagination. For the open source HSQLDB database, Hibernate uses top and limit.

If you only have one result in your HQL result set, Hibernate has a shortcut method for obtaining just that object.

Obtaining a Unique Result

HQL's `Query` interface provides a `uniqueResult()` method for obtaining just one object from an HQL query. Although your query may only yield one object, you may also use the `uniqueResult()` method with other result sets if you limit the results to just the first result. You could use the `setMaxResults()` method discussed in the previous section. The `uniqueResult()` method on the `Query` object returns a single object, or `null` if there are zero results. If there is more than one result, the `uniqueResult()` method throws a `NonUniqueResultException`.

The following short example demonstrates having a result set that would have included more than one result, except that it was limited with the `setMaxResults()` method:

```
String hql = "from Product where price>25.0";
Query query = session.createQuery(hql);
query.setMaxResults(1);
Product product = (Product) query.uniqueResult();
//test for null here if needed
```

Unless your query returns one or zero results, the `uniqueResult()` method will throw a `NonUniqueResultException` exception. Do not expect Hibernate just to pick off the first result and return it—either set the maximum results of the HQL query to 1, or obtain the first object from the result list.

Sorting Results with the order by Clause

To sort your HQL query's results, you will need to use the `order by` clause. You can order the results by any property on the objects in the result set: either ascending (`asc`) or descending (`desc`). You can use ordering on more than one property in the query if you need to. A typical HQL query for sorting results looks like this:

```
from Product p where p.price>25.0 order by p.price desc
```

If you wanted to sort by more than one property, you would just add the additional properties to the end of the `order by` clause, separated by commas. For instance, you could sort by product price and the supplier's name, as follows:

```
from Product p order by p.supplier.name asc, p.price asc
```

HQL is more straightforward for ordering than the equivalent approach using the Criteria Query API.

Associations

Associations allow you to use more than one class in an HQL query, just as SQL allows you to use joins between tables in a relational database. Add an association to an HQL query with the `join` clause. Hibernate supports five different types of joins: `inner join`, `cross join`, `left outer`

join, right outer join, and full outer join. If you use cross join, just specify both classes in the from clause (from Product p, Supplier s). For the other joins, use a join clause after the from clause. Specify the type of join, the object property to join on, and an alias for the other class.

You can use inner join to obtain the supplier for each product, and then retrieve the supplier name, product name, and product price, as so:

```
select s.name, p.name, p.price from Product p inner join p.supplier as s
```

You can retrieve the objects using similar syntax:

```
from Product p inner join p.supplier as s
```

We used aliases in these HQL statements to refer to the entities in our query expressions. These are particularly important in queries with associations that refer to two different entities with the same class—for instance, if we are doing a join from a table back to itself. Commonly, these types of joins are used to organize tree data structures.

Notice that Hibernate does not return Object objects in the result set; instead, Hibernate returns Object arrays in the results. You will have to access the contents of the Object arrays to get the Supplier and the Product objects.

If you would like to start optimizing performance, you can ask Hibernate to fetch the associated objects and collections for an object in one query. If you were using lazy loading with Hibernate, the objects in the collection would not be initialized until you accessed them. If you use fetch on a join in your query, you can ask Hibernate to retrieve the objects in the collection at the time the query executes. Add the fetch keyword after the join in the query, like so:

```
from Supplier s inner join fetch s.products as p
```

When you use fetch for a query like this, Hibernate will return only the Supplier objects, not the Product objects. This is because you are specifying the join, so Hibernate knows which objects to fetch (instead of using lazy loading). If you need to get the Product objects, you can access them through the associated Supplier object. You cannot use the properties of the Product objects in expressions in the where clause. Use of the fetch keyword overrides any settings you have in the mapping file for object initialization.

Aggregate Methods

HQL supports a range of aggregate methods, similar to SQL. They work the same way in HQL as in SQL, so you do not have to learn any specific Hibernate terminology. The difference is that in HQL, aggregate methods apply to the properties of persistent objects. The count(...) method returns the number of times the given column name appears in the result set. You may use the count(*) syntax to count all the objects in the result set, or count(product.name) to count the number of objects in the result set with a name property. Here is an example using the count(*) method to count all products:

```
select count(*) from Product product
```

The distinct keyword only counts the unique values in the row set—for instance, if there are 100 products, but 10 have the same price as another product in the results, then a select count(distinct product.price) from Product product query would return 90. In our database, the following query will return 2, one for each supplier:

```
select count(distinct product.supplier.name) from Product product
```

If we removed the distinct keyword, it would return 5, one for each product.

All of these queries return an Integer object in the list. You could use the uniqueResult() method here to obtain the result.

The aggregate functions available through HQL include the following:

- avg(*property name*): The average of a property's value

- count(*property name* or *): The number of times a property occurs in the results

- max(*property name*): The maximum value of the property values

- min(*property name*): The minimum value of the property values

- sum(*property name*): The sum total of the property values

If you have more than one aggregate method, the result set List will contain an Object array with each of the aggregates you requested. Adding another aggregate to the select clause is straightforward:

```
select min(product.price), max(product.price) from Product prodsuct
```

You can also combine these with other projection properties in the result set.

Bulk Updates and Deletes with HQL

Bulk updates are new to HQL with Hibernate 3, and deletes work differently in Hibernate 3 than they did in Hibernate 2. The Query interface now contains a method called executeUpdate() for executing HQL UPDATE or DELETE statements. The executeUpdate() method returns an int that contains the number of rows affected by the update or delete, as follows:

```
public int executeUpdate() throws HibernateException
```

HQL updates look like you would expect them to, being based on SQL UPDATE statements. Do not include an alias with the update; instead, put the set keyword right after the class name, as follows:

```
String hql = "update Supplier set name = :newName ➥
where name = :name";

Query query = session.createQuery(hql);
query.setString("name","SuperCorp");
query.setString("newName","MegaCorp");
int rowCount = query.executeUpdate();
System.out.println("Rows affected: " + rowCount);
```

```
//See the results of the update
query = session.createQuery("from Supplier");
List results = query.list();
```

After carrying out this query, any supplier previously named SuperCorp will be named MegaCorp. You may use a where clause with updates to control which rows get updated, or you may leave it off to update all rows. Notice that we printed out the number of rows affected by the query. We also used named parameters in our HQL for this bulk update.

Bulk deletes work in a similar way. Use the delete from clause with the class name you would like to delete from. Then use the where clause to narrow down which entries in the table you would like to delete. Use the executeUpdate() method to execute deletes against the database as well. Our code surrounding the HQL DELETE statement is basically the same—we use named parameters, and we print out the number of rows affected by the delete:

```
String hql = "delete from Product where name = :name";
Query query = session.createQuery(hql);
query.setString("name","Mouse");
int rowCount = query.executeUpdate();
System.out.println("Rows affected: " + rowCount);

//See the results of the update
query = session.createQuery("from Product");
List results = query.list();
```

■**Caution** Using bulk updates and deletes in HQL works almost the same as in SQL, so keep in mind that these are powerful and can erase the data in your tables if you make a mistake with the where clause.

Named Queries for HQL and SQL

One of Hibernate's best features is the named query, in which your application can store its HQL queries outside the application in the mapping file. This has many benefits for application maintenance. The first benefit is that many objects can share queries—you could set up static final strings on classes with the HQL queries, but Hibernate already provides a nice facility for the same thing. The next benefit is that named queries could also contain native SQL queries—the application calling the named query does not need to know if the named query is SQL or HQL. This has enormous benefits for migrating SQL-based applications to Hibernate. The last benefit is that you can provide your HQL and SQL queries in a configuration file to your database administrators, who will probably find it easier to work with an XML mapping file than with HQL statements embedded in Java code.

Add named queries in the appropriate Hibernate mapping file. HQL queries use the XML <query> element, and SQL queries use the XML <sql-query> element. Both of these XML elements require a name attribute that uniquely identifies the query in the application. With one

simple HQL named query, and one simple SQL query that does the same thing, we have the Hibernate mapping file shown in Listing 9-4: `Product.hbm.xml`.

Listing 9-4. `Product.hbm.xml`

```xml
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping package="com.hibernatebook.criteria">
    <class name="Product">
        <id name="id" type="int">
            <generator class="native"/>
        </id>

        <property name="name" type="string"/>
        <property name="description" type="string"/>
        <property name="price" type="double"/>
        <many-to-one name="supplier" class="Supplier" column="supplierId"/>

    </class>

    <query name="com.hibernatebook.criteria.Product.HQLpricing"><![CDATA[
        select product.price from Product product where product.price > 25.0]]>
    </query>
    <sql-query name="com.hibernatebook.criteria.Product.SQLpricing">
        <return-scalar column="price" type="double"/>
        <![CDATA[
        select product.price from Product as product where product.price > 25.0]]>
    </sql-query>
</hibernate-mapping>
```

Notice that we embedded the SQL and HQL queries in CDATA regions. This protects our SQL queries from interfering with the XML parser—we don't have to worry about special characters breaking the XML. For the native SQL query, we also had to specify a return type, so Hibernate knows what type of result data to expect from the database. When you use HQL, Hibernate handles that mapping behind the scenes, because it knows which objects went in. With SQL, you have to specify the return types yourself. In this case, we used the <return-scalar> XML element to define our return type as a column named price, with a type of double. Hibernate converts the JDBC result set into an array of objects, just like the previous HQL query. Functionally, they are identical. We discuss native SQL in more detail in the next section of the chapter.

You may also specify the flush mode, whether the query is cacheable, the cache region, the fetch size, and the timeout for the HQL and SQL queries. For the SQL query, you may additionally specify whether the SQL query is callable.

Using Native SQL

Although you should probably use HQL whenever possible, Hibernate does provide a way to use native SQL statements directly through Hibernate. One reason to use native SQL is that your database supports some special features through its dialect of SQL that are not supported in HQL. Another reason is that you may want to call stored procedures from your Hibernate application. We discuss stored procedures and other database-specific integration solutions in Appendix A. Rather than just providing an interface to the underlying JDBC connection, like other Java ORM tools, Hibernate provides a way to define the entity (or join) that the query uses. This makes integration with the rest of your ORM-oriented application easy.

You can modify your SQL statements to make them work with Hibernate's ORM layer. You do need to modify your SQL to include Hibernate aliases that correspond to objects or object properties. You can specify all properties on an object with {objectname.*}, or you can specify the aliases directly with {objectname.property}. Hibernate uses the mappings to translate your object property names into their underlying SQL columns. This may not be the exact way you expect Hibernate to work, so be aware that you do need to modify your SQL statements for full ORM support. You will especially run into problems with native SQL on classes with subclasses— be sure you understand how you mapped the inheritance across either a single table or multiple tables, in order that you select the right properties off of the table.

Underlying Hibernate's native SQL support is the org.hibernate.SQLQuery interface, which extends the org.hibernate.Query interface already discussed. Your application will create a native SQL query from the session with the createSQLQuery() method on the Session interface.

```
public SQLQuery createSQLQuery(String queryString) throws HibernateException
```

After you pass a string containing the SQL query to the createSQLQuery() method, you should associate the SQL result with an existing Hibernate entity, a join, or a scalar result. The SQLQuery interface has addEntity(), addJoin(), and addScalar() methods. For the entities and joins, you can specify a lock mode, which we discuss in Chapter 8. The addEntity() methods take an alias argument and either a class name or an entity name. The addJoin() methods take an alias argument and a path to join.

Using native SQL with scalar results is the simplest way to get started with native SQL. Our Java code looks like this:

```
String sql = "select avg(product.price) as avgPrice from Product product";
SQLQuery query = session.createSQLQuery(sql);
query.addScalar("avgPrice",Hibernate.DOUBLE);
List results = query.list();
```

Because we did not specify any entity aliases, Hibernate executes exactly the same SQL that we passed through:

```
select avg(product.price) as avgPrice from Product product
```

The SQL is regular SQL (we did not have to do any aliasing here). We created an SQLQuery object, and then added a scalar mapping with the built-in double type (from the org.hibernate._Hibernate class). We needed to map the avgPrice SQL alias to the object type. The results are a List with one object—a Double.

A bit more complicated than the previous example is the native SQL that returns a result set of objects. In this case, we will need to map an entity to the SQL query. The entity consists of the alias we used for the object in the SQL query and its class. For this example, we used our `Supplier` class:

```
String sql = "select {supplier.*} from Supplier supplier";
SQLQuery query = session.createSQLQuery(sql);
query.addEntity("supplier", Supplier.class);
List results = query.list();
```

Hibernate modifies the SQL and executes the following command against the database:

```
select Supplier.id as id0_ , Supplier.name as name2_0_ from Supplier supplier
```

The special aliases allow Hibernate to map the database columns back to the object properties.

Summary

HQL is a powerful object-oriented query language that provides the power of SQL while taking advantage of Hibernate's object-relational mapping and caching. If you are porting an existing application to Hibernate, you can use Hibernate's native SQL facilities to execute SQL against the database. The SQL functionality is also useful for executing SQL statements that are specific to a given database and have no equivalents in HQL.

You may turn on SQL logging for Hibernate, and Hibernate will log the generated SQL that it executes against the database. If you add a comment to your HQL query object, Hibernate will display a comment in the log next to the SQL statement—this helps with tracing SQL statements back to HQL in your application.

CHAPTER 10

■ ■ ■ ■

Advanced Queries Using Criteria

Hibernate provides three different ways to retrieve data. We have already discussed HQL and the use of native SQL queries—now we add criteria.

The Criteria Query API lets you build nested, structured query expressions in Java, providing a compile-time syntax-checking that is not possible with a query language like HQL or SQL. The Criteria API also includes *query by example* (QBE) functionality—this lets you supply example objects that contain the properties you would like to retrieve instead of having to spell the components of the query out step by step. It also includes projection and aggregation methods, including counts.

In this chapter, we explore the use of the Criteria API using the sample database established in the previous chapter.

Using the Criteria API

The Criteria API allows you to build up a criteria query object programmatically—the `org.hibernate.Criteria` interface defines the available methods for one of these objects. The Hibernate `Session` interface contains several `createCriteria()` methods. Pass the persistent object's class or its entity name to the `createCriteria()` method, and Hibernate will create a `Criteria` object that returns instances of the persistence object's class when your application executes a criteria query.

The simplest example of a criteria query is one with no optional parameters or restrictions—the criteria query will simply return every object that corresponds to the class.

```
Criteria crit = session.createCriteria(Product.class);
List results = crit.list();
```

When you run this example with our sample data, you will get all objects that are instances of the `Product` class—note that this includes any instances of the `Software` class because they are derived from `Product`.

Moving on from this simple example, we will add constraints to our criteria queries so we can winnow down the result set.

Using Restrictions with Criteria

The Criteria API makes it easy to use restrictions in your queries to selectively retrieve objects; for instance, your application could retrieve only products with a price over $30. You may add these restrictions to a `Criteria` object with the `add()` method. The `add()` method takes an `org.hibernate.criterion.Criterion` object that represents an individual restriction. You can have more than one restriction for a criteria query.

Although you could create your own objects implementing the `Criterion` object, or extend an existing `Criterion` object, we recommend that you use Hibernate's built-in `Criterion` objects from your application's business logic. For instance, you could create your own factory class that returns instances of Hibernate's `Criterion` objects appropriately set up for your application's restrictions.

Use the factory methods on the `org.hibernate.criterion.Restrictions` class to obtain instances of the `Criterion` objects. To retrieve objects that have a property value that equals your restriction, use the `eq()` method on `Restrictions`, as follows:

```
public static SimpleExpression eq(String propertyName, Object value)
```

We would typically nest the `eq()` method in the `add()` method on the `Criteria` object. Here is an example of how this would look if we were searching for products with the name "Mouse":

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.eq("name","Mouse"));
List results = crit.list()
```

Next, we search for products that do *not* have the name "Mouse." For this, we would use the `ne()` method on the `Restrictions` class to obtain a not-equal restriction:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.ne("name","Mouse"));
List results = crit.list();
```

■**Tip** You cannot use the not-equal restriction to retrieve records with a NULL value in the database for that property (in SQL, and therefore in Hibernate, NULL represents the absence of data, and so cannot be compared with data). If you need to retrieve objects with NULL properties, you will have to use the `isNull()` restriction, which we discuss further on in the chapter. You can combine the two with an OR logical expression, which we also discuss later in the chapter.

Instead of searching for exact matches, we can also retrieve all objects that have a property matching part of a given pattern. To do this, we need to create an SQL LIKE clause, with either the `like()` or the `ilike()` method. The `ilike()` method is case-insensitive. In either case, we have two different ways to call the method:

```
public static SimpleExpression like(String propertyName, Object value)
```

or

```
public static SimpleExpression like(String propertyName,
                                    String value,
                                    MatchMode matchMode)
```

The first `like()` or `ilike()` method takes a pattern for matching. Use the % character as a wildcard to match parts of the string, like so:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.like("name","Mou%"));
List results = crit.list();
```

The second `like()` or `ilike()` method uses an `org.hibernate.criterion.MatchMode` object to specify how to match the specified value to the stored data. The `MatchMode` object (a type-safe enumeration) has four different matches:

- `ANYWHERE`: Anyplace in the string

- `END`: The end of the string

- `EXACT`: An exact match

- `START`: The beginning of the string

Here is an example that uses the `ilike()` method to search for case-insensitive matches at the end of the string:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.ilike("name","browser", MatchMode.END));
List results = crit.list();
```

The `isNull()` and `isNotNull()` restrictions allow you to do a search for objects that have (or do not have) null property values. This is easy to demonstrate:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.isNull("name"));
List results = crit.list();
```

Several of the restrictions are useful for doing math comparisons. The greater-than comparison is `gt()`, the greater-than-or-equal-to comparison is `ge()`, the less-than comparison is `lt()`, and the less-than-or-equal-to comparison is `le()`. We can do a quick retrieval of all products with prices over $25 like this:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.gt("price",new Double(25.0)));
List results = crit.list();
```

Moving on, we can start to do more complicated queries with the Criteria API. For example, we can combine `AND` and `OR` restrictions in logical expressions. When you add more than one constraint to a criteria query, it is interpreted as an `AND`, like so:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.gt("price",new Double(25.0)));
crit.add(Restrictions.like("name","K%"));
List results = crit.list();
```

If we want to have two restrictions that return objects that satisfy either or both of the restrictions, we need to use the or() method on the Restrictions class, as follows:

```
Criteria crit = session.createCriteria(Product.class);
Criterion price = Restrictions.gt("price",new Double(25.0));
Criterion name = Restrictions.like("name","Mou%");
LogicalExpression orExp = Restrictions.or(price,name);
crit.add(orExp);
List results = crit.list();
```

The orExp logical expression that we have created here will be treated like any other criterion. We can therefore add another restriction to the criteria:

```
Criteria crit = session.createCriteria(Product.class);
Criterion price = Restrictions.gt("price",new Double(25.0));
Criterion name = Restrictions.like("name","Mou%");
LogicalExpression orExp = Restrictions.or(price,name);
crit.add(orExp);
crit.add(Restrictions.ilike("description","blocks%"));
List results = crit.list();
```

If we wanted to create an OR expression with more than two different criteria, we would use an org.hibernate.criterion.Disjunction object to represent a disjunction. You can obtain this object from the disjunction() factory method on the Restrictions class. The disjunction is more convenient than building a tree of OR expressions in code. To represent an AND expression with more than two criteria, you can use the conjunction() method— although you can easily just add those to the Criteria object. The conjunction is also more convenient than building a tree of AND expressions in code. Here is an example that uses the disjunction:

```
Criteria crit = session.createCriteria(Product.class);
Criterion price = Restrictions.gt("price",new Double(25.0));
Criterion name = Restrictions.like("name","Mou%");
Criterion desc = Restrictions.ilike("description","blocks%");
Disjunction disjunction = Restrictions.disjunction();
disjunction.add(price);
disjunction.add(name);
disjunction.add(desc);
crit.add(disjunction);
List results = crit.list();
```

The last type of restriction is the SQL restriction sqlRestriction(). This restriction allows you to directly specify SQL in the Criteria API. This is useful if you need to use SQL clauses that Hibernate does not support through the Criteria API. Your application's code does not need to know the name of the table your class uses—use {alias} to signify the class's table, as follows:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.sqlRestriction("{alias}.name like 'Mou%'"));
List results = crit.list()
```

The other two `sqlRestriction()` methods permit you to pass JDBC parameters and values into the SQL statement. Use the standard JDBC parameter placeholder (?) in your SQL fragment.

Paging Through the Result Set

One common application pattern that criteria can address is pagination through the result set of a database query. When we say pagination, we mean an interface in which the user sees part of the result set at a time, with navigation to go forward and backward through the results. A naive pagination implementation might load the entire result set into memory for each navigation action, and would usually lead to atrocious performance. Both of us have worked on improving performance for separate projects suffering from exactly this problem. The problem appeared late in testing because the sample dataset that developers were working with was trivial, and they did not notice any performance problems until the first test data load.

If you are programming directly to the database, you will typically use proprietary database SQL or database cursors to support paging. Hibernate abstracts this away for you—behind the scenes, Hibernate uses the appropriate method for your database.

There are two methods on the `Criteria` interface for paging: `setFirstResult()` and `setMaxResults()`. The `setFirstResult()` method takes an integer that represents the first row in your result set, starting with row 0. You can tell Hibernate to retrieve a fixed number of objects with the `setMaxResults()` method. Using both of these together, we can construct a paging component in our web or Swing application. We have a very small dataset in our sample application, so here is an admittedly trivial example:

```
Criteria crit = session.createCriteria(Product.class);
crit.setFirstResult(1);
crit.setMaxResults(2);
List results = crit.list();
```

As you can see, this makes paging through the result set easy. You can increase the first result you return (for example, from 1, to 21, to 41, etc.) to page through the result set. If you only have one result in your result set, Hibernate has a shortcut method for obtaining just that object.

Obtaining a Unique Result

Sometimes you know you are only going to return zero or one objects from a given query. This could be because you are calculating an aggregate (like `COUNT`, which we discuss later), or because your restrictions naturally lead to a unique result—when selecting upon a property under a unique constraint, for example. You may also limit the results of any result set to just the first result, using the `setMaxResults()` method discussed earlier. In any of these circumstances, if you want obtain a single `Object` reference instead of a `List`, the `uniqueResult()` method on the `Criteria` object returns an object or `null`. If there is more than one result, the `uniqueResult()` method throws a `HibernateException`.

The following short example demonstrates having a result set that would have included more than one result, except that it was limited with the `setMaxResults()` method:

```
Criteria crit = session.createCriteria(Product.class);
Criterion price = Restrictions.gt("price",new Double(25.0));
crit.setMaxResults(1);
Product product = (Product) crit.uniqueResult();
```

Again, we stress that you need to make sure that your query only returns one or zero results if you use the `uniqueResult()` method. Otherwise, Hibernate will throw a `NonUniqueResultException` exception, which may not be what you would expect—Hibernate does not just pick the first result and return it.

Sorting the Query's Results

Sorting the query's results works much the same way with criteria as it would with HQL or SQL. The Criteria API provides the `org.hibernate.criterion.Order` class to sort your result set in either ascending or descending order, according to one of your object's properties.

Create an `Order` object with either of the two static factory methods on the `Order` class: `asc()` for ascending or `desc()` for descending. Both methods take the name of the property as their only argument. After you create an `Order`, use the `addOrder()` method on the `Criteria` object to add it to the query.

This example demonstrates how you would use the `Order` class:

```
Criteria crit = session.createCriteria(Product.class);
crit.add(Restrictions.gt("price",new Double(25.0)));
crit.addOrder(Order.desc("price"));
List results = crit.list();
```

You may add more than one `Order` object to the `Criteria` object. Hibernate will pass them through to the underlying SQL query. Your results will be sorted by the first order, then any identical matches within the first sort will be sorted by the second order, and so on. Beneath the covers, Hibernate passes this on to an SQL `ORDER BY` clause after substituting the proper database column name for the property.

Associations

To add a restriction on a class that is associated with your criteria's class, you will need to create another `Criteria` object. Pass the property name of the associated class to the `createCriteria()` method, and you will have another `Criteria` object. You can get the results from either `Criteria` object, although you should pick one style and be consistent for readability's sake. We find that getting the results from the top-level `Criteria` object (the one that takes a class as a parameter) makes it clear what type of object is expected in the results.

The association works going from one-to-many as well as from many-to-one. First, we will demonstrate how to use one-to-many associations to obtain suppliers who sell products with a price over $25. Notice that we create a new `Criteria` object for the `products` property, add restrictions to the products' criteria we just created, and then obtain the results from the supplier `Criteria` object:

```
Criteria crit = session.createCriteria(Supplier.class);
Criteria prdCrit = crit.createCriteria("products");
prdCrit.add(Restrictions.gt("price",new Double(25.0)));
List results = crit.list();
```

Going the other way, we obtain all the products from the supplier MegaInc using many-to-one associations:

```
Criteria crit = session.createCriteria(Product.class);
Criteria suppCrit = crit.createCriteria("supplier");
suppCrit.add(Restrictions.eq("name","MegaInc"));
List results = crit.list();
```

Although we can use either `Criteria` object to obtain the results, it makes a difference which criteria we use for ordering the results. In the following example, we are ordering the supplier results by the supplier names:

```
Criteria crit = session.createCriteria(Supplier.class);
Criteria prdCrit = crit.createCriteria("products");
prdCrit.add(Restrictions.gt("price",new Double(25.0)));
crit.addOrder(Order.desc("name"));
List results = prdCrit.list();
```

If we wanted to sort the suppliers by the descending price of their products, we would use the following line of code. This code would have to replace the previous `addOrder()` call on the supplier `Criteria` object.

```
prdCrit.addOrder(Order.desc("price"));
```

Although the products are not in the result set, SQL still allows you to order by those results. If you get mixed up with which `Criteria` object you are using and pass the wrong property name for the sort-by order, Hibernate will throw an exception.

Distinct Results

If you would like to work with distinct results from a criteria query, Hibernate provides a result transformer for distinct entities, `org.hibernate.transform.DistinctRootEntityResultTransformer`, which ensures that no duplicates will be in your query's result set. Rather than using `SELECT DISTINCT` with SQL, the distinct result transformer compares each of your results using their default `hashCode()` methods, and only adds those results with unique hash codes to your result set. This may or may not be the result you would expect from an otherwise equivalent SQL `DISTINCT` query, so be careful with this. An additional performance note: the comparison is done in Hibernate's Java code, not at the database, so non-unique results will still be transported across the network.

Projections and Aggregates

Instead of working with objects from the result set, you can treat the results from the result set as a set of rows and columns. This is similar to how you would use data from a `SELECT` query with JDBC; also, Hibernate supports properties, aggregate functions, and the `GROUP BY` clause.

To use projections, start by getting the `org.hibernate.criterion.Projection` object you need from the `org.hibernate.criterion.Projections` factory class. The `Projections` class is similar to the `Restrictions` class in that it provides several static factory methods for obtaining `Projection` instances. After you get a `Projection` object, add it to your `Criteria` object with the `setProjection()` method. When the `Criteria` object executes, the list contains object references that you can cast to the appropriate type.

The row-counting functionality provides a simple example of applying projections. The code looks similar to the restrictions examples we were working with earlier in the chapter:

```
Criteria crit = session.createCriteria(Product.class);
crit.setProjection(Projections.rowCount());
List results = crit.list();
```

The results list will contain one object, an `Integer` that contains the results of executing the `COUNT` SQL statement. Other aggregate functions available through the `Projections` factory class include the following:

- `avg(String propertyName)`: Gives the average of a property's value

- `count(String propertyName)`: Counts the number of times a property occurs

- `countDistinct(String propertyName)`: Counts the number of unique values the property contains

- `max(String propertyName)`: Calculates the maximum value of the property values

- `min(String propertyName)`: Calculates the minimum value of the property values

- `sum(String propertyName)`: Calculates the sum total of the property values

We can apply more than one projection to a given `Criteria` object. To add multiple projections, get a projection list from the `projectionList()` method on the `Projections` class. The `org.hibernate.criterion.ProjectionList` object has an `add()` method that takes a `Projection` object. You can pass the projections list to the `setProjection()` method on the `Criteria` object because `ProjectionList` implements the `Projection` interface. The following example demonstrates some of the aggregate functions, along with the projection list:

```
Criteria crit = session.createCriteria(Product.class);
ProjectionList projList = Projections.projectionList();
projList.add(Projections.max("price"));
projList.add(Projections.min("price"));
projList.add(Projections.avg("price"));
projList.add(Projections.countDistinct("description"));
crit.setProjection(projList);
List results = crit.list();
```

When you execute multiple aggregate projections, you get a `List` with an `Object` array as the first element. The `Object` array contains all of your values, in order.

Another use of projections is to retrieve individual properties, rather than entities. For instance, we can retrieve just the name and description from our product table, instead of faulting the classes into memory. Use the `property()` method on the `Projections` class to

create a `Projection` for a property. When you execute this form of query, the `list()` method returns a `List` of `Object` arrays. Each `Object` array contains the projected properties for that row. The following example returns just the contents of the name and description columns from the `Product` data. Remember, Hibernate is polymorphic, so this also returns the name and description from the `Software` objects that inherit from `Product`.

```
Criteria crit = session.createCriteria(Product.class);
ProjectionList projList = Projections.projectionList();
projList.add(Projections.property("name"));
projList.add(Projections.property("description"));
crit.setProjection(projList);
List results = crit.list();
```

Use this query style when you want to cut down on network traffic between your application servers and your database servers. For instance, if your table has a large number of columns, this can slim down your results. In other cases, you may have a large set of joins that would return a very wide result set, but you are only interested in a few columns. Lastly, if your clients have limited memory, this can save you trouble with large datasets. But make sure you don't have to retrieve additional columns for the entire result set later, or your optimizations may actually decrease performance.

You can group your results (using SQL's `GROUP BY` clause) with the `groupProperty` projection. The following example groups the products by name and price:

```
Criteria crit = session.createCriteria(Product.class);
ProjectionList projList = Projections.projectionList();
projList.add(Projections.groupProperty("name"));
projList.add(Projections.groupProperty("price"));
crit.setProjection(projList);
List results = crit.list();
```

As you can see, projections open up aggregates to the Criteria API, which means that developers do not have to drop into HQL for aggregates. Projections offer a way to work with data that is closer to the JDBC result set style, which may be appropriate for some parts of your application.

Query By Example (QBE)

In this section, because of the confusing terminology, we will refer to excerpts from our demonstration code as "samples" rather than "examples," reserving "example" for its peculiar technical meaning in the context of QBE.

In QBE, instead of programmatically building a `Criteria` object with `Criterion` objects and logical expressions, you can partially populate an instance of the object. You use this instance as a template and have Hibernate build the criteria for you based upon its values. This keeps your code clean and makes your project easier to test. The `org.hibernate.criterion.Example` class contains the QBE functionality. Note that the `Example` class implements the `Criterion` interface, so you can use it like any other restriction on a criteria query.

For instance, if we have a user database, we can construct an instance of a user object, set the property values for type and creation date, and then use the Criteria API to run a

QBE query. Hibernate will return a result set containing all user objects that match the property values that were set. Behind the scenes, Hibernate inspects the `Example` object and constructs an SQL fragment that corresponds to the properties on the `Example` object.

To use QBE, we need to construct an `Example` object first. Then we need to create an instance of the `Example` object, using the static `create()` method on the `Example` class. The `create()` method takes the `Example` object as its argument. You add the `Example` object to a `Criteria` object just like any other `Criterion` object.

The following basic example searches for suppliers that match the name on the example `Supplier` object:

```
Criteria crit = session.createCriteria(Supplier.class);
Supplier supplier = new Supplier();
supplier.setName("MegaInc");
crit.add(Example.create(supplier));
List results = crit.list();
```

When Hibernate translates our `Example` object into an SQL query, all the properties on our `Example` objects get examined. We can tell Hibernate which properties to ignore; the default is to ignore null-valued properties. To search our products or software in the sample database with QBE, we need to either specify a price or tell Hibernate to ignore properties with a value of zero, because we used a double primitive for storage instead of a `Double` object. The double primitive initializes to zero, while a `Double` would have been null; and so, left to its own devices, the QBE logic will assume that we are specifically searching for prices of zero, whereas we want it to ignore this default value.

We can make the Hibernate `Example` object exclude zero-valued properties with the `excludeZeroes()` method. We can exclude properties by name with the `excludeProperty()` method, or exclude nothing (compare for null values and zeroes exactly as they appear in the `Example` object) with the `excludeNone()` method. This sample applies the `excludeZeroes()` method to ignore the default zero prices:

```
Criteria crit = session.createCriteria(Product.class);
Product exampleProduct = new Product();
exampleProduct.setName("Mouse");
Example example = Example.create(exampleProduct);
example.excludeZeroes();
crit.add(example);
List results = crit.list();
```

Other options on the `Example` object include ignoring the case for strings with the `ignoreCase()` method, and enabling use of SQL's `LIKE` for comparing strings, instead of just using `equals()`.

We can also use associations for QBE. In the following sample, we create two examples: one for the product and one for the supplier. We use the technique explained in the "Associations" section of this chapter to retrieve objects that match both criteria.

```
Criteria prdCrit = session.createCriteria(Product.class);
Product product = new Product();
product.setName("M%");
Example prdExample = Example.create(product);
prdExample.excludeProperty("price");
prdExample.enableLike();
Criteria suppCrit = prdCrit.createCriteria("supplier");
Supplier supplier = new Supplier();
supplier.setName("SuperCorp");
suppCrit.add(Example.create(supplier));
prdCrit.add(prdExample);
List results = prdCrit.list();
```

We also ignore the price property for our product, and we use LIKE for object comparison, instead of equals.

The QBE API works best for searches in which you are building the search from user input. The Hibernate team recommends using QBE for advanced searches with multiple fields, because it's easier to set values on business objects than to manipulate restrictions with the Criteria API.

Summary

Using the Criteria API is an excellent way to get started developing with HQL. The developers of Hibernate have provided a clean API for adding restrictions to queries with Java objects. Although HQL isn't too difficult to learn, some developers prefer the Criteria Query API, as it offers compile-time syntax checking—although column names and other schema-dependent information cannot be checked until run time.

In the next chapter, we discuss the use of Hibernate filters to restrict the range of data against which queries are applied.

Filtering the Results of Searches

Your application will often need to process only a subset of the data in the database tables. In these cases, you can create a Hibernate *filter* to eliminate the unwanted data. Filters provide a way for your application to limit the results of a query to data that passes the filter's criteria. Filters are not a new concept—you can achieve much the same effect using SQL database views—but Hibernate offers a centralized management system for them.

Unlike database views, Hibernate filters can be enabled or disabled during a Hibernate session. In addition, Hibernate filters can be parameterized, which is particularly useful when you are building applications on top of Hibernate that use security roles or personalization.

When to Use Filters

As an example, consider a web application that manages user profiles. Currently, your application presents a list of all users through a single web interface, but you receive a change request from your end user to manage active users and expired users separately. For this example, assume that the status is stored as a column on the user table.

One way to solve this problem is to rewrite every HQL SELECT query in your application, adding a WHERE clause that restricts the result by the user's status. Depending on how you built your application, this could be an easy undertaking or it could be complex, but you still end up modifying code that you have already tested thoroughly, potentially changing it in many different places.

With Hibernate 3, you can create a filter restriction for the user status. When your end user selects the user type (active or expired), your application activates the user status filter (with the proper status) for the end user's Hibernate session. Now, any SELECT queries will return the correct subset of results, and the relevant code for the user status is limited to two locations: the Hibernate session and the user status filter.

The advantage of using Hibernate filters is that you can programmatically turn filters on or off in your application code, and your filters are defined in your Hibernate mapping documents for easy maintainability. The major disadvantage of filters is that you cannot create new filters at run time. Instead, any filters your application requires need to be specified in the proper Hibernate mapping document. Although this may sound somewhat limiting, the fact that filters can be parameterized makes them pretty flexible. For our user status filter example, only one filter would need to be defined in the mapping document (albeit in two parts). That

filter would specify that the status column must match a named parameter. You would not need to define the possible values of the status column in the Hibernate mapping document—the application can specify those parameters at run time.

Although it is certainly possible to write applications with Hibernate that do not use filters, we find them to be an excellent solution to certain types of problems—notably security and personalization.

Defining Filters

Your first step is to define filters in your application's Hibernate mapping documents, using the `<filter-def>` XML element. These filter definitions must contain the name of the filter and the names and types of any filter parameters. Specify filter parameters with the `<filter-param>` XML element. Filter parameters are similar to named parameters for HQL queries. Both require a `:` before the parameter name. Here is an excerpt from a mapping document with a filter called `latePaymentFilter` defined:

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping>
  <class ...

  </class>
  <filter-def name="latePaymentFilter">
    <filter-param name="dueDate" type="date"/>
  </filter-def>
</hibernate-mapping>
```

Once you have created the filter definitions, you need to attach the filters to class or collection mapping elements. You can attach a single filter to more than one class or collection. To do this, you add a `<filter>` XML element to each class and/or collection. The `<filter>` XML element has two attributes: `name` and `condition`. The name references a filter definition (for instance: `latePaymentFilter`). The condition represents a `WHERE` clause in HQL. Here's an example:

```
  <class ...
    <filter name="latePaymentFilter" condition=":dueDate = paymentDate"/>
  </class>
```

Each `<filter>` XML element must correspond to a `<filter-def>` element. You may have more than one filter for each filter definition, and each class can have more than one filter. This is a little confusing—the extra level of abstraction allows you to define all the filter parameters in one place and then refer to them in the individual filter conditions.

Using Filters in Your Application

Your application programmatically determines which filters to activate or deactivate for a given Hibernate session. Each session can have a different set of filters with different parameter values. By default, sessions do not have any active filters—you must explicitly enable filters programmatically for each session. The `Session` interface contains several methods for working with filters, as follows:

- `public Filter enableFilter(String filterName)`

- `public Filter getEnabledFilter(String filterName)`

- `public void disableFilter(String filterName)`

These are pretty self-explanatory—the `enableFilter(String filterName)` method activates the specified filter, the `disableFilter(String filterName)` method deactivates the method, and if you have already activated a named filter, `getEnabledFilter(String filterName)` retrieves that filter.

The `org.hibernate.Filter` interface has six methods. You are unlikely to use `validate()`; Hibernate uses that method when it processes the filters. The other five methods are as follows:

- `public Filter setParameter(String name, Object value)`

- `public Filter setParameterList(String name, Collection values)`

- `public Filter setParameterList(String name, Object[] values)`

- `public String getName()`

- `public FilterDefinition getFilterDefinition()`

The `setParameter()` method is the most useful. You can substitute any Java object for the parameter, although its type should match the type you specified for the parameter when you defined the filter. The two `setParameterList()` methods are useful for using `IN` clauses in your filters. If you want to use `BETWEEN` clauses, use two different filter parameters with different names. Finally, the `getFilterDefinition()` method allows you to retrieve a `FilterDefinition` object representing the filter metadata (its name, its parameters' names, and the parameter types).

Once you have enabled a particular filter on the session, you do not have to do anything else to your application to take advantage of filters, as we demonstrate in the following example.

A Basic Filtering Example

Because filters are very straightforward, a basic example allows us to demonstrate most of the filter functionality, including activating filters and defining filters in mapping documents.

In the following Hibernate XML mapping document (`User.hbm.xml`), we created a filter definition called `activatedFilter`. The parameters for the filter must be specified with `<filter-param>` XML elements (as shown in Listing 11-1), which use the `<activatedParam>` XML element. You need to specify a type for the filter parameter so that Hibernate knows how

to map values to parameters. Once you have defined your filter, you need to attach the filter definition to a class. At the end of our User class definition, we specify that it uses a filter named activatedFilter. We then need to set a condition corresponding to an HQL WHERE clause for the attached filter. In our case, we used :activatedParam = activated, where :activatedParam is the named parameter specified on the filter definition, and activated is the column name from the user table. You should ensure that the named parameter goes on the left-hand side so that Hibernate's generated SQL doesn't interfere with any joins.

Listing 11-1. *Hibernate XML Mapping for* User

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping>
  <class name="com.hibernatebook.filters.User">
    <id name="id" type="int">
      <generator class="native"/>
    </id>

    <property name="username" type="string" length="32"/>
    <property name="activated" type="boolean"/>
    <filter name="activatedFilter" condition=":activatedParam = activated"/>
  </class>
  <filter-def name="activatedFilter">
    <filter-param name="activatedParam" type="boolean"/>
  </filter-def>
</hibernate-mapping>
```

With the filter definition created and attached to a class with a suitable condition, we need to activate the filter. The next class, SimpleFilterExample, inserts several user records into the database, and then immediately displays them to the screen. The class uses a very simple HQL query (from User) to obtain the result set from Hibernate. The displayUsers() method writes the usernames and activation status out to the console. Before you have enabled any filters on the database, this method will return the full list of users. Once you have enabled the first filter (activatedFilter) to show only activated users, call the same displayUsers() method—the results of the query are the same as if you had added a WHERE clause containing an "activated=true" clause. You can just as easily change the filter's parameter value to show inactive users, as shown in Listing 11-2.

Listing 11-2. *Invoking Filters from Code*

```java
package com.hibernatebook.filters;

import java.util.Iterator;

import org.hibernate.Filter;
import org.hibernate.Query;
import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.Transaction;

import org.hibernate.cfg.Configuration;

public class SimpleFilterExample {
    public static void main (String args[]) {
        SessionFactory factory =
            new Configuration().configure().buildSessionFactory();
        Session session = factory.openSession();

        //insert the users
        insertUser("ray",true,session);
        insertUser("jason",true,session);
        insertUser("beth",false,session);
        insertUser("judy",false,session);
        insertUser("rob",false,session);

        //Show all users
        System.out.println("===ALL USERS===");
        displayUsers(session);

        //Show activated users
        Filter filter = session.enableFilter("activatedFilter");
        filter.setParameter("activatedParam",new Boolean(true));
        System.out.println("===ACTIVATED USERS===");
        displayUsers(session);

        //Show nonactivated users
        filter.setParameter("activatedParam",new Boolean(false));
        System.out.println("===NON-ACTIVATED USERS===");
        displayUsers(session);
```

```
            session.close();
        }

    public static void displayUsers(Session session) {
        session.beginTransaction();
        Query query = session.createQuery("from User");
        Iterator results = query.iterate();
        while (results.hasNext())
        {
            User user = (User) results.next();
            System.out.print(user.getUsername() + " is ");
            if (user.isActivated())
            {
                System.out.println("activated.");
            }
            else
            {
                System.out.println("not activated.");
            }
        }

        session.getTransaction().commit();
    }

    public static void insertUser(String name, boolean activated, Session session) {
        session.beginTransaction();

        User user = new User();
        user.setUsername(name);
        user.setActivated(activated);
        session.save(user);

        session.getTransaction().commit();
    }
}
```

The output of SimpleFilterExample is as follows:

```
===ALL USERS===
ray is activated.
jason is activated.
beth is not activated.
judy is not activated.
rob is not activated.
```

```
===ACTIVATED USERS===
ray is activated.
jason is activated.
===NON-ACTIVATED USERS===
beth is not activated.
judy is not activated.
rob is not activated.
```

Listing 11-3 gives the User class used for this chapter's examples. The only fields it contains are id, username, and activated.

Listing 11-3. *The Source Code for the* User *Class*

```java
package com.hibernatebook.filters;

public class User {
    private int id;
    private String username;
    private boolean activated;


    public boolean isActivated() {
        return activated;
    }

    public void setActivated(boolean activated) {
        this.activated = activated;
    }

    public int getId() {
        return id;
    }

    public void setId(int id) {
        this.id = id;
    }

    public String getUsername() {
        return username;
    }

    public void setUsername(String username) {
        this.username = username;
    }
}
```

Because filters do not use any database-specific functionality beyond the Hibernate configuration, you should not encounter any difficulty running this example on databases other than HSQLDB. The Hibernate configuration file defines the database configuration and connection information, along with the XML mapping document for the User class (see Listing 11-4).

Listing 11-4. *The Hibernate XML Configuration File for the Example*

```xml
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <property name="hibernate.connection.driver_class">
            org.hsqldb.jdbcDriver
        </property>
        <property name="hibernate.connection.url">
            jdbc:hsqldb:file:filterdb;SHUTDOWN=true
        </property>
        <property name="hibernate.connection.username">sa</property>
        <property name="hibernate.connection.password"></property>
        <property name="hibernate.connection.pool_size">0</property>
        <property name="dialect">
            org.hibernate.dialect.HSQLDialect
        </property>

        <!-- Mapping files -->
        <mapping resource="com/hibernatebook/filters/User.hbm.xml"/>
    </session-factory>
</hibernate-configuration>
```

The source code for this chapter includes the schema we used for the HSQL database to create the table for the filterdb database.

Summary

Filters are a useful way to separate some database concerns from the rest of your code. A set of filters can cut back on the complexity of the HQL queries used in the rest of your application, at the expense of some runtime flexibility. Instead of using views (which must be created at the database level), your applications can take advantage of dynamic filters that can be activated as and when they are required.

■ ■ ■

More Advanced Features

In this appendix, we discuss some of the features that, strictly speaking, lie outside the scope of this book, but that you should be aware of if you go on to use Hibernate in more depth.

EJB 3 and the EntityManager

The third version of the Enterprise Java Beans specification, generally known as EJB 3, has recently been finalized. Among other features, EJB 3 includes a standard ORM technology that was significantly influenced by the design of Hibernate.

You encountered this close relationship in Chapter 6 when we discussed Hibernate's use of the EJB 3 annotations for creating entity mappings. Annotations can be used throughout your EJB 3 applications to denote various settings. They are also used to mark for injection of resources from the container in a manner very like that of Spring's dependency injection (see Appendix C). HQL, which was discussed in Chapter 9, is very similar to the EJB QL used in EJB 3 environments—generally speaking, your HQL queries can be used as EJB QL queries without change.

Given these similarities, a Hibernate application can be converted into a portable EJB 3 application with surprisingly few changes. EJB 3 now supports both J2SE environments and those hosted within J2EE application servers; so even a stand-alone application can be written to take advantage of the EJB 3 features.

The standard way to access the ORM components of an EJB 3 application is through the `EntityManager`. The Hibernate team provides appropriate libraries for download on their EntityManager site at `http://entitymanager.hibernate.org`.

The `EntityManager` is configured through a standard file called `persistence.xml`, which must be provided in a `META-INF` directory accessible from the classpath (or, in a J2EE environment, from the root of the deployed archive). This file serves the same purpose as a conventional Hibernate configuration file (`hibernate.cfg.xml`), although its syntax is somewhat different. An example file is given in Listing A-1.

Listing A-1. *An EJB 3* `persistence.xml` *Configuration File*

```
<persistence xmlns="http://java.sun.com/xml/ns/persistence"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/persistence
http://java.sun.com/xml/ns/persistence/persistence_1_0.xsd"
  version="1.0">
```

```
    <persistence-unit name="sampleManager" transaction-type="RESOURCE_LOCAL">
        <class>com.hibernatebook.advanced.Sample</class>
        <properties>
            <property name="hibernate.connection.driver_class"
                value="org.hsqldb.jdbcDriver"/>
            <property name="hibernate.connection.url"
                value="jdbc:hsqldb:file:/advanced/db/advanceddb;SHUTDOWN=true"/>
            <property name="hibernate.connection.username" value="sa"/>
            <property name="hibernate.connection.password" value=""/>
            <property name="hibernate.connection.pool_size" value="0"/>
            <property name="hibernate.show_sql" value="false"/>
            <property name="hibernate.dialect"
                value="org.hibernate.dialect.HSQLDialect"/>
        </properties>
    </persistence-unit>
</persistence>
```

The configuration file can contain multiple named `<persistence-unit>` elements, each corresponding to a different configuration of the ORM environment. In the example in Listing A-1, we have created a single annotation-mapped entity. The `<properties>` element then configures the implementation-specific (i.e., Hibernate-specific) properties. In Listing A-1, we have configured a database connection and dialect. When configuring a J2EE environment, the connection would usually be provided through generic elements of the persistence unit; Listing A-2 shows a configuration that takes advantage of this approach.

Listing A-2. *An EJB 3* `persistence.xml` *Configuration File Using a JTA Data Source*

```
<persistence xmlns="http://java.sun.com/xml/ns/persistence"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xsi:schemaLocation="http://java.sun.com/xml/ns/persistence
http://java.sun.com/xml/ns/persistence/persistence_1_0.xsd"
    version="1.0">
    <persistence-unit name="sampleManager" transaction-type="JTA">
        <jta-data-source>java:comp/env/jdbc/advanced</jta-data-source>
        <class>com.hibernatebook.advanced.Sample</class>
        <properties>
            <property name="hibernate.dialect"
                value="org.hibernate.dialect.HSQLDialect"/>
        </properties>
    </persistence-unit>
</persistence>
```

Listing A-2 shows that the configuration file requires very little Hibernate-specific information. In a J2EE environment, it is possible at deployment time to substitute alternative EJB 3 providers for the providers indicated in the application's metadata.

In a J2SE environment, the configuration information is accessed by creating an `EntityManagerFactory` class by calling the `createEntityManagerFactory()` method of the `Persistence` class, with the configured name of the persistence unit (shown in bold in

Listings A-1 and A-2) containing the appropriate configuration information. From the
EntityManagerFactory class, you can request EntityManager instances that are used to access
the entities. You have probably already spotted that the EJB 3 Persistence class corresponds
roughly to Configuration, that EntityManagerFactory is a dead ringer for SessionFactory, and
that EntityManager is the analog of Session.

The example code in Listing A-3 pushes this point home. The EntityManager instance is
used in a very similar way to the Session class shown throughout this book (although some of
the method names are slightly different—persist() in this example corresponds to Session's
save() method).

Listing A-3. *Using EJB 3 Persistence in J2SE Code*

```java
package com.hibernatebook.advanced;

import java.util.List;

import javax.persistence.EntityManager;
import javax.persistence.EntityManagerFactory;
import javax.persistence.Persistence;
import javax.persistence.Query;

public class Ejb3Example {
    @SuppressWarnings("unchecked")
    public static void main(String[] args) {

        EntityManagerFactory factory =
            Persistence.createEntityManagerFactory("sampleManager");

        EntityManager manager = factory.createEntityManager();

        manager.getTransaction().begin();
        manager.persist(new Sample("FAQ"));
        manager.persist(new Sample("RTFM"));
        manager.persist(new Sample("PDQ"));
        manager.persist(new Sample("ASAP"));
        manager.getTransaction().commit();

        Query query = manager.createQuery("from Sample");

        manager.getTransaction().begin();
        List<Sample> list = (List<Sample>)query.getResultList();
        manager.getTransaction().commit();

        for(Sample sample : list) {
            System.out.println(sample.getContent());
        }
```

```
      manager.close();
      factory.close();
   }
}
```

While the configuration of an EJB 3 application server falls well outside the scope of this book (which is a shame, because the topic is interesting—see *Pro EJB: Java Persistence API*, by Mike Keith and Merrick Schincariol (Apress, 2006), for a good introduction to the subject), the *use* of an `EntityManager` deployed into an EJB 3 application server is straightforward. Typically in such an environment, the container manages the `EntityManager`. Listing A-4 demonstrates how to obtain a reference to an `EntityManager` in such an environment—only very simple changes would be necessary in Listing A-3 to support this. Note that in this environment, there is no need to interact with the `EntityManagerFactory`—the container manages the appropriate interaction with the factory in a way that is transparent to the user code.

Listing A-4. *Obtaining an* `EntityManager` *from the Container by Injection*

```
public class Ejb3Example {
@PersistenceContext(unitName="sampleManager",type=EXTENDED)
EntityManager manager;
   // …
}
```

As Listing A-4 demonstrates, the combination of container-managed `EntityManagers`, annotations, and resource injection makes the acquisition of an `EntityManager` object trivially simple (and remember, the `EntityManager` is essentially the same as a Hibernate `Session` object).

Hibernate provides a couple of additional features to facilitate the transition of Hibernate 3 code to EJB 3. Where your application uses a `Configuration` (or an `AnnotationConfiguration`) object to programmatically configure the Hibernate application, there is now an alternative `Ejb3Configuration` class that can be used in a similar manner to provide the configuration information for the EJB 3 objects without the need for a `configuration.xml` file.

The `<jb3configuration>` element of the Hibernate Tools Ant task conversely allows the configuration of the tasks to be drawn from the classpath's `META-INF/configuration.xml` file, instead of from an explicitly identified Hibernate configuration or properties file.

Managed Versioning and Optimistic Locking

While we have saved versioning for this appendix's discussion of advanced features, it is actually quite straightforward to understand and apply. Consider the following scenario:

- Client A loads and edits a record.

- Client B loads and edits the same record.

- Client A commits its edited record data.

- Client B commits its differently edited record data.

While the scenario is simple, the problems it presents are not. If Client A establishes a transaction, then Client B may not be able to load and edit the same record. Yet in a web

environment, it is not unlikely that Client A will close a browser window on the open record, never committing or canceling the transaction, so that the record remains locked until the session times out. Clearly this is not a satisfactory solution. Usually, you will not want to permit the alternative scenario, in which no locking is used, and the last person to save a record wins!

The solution, versioning, is essentially a type of optimistic locking (see Chapter 8). When any changes to an entity are stored, a version column is updated to reflect the fact that the entity has changed. When a subsequent user tries to commit changes to the same entity, the original version number will be compared against the current value—if they differ, the commit will be rejected.

The Hibernate/EJB 3 annotation mappings and the Hibernate XML-based mappings both provide a simple syntax for indicating which field should be used for storing the managed version information. The annotation for this field is shown in Listing A-5.

Listing A-5. *Marking the Version Attribute Using Annotations*

```
@Version
protected int getVersionNum() {
    return versionNum;
}
```

The default optimistic locking strategy for Hibernate is versioning, so if you provide a `<version>` element in your XML configuration, this will be used as long as you have enabled dynamic updates (as shown in Listing A-6).

Listing A-6. *Marking the Version Attribute Using XML Mappings*

```
<class dynamic-update="version" optimistic-lock="version" ... >
    ...
    <version name="versionNum"/>
</class>
```

The version attribute is defined in a very similar way to the normal property attribute configuration. The version can be of type `long`, `integer`, `short`, `timestamp`, or `calendar` (note that using the `<timestamp ... />` element is an equivalent alternative to the use of the `<version type="timestamp" ... />` element syntax).

The `<class>` element's `optimistic-lock` attribute can be used to override the default versioning-based optimistic locking strategy. You can disable it entirely (despite the presence of a version field) using a value of `none`. You can explicitly state that versioning should be used with a value of `version`. You can elect to use dirty checking, with the `dirty` and `all` options.

If you elect not to use versioning, dirty checking offers an alternative form of optimistic locking. Here, the values of the entities are themselves checked to see if they have changed since the entity was originally obtained. As with versioning-based optimistic locking, the check against the database is carried out when the changes are committed. If an optimistic lock type of `dirty` is selected, then only those fields that have changed since the persistent entity was obtained will be checked (the `Session` keeps track of the appropriate state information). If an optimistic lock type of `all` is selected, then all the fields comprising the entity will

be checked for changes. If the fields being checked have changed prior to the commit, then the commit will fail.

Versioning is generally a simpler and more reliable approach, so we suggest that you use this whenever you need optimistic locking features.

XML Relational Persistence

Hibernate provides a feature that allows XML data to be mapped into the entity model for access using the normal session methods. This functionality is provided primarily so that data can be imported into and exported from the underlying relational data store—it is not intended as a replacement for relational databases!

The feature can be used for various purposes—archiving data, implementing SOAP interfaces, and so on—but the most common use is for the purposes of processing (and providing) external data feeds such as product catalogs. We show here how the example application in Chapter 3 (an advertisements database) can be configured to read and write appropriate XML feeds for the mapped entities.

Hibernate requires the use of Dom4J as the API for XML access because Hibernate's internals already rely upon Dom4J to read configuration and mapping files.

Adding Node Information to Mappings

Two attributes are used to add all the XML-specific information to your existing mapping files: `node` and `embed-xml`.

The `node` attribute applies to most tags that correspond to tables or columns in the database. The value can be a single string, in which case it represents an element name in the XML markup; or it can be preceded by a commercial at symbol (@), in which case it represents the attribute of an element. Paths can be indicated using forward slashes. This is the standard XPath syntax for identifying elements in an XML document.

If set to `true`, the `embed-xml` attribute indicates that the property or entity referenced should be included inline as XML. If set to `false`, it indicates that a reference to the primary key should be substituted instead. This is necessary because the DOM generation does not perform reference handling automatically—so loops in the entity model references would result in infinite loops in XML generation if this option could not be set to `true`.

Note that if `embed-xml` is set to `false` for an association, generating XML output that references another entity will *not* automatically include a representation of the entity elsewhere in the generated document. This is your responsibility. Listing A-7 shows how our example `Advert` class from Chapter 3 might be marked up with `node` and `embed-xml` attributes.

Listing A-7. *The* `Advert` *Example Classes from Chapter 3 Marked Up for XML Persistence*

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-mapping
    PUBLIC "-//Hibernate/Hibernate Mapping DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">
```

```
<hibernate-mapping>
    <class name="sample.entity.User" table="aduser"
           node="user">

        <id name="id" type="long" column="id"
            node="@id">
            <generator class="native"/>
        </id>

        <property name="name" column="name" type="string" unique="true"
                  node="@name"/>

        <property name="password" column="password" type="string"
                  node="@password"/>
    </class>

    <class name="sample.entity.Category" table="category"
           node="category">

        <id name="id" type="long" column="id"
            node="@id">
            <generator class="native"/>
        </id>

        <property name="title" column="title" type="string" unique="true"
                  node="@name"/>

        <set name="adverts" table="link_category_advert"
             node="." embed-xml="true">
            <key column="category" foreign-key="fk_advert_category"/>
            <many-to-many class="sample.entity.Advert"
                          column="advert" foreign-key="fk_category_advert"
                          embed-xml="false" node="advert"/>
        </set>
    </class>

    <class name="sample.entity.Advert" table="advert"
           node="advert">

        <id name="id" type="long" column="id"
            node="@id">
            <generator class="native"/>
        </id>
```

```
      <property name="message" column="message" type="string"
              node="content"/>

      <property name="title" column="title" type="string"
              node="title"/>

      <many-to-one name="user" column="aduser"
         class="sample.entity.User" not-null="true"
         foreign-key="fk_advert_user"
         node="@user" embed-xml="false"/>
   </class>
</hibernate-mapping>
```

Exporting XML Entities

With the entity mappings marked up for XML relational persistence, the generation of output from an existing set of persistent entities is fairly simple. By obtaining a session with the entity mode configured for Dom4J, the entity proxies retrieved from that session will be org.dom4j.Node instances that can be manipulated with the normal Dom4J API.

To access a Session object in this mode, you first open a conventional session, and then invoke the getSession() method on it, passing EntityMode.DOM4J as the sole parameter thus:

```
Session = sessionFactory.openSession();
Session xmlSession = session.getSession(EntityMode.DOM4J);
```

Once an appropriate Dom4J document has been populated with the entities (extracted from Hibernate in the normal way), the session can be closed, and the Dom4J document can be treated as a self-contained entity, as shown in Listing A-8.

Listing A-8. *Exporting the Advert Entities Using Dom4J*

```
package sample.xml;

import java.io.*;
import java.util.*;
import org.dom4j.*;
import org.hibernate.*;
import org.hibernate.cfg.Configuration;

public class ExportXML {
    private static final SessionFactory sessionFactory = new Configuration()
            .configure().buildSessionFactory();

    public static void main(String[] args)
        throws Exception
    {
        System.out.println("Preparing the DOM Document");
        Document document = DocumentHelper.createDocument();
        Element root = document.addElement("catalog");
```

```
        System.out.println("Preparing the Session objects");
        Session session = sessionFactory.openSession();
        Session xmlSession = session.getSession(EntityMode.DOM4J);

        System.out.println("Reading the catalog from the database");
        session.beginTransaction();

        export(xmlSession,root,"from User");
        export(xmlSession,root,"from Advert");
        export(xmlSession,root,"from Category");

        session.getTransaction().commit();
        session.close();

        System.out.println("Dumping the catalog to a file");
        BufferedWriter writer = new BufferedWriter(new FileWriter("catalog.xml"));
        document.write(writer);
        writer.flush();
        writer.close();

        System.out.println("Done.");
    }

    public static void export(Session xmlSession,Element root,String hql) {
        Query query = xmlSession.createQuery(hql);
        List categories = query.list();
        Iterator cit = categories.iterator();
        while(cit.hasNext()) {
            Element element = (Element)cit.next();
            root.add(element);
        }
    }
}
```

The code in Listing A-8 generates the catalog.xml file shown in Listing A-9. Note how the node attributes in the mapping file correspond to the positions of the elements and attributes in the exported XML file. Note also how setting the embed-xml attribute to false substitutes the id value for the generated XML in the elements at the catalog/category/advert path.

Listing A-9. *The XML Exported by Hibernate*

```
<?xml version="1.0" encoding="UTF-8"?>
<catalog>
    <user id="1" name="dave" password="dodgy"/>
    ...
    <advert id="1" user="1">
        <content>48k original box and packaging</content>
        <title>Sinclair Spectrum for sale</title>
    </advert>
    ...
```

```
    <category id="1" name="retro">
        <advert>1</advert>
        <advert>3</advert>
    </category>
    ...
</catalog>
```

Importing XML Entities

The session in Dom4J entity mode can be used to import data into Hibernate as well as export it. The process is really just the export process in reverse—although there is one gotcha to be aware of: by default, the entity names will be assumed to be the same as the element nodes that you attempt to persist from the Dom4J document. Unless your node attributes in the mapping files correspond exactly with the entity names used in the mapping files, calls to the session will need to include the explicit entity name being saved.

Listing A-10 shows an example of a workaround for this issue, in which a hard-coded map translates short element names into full entity names, and the version of the save() method that takes an entity name string is invoked. When the class elements' node attributes correspond exactly with entity names, this sort of approach becomes unnecessary.

Listing A-10. *Importing Entities from an XML Document*

```java
package sample.xml;

import java.util.*;
import org.dom4j.*;
import org.dom4j.io.SAXReader;
import org.hibernate.*;
import org.hibernate.cfg.Configuration;

public class ImportXML {
    private static final SessionFactory sessionFactory = new Configuration()
            .configure().buildSessionFactory();

    public static void main(String[] args) throws Exception {

        System.out.println("Preparing the DOM Document");
        SAXReader reader = new SAXReader();
        Document document = reader.read("catalog.xml");

        System.out.println("Preparing the Session objects");
        Session session = sessionFactory.openSession();
        Session xmlSession = session.getSession(EntityMode.DOM4J);

        System.out.println("Importing the catalog from the document");
        session.beginTransaction();
```

```
    Map names = new HashMap();
    names.put("user","sample.entity.User");
    names.put("advert","sample.entity.Advert");
    names.put("category","sample.entity.Category");

    List entities = document.getRootElement().content();
    Iterator eit = entities.iterator();
    while(eit.hasNext()) {
       Node item = (Node)eit.next();
       String entityName = (String)names.get(item.getName());
       xmlSession.save(entityName,item);
    }

    session.getTransaction().commit();
    session.close();

    System.out.println("Done.");
  }
}
```

Unfortunately, the default value for the node attribute, if it is not explicitly applied to the class elements in the mapping files, is the short (unqualified) name of the class—not the entity name, which defaults to the long (fully qualified) class name.

Other Considerations When Using XML Entities

The objects retrieved from Hibernate in the EntityMode.DOM4J session mode are Dom4J objects, but they are *still* Hibernate persistence entities. You can therefore manipulate these entities through the Dom4J API and persist these changes to the database without needing to access the entity as a POJO.

Tools and APIs that process an XML document can therefore be applied to a Dom4J document extracted from Hibernate, and as long as it is still associated with the originating session, changes made to the document will be reflected in the database when the session is flushed or the transaction is committed.

Hibernate and EJB 3 annotations do not provide support for the use of XML relational persistence. If you want to use this feature to import or export data for an annotated application, you will first need to generate appropriate XML-based mapping files from the annotated classes.

Maps

In addition to the default mode (POJO) and the XML mode (Dom4J) described previously, the Hibernate session can be accessed in one more way: as a map of name/value pairs. This mode is accessed by calling the getSession() method with a parameter of EntityMode.MAP (see Listing A-11).

Listing A-11. *Accessing a Hibernate Session in Map Mode*

```
package sample.map;

import java.util.*;
import org.hibernate.EntityMode;
import org.hibernate.*;
import org.hibernate.cfg.Configuration;

public class AccessAsMap {
   private static final SessionFactory sessionFactory = new Configuration()
         .configure().buildSessionFactory();

   public static void main(String[] args) throws Exception {

      System.out.println("Preparing the Session objects");
      Session session = sessionFactory.openSession();
      Session mapSession = session.getSession(EntityMode.MAP);

      System.out.println("Reading the map entries for XXX");
      session.beginTransaction();

      Map entity = (Map)mapSession.get("sample.entity.Category",new Long(2));
      System.out.println("Category Title: " + entity.get("title"));

      System.out.println("Contains Adverts:");
      Set adverts = (Set)entity.get("adverts");
      Iterator adIt = adverts.iterator();
      while(adIt.hasNext()) {
         Map advert = (Map)adIt.next();
         System.out.println(advert.get("title"));
      }

      session.getTransaction().commit();
      session.close();

      System.out.println("Done.");
   }
}
```

This mode works much the same as the Dom4J mode—changes written to the Map objects will be persisted exactly as if a normal persistent POJO object had been updated. Note that only the entities themselves will be represented as Maps—not any of their attributes having a value type, or associations using Collection types. For example, in Listing A-11, the Category entity is represented as a Map, but its title attribute is represented as a String and its adverts attribute is represented as a Set—however, the Set itself contains Advert entities represented as Maps.

Limitations of Hibernate

First and foremost, Hibernate wants every entity to be identifiable with a primary key. Ideally, it would like this to be a *surrogate key* (a single column distinct from the fields of the table). Hibernate will accept a primary key that is not a surrogate key. For example, the username column might be used to uniquely identify an entry in the user table. Hibernate will also accept a composite key as its primary key, so that the username and hostname might be used to form the primary key if the username alone does not serve to identify the row.

In the real world, things do not really work like that. Any database that has been around the block a few times is likely to have at least one table for which the primary key has been omitted. For instance, the contents of the table may not have needed to be involved in any relations with other tables. While this is still bad database design, the error is only exposed when Hibernate tries to map objects to data. It may be that adding a suitable surrogate key column is an option—when this is the case, we urge you to do so. In practice, however, the fundamental schema may not be under the developer's control, or other applications may break if the schema is radically changed.

In most scenarios, a developer will be able to arrange the creation of views or stored procedures. It may be possible to create the appearance of a suitable primary key using these if no other options present themselves, but you should consult with your database administrators, since a table for which no true primary key can be obtained is likely to cause long-term corruption of your data.

Finally, if you can neither change a broken schema nor add views or stored procedures to ameliorate its effects, you have the option of obtaining a pure JDBC connection (see Listing A-12) from the session to the database, and carrying out traditional connected database access. This is the option of last resort, and is only truly of value when you anticipate being able to correct the faulty schema at some future time.

Listing A-12. *Obtaining a JDBC Connection from Hibernate*

```
SessionFactory factory =
        new Configuration().configure().buildSessionFactory();
Session session = factory.openSession();
Connection connection = session.getConnection();
```

Hand-Rolled SQL

While Hibernate cannot operate upon entities that lack primary keys, it is also extremely awkward to use Hibernate when there is a poor correspondence between the tables and the classes of your object model.

Using a Direct Mapping

Figure A-1 presents a fairly typical example of a valid database model that may be painful to represent in our mapping files.

Figure A-1. *A problematic but legal schema*

Here, the `product` table represents a product (for example, a flashlight). The `color` table represents the colors in which it is sold. The link table named `product_color` then allows us to identify a product by stock keeping unit (SKU), and identify the colors in which it is available.

If we do not mind the `Product` object retaining a set of colors (representing the colors in which it can be sold), then we have no problem; but if we want to distinguish between a red flashlight and a green one, things become more difficult (see Listing A-13).

Listing A-13. *A Fairly Direct Representation of the Product*

```
<class name="com.hibernatebook.legacy.Product" table="product_color">

    <composite-id
        class="com.hibernatebook.legacy.ProductKey"
        name="key">

        <key-property type="int" name="id" column="product_id"/>
        <key-property type="int" name="colorId" column="color_id"/>

    </composite-id>


    <many-to-one
        name="color"
        class="com.hibernatebook.legacy.Color"
        column="color_id"
        insert="false"
        update="false"/>

    <many-to-one
        name="data"
        class="com.hibernatebook.legacy.ProductData"
        column="product_id"
        insert="false"
        update="false"/>

</class>
```

There are several dissatisfying aspects to the mapping in Listing A-13. First, rather than mapping our product table, we have mapped the link table. This makes sense when you consider that the primary key formed from the two columns of this table uniquely identifies a "colored product," which the `product` table alone cannot do.

Second, we are obliged to create a number of distinct objects to represent the class: the `Product` class itself, a class to represent the primary key (inevitable where a composite `id` occurs), a class to represent the other attributes of the product, and the `Color` class.

Last, the use of the columns more than once within the mapping requires us to flag them so that they cannot be written—this is a *read-only mapping*.

Using a View

Fortunately, most databases provide a simple mechanism for manipulating a schema so that it better matches the business requirements. A database view will allow you to put together a join that appears to be a table. By a suitable choice of columns from the existing tables, you can construct a view that is much easier to map (see Listing A-14).

Listing A-14. *A View on the Product Tables*

```
create view vwProduct (ProductKey,ColorKey,Id,SKU,ColorId)
AS
   select
      p.id as ProductKey,
      c.id as ColorKey,
      p.id as Id,
      p.sku as SKU,
      c.id as ColorId
   from
      product p,
      product_color pc,
      color c
   where
      p.id = pc.product_id
   and
      pc.color_id = c.id;
```

This view effectively reformats our table so that it has a correct (composite) primary key formed from the link table's two columns. It makes the SKU data available directly, and it retains the foreign key into the `color` table.

Listing A-15 is a much more natural mapping.

Listing A-15. *The Revised Mapping*

```
<class name="com.hibernatebook.legacy.Product" table="vwProduct">
  <composite-id
      class="com.hibernatebook.legacy.ProductKey"
      name="key">
      <key-property
         type="int"
         name="id"
         column="ProductKey"/>
      <key-property
         type="int"
         name="colorId"
         column="ColorKey" />
  </composite-id>

  <property
      name="id"
      type="int"
      column="id"
      insert="false"
      update="false"
      unique="true"/>

  <property
      name="SKU"
      type="int"
      column="sku"
      insert="false"/>

  <many-to-one
      name="color"
      class="com.hibernatebook.legacy.Color"
      column="ColorId"/>
</class>
```

The behavior of the composite primary key is unchanged, but the SKU now becomes a simple property. The color entity is mapped as before.

The caveat for this approach is the problem of writing data to the mapping. Some databases (for example, versions 4 and lower of MySQL) do not support writable views, and others may have only limited support for them. To avoid views in these circumstances, we must abandon complete portability in favor of database-specific SQL inserted directly into the mapping file.

Putting SQL into a Mapping

Hibernate provides three tags that can be used to override the default behavior when writing to the database. Instead of accepting the SQL generated by Hibernate from the information in

the mapping file, you can dictate exactly how changes to an entity should be enforced. The disadvantage is that you will lose Hibernate's guarantee of cross-database platform portability. The advantage is that you can carry out operations that are not explicitly described in the mapping, such as calculating and inserting values in the process of carrying out an insert.

The tags are `<sql-insert>`, `<sql-update>`, and `<sql-delete>`. All three work in the same way.

If you take a look at the DDL script for this appendix, you will see that our client table includes seven fields, the last of which is the `country` field, as shown in Listing A-16.

Listing A-16. *The DDL Script to Create the Client Table*

```
create table client (
    id int not null primary key,
    name varchar(32) not null,
    number varchar(10),
    streetname varchar(128),
    town varchar(32),
    city varchar(32),
    country varchar(32)
);
```

We will, however, ignore the `country` field in our mapping file. We would like this to be automatically set to UK whenever a client entity is persisted.

Depending on the database, this could be implemented as part of the view's support for writing operations, or as a trigger invoked when the other fields are written—but we use the `<sql-insert>` tag to specify the operation to perform.

The necessary ordering of the parameters can be determined by running Hibernate with logging enabled for the `org.hibernate.persister.entity` level. You must do this before you add the mapping. Listing A-17 shows a suitably formatted `<sql-insert>` element with the parameters suitably ordered. Note that the identifier field `id` is in the last position—not the first, as you might have expected.

Listing A-17. *The Mapping File Using Explicit SQL to Update the Tables*

```
<class name="com.hibernatebook.legacy.Client" table="Client">
    <id type="int" name="id" column="id">
        <generator class="native"/>
    </id>
    <property type="text" name="name" column="name"/>
    <property type="text" name="number" column="number"/>
    <property type="text" name="streetname" column="streetname"/>
    <property type="text" name="town" column="town"/>
    <property type="text" name="city" column="city"/>

    <sql-insert>
    insert into client(name,number,streetname,town,city,id,country)
    values (?,?,?,?,?,?,'UK');
    </sql-insert>
</class>
```

In addition to the three SQL terms for writing to the database, you can specify hand-rolled SQL for reading. This is appended as `<sql-query>` tags outside the `class` tag (see Listing A-18). They are not intrinsically a part of the mapping. However, you *can* specify that one of them should be used as the default loader for your class.

Listing A-18. *An Alternative Mapping File Defining a Default Loader*

```
...
<hibernate-mapping>
   <class name="com.hibernatebook.legacy.Client"
          table="Client">
      <id type="int" name="id" column="id">
         <generator class="native"/>
      </id>
      <property name="name"/>
      <property name="number"/>
      <property name="streetname"/>
      <property name="town"/>
      <property name="city"/>

      <loader query-ref="DefaultQuery"/>
   </class>

<sql-query name="DefaultQuery">
   <return alias="c"
           class="com.hibernatebook.legacy.Client"/>
   SELECT
      id as {c.id},
      'NOT SPECIFIED' as {c.name},
      number as {c.number},
      streetname as {c.streetname},
      town as {c.town},
      city as {c.city}
   FROM
      Client
   WHERE
      id = ?
</sql-query>

</hibernate-mapping>
```

■ **Tip** Unfortunately, this technique is not quite as sophisticated as you might hope—the custom SQL will not be invoked in general terms. Only if the `id` is explicitly supplied, as is the case when calling the `Session` class's `get()` method, will the default handling be overridden in favor of the loader query.

Invoking Stored Procedures

Data outlives application logic. This is a general rule of thumb, and as we can attest, it holds true in practice. The natural lifespan of a database will tend to see multiple applications. The lifespan of some of these applications will, in turn, tend to overlap, so that at any one time we expect substantially different code bases to be accessing the same data.

To resolve such issues, databases usually provide their own programming language to allow complex business rules to be expressed and enforced within the boundary of the database itself. These languages are expressed in stored procedures—essentially an API to the database. Often, free-form SQL access to such a database is denied, and only access through stored procedures is permitted. Barring errors in the code of the stored procedures themselves, this removes any risk of corruption.

One final advantage of using stored procedures is that when a substantial calculation is required, the use of a stored procedure can reduce the network traffic involved. For example, if you invoke a stored procedure to calculate the grand total of a table of accounts, only the request and the result figure would need to traverse the network. The equivalent client-side implementation would need to acquire the value to be totaled from every row!

Taking the client example from the "Putting SQL into a Mapping" section, we could replace the SQL logic in the `<sql-insert>` tag with a call to a suitable stored procedure. The callable attribute is set to `true` to indicate that Hibernate needs to issue a call to a stored procedure instead of a standard query (see Listing A-19).

Listing A-19. *Mapping a Call to the Stored Procedure*

```
<sql-insert callable="true">
   {call insertClient(?,?,?,?,?,?)}
</sql-insert>
```

In the stored procedure definition (see Listing A-20), you will note that the order of the parameters to be passed in has been tailored to match the order in which they will be provided by Hibernate.

Listing A-20. *The Logic of the Stored Procedure*

```
CREATE PROCEDURE
   insertClient( p_name varchar(32),
                 p_number varchar(10),
                 p_streetname varchar(128),
                 p_town varchar(32),
                 p_city varchar(32),
                 p_id int)
AS
BEGIN
   INSERT INTO client
      (id,name,number,streetname,town,city,country)
   VALUES
      (:p_id,:p_name,:p_number,:p_streetname,:p_town,:p_city,'UK');
END
```

By obtaining a JDBC connection from the session, it is of course possible to invoke stored procedures directly; however, you must be aware that the Hibernate session cache cannot track these updates.

Events

Hibernate 3 actually implements most of its functionality as event listeners. When you register a listener with Hibernate, the listener entirely supplants the default functionality. For example, the EJB 3–specific behavior required when using an EntityManager is achieved by plugging in a set of EJB 3–specific event listeners!

If you look at the methods of the SessionImpl class, which is the internal Hibernate implementation of the Session interface, you'll see why this is the case. Most of the methods have a form very similar to that shown in Listing A-21.

Listing A-21. *The Implementation of a Typical Method in* SessionImpl

```
SaveOrUpdateEvent event = new SaveOrUpdateEvent(entityName, obj, this);
listeners.getSaveOrUpdateEventListener().onSaveOrUpdate(event);
```

The listeners field is an instance of SessionEventListenerConfig, which provides the requested event listener, or the default if none is specified. So, if your event listener is provided and doesn't call the default one, nothing else can.

Event listeners are always registered globally for the event that they handle. You can register them in the configuration file or programmatically. Either way, you will need to map your implementation of one of the interfaces to the associated types, which you can look up in Table A-1. (The names are almost—but not quite—standardized.)

Table A-1. *The Listener Names and Their Corresponding Interfaces*

Type Name	Listener
auto-flush	AutoFlushEventListener
delete	DeleteEventListener
dirty-check	DirtyCheckEventListener
evict	EvictEventListener
flush	FlushEventListener
flush-entity	FlushEntityEventListener
load	LoadEventListener
load-collection	InitializeCollectionEventListener
lock	LockEventListener
merge	MergeEventListener
persist	PersistEventListener
post-delete	PostDeleteEventListener
post-insert	PostInsertEventListener
post-load	PostLoadEventListener
post-update	PostUpdateEventListener

Type Name	Listener
pre-delete	PreDeleteEventListener
pre-insert	PreInsertEventListener
pre-load	PreLoadEventListener
pre-update	PreUpdateEventListener
refresh	RefreshEventListener
replicate	ReplicateEventListener
save-update	SaveOrUpdateEventListener

So, for example, your listener for the SaveOrUpdateEvent is mapped to the type name save-update, must implement the SaveOrUpdateEventListener interface, and would normally have been implemented by the DefaultSaveOrUpdateEventListener class. It is wise to follow a similar convention with your own naming, so your mapping file listener entry might read like this:

```
<listener type="save-or-update"
        class="com.hibernatebook.advanced.BookingSaveOrUpdateEventListener"/>
```

Alternatively, a programmatic registration of the same event would be given thus:

```
Configuration config = new Configuration();
config.setListener("save-update", new BookingSaveOrUpdateEventListener());
```

Because they override the default behavior, events are suitable for situations in which you want to fundamentally change the Session's behavior—particularly if you want to prevent a certain event from being processed. Probably the best example of this requirement is in authorizing access to the database, and, in fact, Hibernate provides a set of event listeners for just this purpose. The four events listeners in question override the PreDelete, PreUpdate, PreInsert, and PreLoad listeners. The logic in each case (in pseudocode) runs something like this:

```
if( user does not have permission ) throw RuntimeException
Invoke default listener…
```

Because events are invoked in the same thread as the user's call to the session, the result of an exception in the first step will be an exception (actually a security exception) as the unprivileged user carries out the relevant operation.

To enable policy configuration of security, you would add the following:

```
<listener type="pre-delete"
        class="org.hibernate.secure.JACCPreDeleteEventListener"/>
<listener type="pre-update"
        class="org.hibernate.secure.JACCPreUpdateEventListener"/>
<listener type="pre-insert"
        class="org.hibernate.secure.JACCPreInsertEventListener"/>
<listener type="pre-load"
        class="org.hibernate.secure.JACCPreLoadEventListener"/>
```

An Example Event Listener

Before we get stuck in a simple example, a word of caution: events are very much an exposed part of the inner workings of the Session. While this is ideal for something requiring the level of interference of a security tool, you will not need this for most purposes. Listing A-22 is more in the nature of an illustrative "hack" than a real solution. In a real application, you would probably solve this particular problem either within the body of the application, by using interceptors, or by using triggers.

Listing A-22 shows how an event listener could be used to prevent the booking of certain seats from being persisted to the database in a concert hall ticket booking application (we revisit this example application in a little more detail with a slightly more realistic scenario in the later section, "Interceptors").

Listing A-22. *Programmatically Installing an Event Listener*

```
public class EventExample {
   public static void main(String[] args) {
      Configuration config = new Configuration();

      // Apply this event listener (programmatically)
      config.setListener("save-update", new BookingSaveOrUpdateEventListener());

      SessionFactory factory = config.configure().buildSessionFactory();
      Session session = factory.openSession();

      Transaction tx = session.beginTransaction();

      // Make our bookings... seat R1 is NOT to be saved.
      session.saveOrUpdate(new Booking("charles","R1"));
      session.saveOrUpdate(new Booking("camilla","R2"));

      // The confirmation letters should not be sent
      // out until AFTER the commit completes.
      tx.commit();
   }
}
```

Our example is only going to implement the SaveOrUpdateEventListener interface. You will notice that in Listing A-22, the original calls to save() have been replaced with calls to saveOrUpdate(). There is a close correspondence between the methods on the Session interface and the event listeners with similar names. A call to save() will not invoke saveOrUpdate(), and vice versa. Try using the save() method in the EventExample, and you will see that the BookingSaveOrUpdateListener is not invoked. In Listing A-23, we present the logic of the listener registered in Listing A-22.

Listing A-23. *The Implementation of an Event Listener*

```java
package com.hibernatebook.advanced.events;

import java.io.Serializable;

import org.hibernate.HibernateException;
import org.hibernate.event.SaveOrUpdateEvent;
import org.hibernate.event.def.DefaultSaveOrUpdateEventListener;


public class BookingSaveOrUpdateEventListener
    extends DefaultSaveOrUpdateEventListener
{
    public Serializable onSaveOrUpdate(SaveOrUpdateEvent event)
         throws HibernateException {
       if( event.getObject() instanceof Booking ) {
          Booking booking = (Booking)event.getObject();
          System.out.println("Preparing to book seat " + booking.getSeat());

          if( booking.getSeat().equalsIgnoreCase("R1")) {
             System.out.println("Royal box booked");
             System.out.println("Conventional booking not recorded.");

             // By returning null instead of invoking the
             // default behavior, we prevent the invocation
             // of saveOrUpdate on the Session from having
             // any effect on the database!
             return null;
          }
       }

       // The default behavior:
       return super.onSaveOrUpdate(event);
    }
}
```

Interceptors

Interceptors are privy to a blow-by-blow account of what is going on as Hibernate carries out its duties. While you can listen in, you can only make limited changes to the way in which Hibernate actually behaves. This is the common requirement; unless you are making substantial changes to the persistence behavior, you will usually want only to track what is going on.

Financial packages often require considerable auditing information to be maintained to prevent fraud and aid accountability. Auditing is a natural candidate for implementation as an interceptor, as it would normally require that no changes be made to the persistence process at all.

The question that usually arises when discussing interceptors is "why not use triggers?" Triggers should never embody application logic, only business logic. If any application is going to have audit-free access to the database, you cannot implement the auditing in triggers. Worse, the triggers may not have access to the user information that's needed. In most multi-tier situations, the need to pool the database connections precludes establishing individual user logins to the database. So, for example, the trigger would only know that a user with the login "MonolithicApplication" carried out an update of last year's sales figures—not that it was carried out by, say, Jim from accounts, which is who the auditors are likely to be interested in! Table A-2 summarizes the points in the application life cycle at which the various methods will be invoked.

Table A-2. *The Interceptor Methods*

Name	When Invoked	Comments
afterTransactionBegin()	Invoked immediately after a call to begin() on a Transaction object retrieved from the Session object.	This method can change the state of the transaction—for example, it can call rollback().
afterTransactionCompletion()	Invoked immediately after the completion of a transaction.	
beforeTransactionCompletion()	Invoked immediately prior to the completion of a transaction. This method can change the state of the transaction—for example, it can call rollback().	
findDirty()	Invoked during calls to flush().	This allows the saving of changes to attributes to be prevented or forced.
getEntity()	Invoked when an entity not in the Session object's own cache is requested by its identifier.	
getEntityName()	Invoked when the Session object needs to determine the name of a given entity.	
instantiate()	Invoked when the Session object needs to create an entity instance.	Because the "empty" object can be created here, this allows Hibernate (in legacy applications, for example) to use entities that do not have a default constructor.
isTransient()	Invoked when the Session object needs to determine whether an entity it has been asked to persist is transient—for example, during calls to saveOrUpdate().	
onDelete()	Invoked before an object is deleted.	The object's state should not be tampered with at this point.

Name	When Invoked	Comments
onFlushDirty()	Invoked during a call to flush() after entities have been determined to be dirty. (If the entities are not dirty, then there are no changes to be persisted and Hibernate has no actions to perform—therefore, there is no general case interceptor, and this interceptor will *not* be invoked if the entities are clean.)	
onLoad()	Invoked immediately before an entity object is populated from the database.	The loading can be overridden (by returning false), and the instantiated but uninitialized object is available if supplementary initialization from the listener is needed.
onSave()	Invoked before an object is saved.	This permits the state of the object to be changed immediately before it is saved.
postFlush()	Invoked after the Session object is flushed, if and only if the Session object had to carry out SQL operations to synchronize state with the database.	
preFlush()	Invoked immediately before the Session object is flushed.	

An Example Interceptor

To illustrate how all this works in practice, we will create a simple interceptor from scratch. While the auditing example is a good one, it is rather too involved for our demonstration. Instead, we will consider a concert hall seat-booking system (the entity to represent an entity is shown in Listing A-24) for which the details of bookings will be sent out to customers as they are pushed into the database.

Listing A-24. *The* Booking *POJO*

```
package com.hibernatebook.advanced.events;

public class Booking {
   public Booking(String name, String seat) {
      this.name = name;
      this.seat = seat;
   }

   Booking() {
   }
```

```
   protected String getName() {
      return name;
   }

   protected void setName(String name) {
      this.name = name;
   }

   protected String getSeat() {
      return seat;
   }

   protected void setSeat(String seat) {
      this.seat = seat;
   }

   private String seat;
   private String name;
}
```

Interceptors have to override the org.hibernate.Interceptor interface. You can set a global interceptor for the configuration (see Listing A-25), or you can apply interceptors on a per-session basis. You have to install the interceptor programmatically—there is no syntax for specifying this in the Hibernate configuration file.

Listing A-25. *Installing a Global Interceptor*

```
package com.hibernatebook.advanced.events;

import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.Transaction;
import org.hibernate.cfg.Configuration;

public class MailingExample {
   public static void main(String[] argv) {
      Configuration config = new Configuration();

      // Apply this interceptor at a global level...
      config.setInterceptor(new BookingInterceptor());

      SessionFactory factory = config.configure().buildSessionFactory();
      Session session = factory.openSession();

      // A local interceptor could alternatively
      // be applied here:
      // session.setInterceptor(new BookingInterceptor());
```

```
      Transaction tx = session.beginTransaction();

      // Make our bookings...
      session.save(new Booking("dave","F1"));
      session.save(new Booking("jeff","C3"));

      // The confirmation letters should not be sent
      // out until AFTER the commit completes.
      tx.commit();
    }
}
```

The interceptor that we are applying is going to capture the information from the `Booking` objects that we are storing in the database. Listing A-26 demonstrates the basic mechanism, but it is only a toy example. We will discuss some of its deficiencies in a moment.

Listing A-26. *An Interceptor Implementation*

```
package com.hibernatebook.advanced.events;

import java.io.Serializable;
import java.util.Collection;
import java.util.HashSet;
import java.util.Iterator;

import org.hibernate.CallbackException;
import org.hibernate.EntityMode;
import org.hibernate.Interceptor;
import org.hibernate.Transaction;
import org.hibernate.type.Type;

public class BookingInterceptor implements Interceptor {

   public BookingInterceptor() {
   }

   private ThreadLocal stored = new ThreadLocal();

   public void afterTransactionBegin(Transaction tx) {
      stored.set(new HashSet());
   }

   public void afterTransactionCompletion(Transaction tx) {
      if (tx.wasCommitted()) {
         Iterator i = ((Collection) stored.get()).iterator();
```

```java
        while (i.hasNext()) {
            Booking b = (Booking) i.next();
            sendMail(b);
        }
    }
    stored.set(null);
}

public boolean onSave(Object entity, Serializable id,
        Object[] state, String[] propertyNames, Type[] types)
        throws CallbackException {
    ((Collection) stored.get()).add(entity);
    return false;
}

private void sendMail(Booking b) {
    // Here we would actually send out the e-mail
    System.out.print("Name: " + b.getName());
    System.out.println(", Seat: " + b.getSeat());
}

public void beforeTransactionCompletion(Transaction tx) {
}

public int[] findDirty(Object entity, Serializable id,
        Object[] currentState, Object[] previousState,
        String[] propertyNames, Type[] types) {
    return null;
}

public Object getEntity(String entityName, Serializable id)
        throws CallbackException {
    return null;
}

public String getEntityName(Object object) throws CallbackException {
    return null;
}

public Object instantiate(String entityName, EntityMode entityMode,
        Serializable id) throws CallbackException {
    return null;
}

public Boolean isTransient(Object object) {
    return null;
}
```

```
    public void onDelete(Object entity, Serializable id, Object[] state,
            String[] propertyNames, Type[] types) throws CallbackException {
    }

    public boolean onFlushDirty(Object entity, Serializable id,
            Object[] currentState, Object[] previousState,
            String[] propertyNames, Type[] types) throws CallbackException {
        return false;
    }

    public boolean onLoad(Object entity, Serializable id,
            Object[] state, String[] propertyNames, Type[] types)
            throws CallbackException {
        return false;
    }

    public void postFlush(Iterator entities) throws CallbackException {
    }

    public void preFlush(Iterator entities) throws CallbackException {
    }
}
```

Our interceptor makes use of the afterTransactionBegin() method to prepare to collect booking details, the onSave() method to collect them, and the afterTransactionCompletion() method to report the successful bookings. This sequence guarantees that bookings will not be reported to the users until after we are confident that they have been retained in the database.

A minor deficiency of this implementation is that the e-mail is sent outside the transaction—a system failure immediately after the commit completes could cause the e-mail not to be sent. In our scenario, this is unimportant, because e-mail is already an unreliable transport mechanism; but there are other situations, such as the auditing example discussed earlier, in which this may be unacceptable. In these cases, interception may not be appropriate, and an integrated solution tied into a two-phase commit transaction may be required.

More importantly, our example assumes that the Booking object will not be altered between its addition to the set of e-mails to be sent and their transmission. This is an extremely dangerous assumption! A safer approach would be to create copies of the Booking objects or, better yet, to copy their data into a more appropriate object, as shown in Listing A-27.

Listing A-27. *A Better Approach to Preparing the Mailshot*

```
    public boolean onSave(Object entity, Serializable id,
            Object[] state, String[] propertyNames, Type[] types)
            throws CallbackException
    {
        if( entity instanceof Booking ) {
            Booking booking = (Booking)entity.
            Mailshot mailshot = new Mailshot(booking.getName(), booking.getSeat() );
```

```
        ((Collection) stored.get()).add(mailshot);
      }
      return false;
  }
```

Finally, we don't necessarily have enough information to prepare our mailshot—the e-mail address may be missing. If the name field actually represents the e-mail address, then we are fine; but if it represents a key into other objects, and hence tables in the database, then we have to be careful. It is possible to write database logic from within an interceptor, but the risk of accidentally recursing back into your interceptor logic is high, so we don't recommend it. It's slightly less tricky if you are only using a session-scoped interceptor, but there are probably safer ways to achieve the same end result.

In this example, the methods that we are not using have been given a default implementation. You will note that some of these methods return null or false. These methods are permitted to change the data that is preserved or returned by the session. Returning to the onSave() method, we will consider another possible implementation, shown in Listing A-28.

Listing A-28. *Changing the Data from Within an Interceptor*

```
    public boolean onSave(Object entity, Serializable id,
          Object[] state, String[] propertyNames, Type[] types)
          throws CallbackException
    {
      if( entity instanceof Booking ) {
          state[1] = "unknown";
      }
      return true;
    }
```

Here we are altering the state array. This contains the values of each of the objects' fields that are to be stored (in the order defined in the mapping file). In our Booking class, field 0 is the id, field 1 is the name, and field 2 is the seat—so here we have changed the name value to be preserved in the database. Returning true causes Hibernate to reflect this change when it saves the data. If we left the return flag as false, nothing would happen when the method was called.

The temptation is to assume that returning false guarantees the safety of the data to be preserved, but, in fact, this is not the case. The state array represents copies of the data to be preserved—but we have also been given access to the actual object (entity) that contains the original values. If you amend the fields of the entity before returning, the flag will not prevent your changes from being made. Listing A-29 illustrates how this might occur.

Listing A-29. *Changing the Data in an Unorthodox Way*

```
public boolean onSave(Object entity, Serializable id,
        Object[] state, String[] propertyNames, Type[] types)
        throws CallbackException
```

```
{
    if( entity instanceof Booking ) {
        Booking booking = (Booking)entity;
        booking.setName("unknown");
    }
    // The flag can't save us from ourselves here!
    return false;
}
```

Again, this is probably not the best way to make the changes, but it can be useful when you already have a considerable body of logic prepared to process the entity type.

Overriding the Default Constructor

Occasionally, you will find that it is necessary to persist a POJO that has no default constructor. Usually you will have access to the source code, and should just make the change directly. Occasionally, however, you may find that you are working with classes for which the source code is not available—or that you are working with large bodies of generated objects for which it is extremely inconvenient to manage changes made to the source code. In these circumstances, it is possible to use an interceptor to replace the default object-creation logic.

This technique can be used as long as you have some way of obtaining or applying default values for the parameters of the POJO's non-default constructor. Listing A-30 shows an example of the use of this technique to instantiate a POJO whose only constructor demands a String parameter.

Listing A-30. *Invoking a Non-Default Constructor*

```java
private static class OverrideCtor implements Interceptor {

    public Object instantiate(
        String entityName,
        EntityMode entityMode,
        Serializable id)
            throws CallbackException
    {
        if( entityName.equals(MissingDefaultCtorClass.class.getName())) {
            // My call to CTor
            return new MissingDefaultCtorClass("NOT SET");
        } else {
            // Some other class - continue to default handling
            return null;
        }
    }

    // ... the remaining default method declarations...
}
```

Summary

In this appendix, we have examined Hibernate's place in the new Enterprise Java Beans standard (EJB 3), and looked at alternative mechanisms for accessing Hibernate entities. We have shown how SQL and stored procedures can be integrated into the Hibernate environment, and we have discussed how events and listeners provide internal access to Hibernate's persistence mechanism.

■ ■ ■

Hibernate Tools

The Hibernate Tools toolset really consists of two quite distinct tools: a set of plug-ins to enhance the Eclipse integrated development environment (IDE), and a set of tasks for the Ant build tool. They are packaged together because they share significant chunks of implementation despite their very different façades. We have already used one of the Ant tasks from Hibernate Tools in earlier chapters to generate our database schemas. In this appendix, we will discuss the other available tasks. First, however, we will discuss the use of the toolset as a plug-in for Eclipse.

■**Caution** At the time of writing, Hibernate Tools is at an advanced beta stage of development. You should be aware that you may encounter rough edges and bugs when using these tools. Currently, the Ant tasks are more polished than the Eclipse plug-ins. Even so, we think that it is well worth familiarizing yourself with all these tools and even using them in production, as most of the problems can be worked around.

It is beyond the scope of this book to attempt to teach you how to use Ant or Eclipse (although we do walk you through some of the less common configuration details). To get the most from this appendix, you should be familiar with both Ant and Eclipse—although it is possible to use both parts of the plug-in independently.

General information on the latest versions of Hibernate Tools, any changes or new features, the online documentation, and the locations of the various downloads are available from the Hibernate web site (`http://tools.hibernate.org`).

The Eclipse Plug-In

Eclipse is one of the best-known and liked Java development environments to emerge in recent years. Eclipse evolved originally as a proprietary component of IBM's WebSphere Application Developer (WSAD) environment. IBM chose to release Eclipse, the IDE, as an open source application. Thus, the open source Eclipse IDE emerged as a mature product virtually overnight.

Eclipse is designed as a core application, the Eclipse platform, extended by various plug-ins, typically including Java Development Tools (JDT). For our purposes, we assume that you will start out with this combination (known rather confusingly as the Software Development

Kit, or SDK) as the basis for installing Hibernate Tools. At the time of writing, the latest version of the SDK is 3.1.2.

Eclipse can be downloaded from the Eclipse Foundation web site (`www.eclipse.org`). You will find a file named `eclipse-SDK-3.1.2-win32.zip` in the Downloads section.

The Hibernate team is now employed by JBoss, who provides a branded version of Eclipse that includes the Hibernate plug-ins, including several JBoss-specific ones. This can be downloaded from the JBoss web site (`http://download.jboss.com/jbosside/builds`). If you choose to use this version of Eclipse, then you can omit the installation steps described in the next section.

Installing the Plug-In

We will now walk you through the process of installing the plug-in using Eclipse's standard updates feature.

Select the Find and Install menu option from the Help ➤ Software Updates menu.

You should then select the option to search for new features, as Hibernate is not a standard Eclipse SDK component.

By default, Eclipse will only be aware of the Eclipse Foundation web site as a source of new Eclipse features. Click the New Remote Site button shown in the upper right of Figure B-1 to add the site from which you will obtain the Hibernate Tools plug-in.
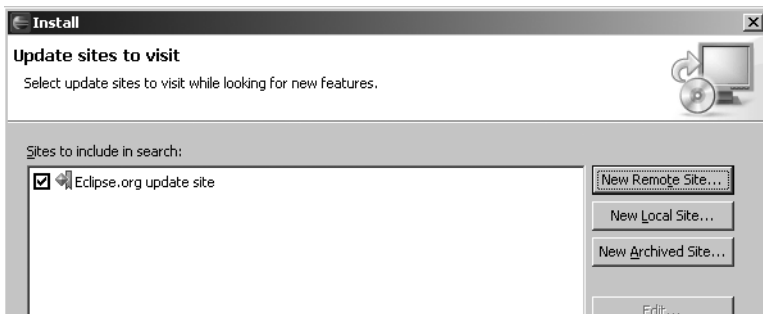
Figure B-1. *By default, only* `www.eclipse.org` *is available as a download site.*

You will now need to enter the URL for the site from which Eclipse will obtain the plug-in, and an informal name for the site, into the dialog shown in Figure B-2. Presently, the download site is at `http://download.jboss.org/jbosside/updates/development`—this is unlikely to change, but if you encounter problems, you should check the `http://tools.hibernate.org` web site to make sure that this is still listed as the update site.

Figure B-2. *Specifying the new download site*

You will then be returned to the list of sites shown in Figure B-3, where you should uncheck the Eclipse download site, as you will only be installing the Hibernate Tools plug-in.

Figure B-3. *Deselect the Eclipse update site.*

Click the Next button, and Eclipse will download the list of available updates from the updates site that you have provided.

If you have not provided the correct site name, or you have other connectivity issues, you will not be able to reach this step. If you are experiencing problems, go back and check that you have configured any necessary proxy settings, and that all the appropriate URLs have been entered correctly.

Once you successfully reach this step, you should select the Hibernate Tools check box and click Next.

Eclipse now lists the features that will be installed and prompts you to accept the license that applies to the plug-in features. You must accept the terms and conditions of the license to proceed beyond this step. Once you have checked the Accept radio button, click the Next button.

You will now be presented confirmation of the features to be installed, and given the option of installing the features to a nonstandard location. (We always accept the default installation into the Eclipse plug-ins directory.) Click Finish to proceed.

Currently, the Hibernate Tools plug-in is not digitally signed. Eclipse warns you of this. In principle, it is possible that a malicious third party with access to servers between you and the Eclipse download site could substitute their own code for the Eclipse tools. In practice this is unlikely, but it is to be hoped that the Hibernate or JBoss teams will start code signing their final releases. To proceed with the plug-in installation, you must accept the verification prompt.

Finally, Eclipse prompts you to restart or apply the changes directly. Generally when installing Eclipse plug-ins, it is safest to select the restart option. Though we know of no specific problems with the Hibernate Tools, we recommend choosing restart here anyway—it won't be necessary to reboot the PC, though!

At this point, you will have successfully completed the installation of the Hibernate Tools plug-in. There will be no immediate visible change to the layout of the workbench—however, as you will see in the next step, there should be some new options for you, accessible via the various menus—including a new Eclipse perspective onto some of the Hibernate workbench views.

■**Note** If you are unable or unwilling to install the software directly from Eclipse, it is also possible to download the tools as a ZIP file from the Hibernate site and copy the plug-ins directory from the archive directly over the top of Eclipse's own plug-ins directory. You should then restart Eclipse for the changes to take effect. We do not recommend that you try this approach unless you are already comfortable with the process of installing Eclipse plug-ins in this manner.

The Eclipse plug-in installation process is now quite streamlined, so it is unlikely that you will encounter any serious problems. However, if you do have problems, first check that you have the correct versions of the downloads described here, and check that you have followed the installation steps as given previously. If you are still encountering problems, you should search the Tools forum on the Hibernate Forums page (`http://forum.hibernate.org`) to see if other users have encountered the same problems. You should also check the Hibernate bug database (`www.hibernate.org/217.html`) to see if there are known problems that match yours.

The Boilerplate Project Configuration

Now that the tools plug-in has been installed, you are ready to set up a project that will take advantage of it. As a first step, we recommend configuring a user library to make adding the appropriate JAR files to new projects more straightforward, and to make the resulting project view less cluttered. Figure B-4 compares the view of the project's contents with and without the use of user libraries.
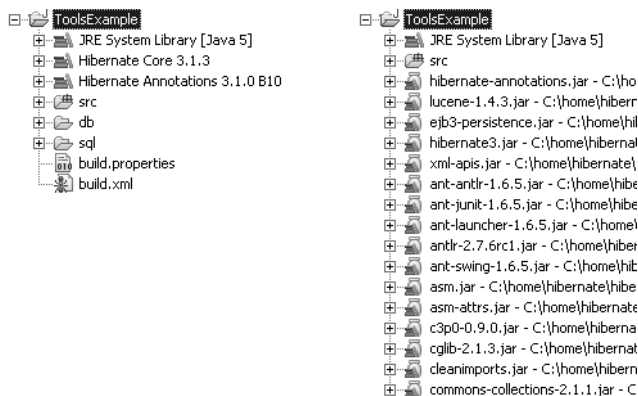
Figure B-4. *An Eclipse project using user libraries (left) compared with one using JAR files directly (right)*

You should be aware that Eclipse user libraries are not a Java standard, and do not exist as independent entities. They are merely a grouping of paths to JAR files to make configuring Eclipse projects more convenient. The user library will not reflect changes made to the underlying JAR files themselves (moving or deleting them, for instance).

Because user libraries belong to the Eclipse workspace, rather than to the individual projects, you can create a library from the Window ➤ Preferences menu before you have created any projects.

You should then drill down through the tree view to select the Java ➤ Build Path ➤ User Libraries node.

Initially there are no user libraries configured. You will need to click the add button to create your own.

You will be prompted for a name for the library. The first library will contain the JAR files for the Hibernate core. We recommend including the full version number in the library name so that you will be able to readily distinguish between versions if you are managing more than one Eclipse project in the same workspace over an extended period of time. (Hibernate updates come thick and fast!)

The library name will be added to the list of libraries, but it does not yet contain any JAR files. You should select the library name in the list, and then click the Add Jars button.

You now want to add all the core Hibernate 3 JAR files to the library. You can add multiple files, but not directories, so this must be carried out in two steps.

You will be presented with a normal file-selection dialog. You should navigate to your core Hibernate install directory and select the `hibernate3.jar` file.

You should select the library name and again click Add Jars, but now navigate to the `lib` subdirectory beneath the core Hibernate install directory. Select all the JAR files in this directory and add them to the library.

Having created the core Hibernate user library, you should repeat the process to create the Hibernate Annotations user library with an appropriate version number.[1] Figure B-5 shows part of an exploded view of the resulting pair of user libraries.
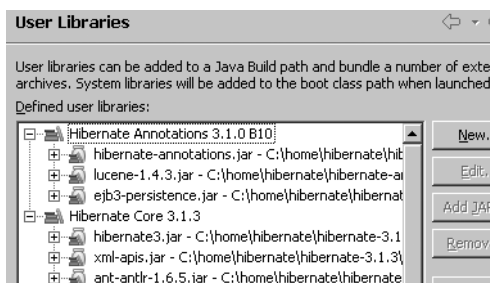
Figure B-5. *The two user libraries*

---

1. Naturally, this step is not essential if you do not intend to use Hibernate annotations in your project. Our example assumes throughout that you will be using annotations for the project's mappings.

Having created it initially, Hibernate allows you to export your user library configuration so that it does not need to be created afresh on each machine that you work on, or when creating a new Eclipse workspace.

You should now create a new Java project called ToolsExample. The source code for the example is available from the Apress web site (www.apress.com). Copy the files into the ToolsExample folder in the workspace directory, and refresh the project.

■**Tip** If Eclipse is running under a 1.4 JVM or earlier, and you want to take advantage of annotations, you must reconfigure it to run under Java 5. Add the -vm flag to the Eclipse command line to force the named JVM to load Eclipse. For example, the shortcut to Eclipse on my Windows desktop contains the following target:

```
C:\eclipse\eclipse.exe -vm c:\home\jdk1.5.0_03\bin\javaw
```

At the time of writing, Hibernate Tools contains a minor bug that causes some problems if the configuration file is in the project root. We have therefore configured the project to contain separate src and bin directories. You should configure the project to use these as the source folder and default output folder, respectively. We expect this bug to be fixed by the time you read this book, but it is good practice to separate the directories in this way anyway.

You now need to add the two user libraries to the project. Select the Build Path ➤ Configure Build Path context menu option on the project. Select the Libraries tab on the resulting dialog.

Click the Add Library button, and you will be presented with a list of library types. Select the User Library option from the list, and click Next.

This will in turn present you with a list of the available user libraries. Check both of the Hibernate libraries that you configured previously, and click Finish.

The libraries will be added to the list in the Libraries tab of the Java Build Path dialog; you should click OK to accept the changes to the path. Figure B-6 shows the resulting uncluttered project view.

Figure B-6. *The configured Java project*

So far in this section, we have configured a Java project in a conventional way. None of the steps we have taken so far would be unusual in preparing to use a third-party library such as the Spring Framework. The next few steps open a Hibernate Console perspective that will allow you to manage the Hibernate-specific aspects of your project.

To open a new Eclipse perspective, you can either select the Window ➤ Open Perspective menu option or click the Open Perspective icon.

Select the Other option from the resulting menu. This will present you with the list of available perspectives. Select the Hibernate Console option from the resulting dialog.

Much as when entering the Eclipse debugging perspective, the layout of Eclipse will change, with various Hibernate-specific views being opened (although you will still have access to the Package Explorer as one of the tabbed options behind the new Hibernate configuration view).

To switch between the perspectives, you can select the Java or Hibernate Console perspective icons on the top right of the main Eclipse window (you may need to adjust the menu layout slightly if these are obscured). Figure B-7 shows the icons for the Hibernate Console and Java perspectives, with the Java perspective selected.

Figure B-7. *The Java and Hibernate Console perspective icons, with the Java perspective currently selected*

Using the Hibernate Console

With the project set up and the Hibernate Console perspective added to Eclipse, you can now start using some of the Hibernate Tools plug-in features more directly.

Creating a Hibernate Console Configuration

The Hibernate Console represents a view onto the Hibernate object store. It therefore needs the same basic configuration information as any other client of Hibernate. You will need a Hibernate XML configuration file and/or a properties file (see Listing B-1).

Listing B-1. *The Configuration File Used in Our Example:* hibernate.cfg.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <!-- The database settings -->
        <property name="hibernate.connection.driver_class">
            org.hsqldb.jdbcDriver
        </property>
        <property name="hibernate.connection.url">
jdbc:hsqldb:file:/workspace/ToolsExample/db/annotationsdb;shutdown=true
        </property>
```

```
        <property name="hibernate.dialect">
            org.hibernate.dialect.HSQLDialect
        </property>
        <property name="hibernate.connection.username">sa</property>
        <property name="hibernate.connection.password"/>
        <property name="hibernate.connection.pool_size">0</property>
        <property name="show_sql">false</property>

        <!-- The class mappings for our example -->
        <mapping class="com.hibernatebook.tools.Note"/>
        <mapping class="com.hibernatebook.tools.Notepad"/>
    </session-factory>
</hibernate-configuration>
```

If you are using the HSQLDB database in conjunction with Hibernate Tools, we recommend that you use fully qualified path names (shown in bold in the preceding code) to avoid possible clashes between the working directories and classpaths used by Ant, the Eclipse IDE, and the Hibernate Console. In our example, we will be working with the annotations-based mappings, so we use the mapping element with a class attribute to determine the mapped classes. However, the console configuration will automatically find any correctly annotated entity classes using reflection, so it is possible to omit these entries.

If you choose to use XML-based mappings, you can use the file attribute to specify the mapping file instead, or specify the paths to the mapping files in the Hibernate configuration entry used by the Hibernate Console.

Hibernate Tools provides a wizard to create the Hibernate configuration entry. To use the wizard, select the File ➤ New ➤ Hibernate Console Configuration menu option, as shown in Figure B-8.
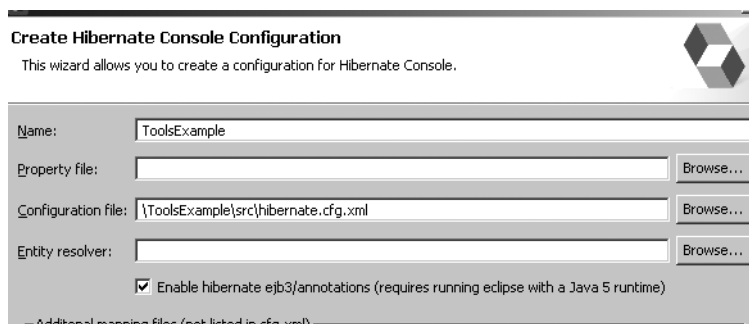
Figure B-8. *Creating a new console configuration*

This will bring up the Create Hibernate Console Configuration dialog page. This consists of a simple form with three basic sections. In the first part of the form shown in Figure B-9, set a suitable name for this configuration and the path to the Hibernate configuration file. You should also tick the check box, as you will be using annotations for your mappings.

Figure B-9. *The Name and Configuration file settings*

The middle section of the form shown in Figure B-10 allows you to specify mapping files for inclusion in the console configuration that are not explicitly included in the configuration file. You would add entries here if your application added these options to the configuration at run time instead of drawing the information from the configuration file (for example, if it only used the properties file for its other configuration settings).
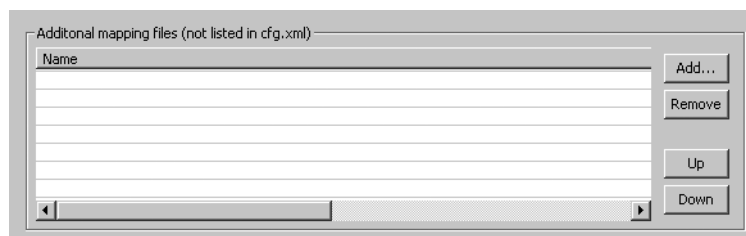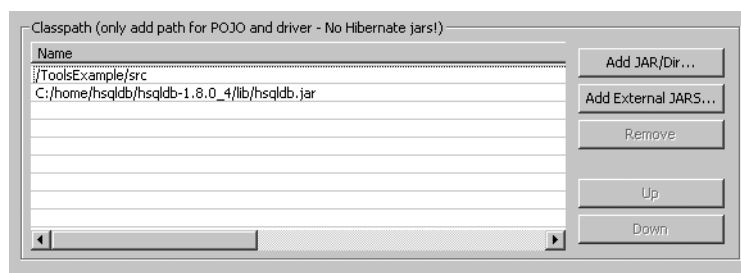
Figure B-10. *The additional mapping file settings*

The last part of the form shown in Figure B-11 allows you to set the classpath entries that will be used by the Hibernate Console. This must include the JDBC driver to access your database and the location of your mapped POJO classes.

Figure B-11. *Setting the classpath*

Click the Finish button, and you should now be able to see the configuration name in the console configuration view.

Generating the Schema

Expand the ToolsExample entry that you created in the console configuration window. There will be three nodes beneath it, as shown in Figure B-12.

Figure B-12. *The three nodes beneath ToolsExample*

The third of these is a view of the tables in the database that correspond to the entities mapped in your console configuration (either explicitly or via the `hibernate.cfg.xml` file). Unless your database contains preexisting tables, the database view will be empty when you attempt to drill down into it, as shown in Figure B-13.

Figure B-13. *The empty database view*

The Hibernate Console now has all of the entity mapping and database connection information. It therefore offers a Run SchemaExport context menu option (accessed from the Configuration node, *not* the Database node) to generate the appropriate tables (see Figure B-14).

Figure B-14. *The Run SchemaExport context menu option*

■**Caution** There are currently no safeguards built in here. If the configured database user has table dele-tion privileges, running SchemaExport in this way will drop and re-create all of your tables with no further warning!

If you select Run SchemaExport from this menu, the database view will be updated to reflect the changes, and you will be able to drill down to view the individual table column names and data types, as shown in Figure B-15.

Figure B-15. *A column-level view of the newly generated database*

■**Note** When you run SchemaExport for the first time, you may see error messages in the event log view related to attempts to drop nonexistent tables. These result from unconditional DROP commands to remove preexisting tables. They should go away if you run SchemaExport a second time—and assuming this is the case, they can be safely ignored. As always, given that the tools are still beta versions at the time of writing, this behavior may have improved by the time you read this book.

This view does not offer the sort of comprehensive access to the database metadata that is available from some tools, and it does not permit you to directly change any of the features. It does, however, offer a useful visual confirmation that the appropriate tables and columns have been created.

Running the Test Program

For the purpose of generating some test data for manipulation in the remainder of this sec-tion, we have created a test application (available with the other source code for this book on the Apress web site) to populate the database with some suitable records. Now that you have created an appropriate database schema, it is possible to run this application.

Switch back to the Java perspective and add a new Java Application run configuration for the PopulateNotes class. You will need to add the path to your JDBC driver to the configuration on the Classpath tab.

Now run the application. You should see the following output, which confirms that a `Notepad` entity and 100 associated `Note` entities have been persisted to the database:

```
Creating test notepad...
Test notepad created.
```

You can now switch back to the Hibernate Console perspective in order to browse this data.

Browsing the Model

Back in the Hibernate Console perspective, you can access the contents of the object model using the Session Factory node of the console configuration view. This node is named after the `SessionFactory` object created and maintained internally by the Hibernate Eclipse plug-in tools when you drill down into the console configuration view.

■**Tip** If you manage to put the internal `SessionFactory` object into a bad state, it is possible to close and re-create it from the configuration view's context menu when the configuration view is collapsed to a single node.

Opening the Session factory node, you will see nodes representing each of the mapped classes. Drilling down further, you can see a representation of the mapped fields and attributes (along with a graphical representation of the association rules, primary keys, and so on, where appropriate).

If you double-click the mapped class, all the instances of the class in the database will be retrieved, and a `toString` representation will be displayed in the Hibernate Query Result view.

As you can see from Figure B-16, the generated HQL query to produce this information is shown in a second tabbed pane (further queries will be added as new tabs within this window). The column is numbered, as no column name can be specified. Note that we have overridden the `Notepad` class's `toString()` method to ensure that a human-readable representation of the class contents is displayed in this view.

Figure B-16. *The Hibernate Query Result view of the mapped* `Notepad` *class*

All results of queries generated manually or automatically through the Hibernate Console perspective will be displayed in the Hibernate Query Result view.

Testing HQL Queries

HQL is a powerful way to interact with the database; but traditionally, debugging errors in these queries has been a long-winded process. Typically, a developer would end up creating a unit test for the query and then spending considerable time tweaking it into working correctly. Hibernate Tools now provides a query editor that allows you to prototype your HQL queries and run them against the database without needing to write any code.

To run an HQL query, select HQL Scratchpad from the context menu for the project in the configuration view (by default, this is to the left of the main window in the Hibernate Console perspective).

A new editor window will be created and labeled with HQL and the name of the project—in our example, it's "HQL: ToolsExample." This is an HQL editor with some context sensitivity. Context sensitivity is activated with the Ctrl+Space key combination, as is usual for Eclipse editors. For example, enter the following query in the scratch pad editor and press Ctrl+Space:

```
from Note
```

This will display a context menu offering the two class names that are legal at this point in the query (as shown in Figure B-17).

Figure B-17. *The context-sensitive scratch pad*

■**Caution** At the time of writing, there are various bugs related to saving scratch pads to physical files, and to running HQL queries loaded-in directly from text files. As a workaround, it is possible to copy and paste between a scratch pad window and a normal text editor view in order to run queries. It is likely that these problems will have been addressed by the time you read this.

As you create the query, the corresponding SQL query against the underlying relational database is updated in the Hibernate Dynamic SQL Preview window shown in Figure B-18.
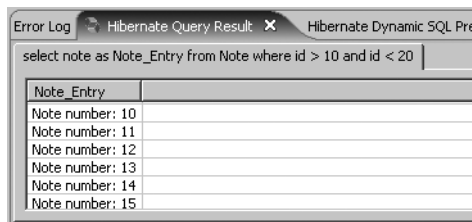
Figure B-18. *The Hibernate Dynamic SQL Preview window*

The generated SQL will be visible whenever the HQL query is syntactically valid, allowing you to confirm the basic correctness of the HQL and confirm the sanity of the underlying query at the same time.

Once you have a satisfactory HQL query, it can be executed against the database to validate the returned results. Enter the following query in the scratch pad window, and then click the green play button above the editor window.

```
select note as Note_Entry
from Note
where id > 10 and id < 20
```

The view shown in Figure B-19 will list the results of the query. Note that the results have been given the correct column heading.

Figure B-19. *The output from the query*

Query Parameters

While the scratch pad alone is adequate for running stand-alone queries, much of the HQL that you will use in practice in your applications will be parameterized. Consider the following typical query:

```
select owner from Notepad where owner = :username
```

This query uses a named parameter, `username`, as a placeholder for the actual value. When converted to SQL, this will become a parameter in a prepared statement, and will not be specified until the application is actually running.

The Hibernate Console provides the novel feature of allowing you to specify the values of parameters in Hibernate queries of this type.[2] This means that you can cut and paste an HQL query from your code into the scratch pad, and test it there without having to edit any of the parameters themselves.

If you paste the preceding query into a scratch pad window and try to run it as is, the console will log an appropriate error, warning you that the query parameters have not all been set (currently this appears in the Hibernate Query Result view). On the right-hand side of the perspective's default layout, there is a view called Query Parameters. Adjust the layout so that you can see the three columns within this view, labeled Name, Type, and Value.

---

2. If you've ever encountered problems debugging complex SQL queries, in which the process of manually substituting in parameters disguised the cause of the problem during testing, you will understand why we are extremely enthusiastic about this ingenious feature!

With the query still in the scratch pad, click the icon shaped like the letter *P* (shown in Figure B-20, just to the right of the Outline tab). A new row will be added to the view, representing the parameter in the query. By editing the values in this row, you can change the details of the parameter (both the value assigned to it and the type of that value).

Figure B-20. *Assigning a typed value to the* id *query parameter*

If you edit the query parameter row so that it contains the data shown in Figure B-20, running the query will return results equivalent to those produced by the following code fragment:

```
Query query = session.createQuery(
    "select owner from Notepad where owner = :username");
query.setString("username","Dave Minter");
List results = query.list();
```

All the standard Hibernate types are available to you, and the console will log appropriate errors when you try to execute queries that combine types in inappropriate ways for the object model.

The Entity Model

If you select the Hibernate Entity Model view (usually located as a tab option in the view at the bottom right-hand corner of the Hibernate Console perspective) for the project, and then select the Configuration node of the configuration view, a graphical representation of the entities in the model will be displayed in the Entity Model view, as shown in Figure B-21.

Figure B-21. *The Entity Model view of our model*

The view allows you to place the entities using the mouse or lay them out automatically. You can zoom into or out of the view, and you can print its contents. At the time of writing, this feature is quite primitive and is therefore of limited use.

Generating a Mapping File

Hibernate Tools provides a wizard to make the creation of the boilerplate XML mapping file contents simpler. You can access the wizard via the File ➤ New ➤ Hibernate XML Mapping File menu option from within the Hibernate Console perspective.

This brings up a standard Eclipse view in which you select the name and location of the generated mapping file.

Clicking Next then brings you to a dialog that requires you to enter the fully qualified name of the class that the mapping file will represent.

If you cannot recall the specific package and class name details, you can use the Browse button to select the class from the class hierarchy.

The boilerplate mapping file will be created in the appropriate location. The wizard does not currently populate this with any of the property or association mapping details. The generated file is therefore incomplete and can only be used as a skeleton to be filled in manually. The sample output for our example Bookshelf class (which is not annotated) follows.

```xml
<?xml version="1.0"?>
<!DOCTYPE hibernate-mapping PUBLIC
    "-//Hibernate/Hibernate Mapping DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">
<hibernate-mapping package="com.hibernatebook.tools">
    <class name="Bookshelf">
    </class>
</hibernate-mapping>
```

The Hibernate Tools plug-in for Eclipse includes an editor for its various mapping files that provides context-sensitive completion—this is based upon the Web Tools Platform's (WTP) support for XML files, but adds awareness of mapping file details beyond the support for the basic DTD description that WTP provides.

Generating a Configuration File

Hibernate Tools provides another wizard to aid in the creation of a Hibernate configuration XML file. This wizard is accessed from the File ➤ New ➤ Hibernate Configuration File menu option from within the Hibernate Console perspective.

This wizard prompts you (as with the mapping file wizard) to specify a path and file name for the generated configuration file. You don't want to overwrite the configuration file used by our sample application, so for this exercise, specify a file name of example.hbm.xml in the root of the project directory.

Clicking the Next button will then take you to the dialog shown in Figure B-22, in which all the basic configuration properties can be specified. When possible, a set of default options are provided as combo box lists. The "Database dialect" field presents the dialects in a more human-readable format, and also filters the Driver class and Connection URL fields to the likely options.
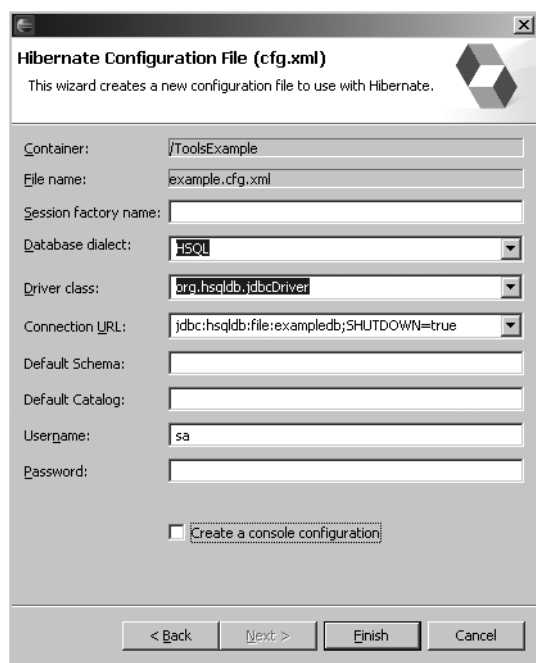
Figure B-22. *Selecting default values for the configuration file*

All the values selected in the drop-down menus can be overtyped if the option you need is not listed. (Note that we have amended the connection URL with our preferred HSQLDB connection details.) You can type in the fully qualified dialect class name instead of choosing from the short names available from the drop-down list.

At this point, you also have the option of selecting the Create a console configuration check box. If you do so, the Hibernate Configuration File wizard will pass you to the Hibernate Console Configuration wizard (and will automatically populate the configuration file field). However, since we've already discussed this earlier in the chapter, leave the check box empty and click Finish.

The resulting example.cfg.xml file is shown in the following code (slightly reformatted to ensure that it fits on the page). Unlike the generated mapping files, this configuration file is pretty much ready to use if you plan to use programmatic configuration of your application, or if you are using it purely as the basis of a Hibernate Console configuration to connect to an existing Hibernate database.

```xml
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
   <session-factory>
      <property name="hibernate.connection.driver_class">
         org.hsqldb.jdbcDriver
```

```
        </property>
        <property name="hibernate.connection.url">
            jdbc:hsqldb:file:exampledb;SHUTDOWN=true
        </property>
        <property name="hibernate.connection.username">sa</property>
        <property name="hibernate.dialect">
            org.hibernate.dialect.HSQLDialect
        </property>
    </session-factory>
</hibernate-configuration>
```

The Reverse Engineering and Code Generation Tools

The last Hibernate Tools wizard is for the reverse engineering tool. This allows you to extract the schema information from the database into an XML file. Additional information can be added to this file using a tabbed control panel, which allows it to be used in the generation of source code using the Hibernate Code Generation tool (accessed separately from the wizard).

This wizard is accessed from the File ➤ New ➤ Hibernate Reverse Engineering File menu option from within the Hibernate Console perspective.

This then prompts you to specify a path and file name for the generated reverse engineering file. Select the default file name and location within the ToolsExample project.

When you are done, you have the option of selecting either Finish or Next. If you select Finish, the wizard will assume that you want to use the details of all the tables in your database when creating the reverse engineering file. Alternatively, if you click Next, you have the opportunity to specify the individual tables and columns that should be included (as shown in Figure B-23).
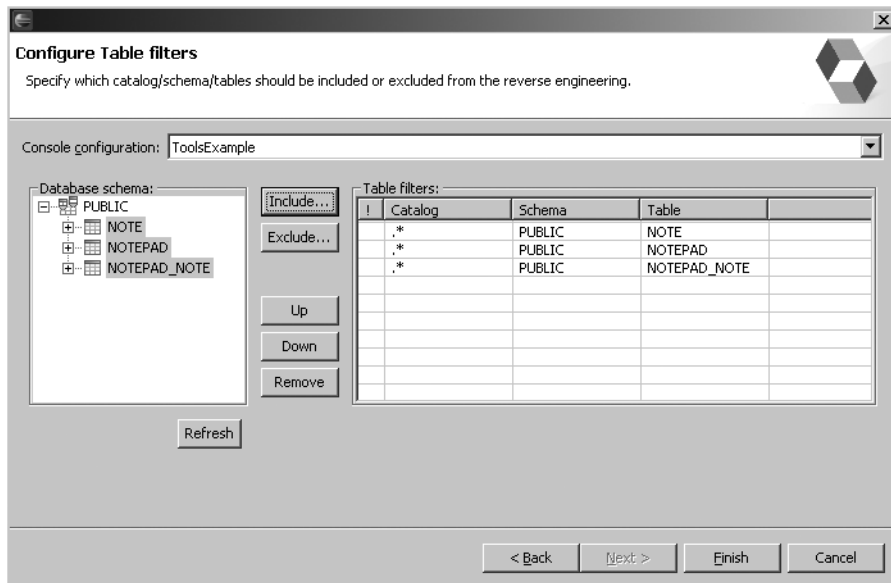
Figure B-23. *Selecting the tables to reverse engineer*

In Figure B-23, we have selected all the tables anyway. Once you click Finish, the reverse engineering file will be generated, and an editor will be opened for it. This is a tabbed view of the XML file (the last tab shows the XML source code). This allows you to change the default mappings between the database and Hibernate types, to alter the table and column details, and to apply *additional* filters to the file itself that dictate which details will be ignored when generating output from the reverse engineering file.

We would not recommend trying to get to grips too closely with the reverse engineering tool until you have some experience in creating mappings manually—the various settings will seem quite opaque when taken out of context. Once you have created a few simple entity models from scratch, however, the need for the various options should become clearer.

In order to actually generate output from the tool, you will need to create a Hibernate code generation configuration by selecting the option in the toolbar of either perspective, as shown in Figure B-24.

Figure B-24. *Selecting the Hibernate Code Generation configuration tool*

■**Note** At the time of writing, the Hibernate Code Generation icon does not seem to be added to the Eclipse toolbar when you install the plug-in. This will probably be fixed by the time you read this book, but if not, you may have to use the JBoss IDE download or an Ant task if you want to take advantage of these features.

The Code Generation tool allows you to create various different configurations for generating files based upon your reverse engineering XML file(s). The tool can output the following:

- Ordinary POJOs

- POJOs using generics

- POJOs using EJB 3 annotations (with or without generics)

- DAO objects

- XML mapping files

- Hibernate configuration files

- HTML documentation of the database schema

It also permits you to create a JBoss Seam skeleton application, but this functionality lies well outside the scope of this book.

The only missing output is the database schema. The reverse engineering file can be generated from the database as we have described here, but it can also be created manually. We have already shown that it is possible to use the Hibernate Console's SchemaExport function to export an annotated POJO or mapping file into a database schema. So, using this functionality, it is possible to create pretty much everything from scratch, starting from a mapping, a database, or a set of POJOs.

When you create your code generation configuration, you need to specify the location of your reverse engineering file, and the console configuration that will be used to establish the appropriate connection details. You need to specify the output directory into which the generated files should be written, and you can also specify a package that any generated classes should belong to.

You will then need to select the types of output to be generated, as shown in Figure B-25.

Name: Example Code Gen Config

Main Exporters Refresh Common

☑ Generate domain code (.java)
  ☑ JDK 1.5 Constructs (generics, etc.)
  ☐ EJB3/JSR-220 annotations
☑ Generate DAO code (.java)
☑ Generate mappings (hbm.xml)
☐ Generate hibernate configuration (hibernate.cfg.xml)
☑ Generate schema html-documentation
☐ Generate JBoss Seam skeleton app (beta)
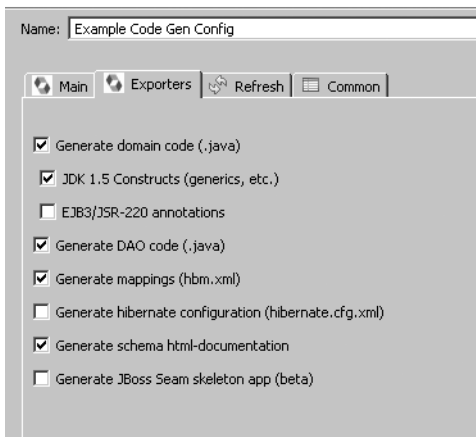
Figure B-25. *Selecting the types of output to be generated*

In Figure B-25, we have selected everything except the Hibernate configuration file (we've already got one), the JBoss Seam application (we're not interested in that), and—because of a minor bug—the EJB 3 annotations option. You would not normally elect to generate both annotations and XML mapping files anyway.

■**Note** At the time of writing, the EJB3/JSR-220 annotations exporter option does not work. If this problem has not been fixed by the time you read this book, you can use an Ant task instead.

The Ant Tasks

As you will have noticed if you followed through the Eclipse discussion, our example project for this appendix includes `build.xml` and `build.properties` files. These are not strictly necessary when working with Eclipse, but it is often desirable to be able to build all the components

of a project outside the IDE. This can be useful in maintaining automated builds for regression testing—and of course, not all Java IDEs offer integrated support for Hibernate anyway, while most of them do offer support for the Ant build tool. And besides, you might not be using an IDE in the first place!

The Ant tasks are part of the Hibernate Tools download, which is largely oriented toward use as an Eclipse plug-in. The Ant tools themselves rely upon a set of three JAR files: `hibernate-tools.jar`, `freemarker.jar`, and `jtidy-r8-21122004.jar`. The `hibernate-tools.jar` file is currently only available as a download with the rest of the plug-in files. You will find the JAR file within the ZIP file in a directory called `plugins\org.hibernate.eclipse_X.x.x.x\lib\tools`, where the *X*s are substituted for the Hibernate Tools version number. The `freemarker` and `jtidy` JAR files will also be available from this directory. Since these JAR files have no dependencies upon the other parts of the plug-in, you can copy them to a directory with a less unwieldy name without further consequences.

How the Ant Tasks Work

Despite their diverse set of outputs, the Hibernate tasks actually all work in the same way. The task is imported into Ant using the standard `TaskDef` element. This makes the Hibernate tools libraries available to Ant itself, and allows you to select an appropriate element representing the Hibernate tools to use in your script.

```
<taskdef
    name="htools"
    classname="org.hibernate.tool.ant.HibernateToolTask"
    classpathref="classpath.tools"/>
```

The Hibernate Tools JAR file must be made available to Ant on the classpath—our example uses a preexisting classpath declaration (referenced through its `id` of `classpath.tools`—see the "Configuring the Classpath" subsection later in this section).

A standard Ant target is declared to contain the set of operations that you want to perform.

```
<target name="exportDDL" depends="compile">
…
</target>
```

Within the target with any other standard (or imported) Ant tasks, you can then include the element that you declared using the `<taskdef>` element. The other Hibernate task elements are only applicable within this task element.

```
<target name="exportDDL" depends="compile">
    <htools destdir="${sql}">
    …
    </htools>
</target>
```

This outermost task element accepts three attributes, as listed in Table B-1.

Table B-1. *The Daughter Elements of the Hibernate Tools Task*

Attribute	Description
classpath	The path to use when locating libraries and configuration files.
destDir	The base directory (relative to the build script's own base directory) into which any generated output will be written.
templatePath	The path containing user-created template files (see the further discussion of templates in the "Templates" section later in the chapter.)

Within the declared Hibernate task, a number of additional standard elements can be created, consisting of a classpath declaration (an alternative to using the classpath attribute), a set of configuration elements, and a set of exporter elements.

The classpath element and attribute are standard Ant features that allow you to bring in any necessary resources used by the Hibernate tasks.

The clever bit of the Hibernate Ant task lies in the configuration elements. Declaring a configuration task causes an appropriate configuration object to be built in memory. These in-memory configuration objects all extend the standard org.hibernate.cfg.Configuration class. The Configuration class represents the mapping relationships between entities (combined with information from any configuration, properties, or reverse engineering files), and it is this information, the metamodel, that is then used to generate the various output files. The provided configuration elements can conjure up a Configuration object from the standard mapping files, from the metadata information gathered over a JDBC connection, and from the annotations discussed in Chapter 6.

```
<target name="exportDDL" depends="compile">
    <htools destdir="${sql}">
        <annotationconfiguration
            configurationfile="${src}/hibernate.cfg.xml"/>
        …
    </htools>
</target>
```

Within any given Hibernate Tools task, you can only have one configuration element configured—normally, you would not want to generate output from two distinct representations of the mapping information, so the single declaration is shared between the generation tasks enclosed within the toolset task elements. The following list describes the configuration tasks and the attributes of each:

- <configuration>: Mapping relationships are generated from conventional XML-based mapping files and information in a *.cfg.xml or *.properties file.

    - configurationfile: The name of the XML configuration file being used.

    - propertyfile: The name of the properties file being used.

    - entityresolver: The name of the SAX EntityResolver to use when resolving "external" XML entities (rarely used).

    - namingstrategy: A naming strategy to use (see Chapter 3) to establish table names from entity names.

- `<annotationconfiguration>`: Mapping relationships are generated from the EJB 3 and Hibernate 3 annotations in conjunction with a `*.cfg.xml` or `*.properties` file.

    - Identical to `<configuration>`.

- `<jdbcconfiguration>`: Mapping relationships are generated from the schema metadata obtained over a JDBC connection. The connection details are configured from a properties file.

    - All those from `<configuration>`, plus the following:

    - packagename: The name of the package that entities should belong to.

    - reversestrategy: The fully qualified name of a class implementing the `org.hibernate.cfg.reveng.ReverseEngineeringStrategy` interface. This is the programmatic equivalent of the `reveng.xml` file approach.

    - revengfile: The name of a reverse engineering file to use when processing metadata information. See the discussion later in this section.

- `<ejb3configuration>`: Mapping relationships are generated from the EJB 3 and Hibernate 3 annotations in conjunction with an EJB 3–compliant `persistence.xml` file.

    - entityresolver: The name of the SAX EntityResolver to use when resolving "external" XML entities.

    - namingstrategy: A naming strategy to use (see Chapter 3) to establish table names from entity names.

The `<configuration>` element also allows you to specify a standard Ant `<fileset>` of `*.hbm.xml` mapping files. If you use this in conjunction with a `*.cfg.xml` configuration file, you must not permit any mapping resources to be duplicated, as this will result in duplicate import mapping exceptions.

Your choice of configuration element will be driven by the data sources that you have available to you. For example, if you have created your XML mapping files, you will want to use the standard configuration element, but if you have only a normalized database, you will want to generate the mapping information from this using the JDBC configuration (although you may well choose to create a reverse engineering file to control this).

Once you have correctly configured an annotation object, however, you can generate any of the other resources that you might need using one or more exporter elements. These are listed in Table B-2.

Table B-2. *The Available Exporter Elements*

Element	Description
`<hbm2ddl>`	Generates tables from the metamodel
`<hbm2cfgxml>`	Generates a `*.cfg.xml` configuration file from the metamodel
`<hbm2java>`	Generates entity POJOs from the metamodel
`<hbm2hbmxml>`	Generates Hibernate `*.hbm.xml` mapping files from the metamodel

Continued

Table B-2. *Continued*

Element	Description
`<hbm2doc>`	Generates HTML documentation for the database schema from the metamodel
`<hbm2dao>`	Generates standard DAOs from the metamodel
`<hbmtemplate>`	Generates arbitrary user-defined output from the metamodel
`<query>`	Runs arbitrary HQL queries against the database using the mapping information in the metamodel

You may notice that the exporters available as Ant tasks correspond fairly closely to the exporters available in the Hibernate Code Generation tool—largely because they rely upon the same underlying implementations.

The two most commonly used tasks are `<hbm2ddl>`, which can generate a database schema directly from the `*.hbm.xml` mapping files, and `<hbm2hbmxml>`, which, conversely, can generate mapping files directly from the database.

The `<hbm2ddl>` element generates DDL scripts from the metamodel. These can be written to a file—or, if the configuration object is provided with database connection details, they can be run directly against the database. Table B-3 shows the attributes that can be supplied.

Table B-3. *The Attributes Available to the* `<hbm2ddl>` *Element*

Property	Default	Description
create	true	If set to `true`, causes the generated DDL to include commands to create database objects. This allows to distinct tasks to be created: one to drop all relevant database objects (using the `drop` attribute) and the other to create them.
console	true	If set to `true`, causes the generated DDL to be displayed on the console.
delimiter	;	Specifies the delimiter to be used to separate DDL statements.
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task.
drop	false	If set to `true`, causes the generated DDL to include commands to drop preexisting database objects before it tries to create them. This may cause warning messages, depending upon the preexisting state of the database; and it of course has the potential to destroy existing data.
export	true	If set to `true`, causes the DDL to be run directly against the database (this has the potential to delete data—do not use carelessly).
format	false	If set to `true`, causes the generated DDL to be formatted using whitespace in a more readable fashion. We recommend using this option if you will be writing the DDL to a file.
haltonerror	false	If set to `true`, causes the script to halt if an error is encountered while generating the DDL (typically, this is used while exporting directly to the database to increase the visibility of any problems encountered while setting up the schema).
outputfilename		Specifies the name of the file name that the generated DDL should be stored in. If left unset, the generated DDL will not be stored.
update	false	Indicates that the tool should attempt to generate the appropriate statements to bring the existing schema inline with the model. We don't recommend using this option.

The task shown in Listing B-2, which completes the simple example that we've been building up in this section, creates a schema generation script from the annotation-based mappings referenced in the project's `hibernate.cfg.xml` file.

Listing B-2. *A Complete Hibernate Mapping Target*

```
<target name="exportDDL" depends="compile">
    <htools destdir="${sql}">
        <annotationconfiguration
            configurationfile="${src}/hibernate.cfg.xml"/>
        <hbm2ddl
            create="true"
            drop="true"
            format="true"
            export="true"
            outputfilename="${ant.project.name}.dll"/>
    </htools>
</target>
```

The `<hbm2cfgxml>` element generates a Hibernate XML configuration file from the meta-model information. Table B-4 shows the attributes that can be supplied.

Table B-4. *The Attributes Available to the* `<hbm2cfgxml>` *Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task.
ejb3	false	By default, causes entities to be mapped using `<mapping resource="..."/>` entries in the configuration file. If set to `true`, the entities will be mapped using the `<mapping class="..."/>` approach to pick up EJB 3 annotations in the mapped classes. This setting does *not* cause a `persistence.xml` file to be generated!

Typically, the `<hbm2cfgxml>` element is used when the configuration task has been configured from a properties file—for example, when using `<jdbcconfiguration>`, you would typically start with a normalized database schema and a properties file containing the connection details, and use this exporter to create a `hibernate.cfg.xml` file containing both the connection details and the details of the mapped entities.

The `<hbm2java>` element generates the Java source code for POJOs for each of the entities held in the metamodel. Table B-5 shows the attributes that can be supplied.

Table B-5. *The Attributes Available to the* `<hbm2java>` *Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task
ejb3	false	If set to `true`, causes the POJOs to be generated with EJB 3 annotations
jdk5	false	If set to `true`, causes the POJOs to be generated with Java 5 constructs (generics, enumerations, etc.)

This exporter can be used to permit the mapping file–based creation of suitable classes, or to create classes from the database schema when `<jdbcconfiguration>` is used.

The `<hbm2hbmxml>` element generates the XML mapping files from the information contained in the metamodel. Table B-6 shows the attributes that can be supplied.

Table B-6. *The Attributes Available to the* `<hbm2hbmxml>` *Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task

This exporter is particularly terse because it only writes out the mapping information stored in the metamodel. This is all handled by the appropriate configuration element. The `<hbm2hbmxml>` exporter just needs to know which path to write the XML files into. And that can be specified at the Hibernate tool level. A fairly typical invocation of this exporter is, in its entirety, the ridiculously simple `<hbm2hbmxml/>`.

The `<hbm2doc>` element generates HTML documentation of the schema and entities in a style similar to the familiar javadoc output. Table B-7 shows the attributes that can be supplied.

Table B-7. *The Attributes Available to the* `<hbm2doc>` *Exporter Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task

The `<hbm2dao>` element generates a set of basic DAO classes—one for each of the entities in the metadata. Table B-8 shows the attributes that can be supplied.

Table B-8. *The Attributes Available to the* `<hbm2dao>` *Exporter Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task

By default, these DAO classes will be named after your entity, suffixed with Home. The generated DAOs provide a set of methods roughly corresponding to the methods available on the Session interface, but strongly tied to the entity type. While the generated DAOs can provide a useful foundation for your own more specific DAOs, we generally find that they offer little value beyond that already offered from the standard Session interactions.

■**Note** Nothing in any of the exporters intrinsically stops you from generating "silly" combinations of output—but this has its advantages; for example, it is possible to use an `<annotationsconfiguration>` configuration element with the `<hbm2java>` exporter to generate POJOs. While that might seem pointless, given that you have to start with POJOs to use an annotations-based configuration in the first place, it actually provides the useful ability to generate Java 1.4–compatible source code from annotated Java 5 class files!

In principle, the `<query>` element allows you to specify an arbitrary HQL query that will be run against the database using the configuration's mapping information. Table B-9 shows the attributes that can be supplied.

Table B-9. *The Properties Available to the* `<query>` *Exporter Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the tools task.
destfile		If set, specifies the file into which the output of the queries will be written. If left unset, the query is carried out, but the output is not saved.

The HQL query itself is included as the body of the `<query>` element.

```
<query destdir="output" destfile="sql.log">
   select n.owner from Notepad n
</query>
```

If you want to include multiple SQL queries in the task, you can include multiple nested `<hql>` elements thus:

```
<query destdir="output" destfile="sql.log">
   <hql>select n.owner from Notepad n</hql>
   <hql>select n.owner from Note n</hql>
</query>
```

Reverse Engineering

As we discussed in the previous section, the `<jdbcconfiguration>` task can be used to create a configuration metamodel directly from a database schema. However, a database schema is not necessarily an exact match for the entity mappings that we want to create. It may contain tables and columns that we would prefer to omit from our entities. It may also lack some of the information that we want to express. For example, database types such as `VARCHAR` do not map exactly to Java types such as `char[]` and `java.lang.String`, and the tables will have names that do not conform to the Java naming conventions.

Used as is, the `<jdbcconfiguration>` task will select sensible defaults for the type names, and will assign the reverse-engineered tables suitable Java names derived from the table names. It also provides an attribute to allow you to directly specify a suitable package name. Even so, we would really like more control over the specifics of the reverse engineering process. Note that while it also provides a naming strategy attribute, this has no effect during the reverse engineering process, as naming strategy classes can only be used to determine schema names from mapping information, not vice versa.

Hibernate provides two ways in which this process can be controlled—you can specify one of two attributes on the `<jdbcconfiguration>` task to override the default behavior.

The `reversestrategy` attribute allows you to specify the fully qualified class name of a custom implementation of the `org.hibernate.cfg.reveng.ReverseEngineeringStrategy` interface. This interface defines methods that can be used to select the following:

- Class names from table details

- Property names from column details

- One-to-many associations from foreign key relationships

- Many-to-one associations from foreign key relationships

- Inverse relationships from foreign key relationships

- Lazy loading details from foreign key relationships

- Collection attribute names from foreign key relationships

- Entity names from foreign key relationships

- Tables and columns to include and exclude

- Primary key columns from table details

- The column to use for optimistic locking strategies

- Composite ID details from table details

It also allows you to provide information about how schemas should be *generated* from the resulting metamodel. This information is as follows:

- Additional foreign key relationships

- The table naming strategy to be used

The disadvantage of this approach to managing the reverse engineering process is that it is not particularly flexible, and it requires a lot of coding. Reverse engineering is often carried out only once to establish the mappings, with the schema thereafter being driven *from* the mappings, rather than being used to create them. The Hibernate tools therefore provide a second mechanism for controlling the reverse engineering process by specifying an XML configuration file using the `revengfile` attribute of the `<jdbcconfiguration>` task. This provides nearly the same degree of control, but is much simpler to create—especially if you intend to manipulate only minor details of the process.

The output of the reverse engineering tool described in the "The Reverse Engineering and Code Generation Tools" section is actually a reverse engineering XML file following this format. A very simple example reverse engineering file is shown in the following code:

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE hibernate-reverse-engineering
PUBLIC "-//Hibernate/Hibernate Reverse Engineering DTD 3.0//EN"
"http://hibernate.sourceforge.net/hibernate-reverse-engineering-3.0.dtd"
>
```

```
<hibernate-reverse-engineering>
    <table-filter match-schema="PUBLIC" match-name="NOTE"/>
    <table-filter match-schema="PUBLIC" match-name="NOTEPAD"/>
    <table-filter match-schema="PUBLIC" match-name="NOTEPAD_NOTE"/>
</hibernate-reverse-engineering>
```

The reverse engineering file in the preceding code limits the tables to be used when generating the mapping information from the schema to the three explicitly named tables (NOTE, NOTEPAD, and NOTEPAD_NOTE) in the database's public schema.

A reverse engineering file always consists of a single top-level `<hibernate-reverse-engineering>` element containing various additional elements. These daughter elements are given in order in Table B-10:

Table B-10. *The Elements Used in Configuring the Reverse Engineering File*

Element	Cardinality	Description
`<schema-selection>`	Zero or more	Allows the reverse engineering process to be limited by catalog, schema, and table name
`<type-mapping>`	Zero or one	Allows you to override the default mapping between database types and Java types
`<table-filter>`	Zero or more	Allows you to include or exclude tables by catalog name, schema name, and table name, and allows them to be grouped into a particular package
`<table>`	Zero or more	Allows you to override the default mappings of tables into entities

Rather than trying to exhaustively specify the syntax of a reverse engineering file, which is anyway available through the DTD at http://hibernate.sourceforge.net/hibernate-reverse-engineering-3.0.dtd, we think it is easier to follow the basic requirements of the file format with some examples of valid `<schema-selection>`, `<type-mapping>`, `<table-filter>`, and `<table>` elements.

Our first example specifies the following rule: tables should only be used for reverse engineering if they are in the public schema and their names begin with NOTE:

```
<schema-selection match-schema="PUBLIC" match-table="NOTE*"/>
```

Our next example enforces a rule that database INTEGER types for which the column is specified as NOT NULL must be represented using Hibernate's int type. It also enforces a rule that database VARCHAR types that have a specified length of 250 should be treated as Hibernate string types.

```
<type-mapping>
    <sql-type jdbc-type="INTEGER" hibernate-type="int" not-null="true"/>
    <sql-type jdbc-type="VARCHAR" hibernate-type="string" length="250"/>
</type-mapping>
```

These type mappings apply throughout the reverse engineering process—you cannot specify them on a per-table basis using the `<type-mapping>` element, but you can using the `<table>` element.

The `<table-filter>` element allows you to include and exclude groups of tables from the mapping process on the basis of pattern matches on their names. Where the `<schema-selection>` element allows you to specify a set of tables matching a single pattern to be reverse engineered, the `<table-filter>` element operates within this and allows multiple patterns to be applied to include and exclude tables. Here's an example:

```
<table-filter match-name="NOTEPAD_ARCHIVE*" exclude="true"/>
```

Although the previous `<schema-selection>` element included all tables within the current schema that matched the pattern NOTE*, this table filter excludes any tables that match the pattern NOTEPAD_ARCHIVE* from reverse engineering. Table filters are applied in order, so using this technique, you can build up a filter that only includes a specific set of tables.

The `<table>` task permits almost total control over the mapping details of the entity. You can select a specific table by catalog, schema, and name. You can specify the class name that it should take, how the primary key column(s) relate to that class's properties, the primary key generator that it should use, the column types and properties that they are associated with, and the details of the associations formed by foreign key relationships. Our simple example places the generated entity into an appropriate class, with a nonstandard primary key property name and a nonstandard type mapping for one of the columns (note). It also excludes one of the columns (audit) from the entity model.

```
<table schema="PUBLIC"
       name="NOTEPAD"
       class="com.hibernatebook.tools.Notepad">
   <primary-key>
      <column name="id"
              jdbc-type="INTEGER"
              property="notepadPk"
              type="int"/>
   </primary-key>
   <column name="note" jdbc-type="VARCHAR" type="char" property="note"/>
   <column name="audit" exclude="true"/>
</table>
```

If it looks like you will have to manage the reverse engineering process to this level of detail, it may in fact be more appropriate to create some or all of the mapping files manually, which gives you total control over the specification of those entities. Complex specification of mapping information in the reverse engineering file is really only appropriate if it is for exceptional classes when the general cases are common; or if you expect to need to regenerate the model from the schema very frequently in response to changes in the schema details.

Templates

With the exception of `<hbmtemplate>` and `<query>`, all the Ant exporter tasks take the meta-model information from the configuration, pass it to a set of FreeMarker templates, and write the output to files. For more information on the FreeMarker template scripting language, see the FreeMarker site at http://freemarker.sourceforge.net.

■**Note** Earlier versions of Hibernate Tools used the Velocity template language instead of FreeMarker, but in other respects, they behaved in the same way.

If the existing exporters do not meet your needs, you can specify your own additional code generation tasks using the `<hbmtemplate>` task. Table B-11 shows the attributes that can be supplied to this task.

Table B-11. *The Properties Available to the* `<hbmtemplate>` *Exporter Element*

Property	Default	Description
destdir		If set, overrides, for this exporter only, the destination directory specified on the Hibernate tool task.
exporterclass		Specifies a class to use to generate output. It will be invoked once only, and the configuration object, output directory, template path and prefix, and any configuration properties will be passed into it.
filepattern		When using templates, represents the FreeMarker macro that should be used to determine the file name for the entity being processed.
template		Specifies the template to use to generate output. It will be invoked for each of the entities in the configuration metamodel.
templatepath		Specifies the path from which your template(s) will be loaded, overriding any default location.
templateprefix		Specifies an optional prefix to your template file name within the template path. It is set by the standard tasks, so overriding this allows you to import your own tasks instead.

Again, you have two options when carrying out this process. You can set the `exporterclass` attribute to the name of the class to be used to carry out the export process. This class must implement the `org.hibernate.tool.hbm2x.Exporter` interface. This is passed a reference to the current configuration object and any other attributes that were set on the `<hbmtemplate>` task.

Alternatively, you can specify the name of a FreeMarker template to be used in processing the configuration object and the name of a prefix.

```
<hbmtemplate
   destdir="generated_txt"
   templateprefix="foo"
   template="template/MyTemplate.ftl"
   filepattern="{package-name}/{class-name}.txt">
```

Note that `filepattern` contains FreeMarker macros that will be expanded at run time to determine the appropriate file names for the tool's output. The task will search for this file on the classpath, and then as a file resource.

If the configuration object does not contain some of the information that you need in order to produce the desired output, you can also specify additional arbitrary details using a standard Ant property or property set. Here's an example:

```
<hbmtemplate …>
   <property key="bar" value="BAR!"/>
</hbmtemplate>
```

In addition to any properties you add to the template task yourself, you will have access to the scripting variables listed in Table B-12.

Table B-12. *The Standard Scripting Variables Available to a Template Task*

Variable	Description
artifacts	An instance of org.hibernate.tool.hbm2x.ArtifactCollector that can be populated with values to reflect the actions taken during output generation
c2h	An instance of the org.hibernate.tool.hbm2x.Cfg2HbmTool class providing helper methods for converting configuration object values into Hibernate mapping files
c2j	An instance of the org.hibernate.tool.hbm2x.Cfg2JavaTool class providing helper methods for converting configuration object values into Java class files
cfg	A reference to the configuration object
outputdir	The path specified as the <hbmtemplate> element's destdir attribute
template_path	A list of the paths to directories containing FreeMarker templates

The standard Hibernate Tools exporter tasks are implemented in much the same way. Although we haven't shown this when discussing their attributes earlier, all the exporter tasks support the templatepath and templateprefix attributes, allowing you to override their default behavior by instructing them to use a different set of FreeMarker macros than those included in the hibernate-tools.jar file. All attributes also support the use of property sets to pass in information that is required by your custom macros but isn't available from the configuration object.

A very simple FreeMarker script is shown in the following code. This is not very useful in itself, as the first four variables simply display their hashcode representations from the default toString() implementation, but it provides you with a syntactically valid starting point for exploration of the code generation tools.

```
Configuration object: ${cfg}
Artifacts object: ${artifacts}
Cfg2Hbm Helper: ${c2h}
Cfg2Java Helper: ${c2j}
Output Directory: ${outputdir}
Template path: ${template_path[0]}
```

Configuring the Classpath

There are two distinct classpaths to consider when setting up the Hibernate Tools Ant tasks: the classpath of the task *definition*, and the classpath to be used by the *tasks*. The task definition needs to have in its classpath the Hibernate Tools JAR file, the Hibernate libraries, the Hibernate Annotations libraries, and the JDBC driver that will be used to access the database. A typical configuration of this classpath is as follows:

```
<path id="classpath.base">
   <pathelement location="${hibernate.path}"/>
   <fileset dir="${hibernate.lib}" includes="**/*.jar"/>
      <pathelement location="${hibernate.annotations.path}"/>
   <fileset dir="${hibernate.annotations.lib}" includes="**/*.jar"/>
</path>
<path id="classpath.tools">
   <path refid="classpath.base"/>
   <pathelement location="${hibernate.tools.path}"/>
   <pathelement location="${jdbc.driver.path}"/>
</path>
```

The task definition (as shown earlier in this section) would use the classpath with the ID `classpath.tools`.

The tasks themselves will need access to two additional sets of resources: the configuration file(s) and the compiled classes.

```
<path id="classpath.apps">
   <path refid="classpath.base"/>
   <pathelement path="${src}"/>
   <pathelement path="${bin}"/>
   <pathelement location="${jdbc.driver.path}"/>
</path>
```

The configuration files will include the `hibernate.cfg.xml` and/or `hibernate.properties` files, along with any log4j configuration files, cache configuration files, and applicable XML mapping files.

If you are using annotations in any of your tasks, you will need to ensure that the task is assigned a dependency upon the compiled POJOs—annotations cannot be read at run time from Java source files, only from compiled classes.

Summary

In this appendix, we have discussed the installation and use of Hibernate Tools, including the Eclipse plug-in and the Ant tasks. Together, these remove most of the need to manually create boilerplate configuration code.

In Appendix C, we discuss how Hibernate can be used as the data access layer within the Spring Framework.

■ ■ ■

Hibernate and Spring

The Spring Application Framework offers developers an environment that ties together numerous APIs into a coherent whole. Spring applies the philosophy of "dependency injection" by providing appropriate configurable wrapper classes for all sorts of popular Java libraries.

The standard Spring API is immense, and its standardized approach to dependency management means that any existing API can in principle become a "Spring" API. If you want a good introduction to using Spring, then we recommend the excellent *Pro Spring*, by Rob Harrop and Jan Machacek (Apress, 2005). For an overview, visit the Spring web site at `http://springframework.org`.

In view of its scope, we cannot and do not make any attempt to teach you even the basics of the Spring Framework in this appendix—instead, we assume that you are already familiar with Spring in general, and offer a focused introduction to the Hibernate-related components.

Throughout this appendix, we refer to a simple sample application that represents a "newsstand" of papers consisting of sets of articles. At the end of this appendix, we include the complete Spring bean configuration file for the example application; and as with all the examples in this book, the entire application itself can be downloaded from the Apress web site (`www.apress.com`).

Spring Libraries

The Spring Framework essentially provides wrappers and utility classes for working with various other frameworks, as well as some of its own implementations. For example, it provides its own model-view-controller (MVC) patterned web application framework, and also supports Struts and vanilla JSPs.

Spring is distributed as a small set of JAR files that contain the Spring-specific functionality. The third-party components are made available with the "including dependencies" distribution of Spring, but can also be downloaded independently from their respective web sites.

It is possible to copy all the libraries that Spring is capable of using into the classpath, but this is a somewhat inelegant approach. We prefer to pick the core Spring JARs and add to them only what is necessary to support the application being built.

For our example application, we have therefore included the following sets of libraries:

- The Hibernate JAR and its required dependencies

- The Hibernate Annotations JAR and its required dependencies

- The Java Standard Template Library JARs

- The HSQLDB driver JAR

- The Spring JAR

This requires the list of JAR files shown in Listing C-1 to be included in the `WEB-INF/lib` directory of our example application.

Listing C-1. *The Required JAR Files*

```
antlr-2.7.6rc1.jar
asm-attrs.jar
asm.jar
cglib-2.1.3.jar
commons-collections-2.1.1.jar
commons-logging-1.0.4.jar
dom4j-1.6.1.jar
ehcache-1.1.jar
ejb3-persistence.jar
hibernate-annotations.jar
hibernate3.jar
hsqldb.jar
jdbc2_0-stdext.jar
jstl.jar
jta.jar
spring.jar
standard.jar
xml-apis.jar
```

Configuring Hibernate from a Spring Application

A conventional Hibernate application needs access to its database and the entity mapping information. The point of access to a fully configured Hibernate environment is the session factory, from which `Session` objects are obtained. Spring provides a bean to represent the session factory, but provides a few additional options in order to configure its resources.

In our example application, we take the line of least resistance and use a Hibernate configuration file (`hibernate.cfg.xml`) to represent both the mapping information and the database configuration. For easy reference when setting up a Spring application, we show a sample Hibernate configuration file in Listing C-2.

Listing C-2. *Familiar Territory: A Standard Hibernate Configuration File Used in Spring*

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
```

```
<hibernate-configuration>
    <session-factory>
        <property name="connection.driver_class">
            org.hsqldb.jdbcDriver
        </property>
        <property name="connection.url">
            jdbc:hsqldb:file:/spring/db/springdb;SHUTDOWN=true
        </property>
        <property name="connection.username">sa</property>
        <property name="connection.password"></property>
        <property name="hibernate.connection.pool_size">0</property>
        <property name="show_sql">true</property>
        <property name="dialect">org.hibernate.dialect.HSQLDialect</property>

        <mapping class="com.hibernatebook.spring.Paper"/>
        <mapping class="com.hibernatebook.spring.Article"/>
    </session-factory>
</hibernate-configuration>
```

Spring represents the configured session factory as a `LocalSessionFactoryBean`. Our example application uses annotations to manage the mappings, so we specify that the Hibernate `AnnotationConfiguration` type should be used in our bean instead of the default `Configuration`.

■Caution Spring maintains two sets of Hibernate packages: `org.springframework.orm.hibernate...` for Hibernate 2 functionality, and `org.springframework.orm.hibernate3...` for Hibernate 3 functionality—a single-character difference. Be careful to select the correct one, as debugging the `ClassNotFound` and similar exceptions that result if you use the wrong one can be extremely time-consuming!

We also specify the location and name of the configuration file relative to the classpath as indicated by the `classpath:` prefix (see Listing C-3).

Listing C-3. *Configuring a Session Factory Bean in Spring*

```
<bean id="sessionFactory"
    class="org.springframework.orm.hibernate3.LocalSessionFactoryBean">
    <property name="configurationClass"
        value="org.hibernate.cfg.AnnotationConfiguration" />
    <property name="configLocation"
        value="classpath:hibernate.cfg.xml" />
</bean>
```

As noted, our simple web application derives its database connection details from the Hibernate configuration file. However, a larger web application typically needs to provide database resources to other applications, in which case a Hibernate-specific configuration file

is not the appropriate location for its details to be stored. Moreover, a well-behaved web application will draw its database configuration from a JNDI-provided `DataSource` object so that connection details can be uniformly managed at deployment time.

Spring allows data sources to be managed centrally as beans, and if a `JndiObjectFactoryBean` bean is used, it can in turn draw its details from JNDI. The `LocalSessionFactoryBean` therefore provides a `dataSource` property into which the appropriate Spring `DataSource` bean can be injected.

Typically, to manage a data source from within the Spring configuration, but without deferring the details to a JNDI resource, you would use the `DriverManagerDataSource` bean (see Listing C-4).

Listing C-4. *Configuring a Typical* `BasicDataSource` *Bean*

```
<bean id="sampleDataSource"
   class="org.springframework.jdbc.datasource.DriverManagerDataSource"
   destroy-method="close">
   <property name="driverClassName">
      <value>org.hsqldb.jdbcDriver</value>
   </property>
   <property name="url">
      <value>
         jdbc:hsqldb:file:/spring/db/springdb;SHUTDOWN=true
      </value>
   </property>
   <property name="username" value="sa"/>
   <property name="password" value=""/>
</bean>
```

Alternatively, if the data source resources are to be drawn from an existing JNDI-accessible data source, then the Spring `JndiObjectFactoryBean` should be used to represent the data source (see Listing C-5).

Listing C-5. *Configuring a Typical* `JndiObjectFactoryBean`

```
<bean id="sampleDataSource"
   class="org.springframework.jndi.JndiObjectFactoryBean">
   <property name="jndiName" value="java:comp/env/jdbc/spring"/>
</bean>
```

It is not just the connection details that can be migrated from the Hibernate configuration file into the Spring configuration. The property attributes and the mappings (class names or mapping file names) can also be assigned during the configuration of a `LocalSessionFactory` bean (see Listing C-6).

Listing C-6. *Configuring Hibernate Purely from Spring*

```
<bean id="sampleSessionFactory"
   class="org.springframework.orm.hibernate3.LocalSessionFactoryBean">
   <property name="dataSource" ref="sampleDataSource"/>
   <property name="mappingResources">
      <list>
         <value>com/hibernatebook/spring/Paper.hbm.xml</value>
         <value>com/hibernatebook/spring/Article.hbm.xml</value>
      </list>
   </property>
   <property name="hibernateProperties">
      <props>
      <prop key="hibernate.connection.pool_size">0</prop>
      <prop key="hibernate.show_sql">true</prop>
      <prop key="hibernate.dialect">org.hibernate.dialect.HSQLDialect</prop>
      </props>
   </property>
</bean>
```

Note that in Listing C-6, purely in order to demonstrate the use of mapping files in a `LocalSessionFactoryBean` configuration, we omit the specification of a Hibernate `AnnotationConfiguration` for the `configurationClass` property, causing it to default to the normal (mapping file–based) Hibernate `Configuration` object.

Typically, the mappings themselves are specified in the conventional Hibernate manner through XML mapping files or Java annotations. It would be entirely possible to arrange to configure these externally, but no default Spring classes are provided to achieve this, and it is difficult to see any obvious benefit that would accrue from such an approach.

Using Hibernate in Your Spring Beans

With your session factory configured as a Spring bean, you can now go on to create DAOs that take advantage of Hibernate's functionality. Here, Spring really starts to come into its own, as it provides you with the invaluable `HibernateDaoSupport` class to form the basis of your DAOs (see Listing C-7).

Listing C-7. *Declaring the Interface for Our DAO*

```
package com.hibernatebook.spring.dao;

import java.util.List;
import com.hibernatebook.spring.Article;
import com.hibernatebook.spring.Paper;
```

```
public interface PaperDao {
    public List<Paper> getAll();
    public void createPaper(final Paper paper);
    public Paper getPaper(final Integer paperId);
    public Paper createArticle(final Integer paperId,final Article article);
}
```

Ideally, you should define an interface to specify the methods that your DAO will contain. Our sample application requires a single DAO with a few simple methods (shown in Listing C-7).

With the interface clearly specified, your DAO class should then extend the HibernateDaoSupport class and implement the interface that you have defined. Extending HibernateDaoSupport provides a number of get/set pairs for necessary Spring attributes (specifically the sessionFactory attribute) and various helper methods for implementing your DAO. A typical DAO method implementation using these methods is shown in Listing C-8.

Listing C-8. *Implementing* getAll() *Using the Methods from the* HibernateDaoSupport *Class*

```
@SuppressWarnings("unchecked")
public List<Paper> getAll() {
    Session session = getSession();
    List<Paper> papers = (List<Paper>)session.createQuery("from Paper").list();
    releaseSession(session);
    return papers;
}
```

The getSession() and releaseSession() methods are derived from the DAO class. They are roughly equivalent to the session factory's openSession() method and the session's close() method, respectively.

The HibernateDaoSupport class also provides access to an appropriately configured helper object: HibernateTemplate. Using this object, the preceding code can be rewritten as shown in Listing C-9.

Listing C-9 *Implementing the* getAll() *Method Using the* HibernateTemplate *Class*

```
@SuppressWarnings("unchecked")
public List<Paper> getAll() {
    return (List<Paper>)getHibernateTemplate().find("from Paper");
}
```

This notably reduces the amount of boilerplate session management code that is required to process this simply query. HibernateTemplate provides a set of methods (shown in Table C-1) that allow you to carry out most of the basic Hibernate operations in a similar single line of code.

Table C-1. *The Core* HibernateTemplate *Methods*

Method	Description
bulkUpdate()	Performs a bulk update or delete on the database, according to the provided HQL query and entities
contains()	Determines whether the given object exists as an entity in the database
delete()	Deletes an entity from the database
find()	Carries out an HQL query
get()	Obtains an entity by its id (primary key)
persist()	Saves an entity to the database
refresh()	Refreshes an entity from the database
save()	Saves an entity to the database
saveOrUpdate()	Saves an entity to the database or updates it as appropriate
update()	Updates an entity in the database

When a more complex set of operations is required than can be achieved in a single line, the execute() method is used to invoke an instance of HibernateCallback. Our example application uses HibernateCallback to implement its createArticle() method. This is shown in Listing C-10.

Listing C-10. *Invoking a* HibernateCallback *Object from the Template's* execute() *Method*

```
public Paper createArticle(final Integer paperId,final Article article) {
   HibernateCallback callback = new HibernateCallback() {
      public Object doInHibernate(Session session) {
         Paper paper = (Paper)session.get(Paper.class,paperId);
         paper.addArticle(article);
         session.update(paper);
         return paper;
      }
   };
   return (Paper)getHibernateTemplate().execute(callback);
}
```

Although this allows us to invoke more complex code from the HibernateTemplate class, and HibernateTemplate obviates the need for specific management of the session, it doesn't make the implementation particularly terse, and probably makes it slightly harder to understand. Listing C-11 shows how the same method can be implemented without using HibernateTemplate.

Listing C-11. *The* createArticle() *Method Without* HibernateTemplate

```
public Paper createArticle(final Integer paperId,final Article article) {
   Session session = getSession();
   Paper paper = (Paper)session.get(Paper.class,paperId);
   paper.addArticle(article);
   session.update(paper);
   releaseSession(session);
   return paper;
}
```

Regardless of how you use it, configuring your HibernateDaoSupport-derived template is extremely simple. The basic requirement is that you provide its sessionFactory property with a session factory bean from which to obtain Hibernate Session objects (see Listing C-12).

Listing C-12. *Configuring a* HibernateDaoSupport-*Derived Bean*

```
<bean id="sampleDao"
   class="com.hibernatebook.spring.dao.PaperDaoImpl">
   <property name="sessionFactory" ref="sessionFactory"/>
</bean>
```

In practice, however, this is usually made somewhat more complex by the need to declare the transactional behavior that applies to the DAO's methods.

Declarative Transaction Management

The getAll() method as implemented in Listings C-8 and C-9 omits any explicit transaction management. It is entirely possible to manage transactions directly—you have access to the Hibernate Session object, and this can be used as shown in Listing C-13.

Listing C-13. *Explicit Transaction Management (Not Recommended)*

```
@SuppressWarnings("unchecked")
public List<Paper> getAll() {
   Session session = getSession();
   session.beginTransaction();
   List<Paper> papers = (List<Paper>)session.createQuery("from Paper").list();
   session.getTransaction().commit();
   releaseSession(session);
   return papers;
}
```

However, while this is possible, it is not recommended—the use of OpenInViewInterceptor or OpenInViewFilter can prevent this code from behaving as you might expect, as can various other indirectly applied beans. This is equally applicable when you are using HibernateTemplate.

Because of these risks, you should favor the use of declarative transaction management. With this, the beans' methods are marked as being the boundaries of transactions, and the appropriate transaction isolation level and propagation behavior can be specified.

To support transactional behavior, a transaction manager bean must first be applied to the session factory. Typically, this would be a `HibernateTransactionManager` in a self-contained application, or a `JtaTransactionManager` in an environment in which the container is managing transactions. Our application uses the Hibernate transaction manager, as declared in Listing C-14.

Listing C-14. *Declaring the* `HibernateTransactionManager` *Bean*

```
<bean id="transactionManager"
    class="org.springframework.orm.hibernate3.HibernateTransactionManager">
    <property name="sessionFactory" ref="sessionFactory"/>
</bean>
```

The transaction manager must be notified of the session factory in use so that it can manage the transactions of the database connection configured in the session factory. If you want to be able to use nested transactions so that multiple calls to transactional methods can be made from a method that is itself enclosed in a transaction, you must set the `nestedTransactionAllowed` property on the `HibernateTransactionManager` bean. Note that Hibernate does not support the use of savepoints within nested transactions because it is unable to rollback the session cache's state.

The transaction boundaries are applied to a bean by wrapping it in a proxy class that then honors the original bean's API as declared in its interface(s). Typically, the basis of the proxy is therefore declared as an abstract bean so that it can be applied to multiple DAO beans as required. For illustrative purposes, our example application also uses this approach (see Listing C-15).

Listing C-15. *Declaring the Default Transactionality for Our DAO Beans*

```
<bean id="daoTxTemplate" abstract="true"
class="org.springframework.transaction.interceptor.TransactionProxyFactoryBean">
    <property name="transactionManager" ref="transactionManager"/>
    <property name="transactionAttributes">
        <props>
            <prop key="create*">
                PROPAGATION_REQUIRED,ISOLATION_READ_COMMITTED
            </prop>
            <prop key="get*">
                PROPAGATION_REQUIRED,ISOLATION_READ_COMMITTED
            </prop>
        </props>
    </property>
</bean>
```

With the transaction manager's template prepared, the declaration of any DAO objects must be wrapped in a bean declaration derived from this template.

Our wrapped bean is shown in Listing C-16. The lines highlighted in bold are the wrapped bean's declaration—since it is only used as a property for the enclosing proxy, it does not need to be assigned an `id`—instead, the `id` of the `paperDao` proxy is used when a DAO reference is required. The proxy will honor the `PaperDao` interface declared in Listing C-7.

Listing C-16. *Wrapping the DAO Implementation Bean with Appropriate Transactionality*

```
<bean id="paperDao" parent="daoTxTemplate">
   <property name="target">
      <bean class="com.hibernatebook.spring.dao.PaperDaoImpl">
         <property name="sessionFactory" ref="sessionFactory"/>
      </bean>
   </property>
</bean>
```

Managing the Session

A familiar problem encountered when using Hibernate is the `LazyInitializationException`. This occurs when you try to access the lazily loaded attributes of a detached entity—typically when the session that loaded it has been closed. This causes problems in Spring when you want to use information obtained from a DAO in a controller in the view to which it forwards the request.

By default, a DAO derived from the `HibernateDaoSupport` class closes the session as soon as any `HibernateTemplate` methods complete (or as soon as the session is released, if you are not using `HibernateTemplate`). Entities that have been retrieved from the DAO will therefore become detached from their session. If they are then passed to a view—typically a JSP—your client code will produce a `LazyInitializationException` when it tries to access any lazy properties of the entity that were not accessed prior to completion of the original DAO method (or forced to load in some other way).

Clearly, marking all the properties of your entities as being eagerly loaded is not practical—and typically, it is not possible to determine in advance exactly which properties of your entity should be actively loaded.

Instead, Spring provides a mechanism to implement the `OpenSessionInView` pattern of behavior. This ensures that the `Session` object is retained until processing of the view is complete. Only then is it closed—it must be closed at some point to ensure that your web applications don't leak a `Session` for every user request!

The effect is that with either an `OpenSessionInViewInterceptor` in your Spring configuration file or an `OpenSessionInViewFilter` configured in your `web.xml` file, you can access lazily loaded attributes of entities acquired from your DAOs without any risk of the dreaded `LazyInitializationException`. Note that only one of the two options is required—they differ only in their internal details, not the outcome of applying them.

Generally, we use the `OpenSessionInViewInterceptor`, as it is configured like any other Spring bean. Our example application makes use of this to ensure that the lazily loaded `articles` attribute of the `Paper` entity can be accessed from the JSP view implementations (see Listing C-17).

Listing C-17 *Declaring the* `OpenSessionInViewInterceptor` *for Use in the Spring View*

```
<bean name="openSessionInViewInterceptor"
class="org.springframework.orm.hibernate3.support.OpenSessionInViewInterceptor">
    <property name="sessionFactory" ref="sessionFactory"/>
    <property name="singleSession" value="true"/>
</bean>
```

The interceptor needs to control the behavior of the session factory, so we must provide it with a reference to the appropriate bean. We can also dictate whether a single `Session` object will be used for the entire duration of the user request. Setting this to `true` is the most efficient approach, and therefore the default; however, it has potential side effects, particularly if declarative transactions are not in use. When set to `false`, individual sessions will be acquired for each DAO operation.

The Sample Configuration File

Listing C-18 shows the full Spring configuration file of our example application as a handy reference to what's required in creating a Hibernate-based Spring application.

Listing C-18. *The Complete Spring Configuration File*

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE beans PUBLIC "-//SPRING//DTD BEAN//EN"
    "http://www.springframework.org/dtd/spring-beans.dtd">

<beans>
    <!-- Hibernate configurations -->

    <bean id="sessionFactory"
class="org.springframework.orm.hibernate3.LocalSessionFactoryBean">
        <property name="configurationClass"
            value="org.hibernate.cfg.AnnotationConfiguration" />
        <property name="configLocation"
            value="classpath:hibernate.cfg.xml" />
    </bean>

    <bean name="openSessionInViewInterceptor"
class="org.springframework.orm.hibernate3.support.OpenSessionInViewInterceptor">
        <property name="sessionFactory" ref="sessionFactory"/>
    </bean>
```

```xml
    <bean id="transactionManager"
class="org.springframework.orm.hibernate3.HibernateTransactionManager">
        <property name="sessionFactory" ref="sessionFactory"/>
    </bean>

    <bean id="daoTxTemplate"
        abstract="true"
class="org.springframework.transaction.interceptor.TransactionProxyFactoryBean">
        <property name="transactionManager" ref="transactionManager"/>
        <property name="transactionAttributes">
            <props>
                <prop key="create*">
                    PROPAGATION_REQUIRED,ISOLATION_READ_COMMITTED
                </prop>
                <prop key="get*">
                    PROPAGATION_REQUIRED,ISOLATION_READ_COMMITTED
                </prop>
            </props>
        </property>
    </bean>

    <!-- DAO configurations (note use of template) -->

    <bean id="paperDao" parent="daoTxTemplate">
        <property name="target">
            <bean class="com.hibernatebook.spring.dao.PaperDaoImpl">
                <property name="sessionFactory" ref="sessionFactory" />
            </bean>
        </property>
    </bean>

    <!-- Basic Spring MVC configurations -->

    <bean id="viewResolver"
class="org.springframework.web.servlet.view.UrlBasedViewResolver">
        <property name="prefix" value="/WEB-INF/jsp/" />
        <property name="suffix" value=".jsp" />
        <property name="viewClass"
            value="org.springframework.web.servlet.view.JstlView"/>
    </bean>
```

```xml
    <bean id="urlMapping"
class="org.springframework.web.servlet.handler.SimpleUrlHandlerMapping">
        <property name="interceptors">
            <list>
                <ref bean="openSessionInViewInterceptor" />
            </list>
        </property>

        <property name="mappings">
            <props>
                <prop key="/viewPapers.do">viewPapersController</prop>
                <prop key="/createPaper.do">createPaperController</prop>
                <prop key="/viewArticles.do">viewArticlesController</prop>
                <prop key="/createArticle.do">createArticleController</prop>
            </props>
        </property>
    </bean>

    <!-- Message resources -->

    <bean id="messageSource"
class="org.springframework.context.support.ResourceBundleMessageSource">
        <property name="basename" value="message"/>
    </bean>

    <!-- Validators -->

    <bean id="createPaperValidator"
class="com.hibernatebook.spring.validator.CreatePaperValidator"/>

    <bean id="createArticleValidator"
class="com.hibernatebook.spring.validator.CreateArticleValidator"/>

    <!-- Controller configurations -->

    <bean id="viewPapersController"
        class="com.hibernatebook.spring.controller.ViewPapersController">
        <property name="paperDao" ref="paperDao" />
    </bean>
```

```
    <bean id="createPaperController"
class="com.hibernatebook.spring.controller.CreatePaperController">
        <property name="commandClass"
           value="com.hibernatebook.spring.form.CreatePaper"/>
        <property name="commandName" value="paperForm"/>
        <property name="formView" value="createPaper"/>
        <property name="successView" value="viewPapers"/>
        <property name="validator" ref="createPaperValidator"/>
        <property name="paperDao" ref="paperDao"/>
    </bean>

    <bean id="viewArticlesController"
class="com.hibernatebook.spring.controller.ViewArticlesController">
        <property name="paperDao" ref="paperDao"/>
    </bean>

    <bean id="createArticleController"
class="com.hibernatebook.spring.controller.CreateArticleController">
        <property name="commandClass"
value="com.hibernatebook.spring.form.CreateArticle"/>
        <property name="commandName" value="articleForm"/>
        <property name="formView" value="createArticle"/>
        <property name="successView" value="viewArticles"/>
        <property name="bindOnNewForm" value="true"/>
        <property name="validator" ref="createArticleValidator"/>
        <property name="paperDao" ref="paperDao"/>
    </bean>
</beans>
```

Summary

The Spring Framework offers excellent support for Hibernate as a persistence mechanism. It offers excellent support for creating Hibernate-based DAOs and various convenient features to smooth over the problems that you would otherwise encounter in building web applications based around Hibernate.

Upgrading from Hibernate 2

Hibernate 3 represents a major change from the ways of doing things in Hibernate 2. On the whole, it is a better product, and we applaud the Hibernate developers for their efforts. One particular group of users will be made nervous by all the changes: the existing users of Hibernate 2.

Well, there is good news, and there is . . . no bad news! Hibernate 3 has gone the extra mile to allow earlier users to get along. In this appendix, we will discuss the differences between the two versions and explain how a Hibernate 2 user can take advantage of them without conducting a major code rewrite.

Hibernate 3 does make changes: the package names have changed, the DTDs have changed, the required libraries are different, and some of the method names and signatures have been altered. Even so, we don't think that these differences will cause you much grief when upgrading to the new version.

You can run Hibernate 3 on a 1.4 or later JVM quite easily. It is possible in principle to run Hibernate 3 on a 1.3 JVM, although this will require you to recompile it from the source code and find older versions of some of the libraries that it depends upon.

Once you have read this appendix, we also recommend that you consult the Hibernate 3 Migration Guide in the Documentation section of the Hibernate web site (`http://hibernate.org`). The Hibernate team maintains and updates this section to reflect users' experiences, so you can find hints and tips gathered from developers at the cutting edge of just this sort of upgrade.

Package and DTD Changes

The package names for Hibernate 2 have changed with Hibernate 3. Hibernate 2 used a base package of `net.sf.hibernate`, while Hibernate 3 uses a base package of `org.hibernate`.

This is, in itself, a completely trivial difference—you might imagine that it is purely the result of a migration from Hibernate's hosting from SourceForge (`http://sf.net` or `http://sourceforge.net`) to their own web site (`http://hibernate.org`); but, in fact, there is another reason for the change.

Because of the package name change, it is possible for an application to use Hibernate 2 and Hibernate 3 simultaneously, allowing legacy Hibernate 2 code to run unmodified within the same application as more recent Hibernate 3–based components. If the same package name had been used, then it would be nearly impossible to achieve this.

This is not a coincidence—in addition to the package name, there are now two versions of the DTDs for the XML configuration files. Unchanged Hibernate 2 code should use the usual mapping DTD reference of

```
http://hibernate.sourceforge.net/hibernate-mapping-2.0.dtd
```

And for your new Hibernate 3 code, you will use

```
http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd
```

Similarly, for the Hibernate configuration file, your version 2 code will use

```
http://hibernate.sourceforge.net/hibernate-configuration-2.0.dtd
```

And your version 3 code will use

```
http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd
```

■**Caution** If you do not update your mapping configuration from the Hibernate 2 form to the Hibernate 3 form, the time it takes to create a configuration and session factory will increase from around 20 seconds to a matter of minutes.

Obviously, it will not be possible for you to have two configuration files with the same (default) name of `hibernate.cfg.xml`, but either version of Hibernate permits you to construct a `Configuration` object and then specify an explicit location for the configuration file using the `configure()` methods (as shown in Listing D-1). If you are using both versions of Hibernate simultaneously, you must make sure that POJOs are not passed from one version's `Session` object to another's. If you need to persist your old POJOs using the new version, you must update the older code to use Hibernate 3. For an explanation of how this sort of upgrade is supported by Hibernate 3, see the description of the "classic" API in the following section, "New Features and Support for Old Ones."

Listing D-1. *Using an Explicitly Named Configuration File in Hibernate 3*

```
File configFile = new File("hibernate3.cfg.xml");
Configuration v3Config = new Configuration();
v3Config.configure(configFile);
SessionFactory sessionFactory =
    v3Config.buildSessionFactory();

Session session = sessionFactory.openSession();
// ...
```

You should be aware that in your Hibernate 3 logic, some of the defaults for entries in the mapping file have changed. If you have logic that relies upon implicit settings, you should review your converted mapping files against the version 3 DTD to check that they will behave as expected. The most significant change is that all mappings now default to lazy loading.

New Features and Support for Old Ones

If you are a Hibernate 2 developer and you have browsed through the earlier chapters, you will have realized that Hibernate 3 offers a lot of new features. You will have also realized that some of the Hibernate 2 features that you rely on may no longer be supported in version 3. For the most part, though, this is not the case.

Changes and Deprecated Features

If you do not plan to take advantage of the Hibernate 3 features in any of your existing code, you can, as discussed, simply run the two versions side by side without concern. If you are prepared to make some changes to your existing code, then it is better to take the opportunity to update your existing code to use Hibernate 3 directly. In order to make this second choice a little easier, Hibernate 3 provides a number of "deprecated" APIs that permit fewer changes than a full-blown conversion.

This reduces the immediate impact of the change, and allows you to conduct the rest of the conversion at your leisure, while still allowing you to remove the legacy Hibernate 2 libraries from your application.

■**Tip** Hibernate exceptions are now thrown as unchecked exceptions. This will not impact existing code, but you may want to revisit APIs that explicitly declare `HibernateException`s. This change is intended to increase the clarity of API signatures by removing the need for the explicit `throws` clause in code, which uses Hibernate but does not catch its exceptions. There are ongoing debates over the relative merits of the two approaches, but certainly the change *from* checked *to* unchecked does not introduce any incompatibilities (whereas the reverse would).

Some changes to HQL have occurred between versions 2 and 3. If you have a substantial body of existing HQL syntax, you can elect to retain the old syntax. The selection is made with the `hibernate.query.factoryclass` configuration attribute, which selects the class to load for translating HQL into database queries. The options are listed in Table D-1.

Table D-1. *The HQL Processing Classes*

Query Factory Class	HQL Version
`org.hibernate.hql.ast.ASTQueryTranslatorFactory`	3 (default)
`org.hibernate.hql.classic.ClassicQueryTranslatorFactory`	2

It is not possible to switch between the two query translators within a `SessionFactory` instance. Because HQL queries are not parsed until run time[1], you will need to run extensive tests to ensure that your modified queries are correct if you decide to convert to the Hibernate 3 syntax.

---

1. Named queries were introduced in Hibernate 3. These are stored in the mapping file or the annotations, and are parsed on application initialization—so while they are still parsed at run time, you will not need to run extensive tests to spot syntactical problems with them.

The object retrieved from the session factory in Hibernate 3 implements both the pure Hibernate 3 `org.hibernate.Session` interface and a Hibernate 2–friendly `org.hibernate.classic.Session` interface. By using a classic `Session` reference instead of the usual one, you will have access to the methods that have now been deprecated. Despite their deprecated status, all of these methods are fully implemented at the present time. Most of them are backward-compatible with their Hibernate 2 counterparts—but the `delete(String)` method has changed slightly in that deletions are no longer cascaded to associated entities.

A fully converted Hibernate 3 application will not need to invoke any of these methods, so you should use a reference to the standard interface unless you're absolutely compelled by existing logic.

Other deprecated features reside in the classic package and its subpackages. Notable examples are listed in Table D-2.

Table D-2. *Feature Replacements in Hibernate 3*

Feature	Location in Hibernate 3	Use in Preference
Life cycle	`org.hibernate.classic`	Interceptor or event
Validatable	`org.hibernate.classic`	Interceptor or event

Some of the changes to Hibernate 3 have not justified the creation of a replacement class. A few methods will have been removed, replaced, or renamed. In these few cases, if you do not want to run Hibernate 2 and 3 side by side, you will be forced to update your code. When compilation produces errors, consult the javadoc API at `http://hibernate.org` to see whether the API has changed, and to determine how to alter your code to work in the new environment. In practice, there are few changes in the core API between versions 2 and 3, and the changes that do exist have well-signposted successors in the new API documents.

Additions

The `Event` class is new to Hibernate 3. If you are familiar with the `Interceptor` class (which is retained), you will have some idea of what to expect. This topic is discussed in Appendix A.

The Criteria and Filter APIs have been extended considerably. These are discussed in detail in Chapter 11.

The flexibility of the mappings has been improved. For example, the `join` element permits a single class to be represented by multiple tables. Support for stored procedures allows better integration with legacy databases. The mapping file format is discussed in Chapter 8, and support for stored procedures and other new features is discussed in Appendix A.

Changes to Tools and Libraries

As you may expect, the libraries upon which Hibernate is based have changed in version 3. Some have been added and others have been brought up to date. Rather than enumerate these here, we refer you to the `lib/README.txt` file in your Hibernate 3 distribution, which explains in detail whether or not individual libraries are required, and what purpose each serves.

Hibernate 2 provided a number of aids to the generation of POJOs, mapping files, and database schemas. Hibernate 3 has started a process of migrating to external support for these processes. Where these tools are retained, they can be found in the `org.hibernate.tool` package and its subpackages. For example, previously a facility provided by the `CodeGenerator` class to generate DDL from your mappings existed in Hibernate 2. This is still provided in Hibernate 3, but the fully qualified name of the `SchemaExport` class is now `org.hibernate.tool.hbm2ddl.SchemaExport`—but even though the tool still exists, the generation of mapping files from POJOs in Hibernate 3 would usually be conducted by an Eclipse plug-in, or the Hibernate Tools Ant tasks. See Appendix B for an extensive discussion of the use of the Eclipse plug-ins and Ant tasks.

Changes with Java 5

The latest release of Java introduced some pretty substantial changes to the language. It also introduced some incompatibilities with the previous class file format.

The only significant addition to Hibernate 3 that relies directly upon a Java 5–specific feature is the annotations support. Otherwise, a 1.4 JVM will work fine. In fact, it is even possible to build Hibernate 3 for a 1.3 JVM, although a suitable binary version is no longer distributed. The use of annotation-based mapping in Hibernate 3 is discussed in depth in Chapter 6.

Summary

In this appendix, we examined some of the changes that have been introduced with Hibernate 3, and showed how code written for Hibernate 2 can be run in parallel with Hibernate 3, or be readily adapted to run directly under Hibernate 3.

Index

P